

Utreeexo - What is it?

Quick intro

Agenda for today

Show & explain

- Smaller full nodes
- Performance wins
- The catch
- Deployment
- Aligning incentives
- How it works

Smaller full nodes

```
[I] calvin@Calvins-MacBook-Pro ~/L/A/Bitcoin> du -h .
111M ./blocks/index
744M ./blocks
7.4G ./chainstate
88K ./coldcardmk4
104K ./coldcardmk3
8.2G .
```

```
[1] calvin@Calvins-MacBook-Pro:~/L/A/Bitcoin> du -h .
```

111M	./blocks/index
744M	./blocks
7.4G	./chainstate
88K	./coldcardmk4
104K	./coldcardmk3
8.2G	.

```
[I] calvin@Calvins-MacBook-Pro ~/L/A/Bitcoin> du -h .
111M    ./blocks/index
744M    ./blocks
7.4G    ./chainstate
  0K    ./coldcardmk4
104K    ./coldcardmk3
8.2G    .
```

UTXO set lives in chainstate/

The UTXO set

Quick overview

- TXO - An Output of a TX (transaction)

Typical TX

[8ad63210d059908229f00ad177469d9f16a5d37a7ea64cede9c92684ff7cf06e](#)

DETAILS



#0 [f758ad929fbefac99ef0d1f8b13990fa75efb517ef6c919178c8af7308f145](#) 0.04852489 BTC
07:1



#0 [17dSwJytZBszAafyWrK8Pue7rsRh9s5xeJ](#) 0.04033963 BTC

#1 [bc1qwqdg6squ38e46795at95yu9atm8azzmyvckulcc7kytlcckxswvvzej](#)

1 CONFIRMATION 0.04792489 BTC

Outputs of a TX

8ad63210d059908229f00ad177469d9f16a5d37a7ea64cede9c92684ff7cf06e

DETAILS +

#0 f758ad929fbefac99ef0d1f8b13990fa75efb517ef6c919178c8af7308f145 0.04852489 BTC
07:1

#0 17dSwJytZBszAafyWrK8Pue7rsRh9s5xeJ 0.04033963 BTC

#1 bc1qwqdg6squ38e46795at95yu9atm8azzmyvckulcc7kytlcckxswvv 0.00758526 BTC
zej

1 CONFIRMATION 0.04792489 BTC

The UTXO set

Quick overview

- TXO - An Output of a TX (transaction)
- UTXO - Unspent TXO

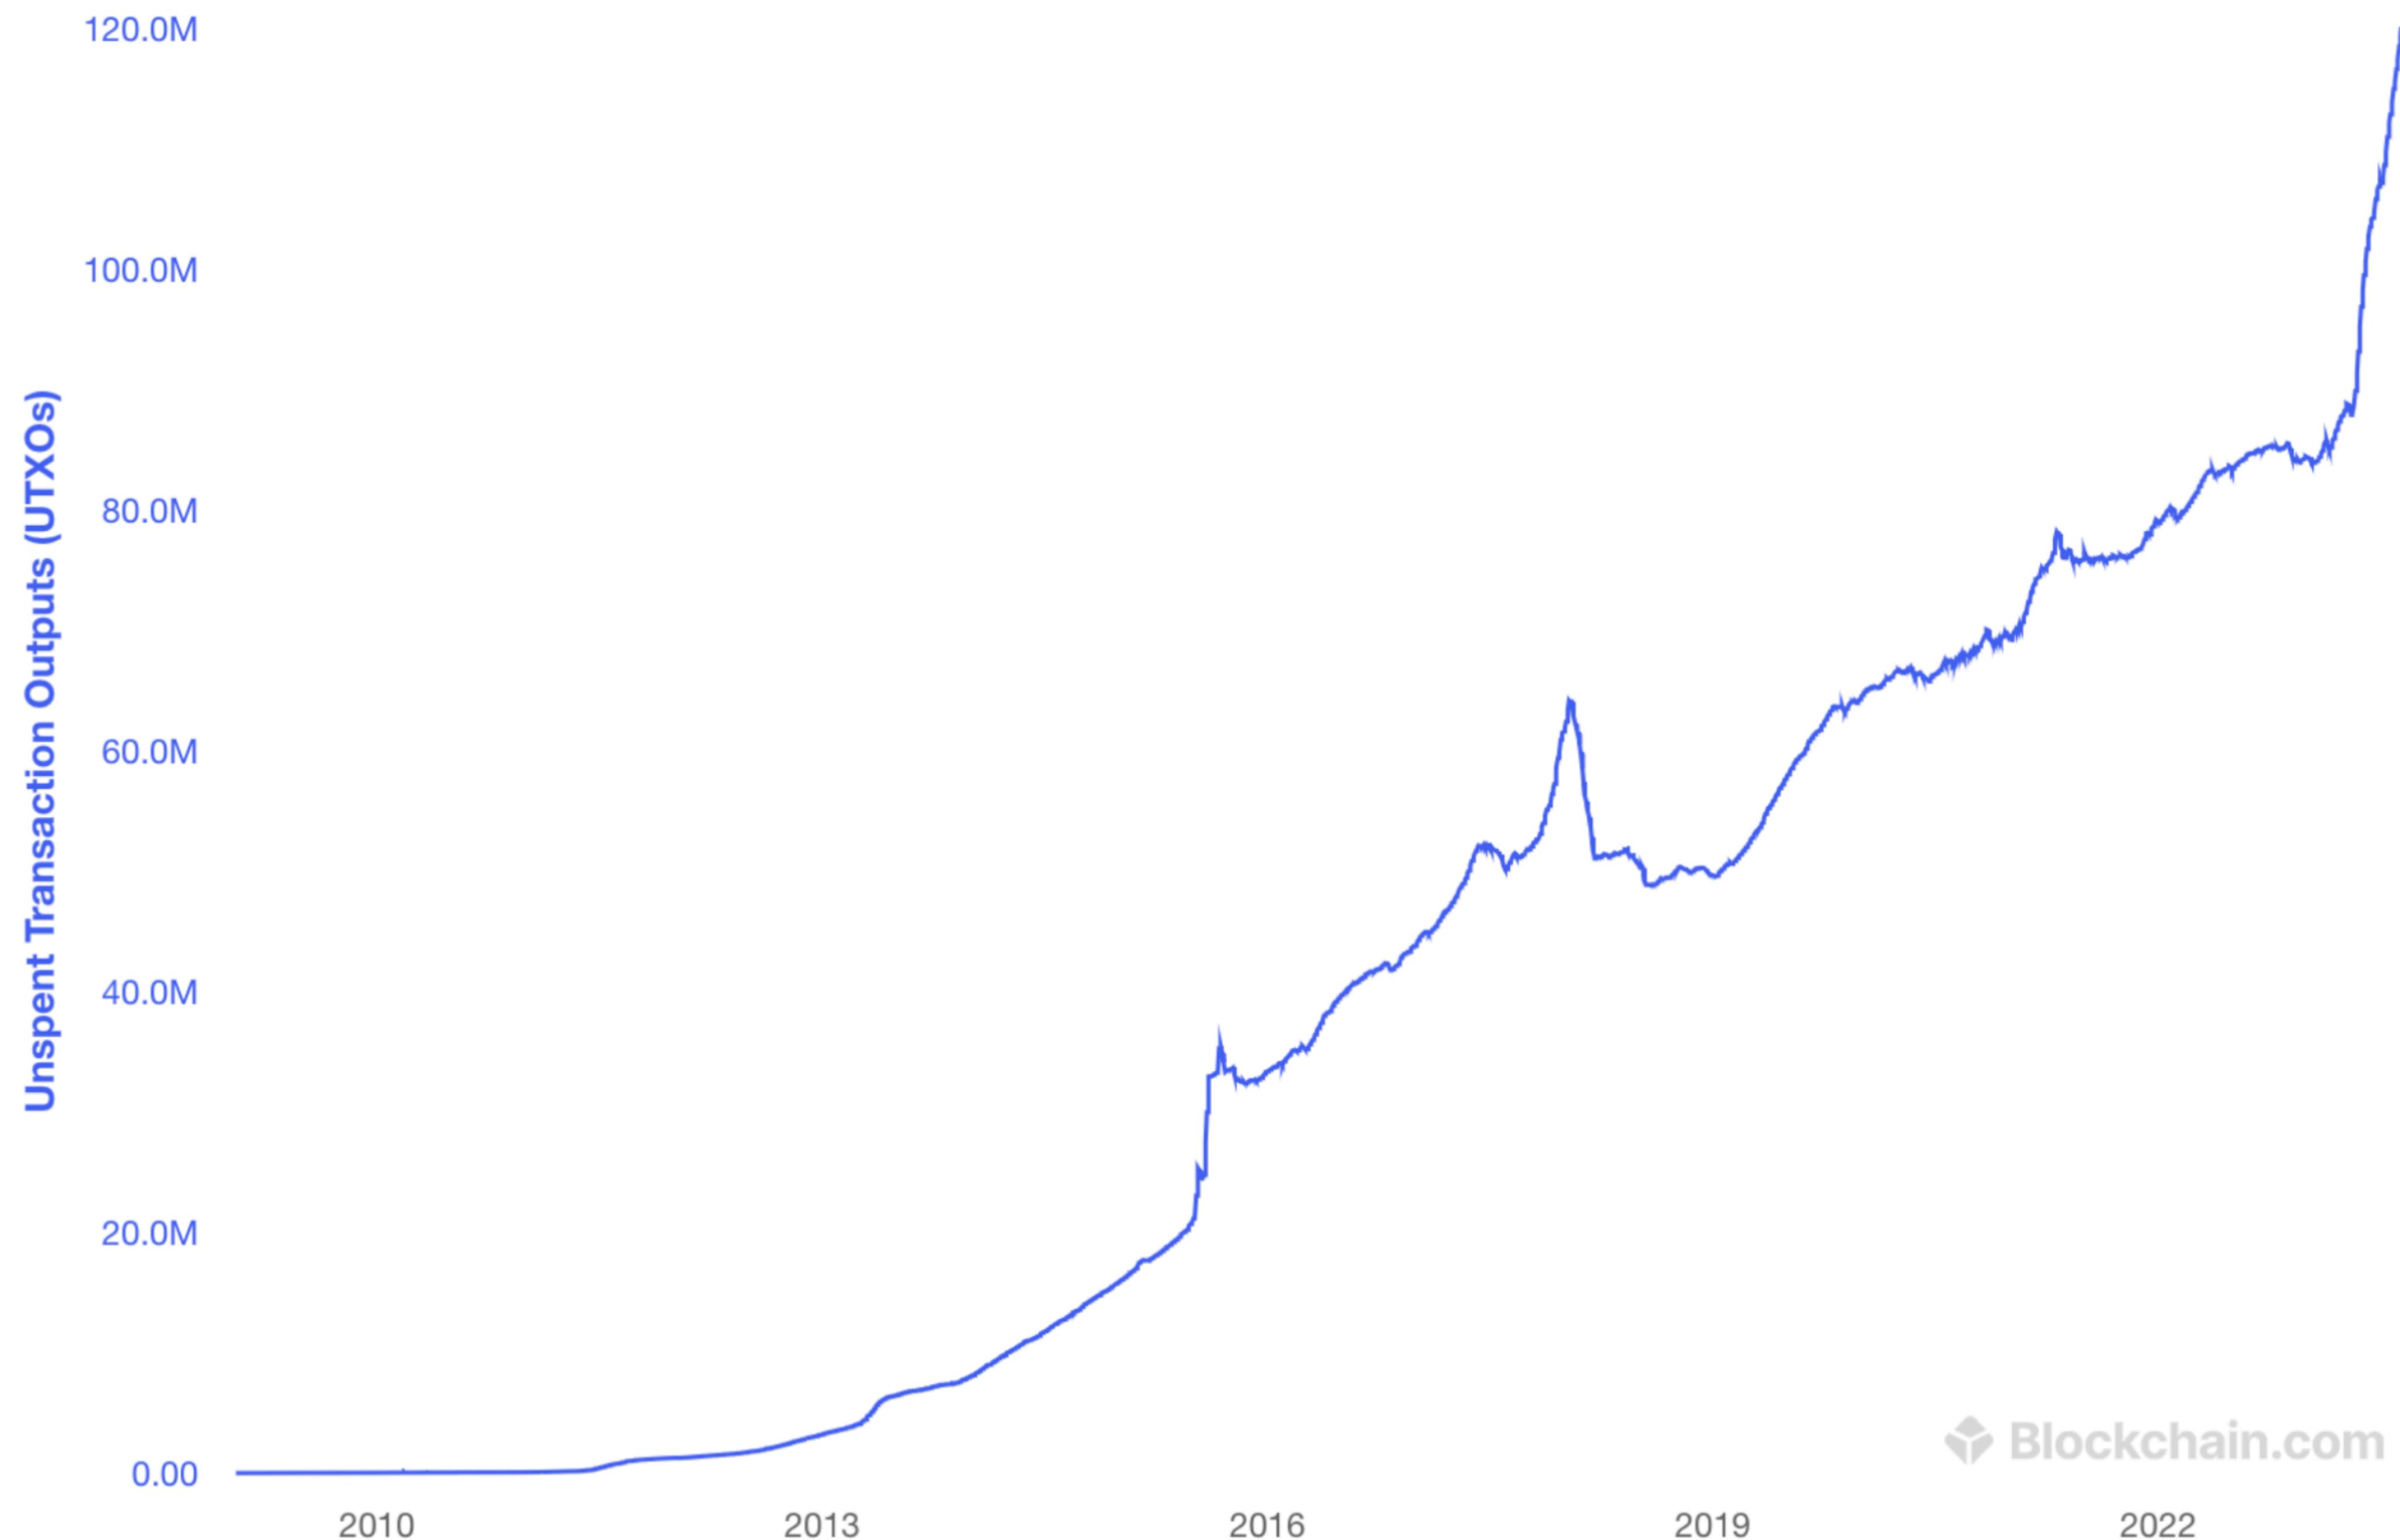
The UTXO set

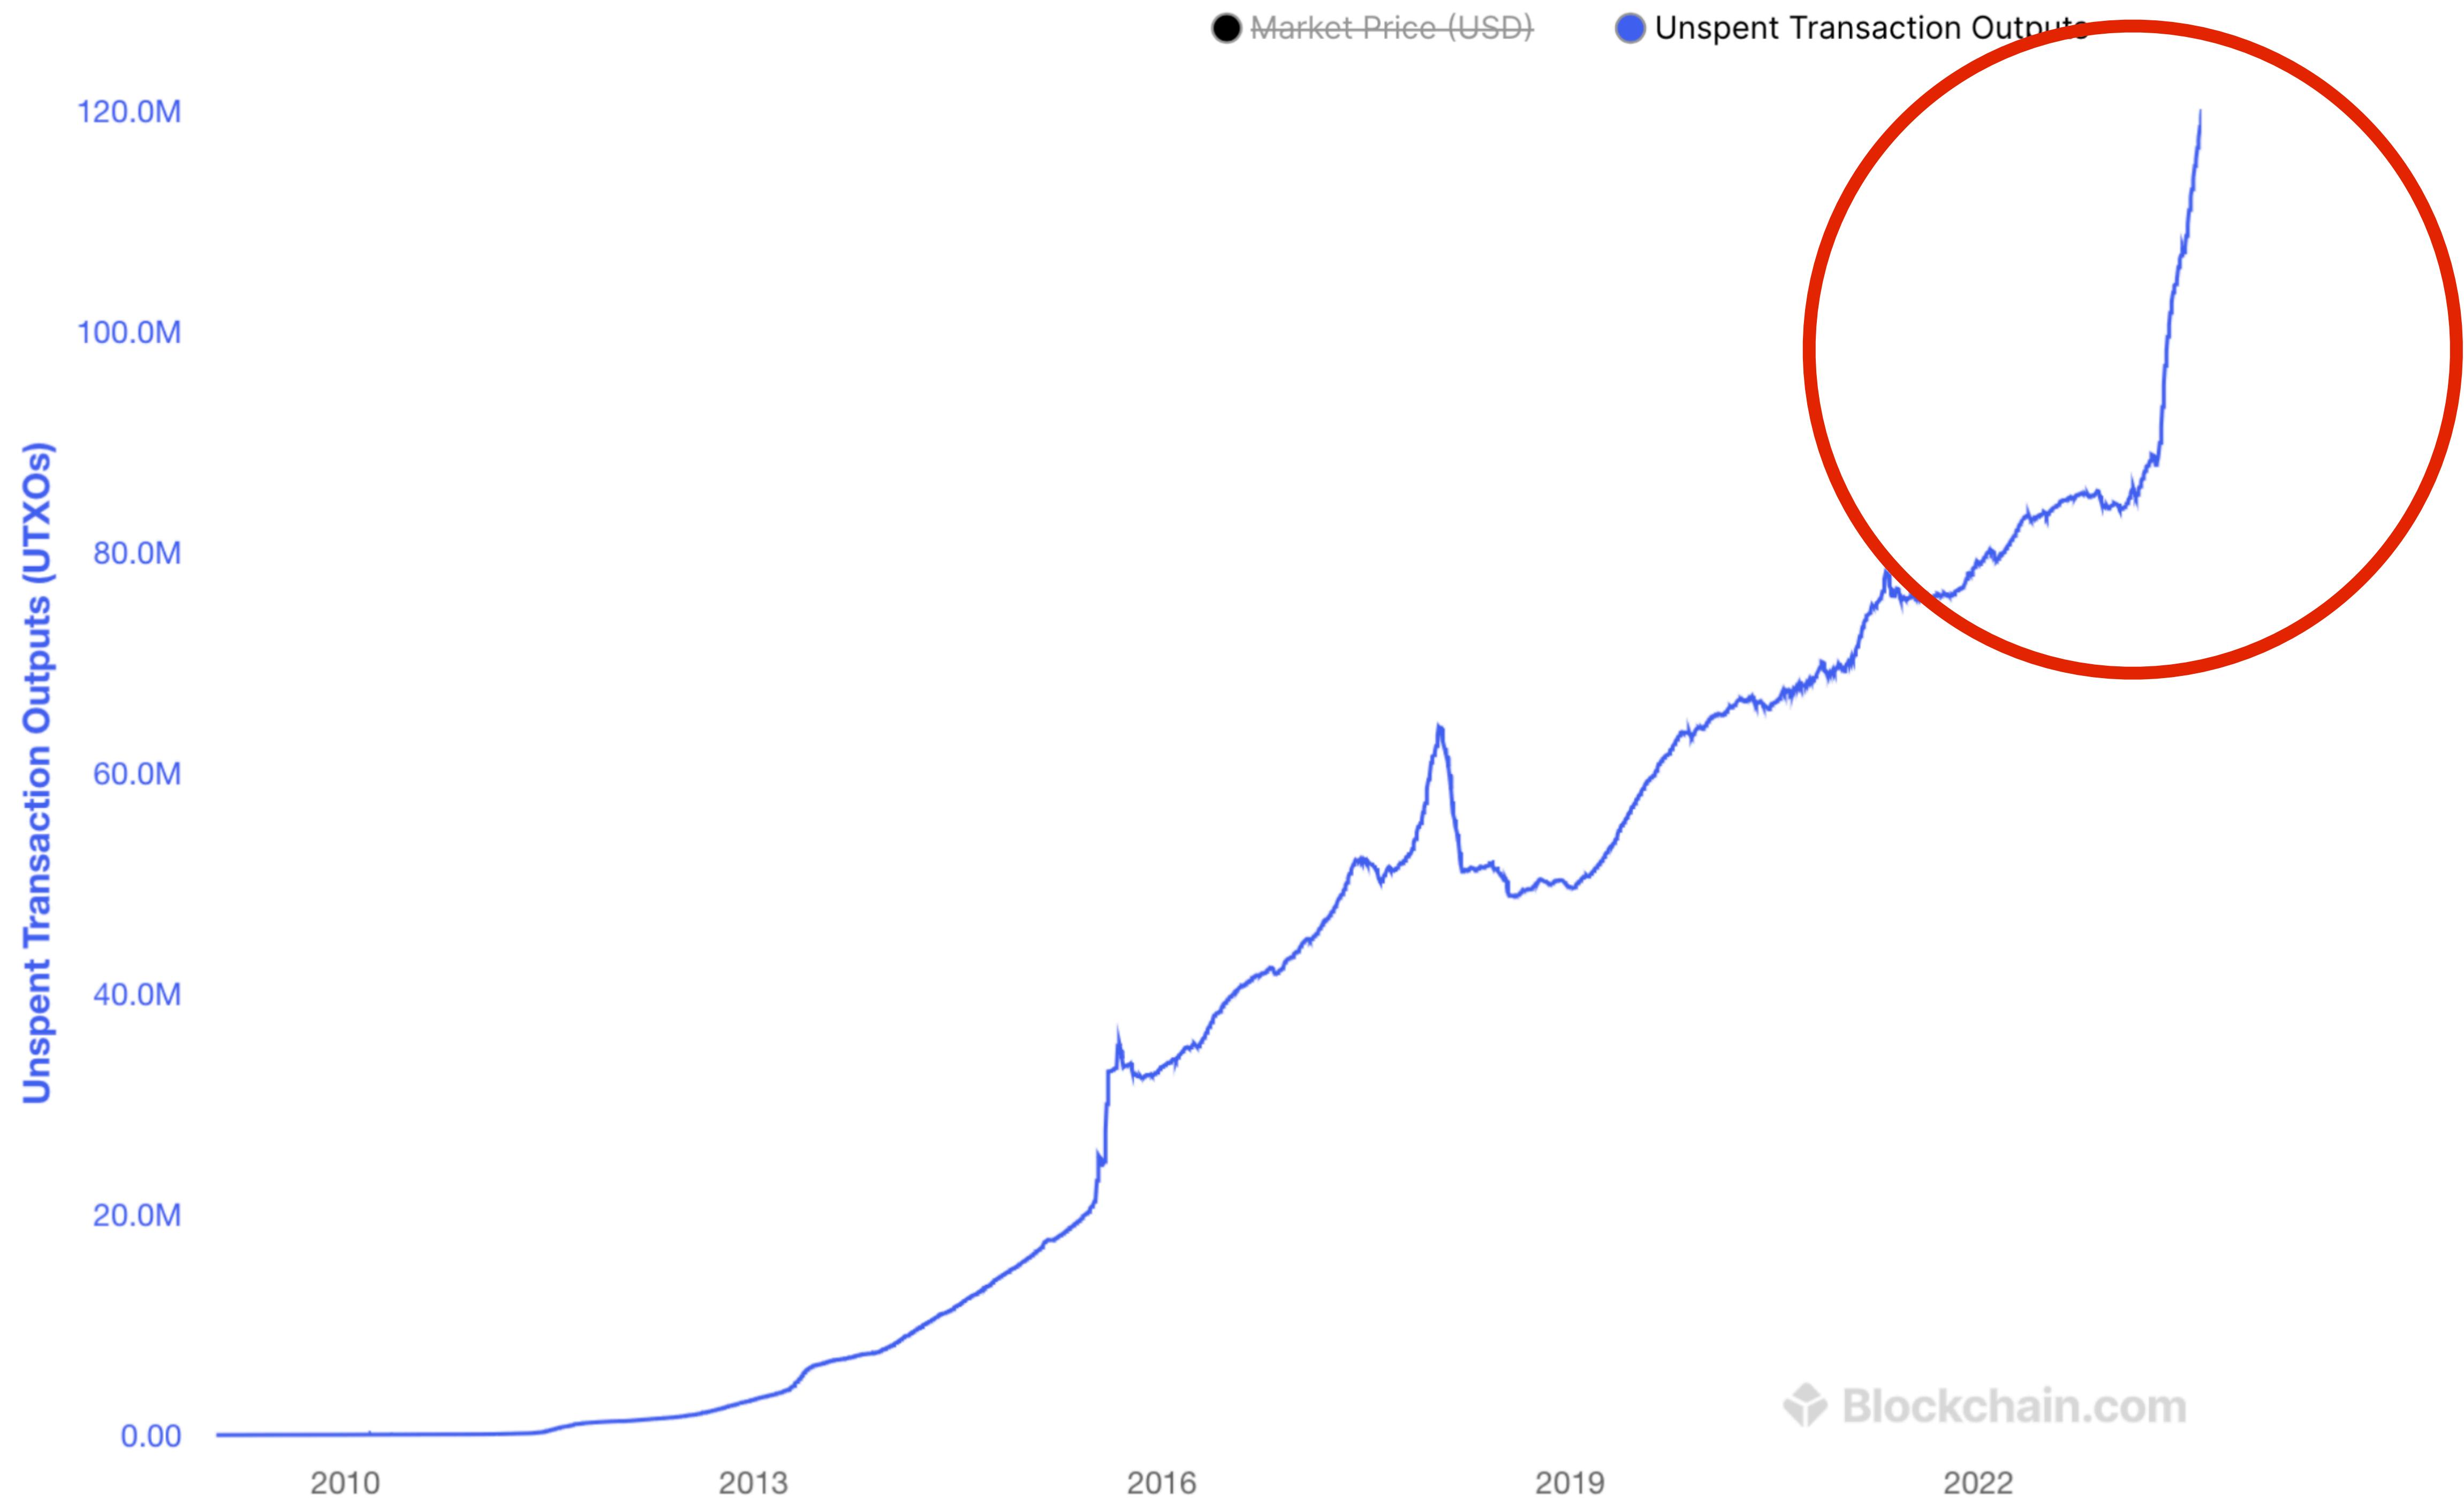
Quick overview

- TXO - An Output of a TX (transaction)
- UTXO - Unspent TXO
- UTXO set - Set of all UTXOs

**UTXO set is
unprunable**

● Market Price (USD) ● Unspent Transaction Outputs





**Turn the UTXO set into a
merkle tree**

Turn the UTXO set into a
fancy merkle tree



Merkle trees



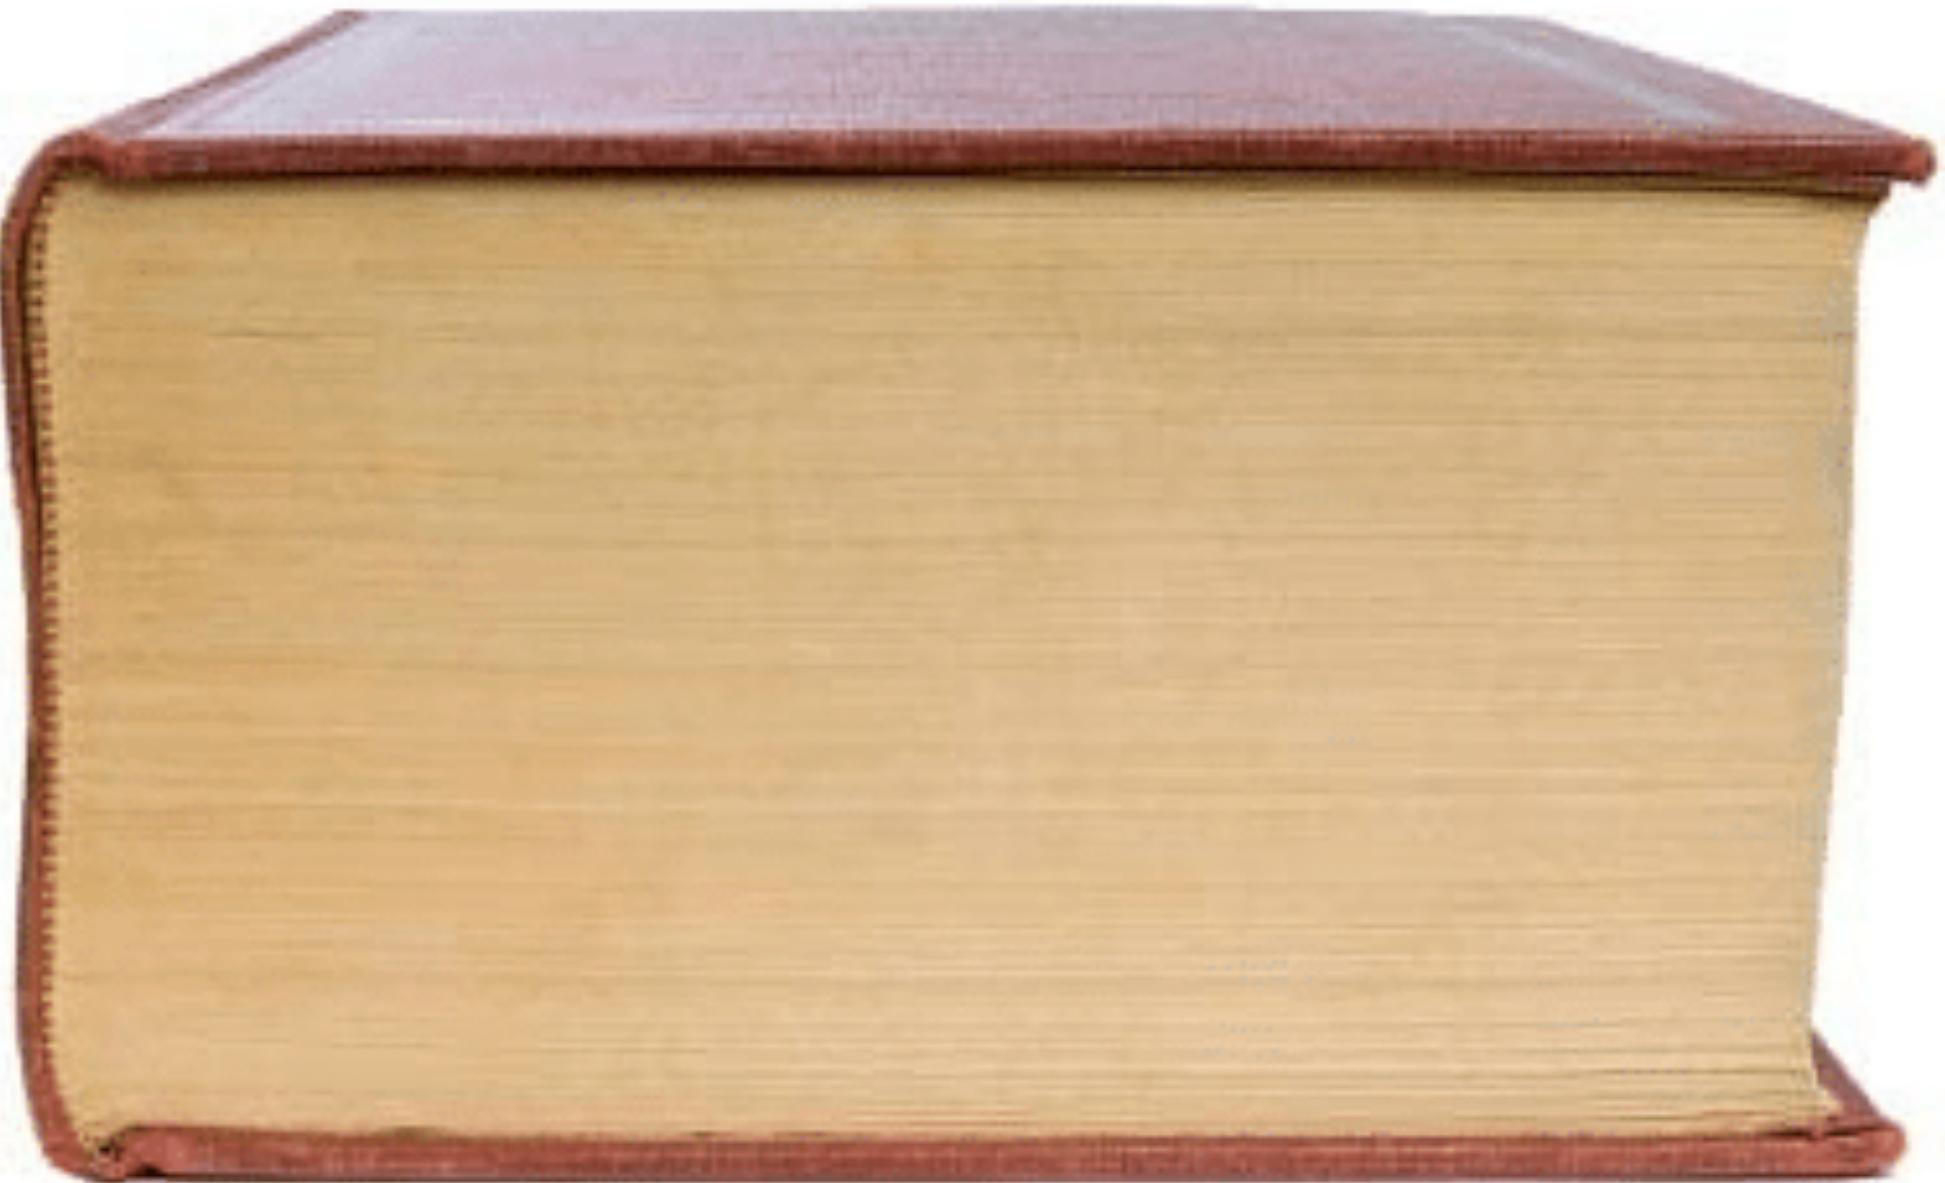
Utreeexo

The UTXO set

Why merkelize it?

- Puts a bound to the UTXO set growth

Pruned nodes



Without
Utreeexo



With
Utreeexo

Pruned nodes



Without
Utreeexo



With
Utreeexo

Pruned nodes



O(N) -> O(log2 N)

N=8 billion UTXOs, at 59B per entry

472GB -> 416B

Performance wins

**~Same Performance
on HDD or SSD**

Avoid disk i/o

No more random reads

Prove it?

Test methodology

- Compare performance degradation of pruned nodes on HDD vs SSD

36% slower

Current nodes

3% slower

Utreeexo nodes

What we flush for utreexo

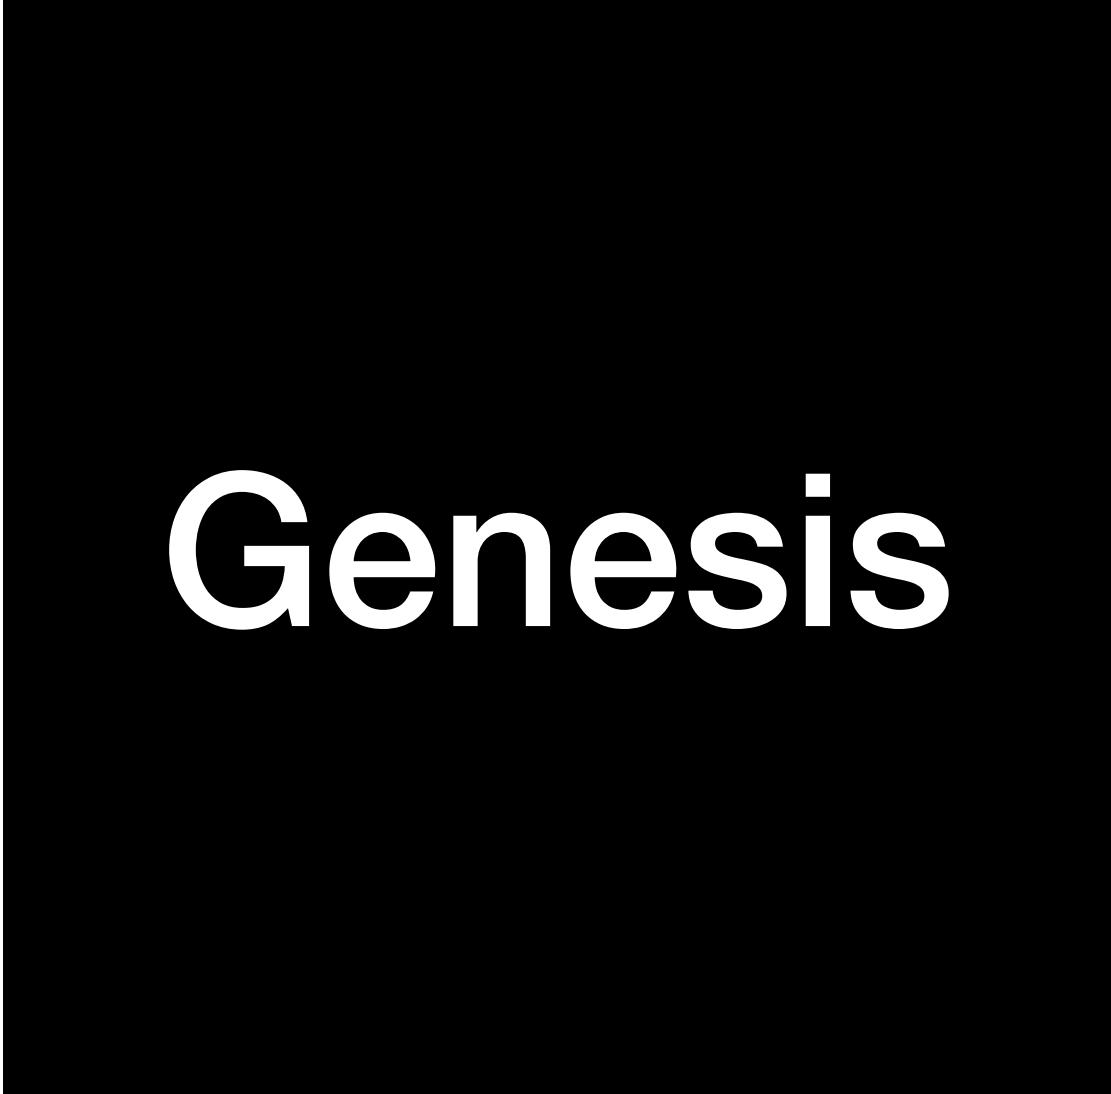
No random i/o

```
{443000, newHashFromStr("000000000000000047efc3c126c3398cfdef1609682d1492cd2909e3c0a17f"),  
    43243051, []*chainhash.Hash{  
        newLeafHashFromStr("df31c496a54a8021c38c77809da15a8bd3968970315f5a56c07f0fd9763262ff"),  
        newLeafHashFromStr("755f4cfbf5727c3e14aa6dcc6e5ab150c61101cead06fade5a886ee8d1dd377"),  
        newLeafHashFromStr("c572a92c7cd9e14ce5d7c94473ef90bee21856654eed5535314b3c117006bc83"),  
        newLeafHashFromStr("9e85567347ef0ea38c1797d8ed8e399ed0d99a07025388f6e06791ff6b88643c"),  
        newLeafHashFromStr("ab4efc738bd7b8281689e899979d44ed0ca8fe1489efa0f5a09f5b1c276fe792"),  
        newLeafHashFromStr("1c763d8f490baef0b756597dfc2da5a31d0aee324da6314b52f21707351cf51a"),  
        newLeafHashFromStr("29fc9989939a17e5d113fb1f247f3bd05606488bcb48c9cfad55a57006f2248"),  
        newLeafHashFromStr("e79f840fe1bcd6b638a53b637885dda7ea3da2b2b5f1f5447ebf4d54fc0a2e"),  
        newLeafHashFromStr("02f33c3ce0d684e644410fa35fee4c7572454facc710389ea0f69a00e27c1780"),  
        newLeafHashFromStr("4e29ad940088ab42991bd99d7943a68ea38d498f8435c4bea21b8252788b5a7d"),  
        newLeafHashFromStr("feabfc6565e918827df8ade12d0885e7da28bb81410c71a0fc92da28e5538e"),  
        newLeafHashFromStr("42275481edcd150c0891f8f7486941e7ecc9dbec0d1554d1ae3027d8523571b4"),  
        newLeafHashFromStr("84dd84bbb564ab9e878ea774d66a2cb818c2b74680974e01a07dde68bd064bfd"),  
        newLeafHashFromStr("19a64f2f77403d4a1e3482b3a1801aa360120f6429f5ed47261fc8676920d9d"),  
    },  
},
```

Parallel Block validation

Current block validation

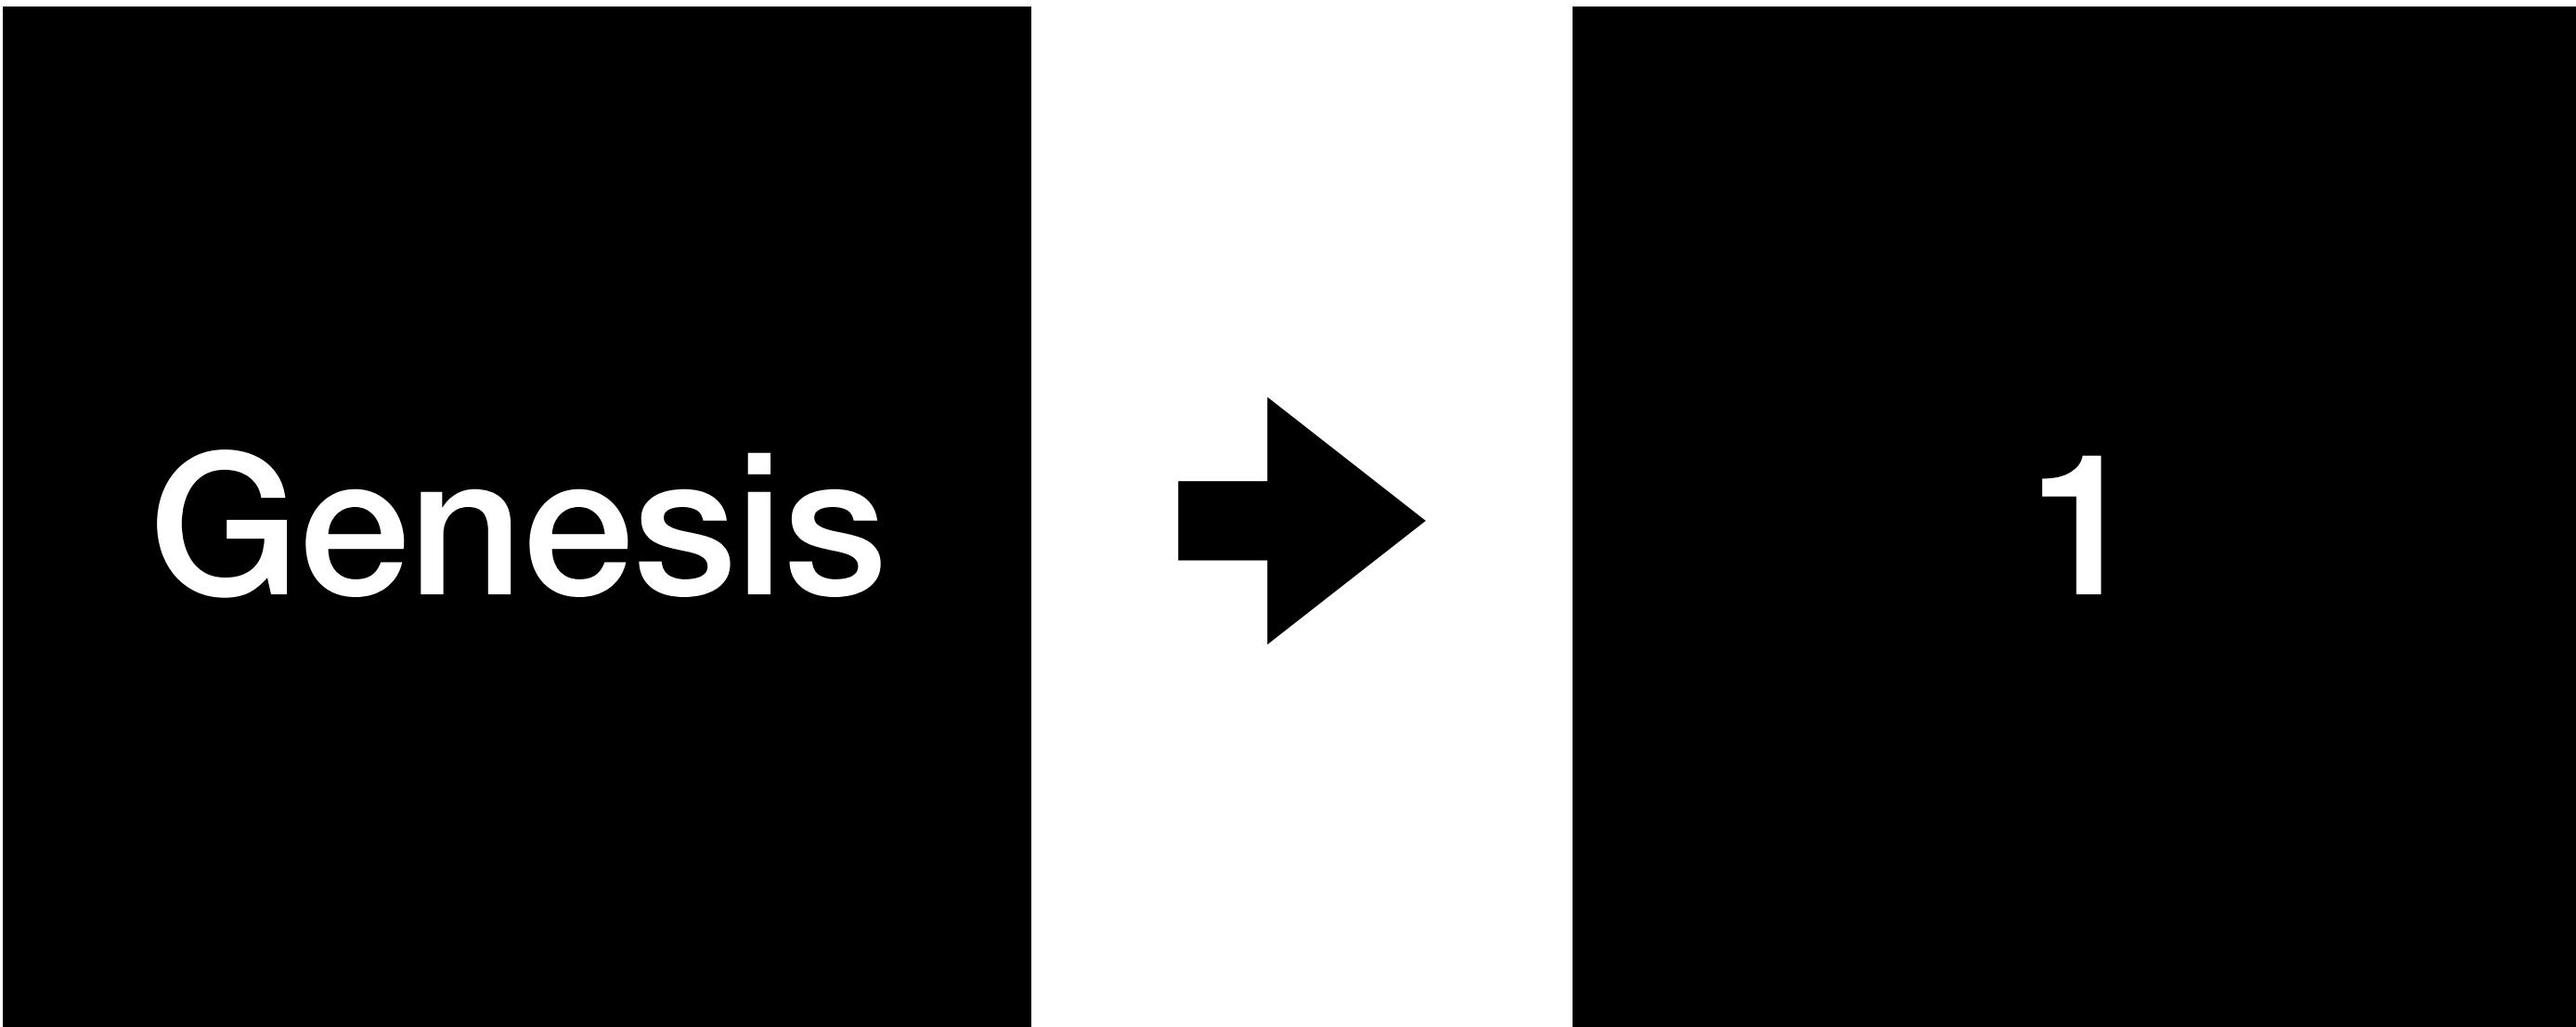
Sequential validation



Genesis

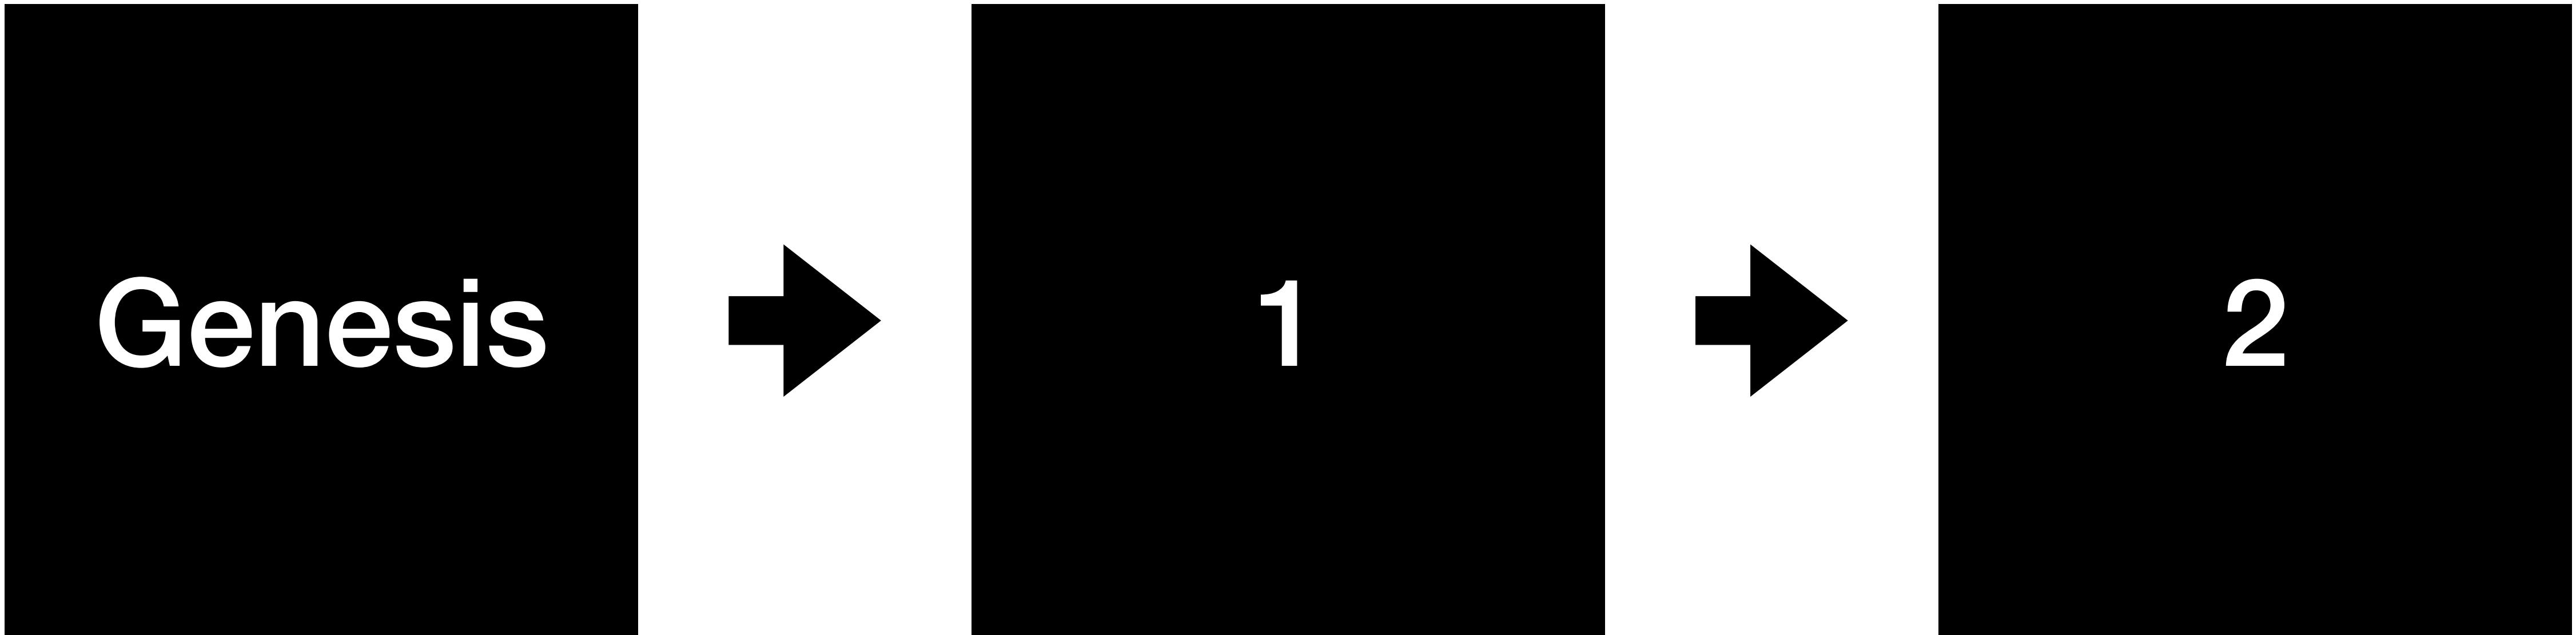
Current block validation

Sequential validation



Current block validation

Sequential validation



Parallel block validation

Replay blocks in any order

- Utreexo roots committed into the binary
- Process blocks starting from any of the roots that are committed

With Utreexo

Efficient parallel validation

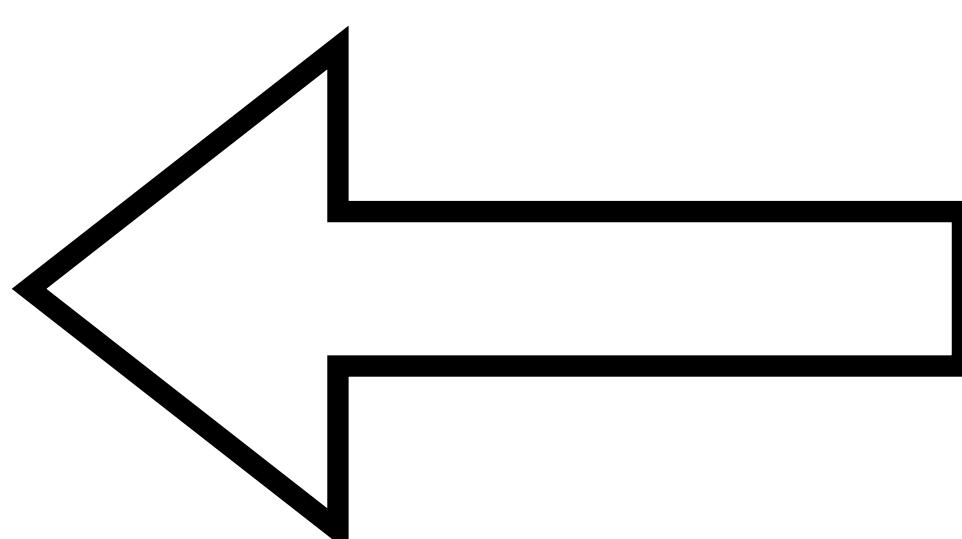
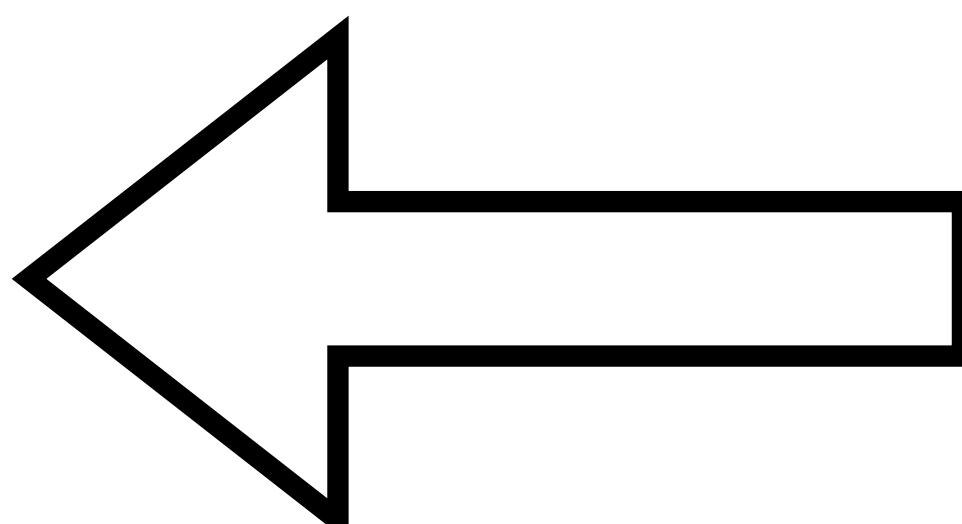
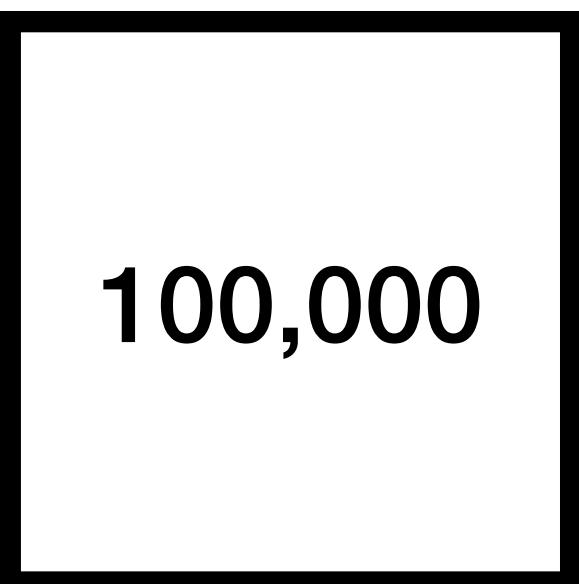
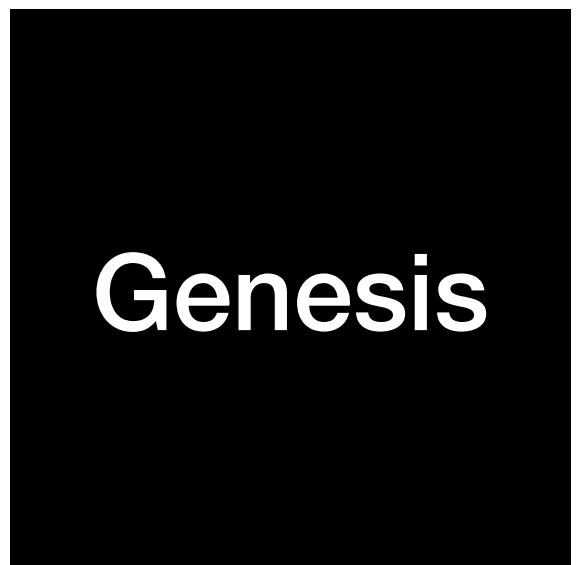
Genesis

100,000

200,000

With Utreexo

Efficient parallel validation



Untrusted roots at
the start

With Utreexo

Efficient parallel validation

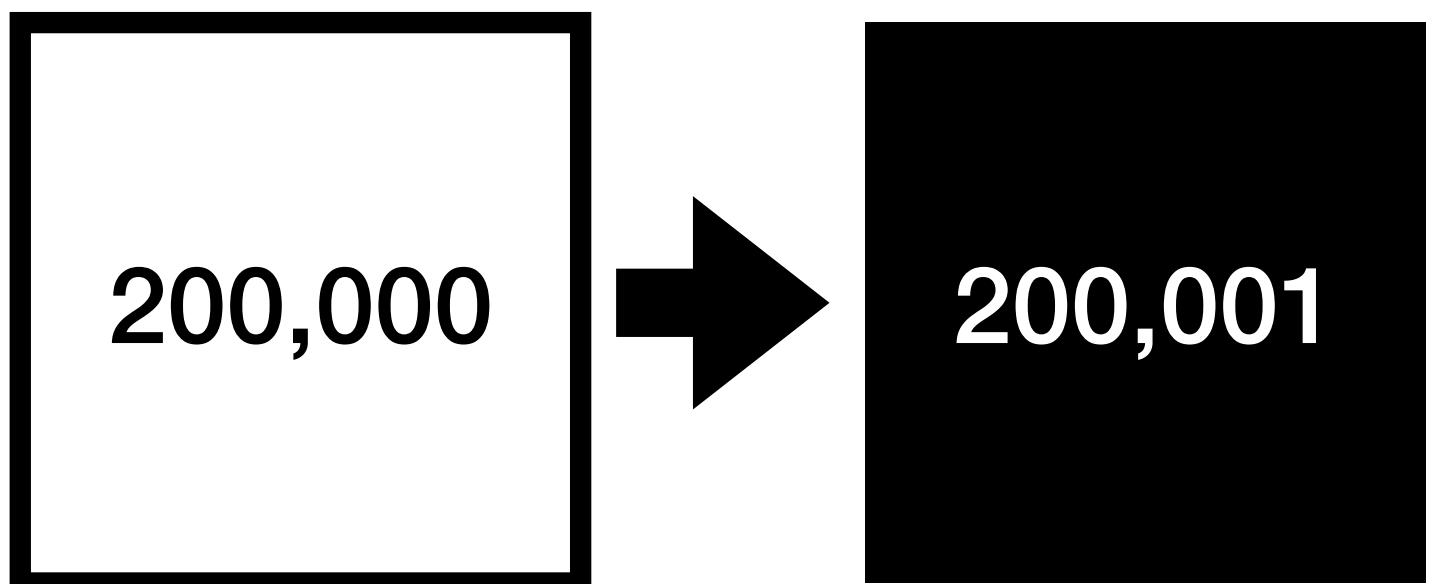
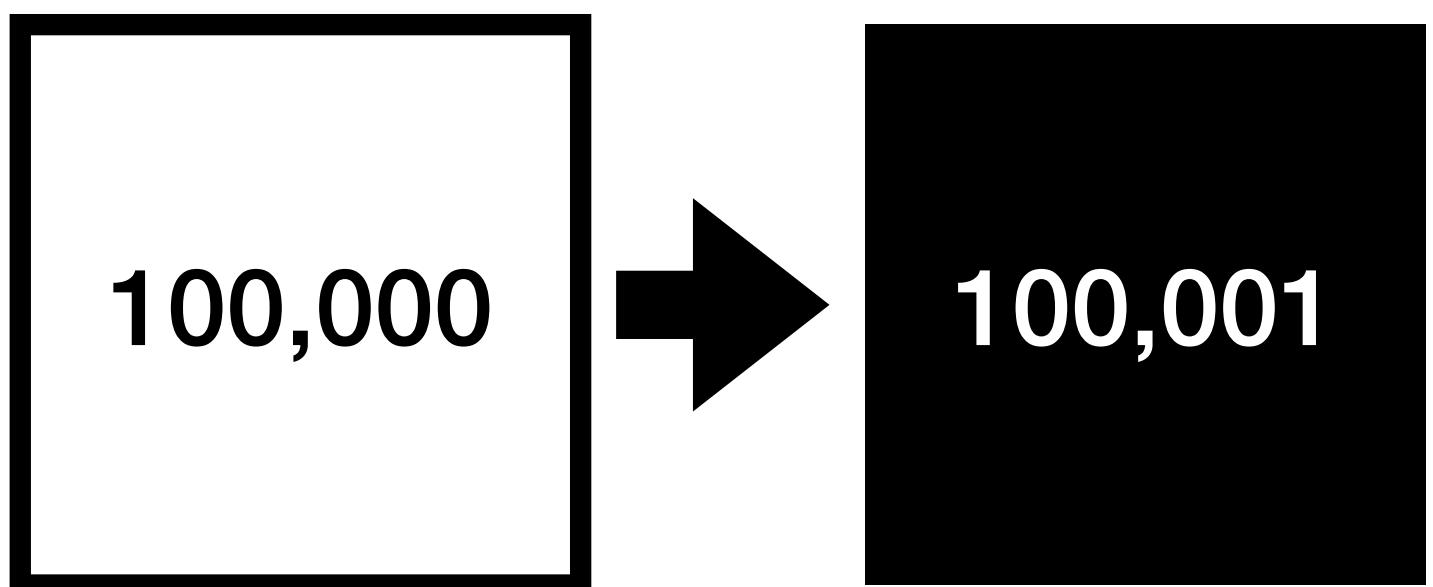
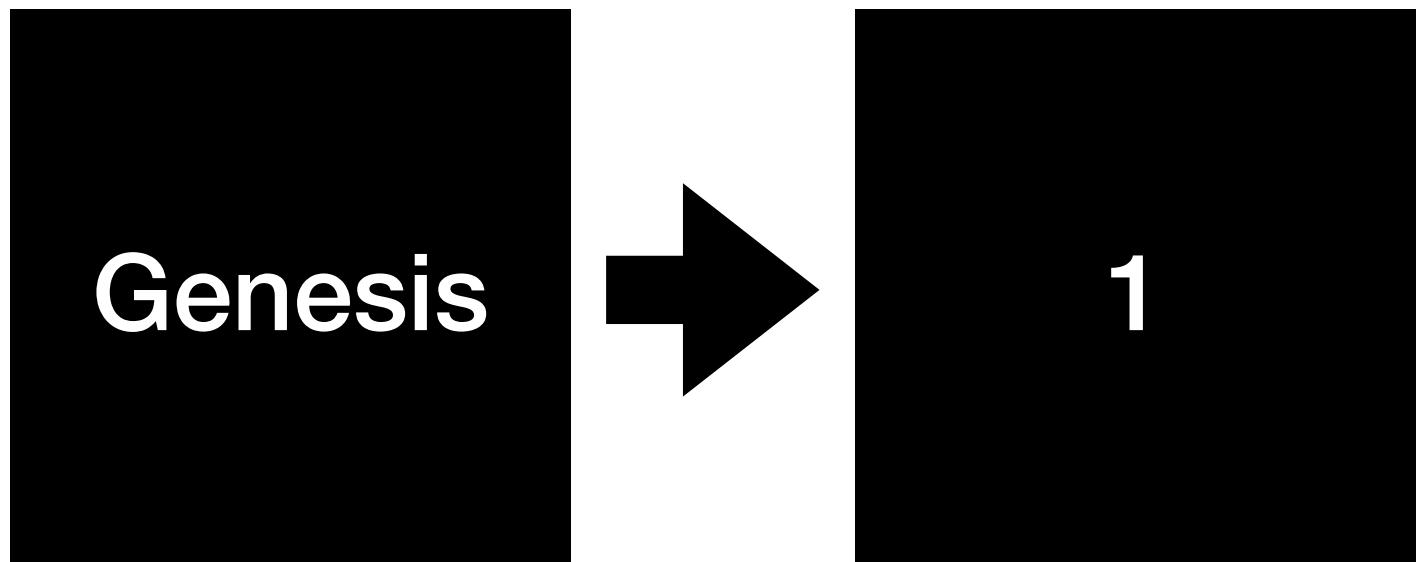
Genesis

100,000

200,000

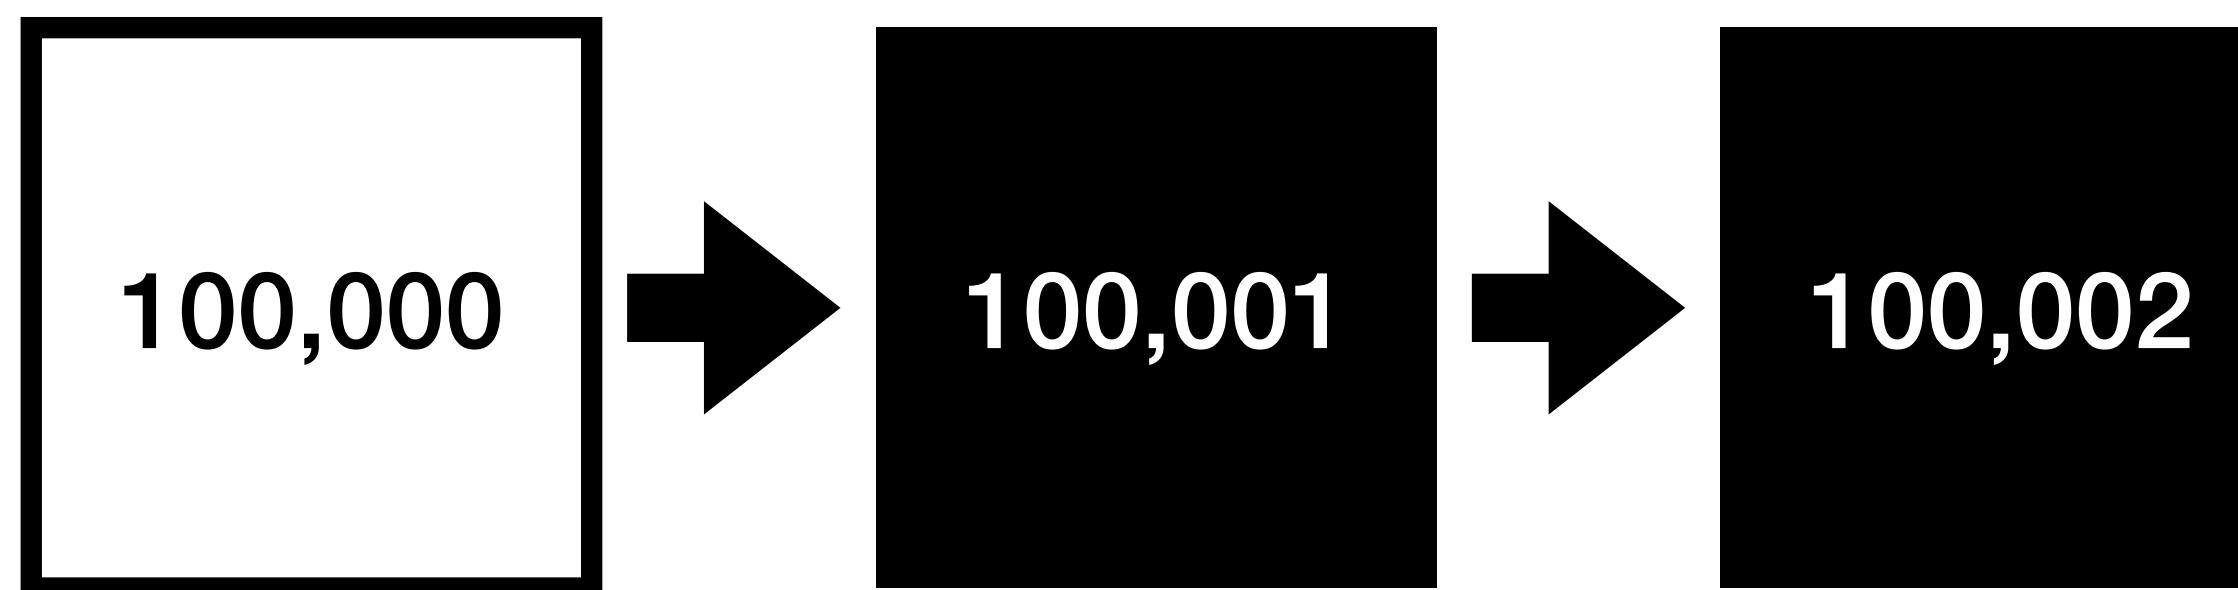
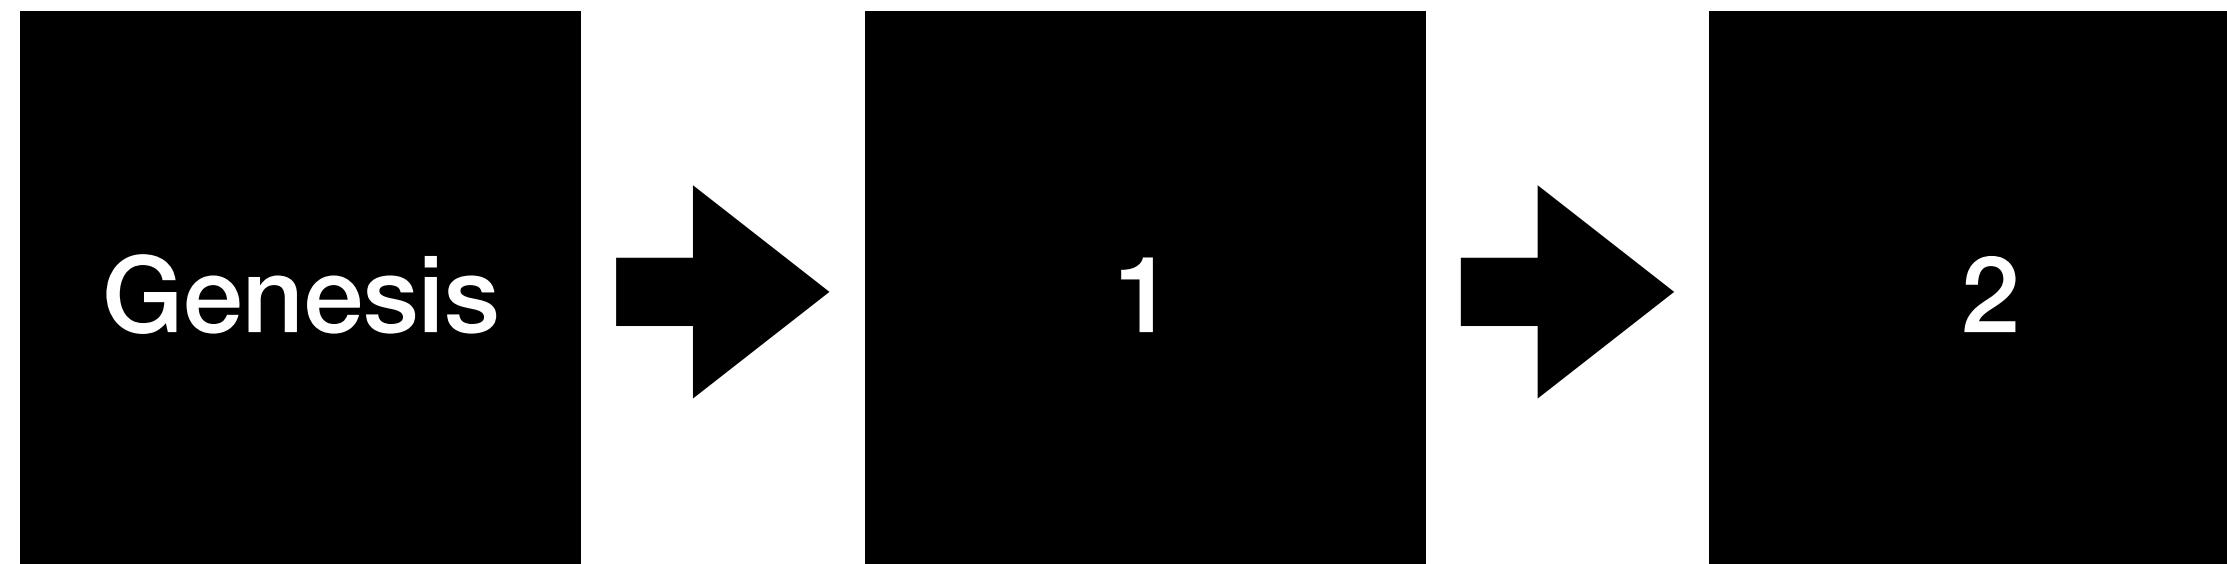
With Utreexo

Efficient parallel validation



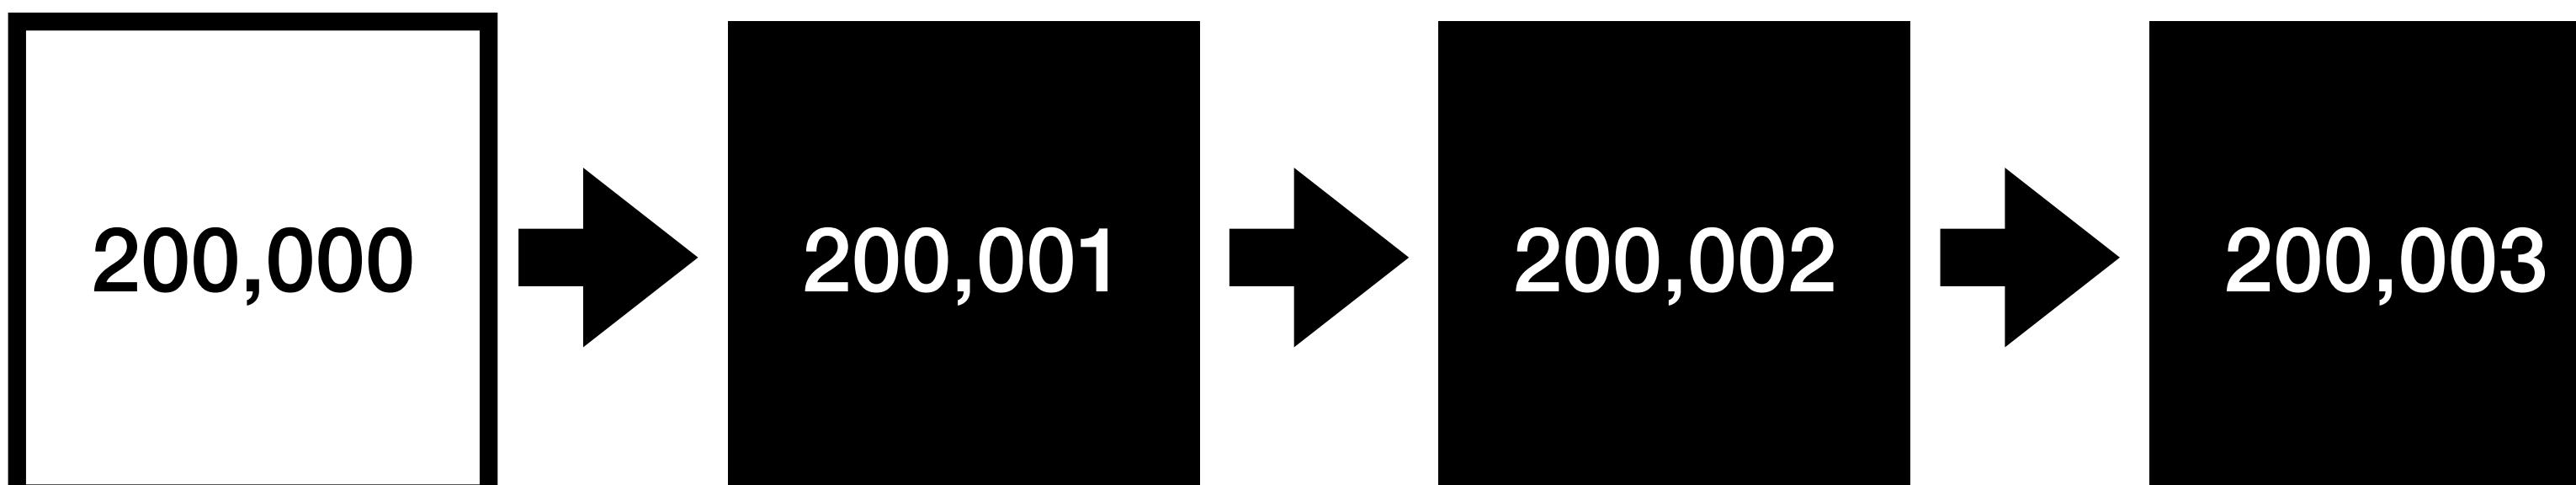
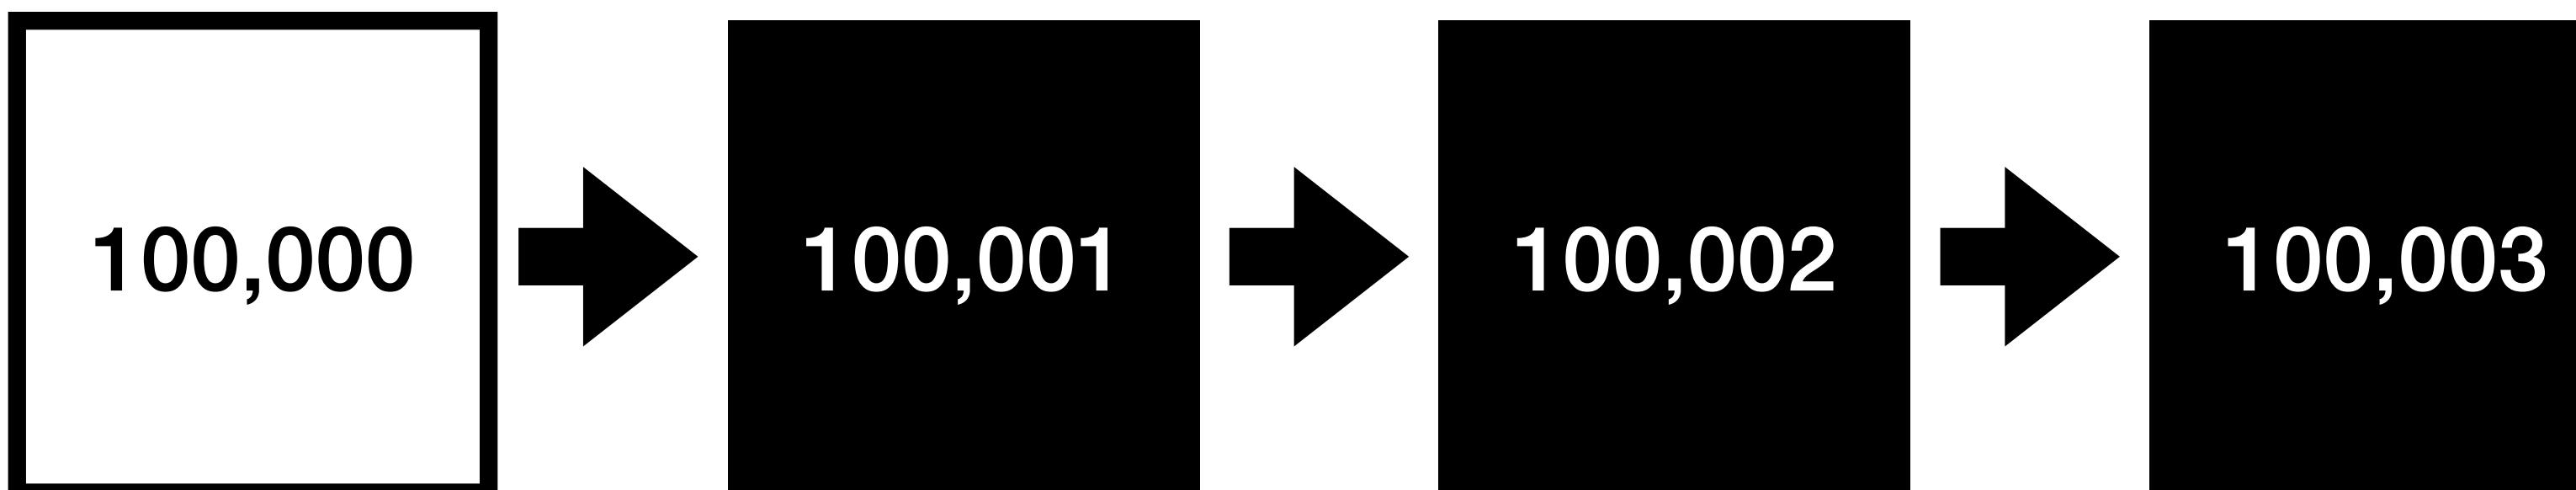
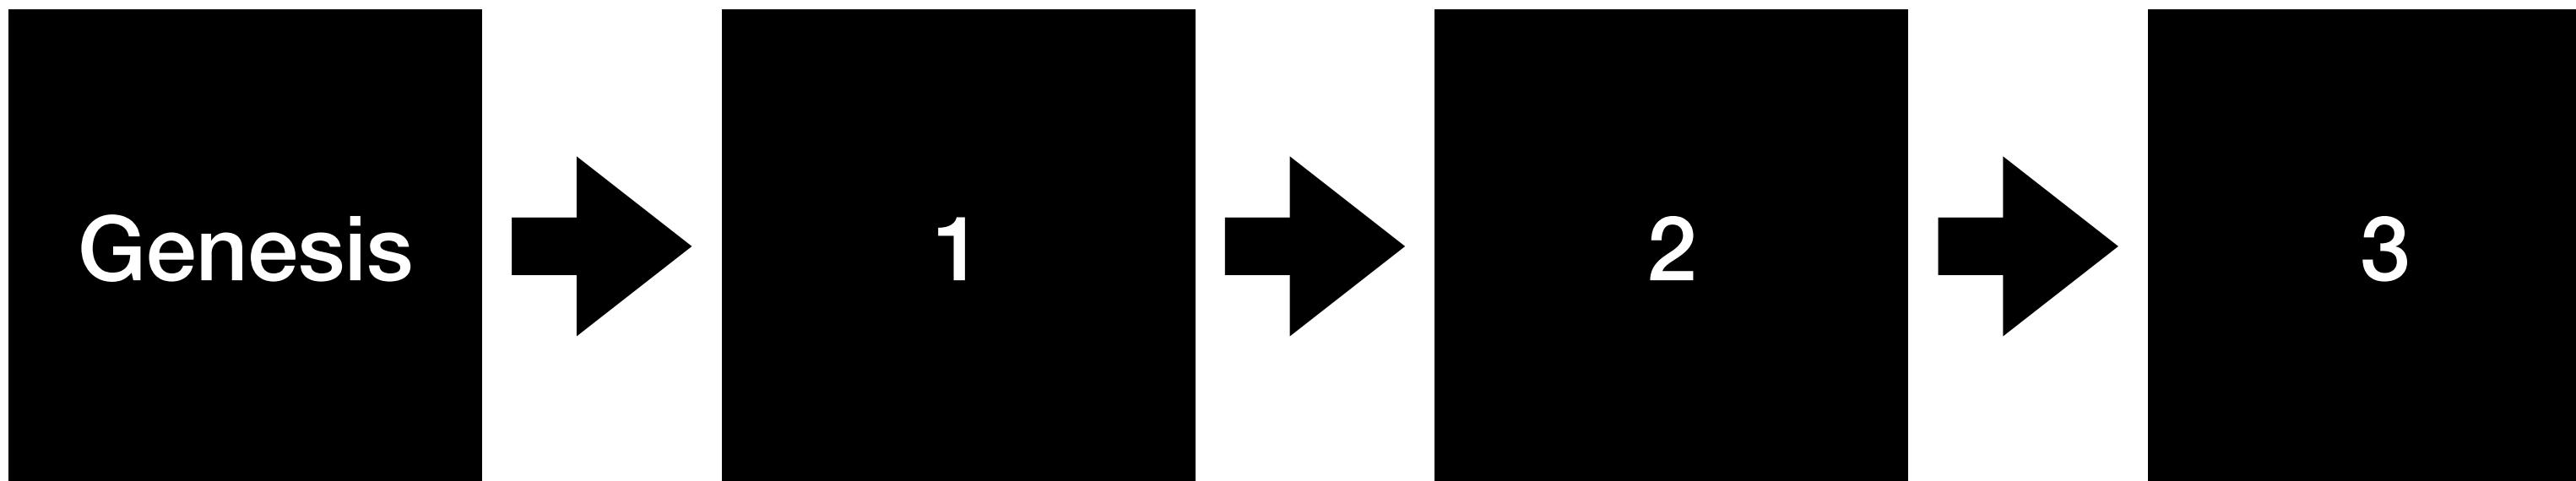
With Utreexo

Efficient parallel validation



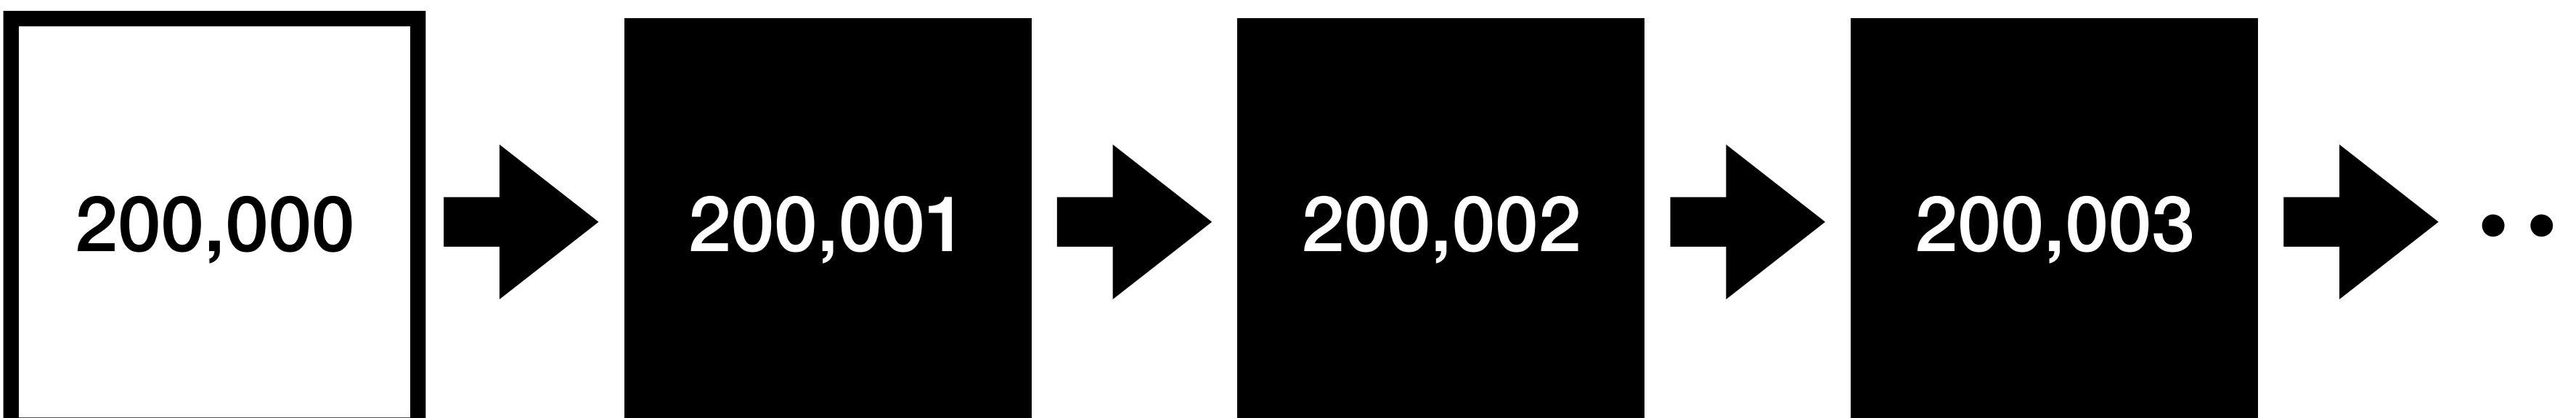
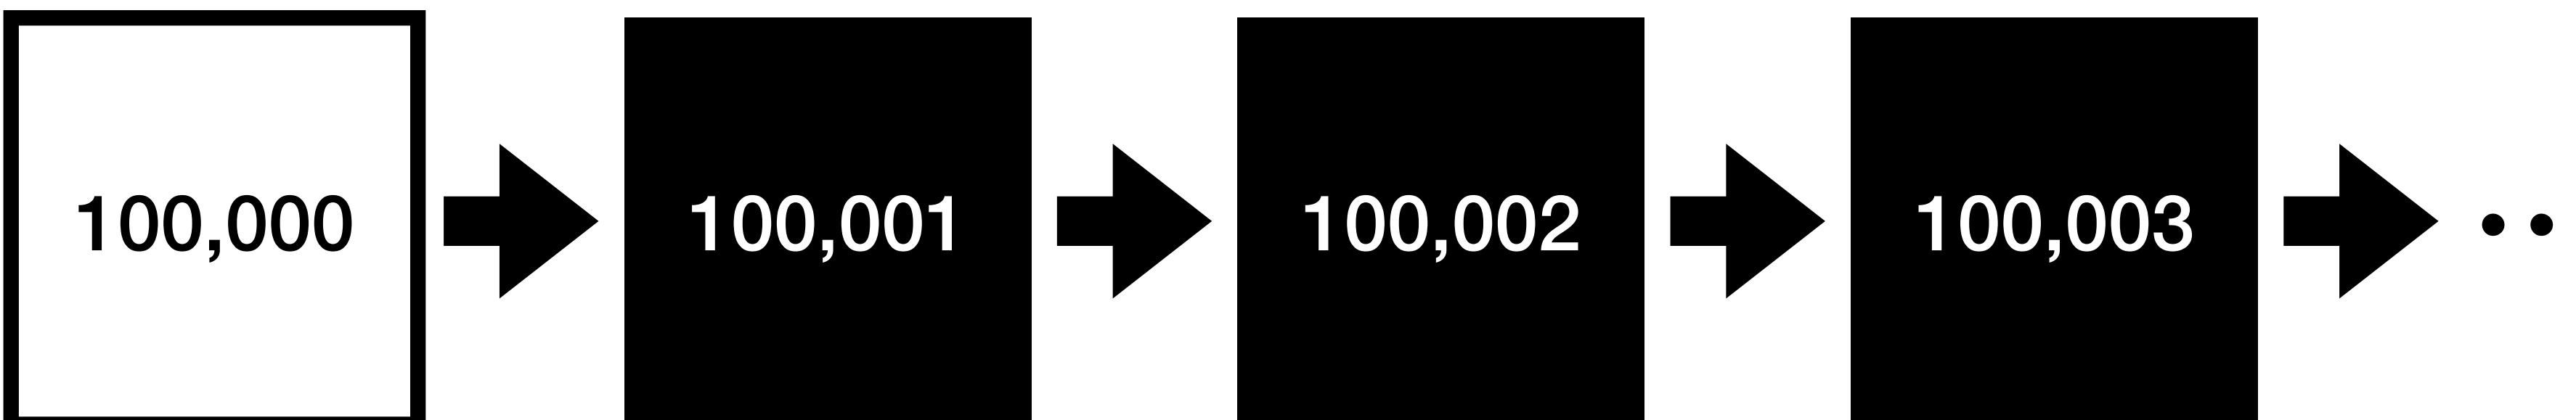
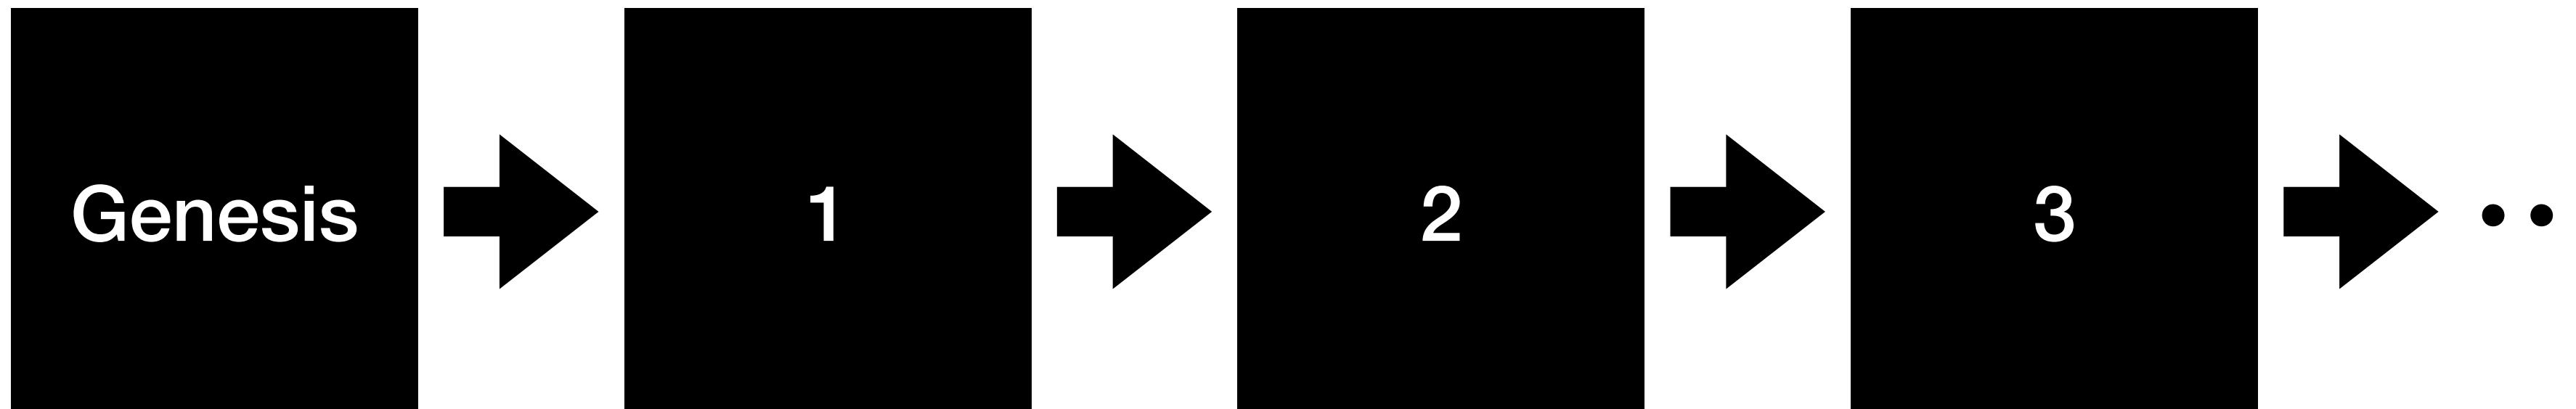
With Utreexo

Efficient parallel validation



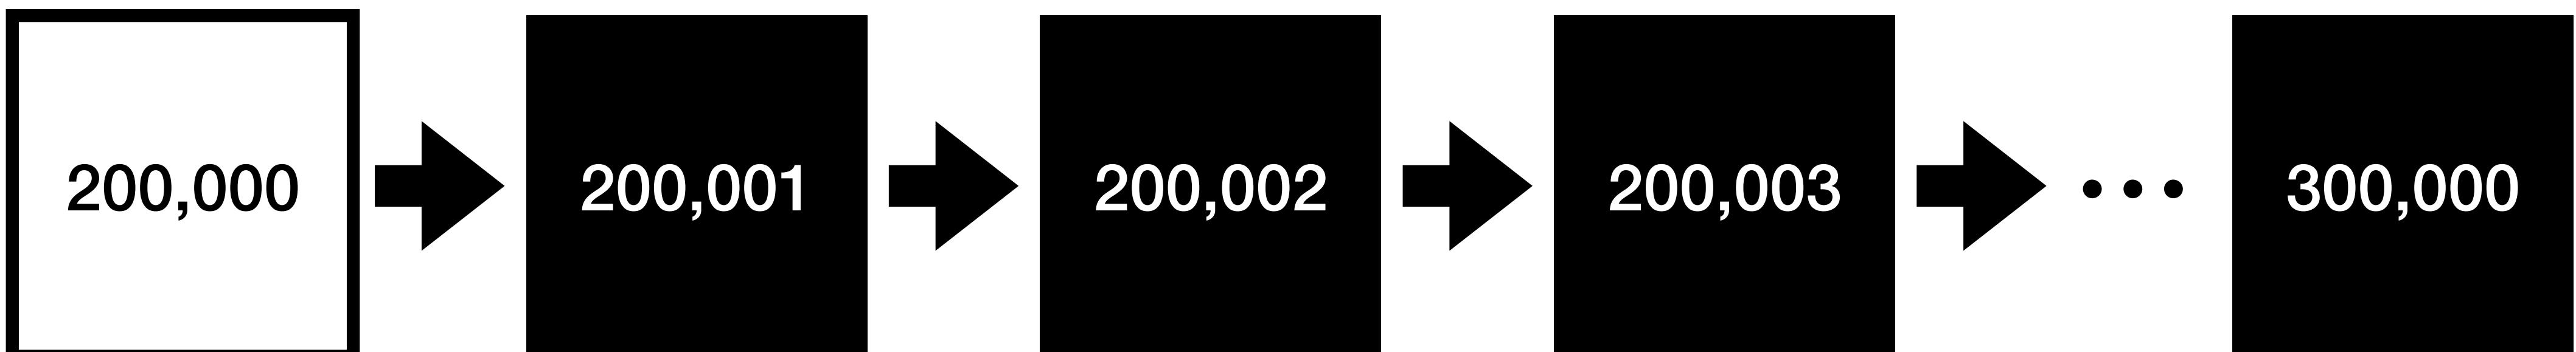
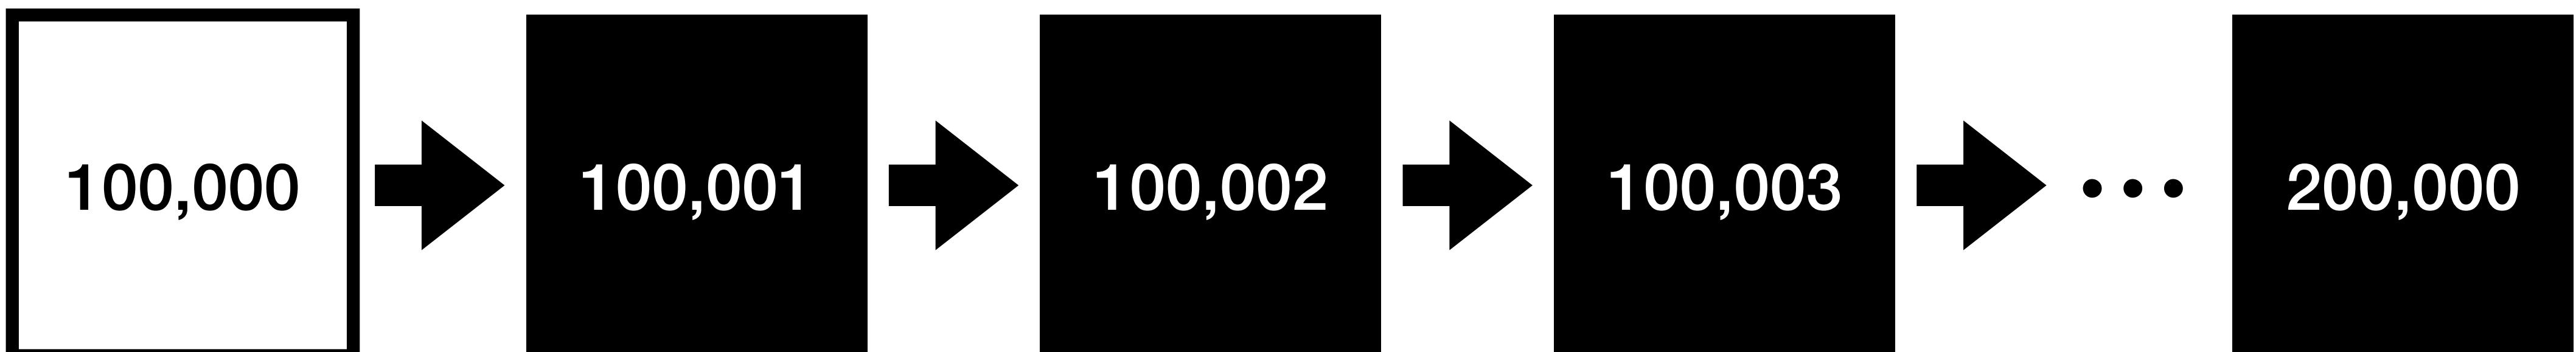
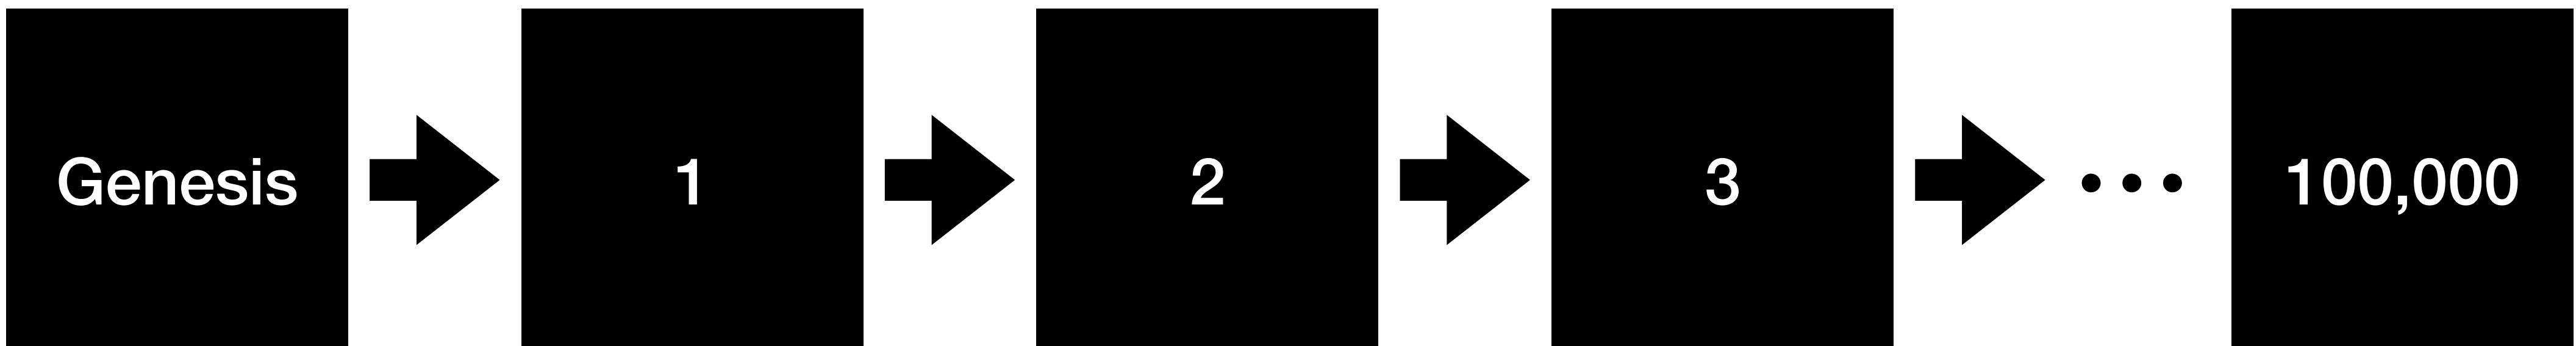
With Utreexo

Efficient parallel validation



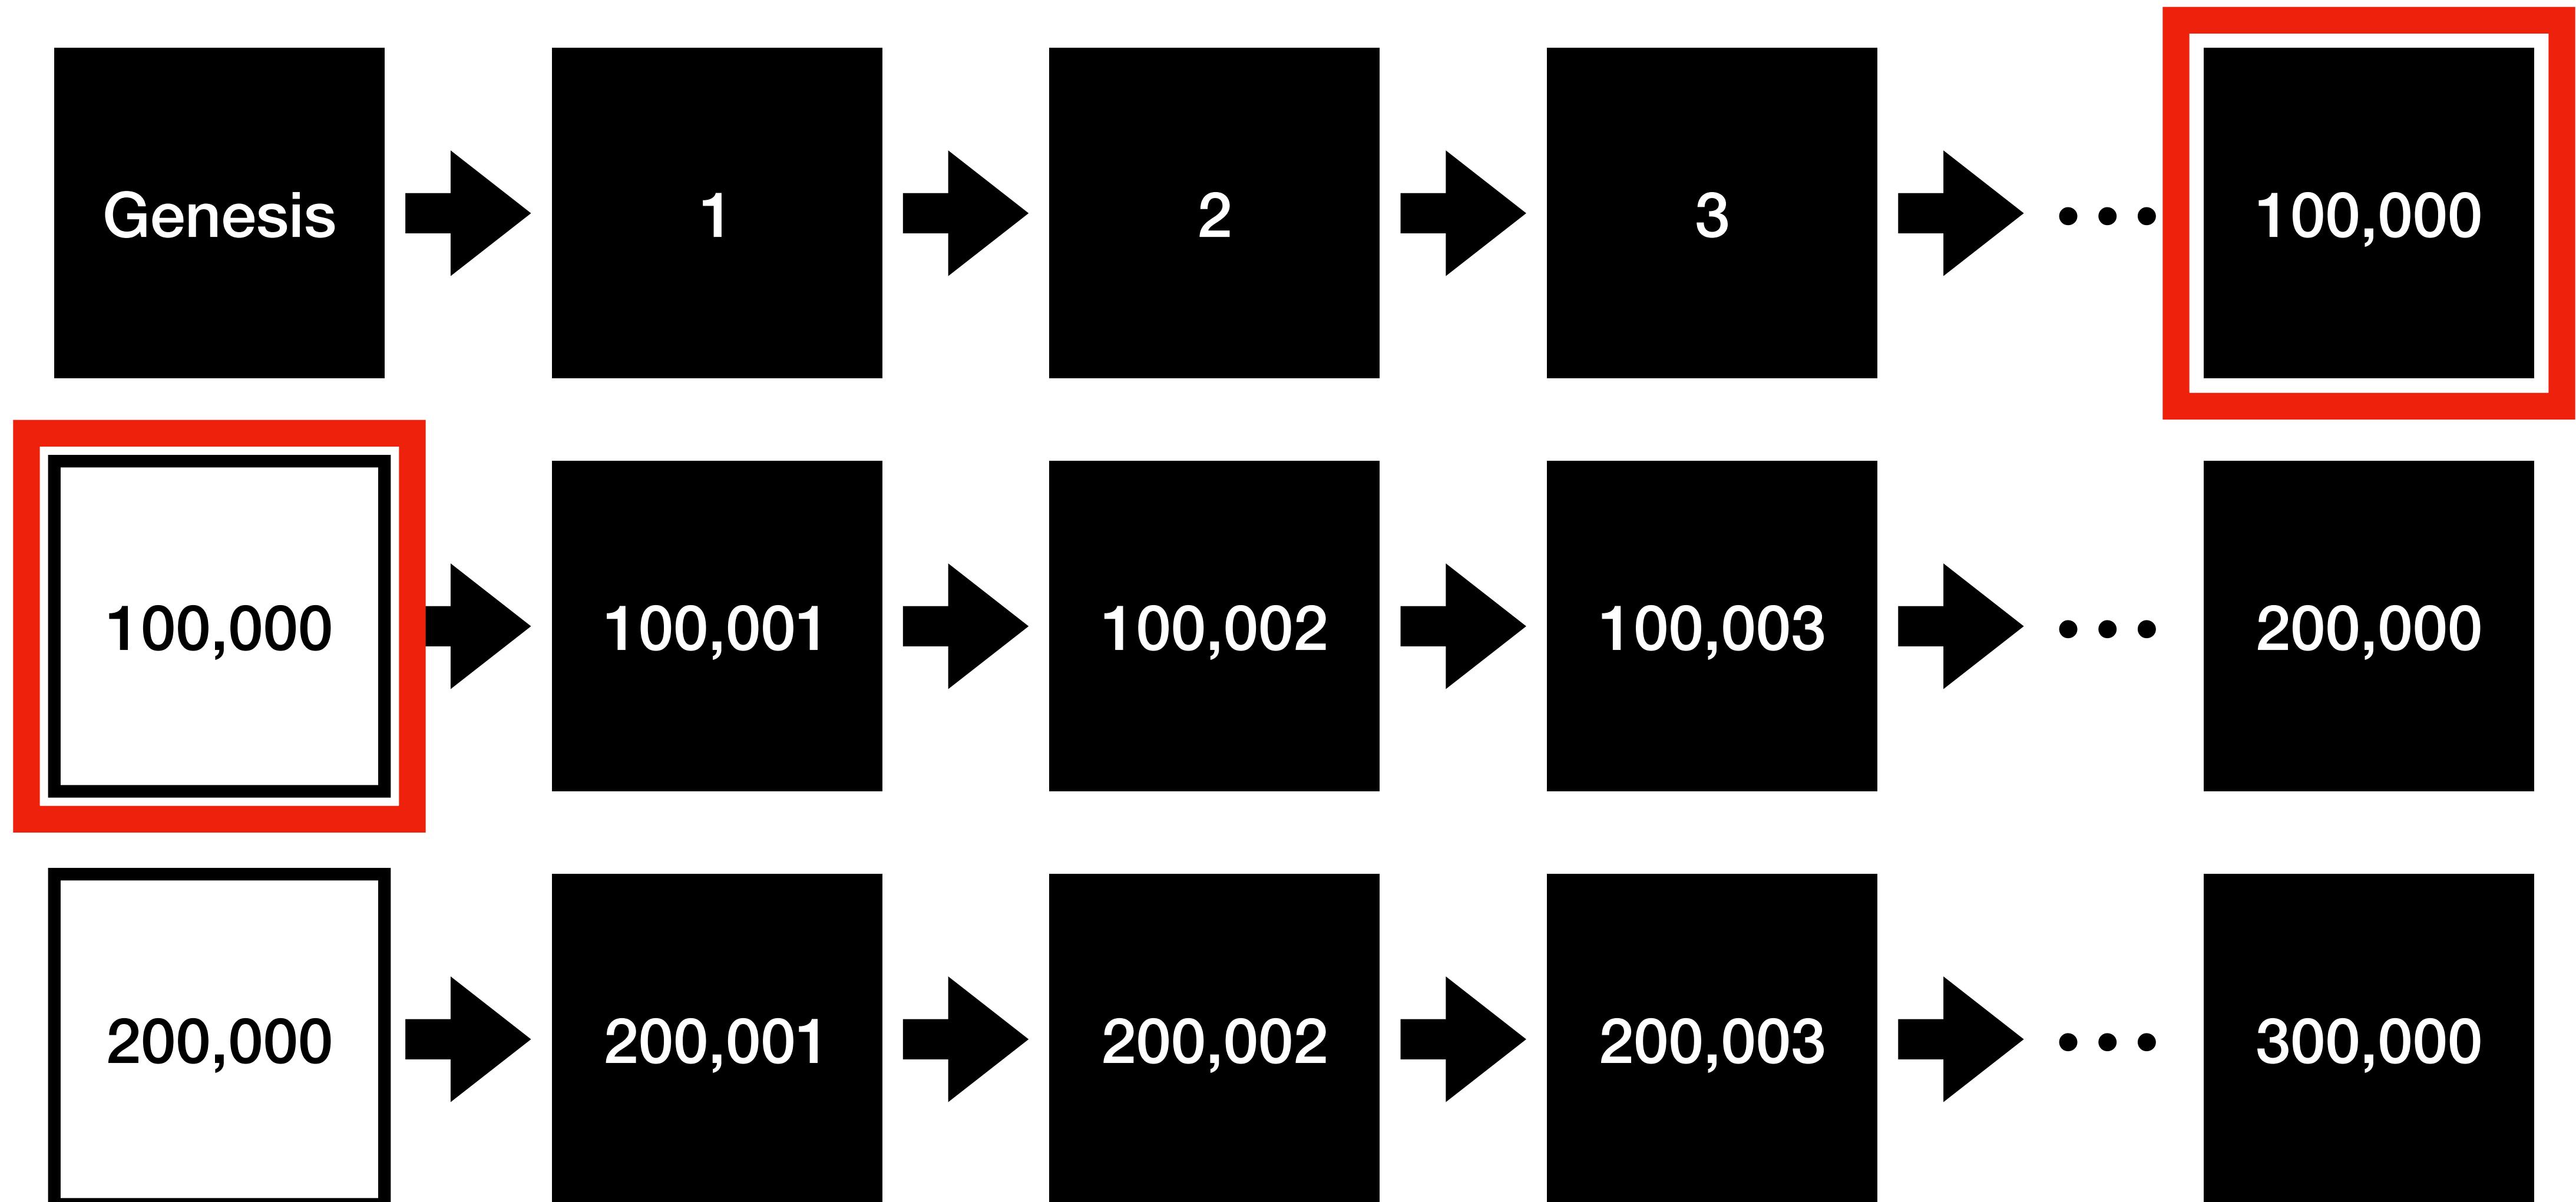
With Utreexo

Efficient parallel validation



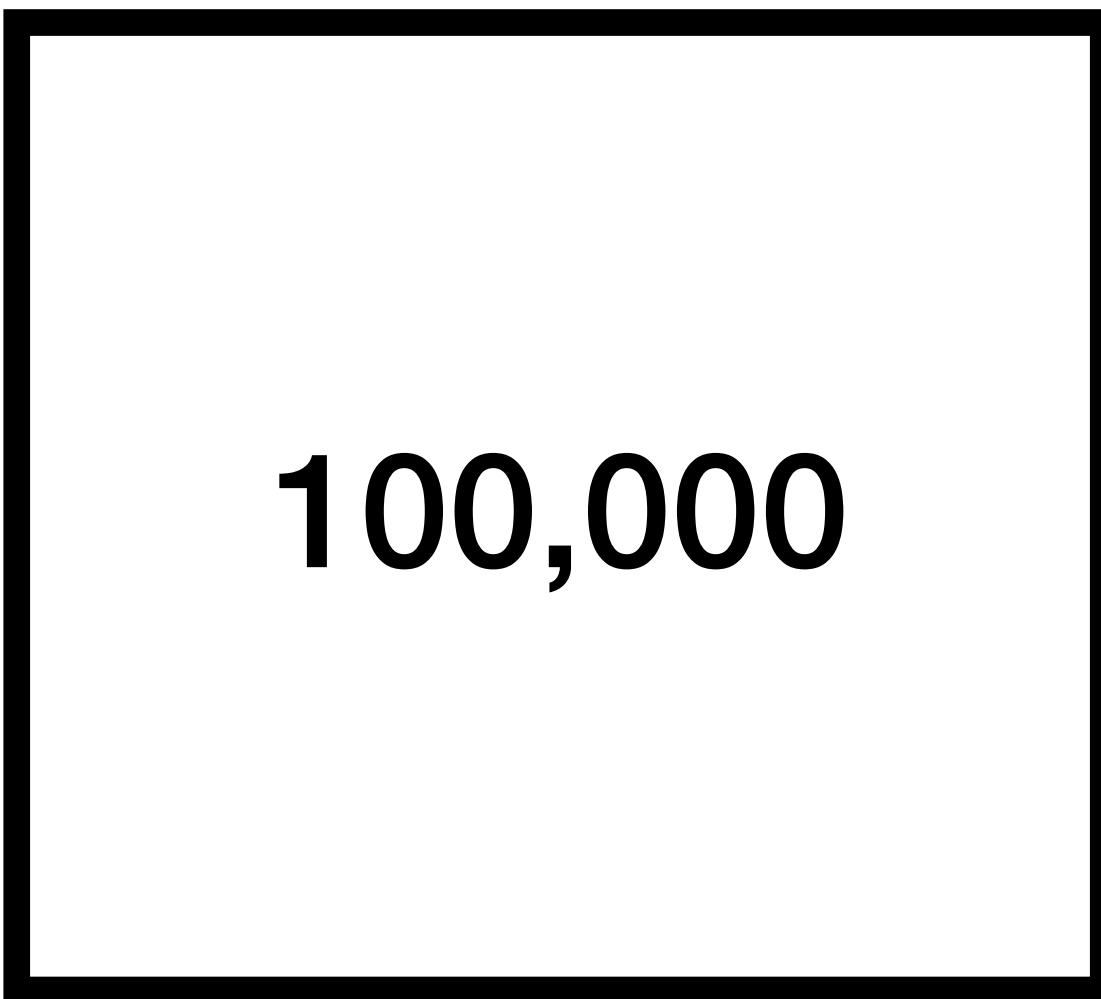
With Utreexo

Efficient parallel validation

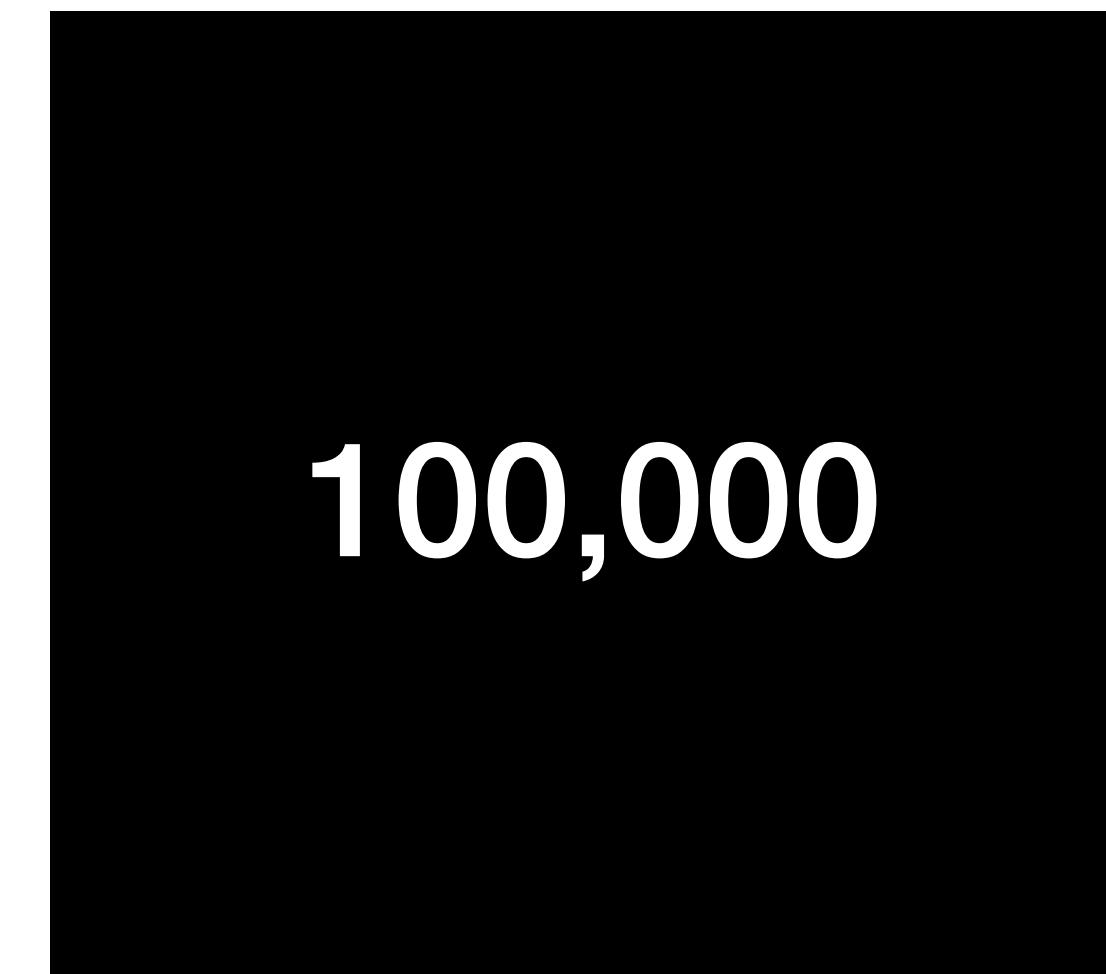


Compare roots

Root at 100,000 becomes trusted if equal

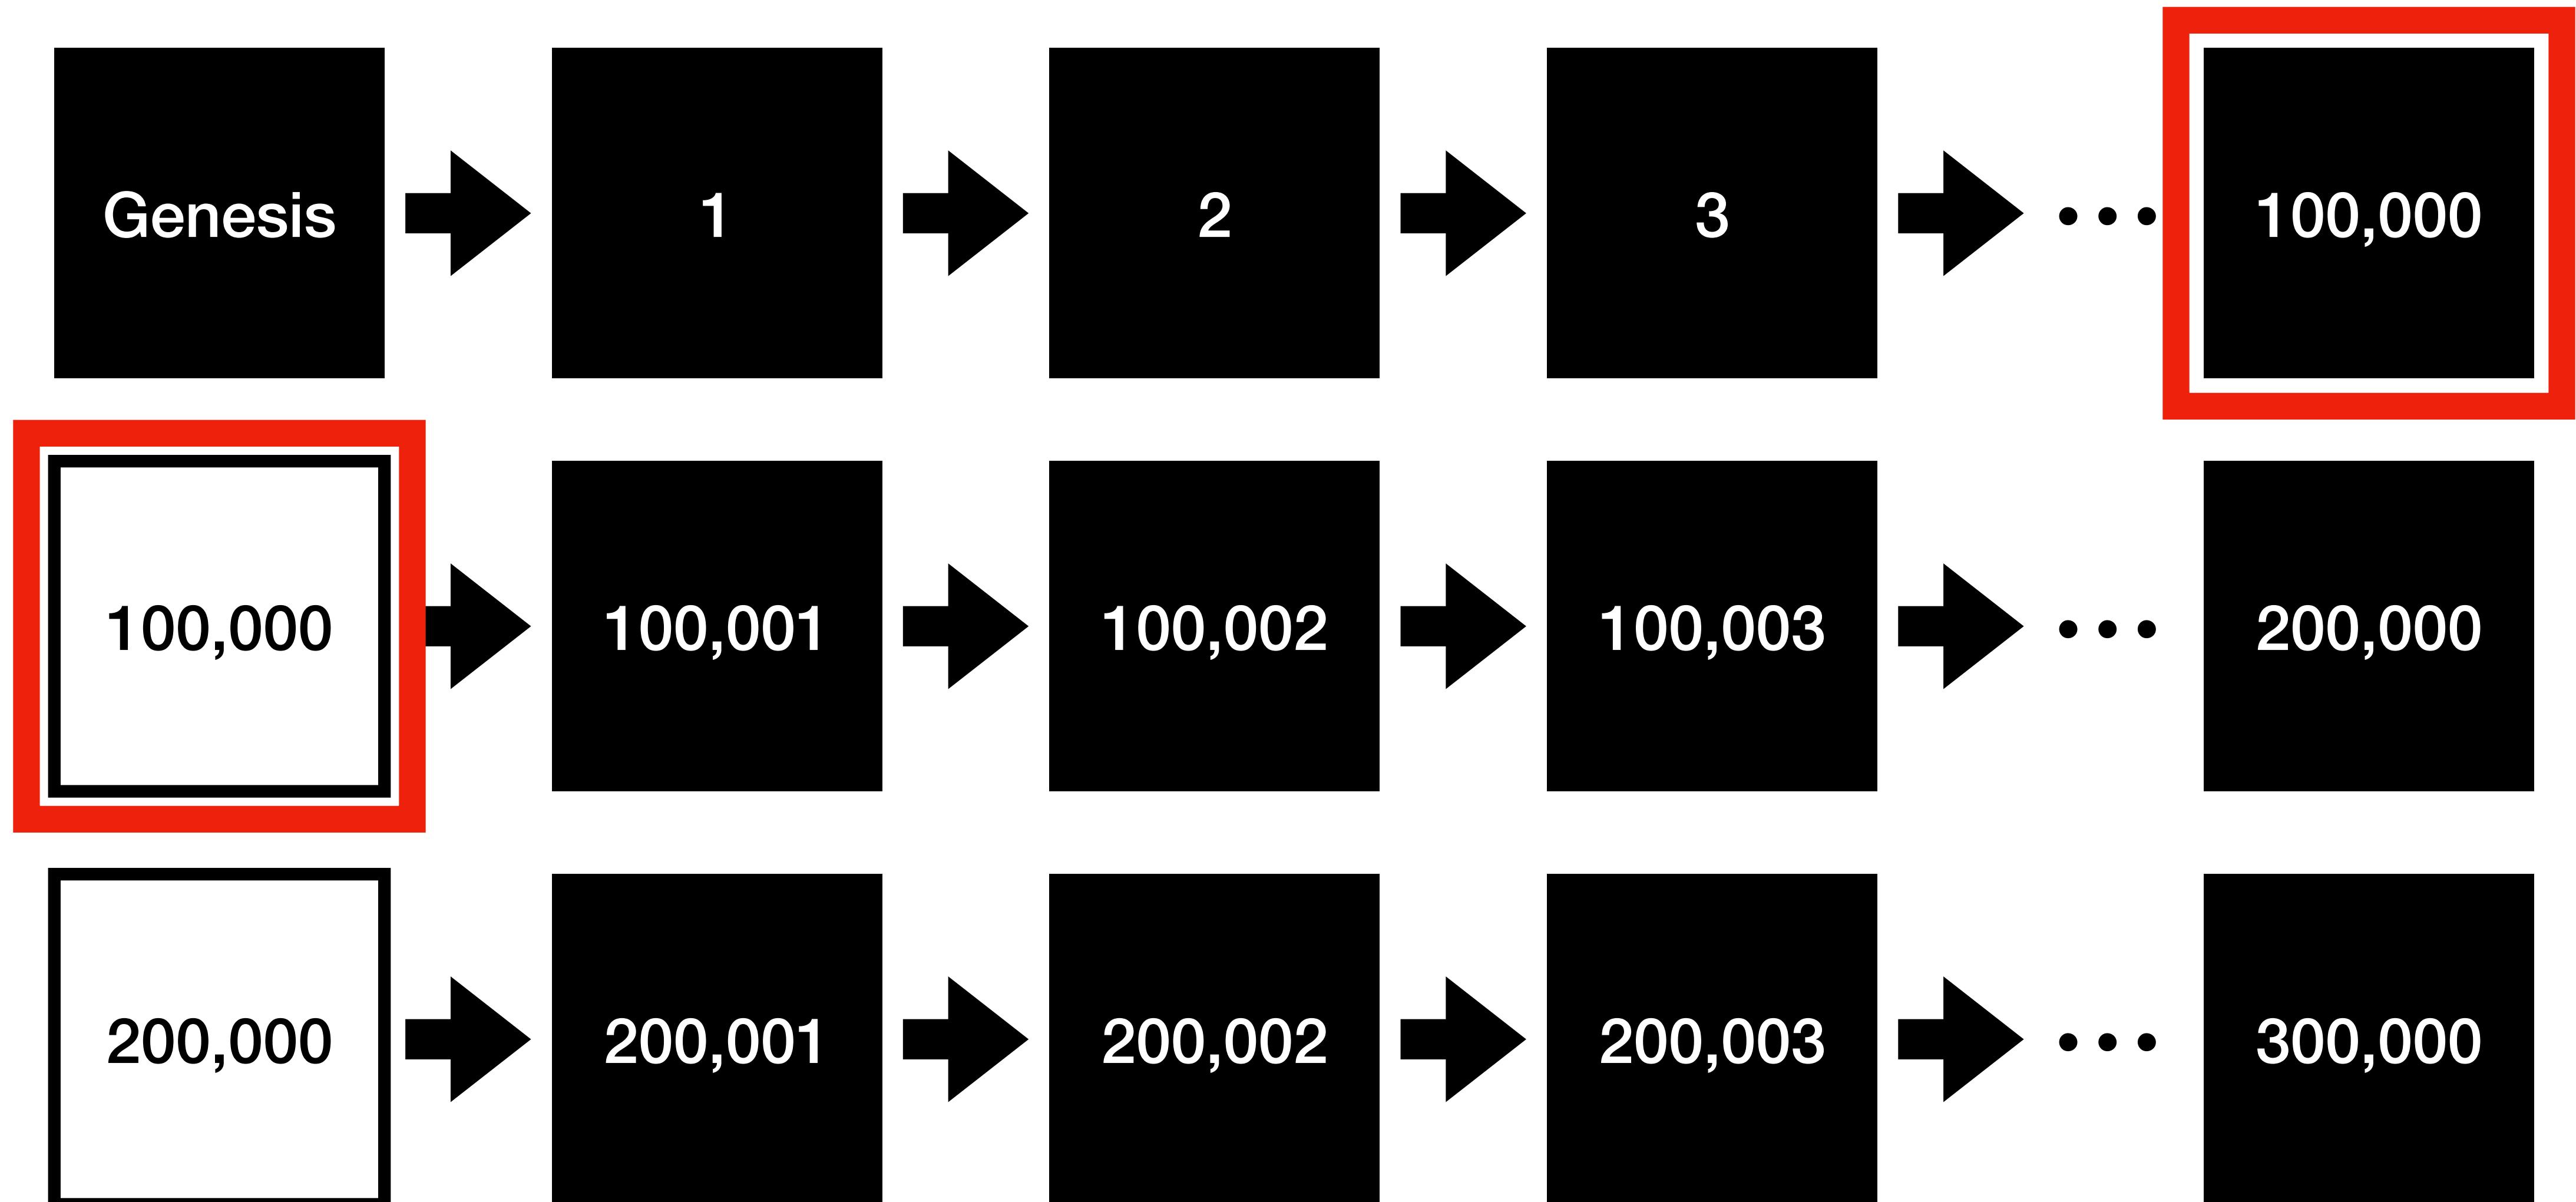


==



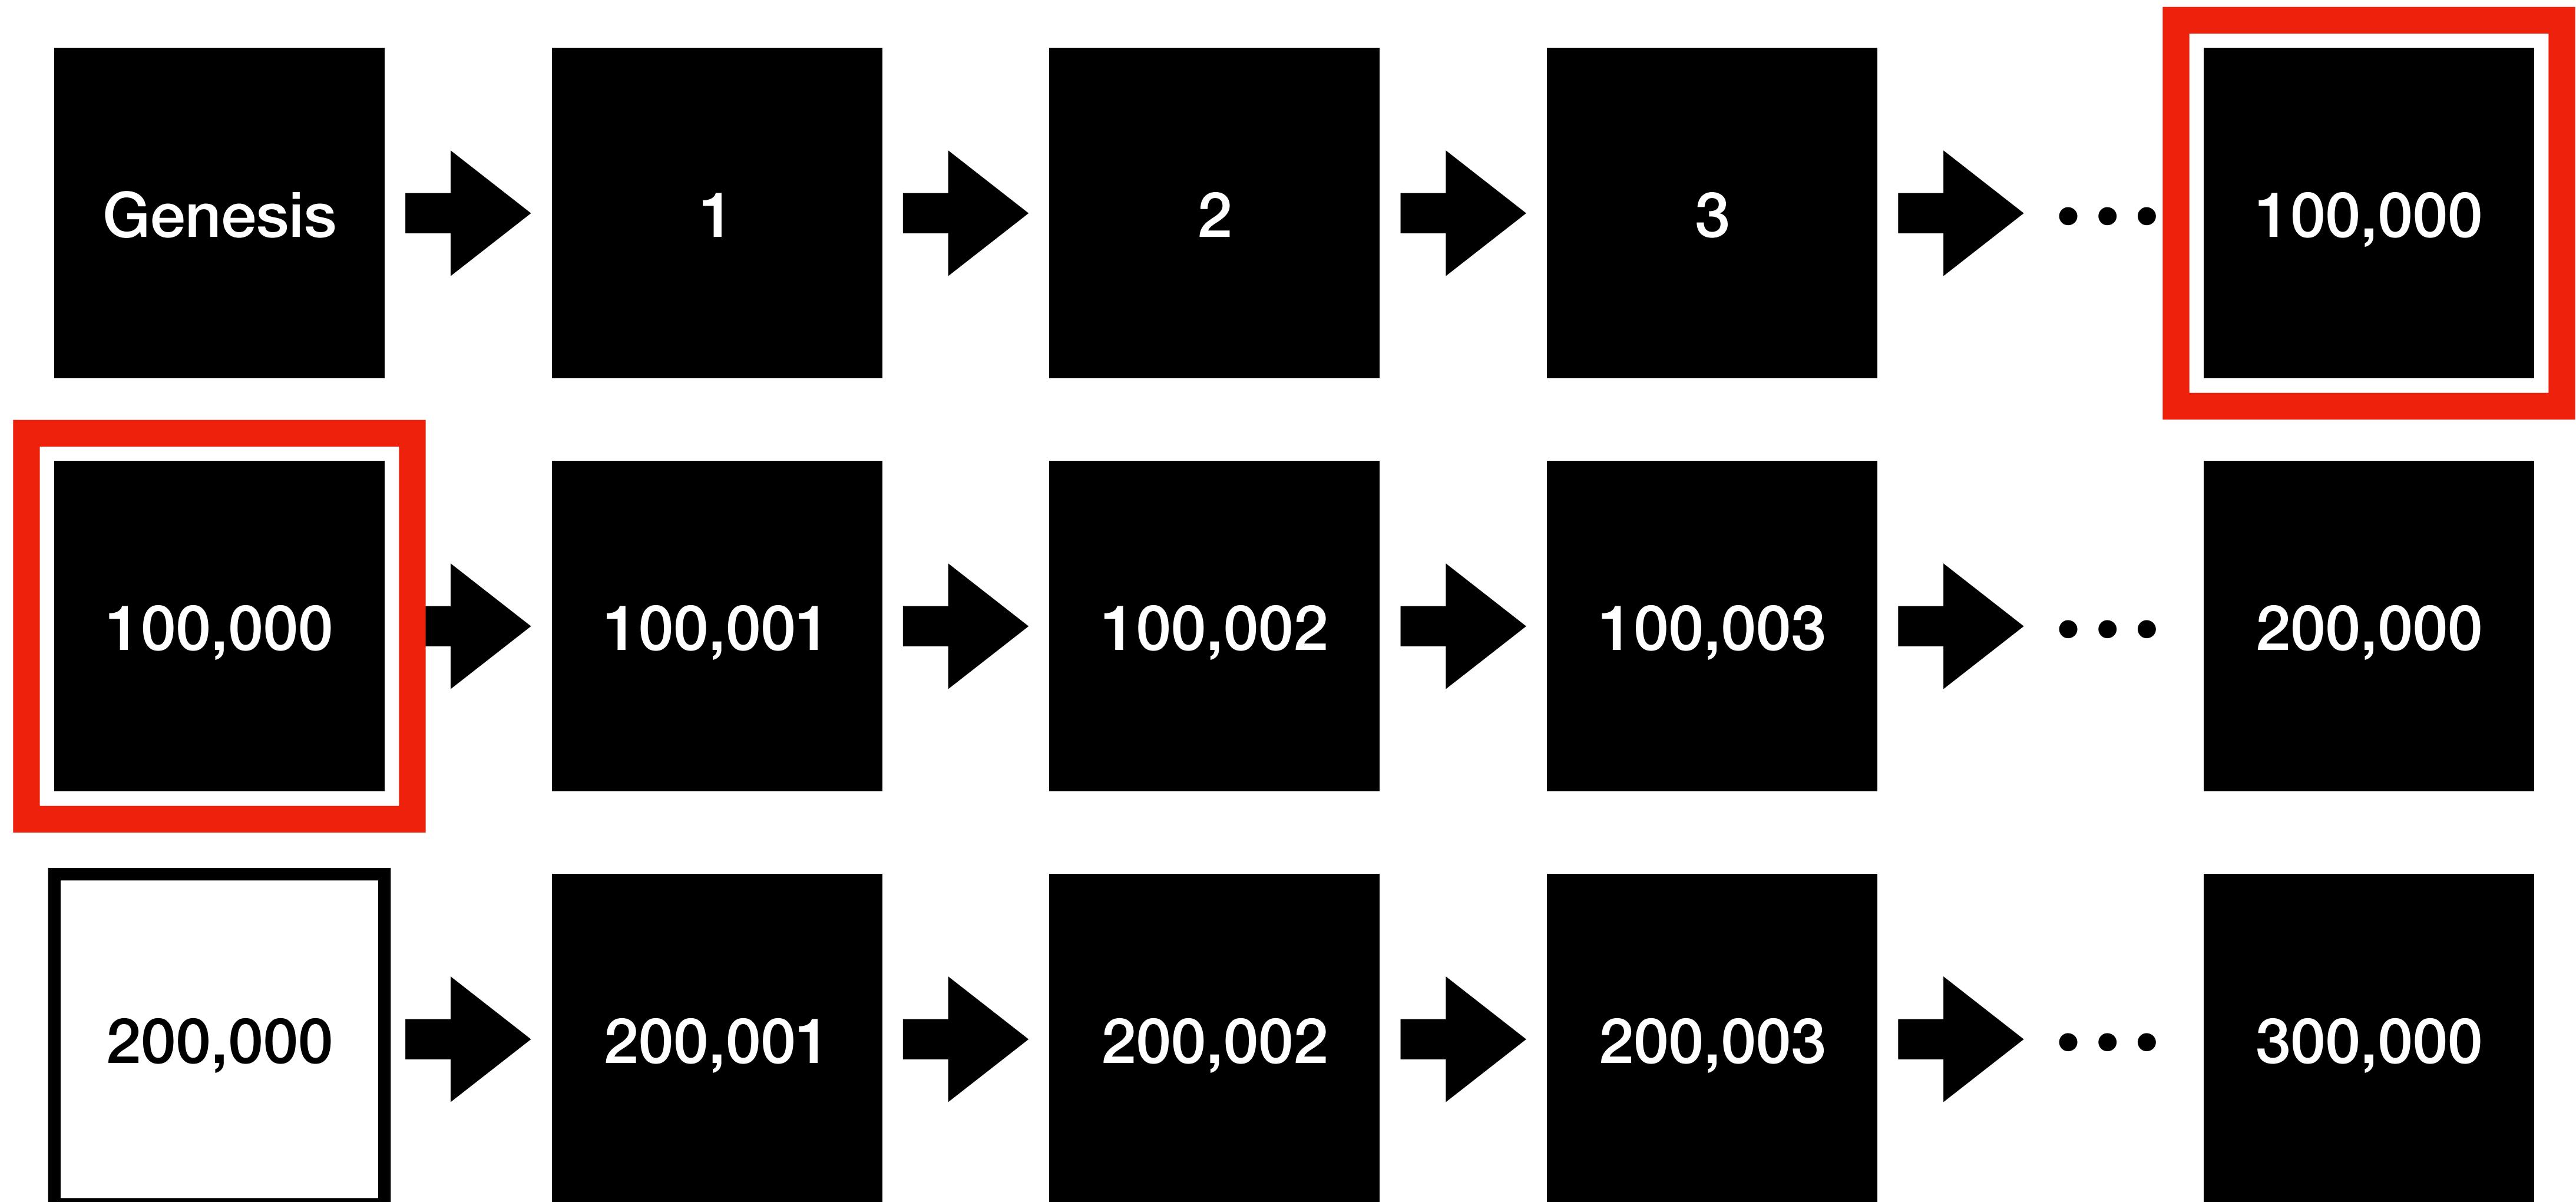
With Utreexo

Efficient parallel validation



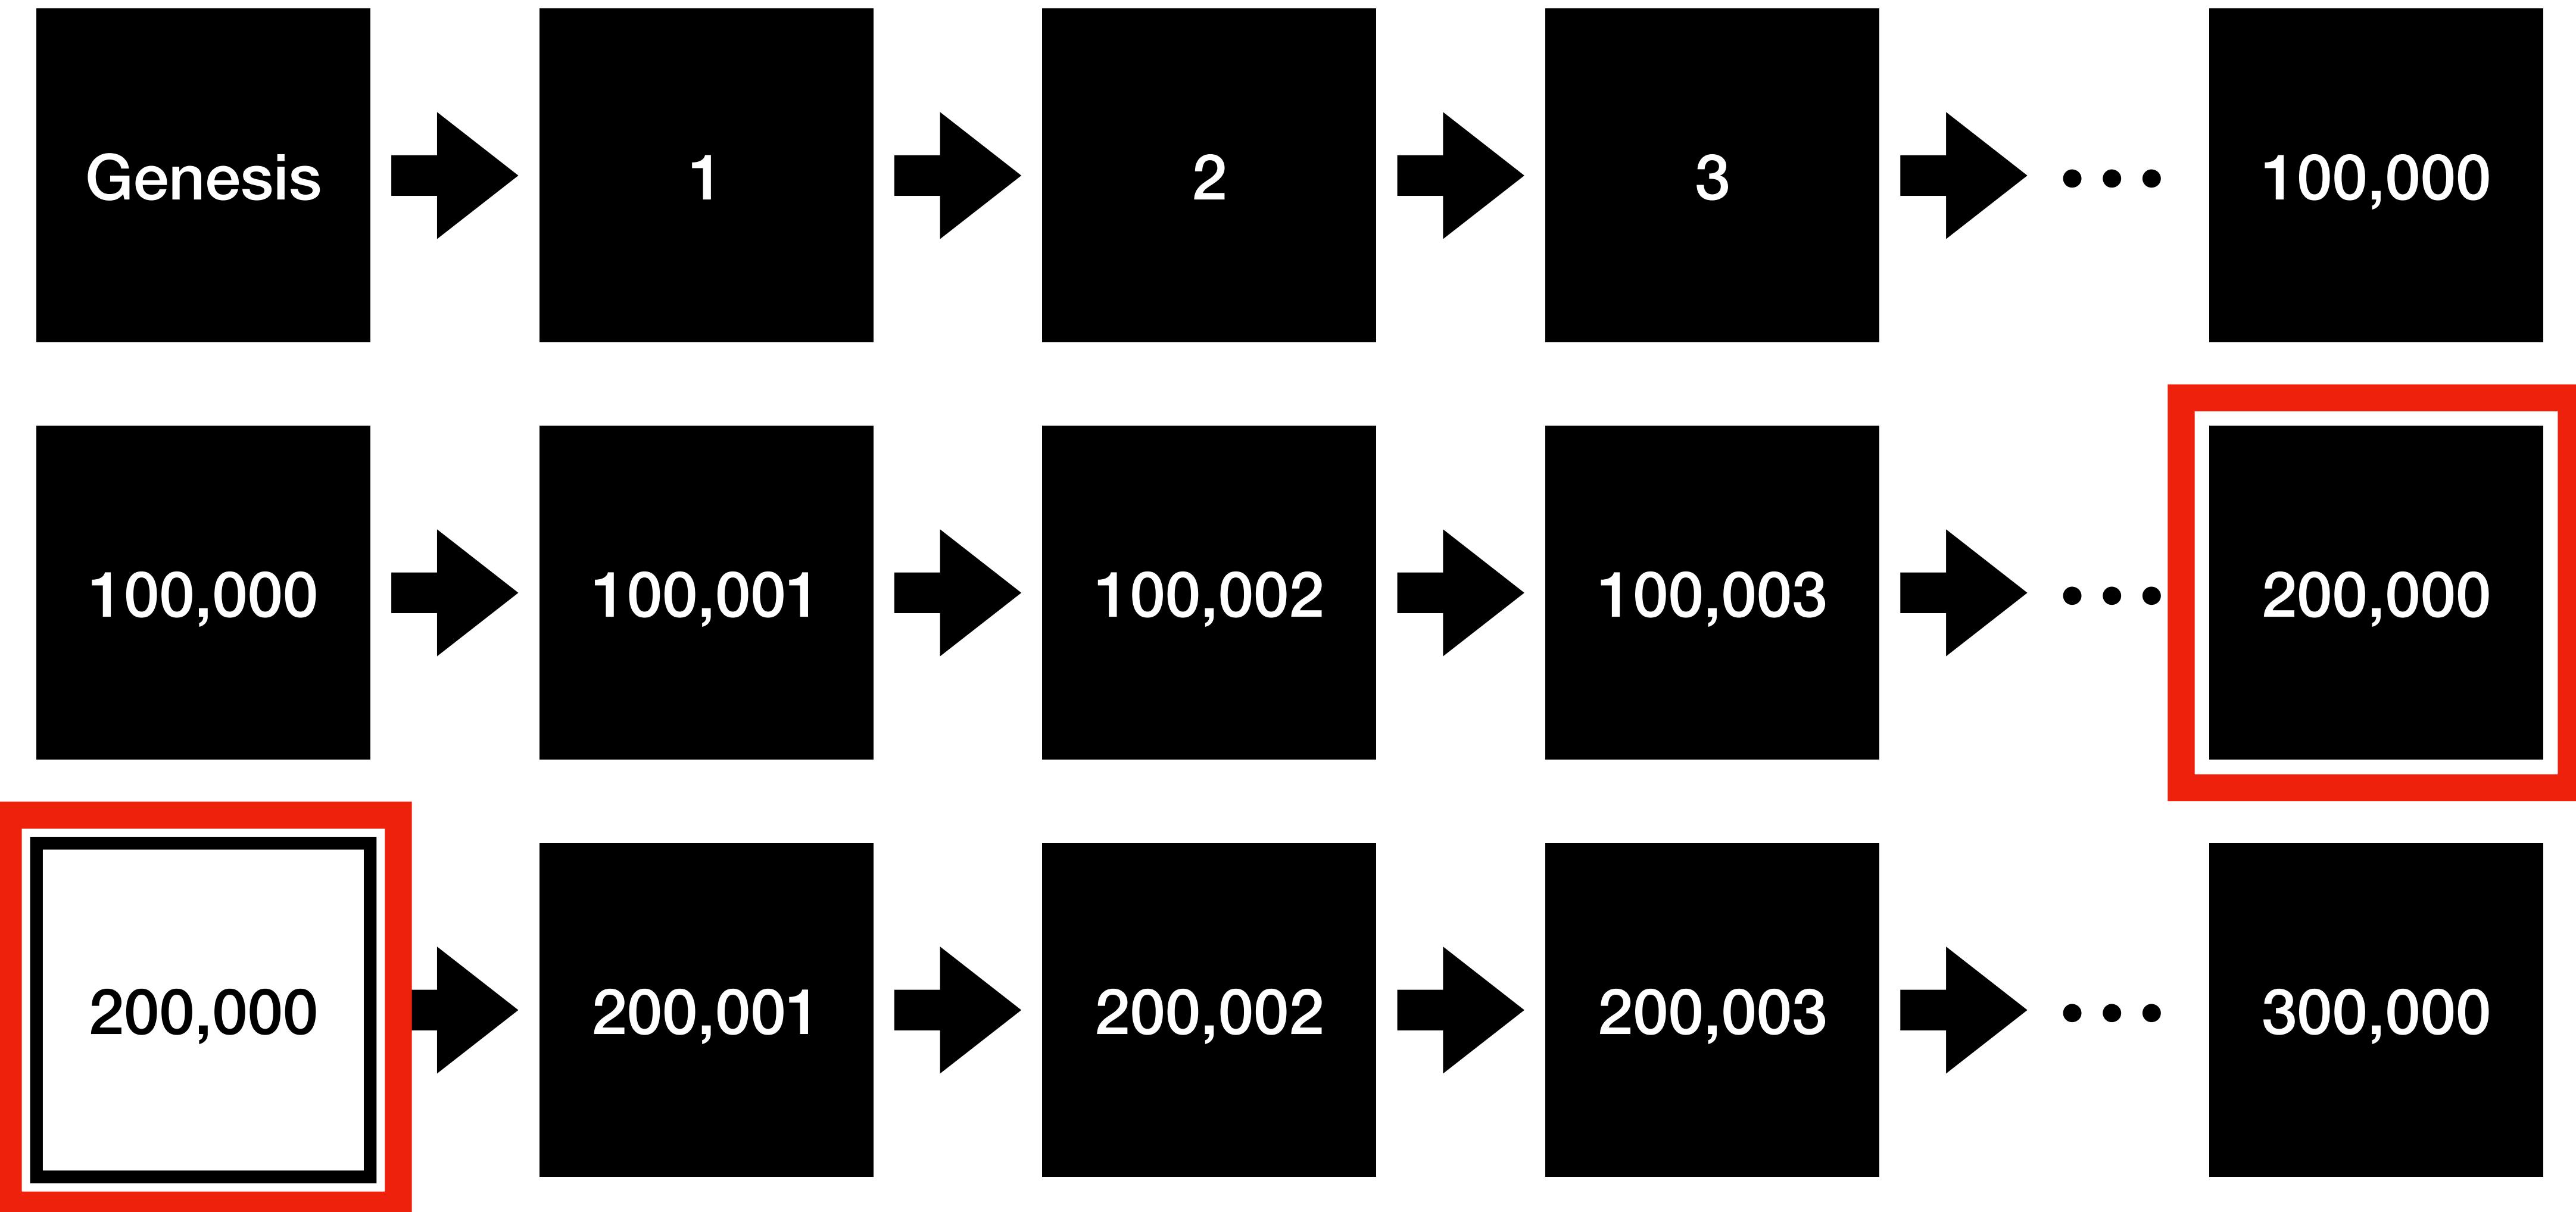
With Utreexo

Efficient parallel validation



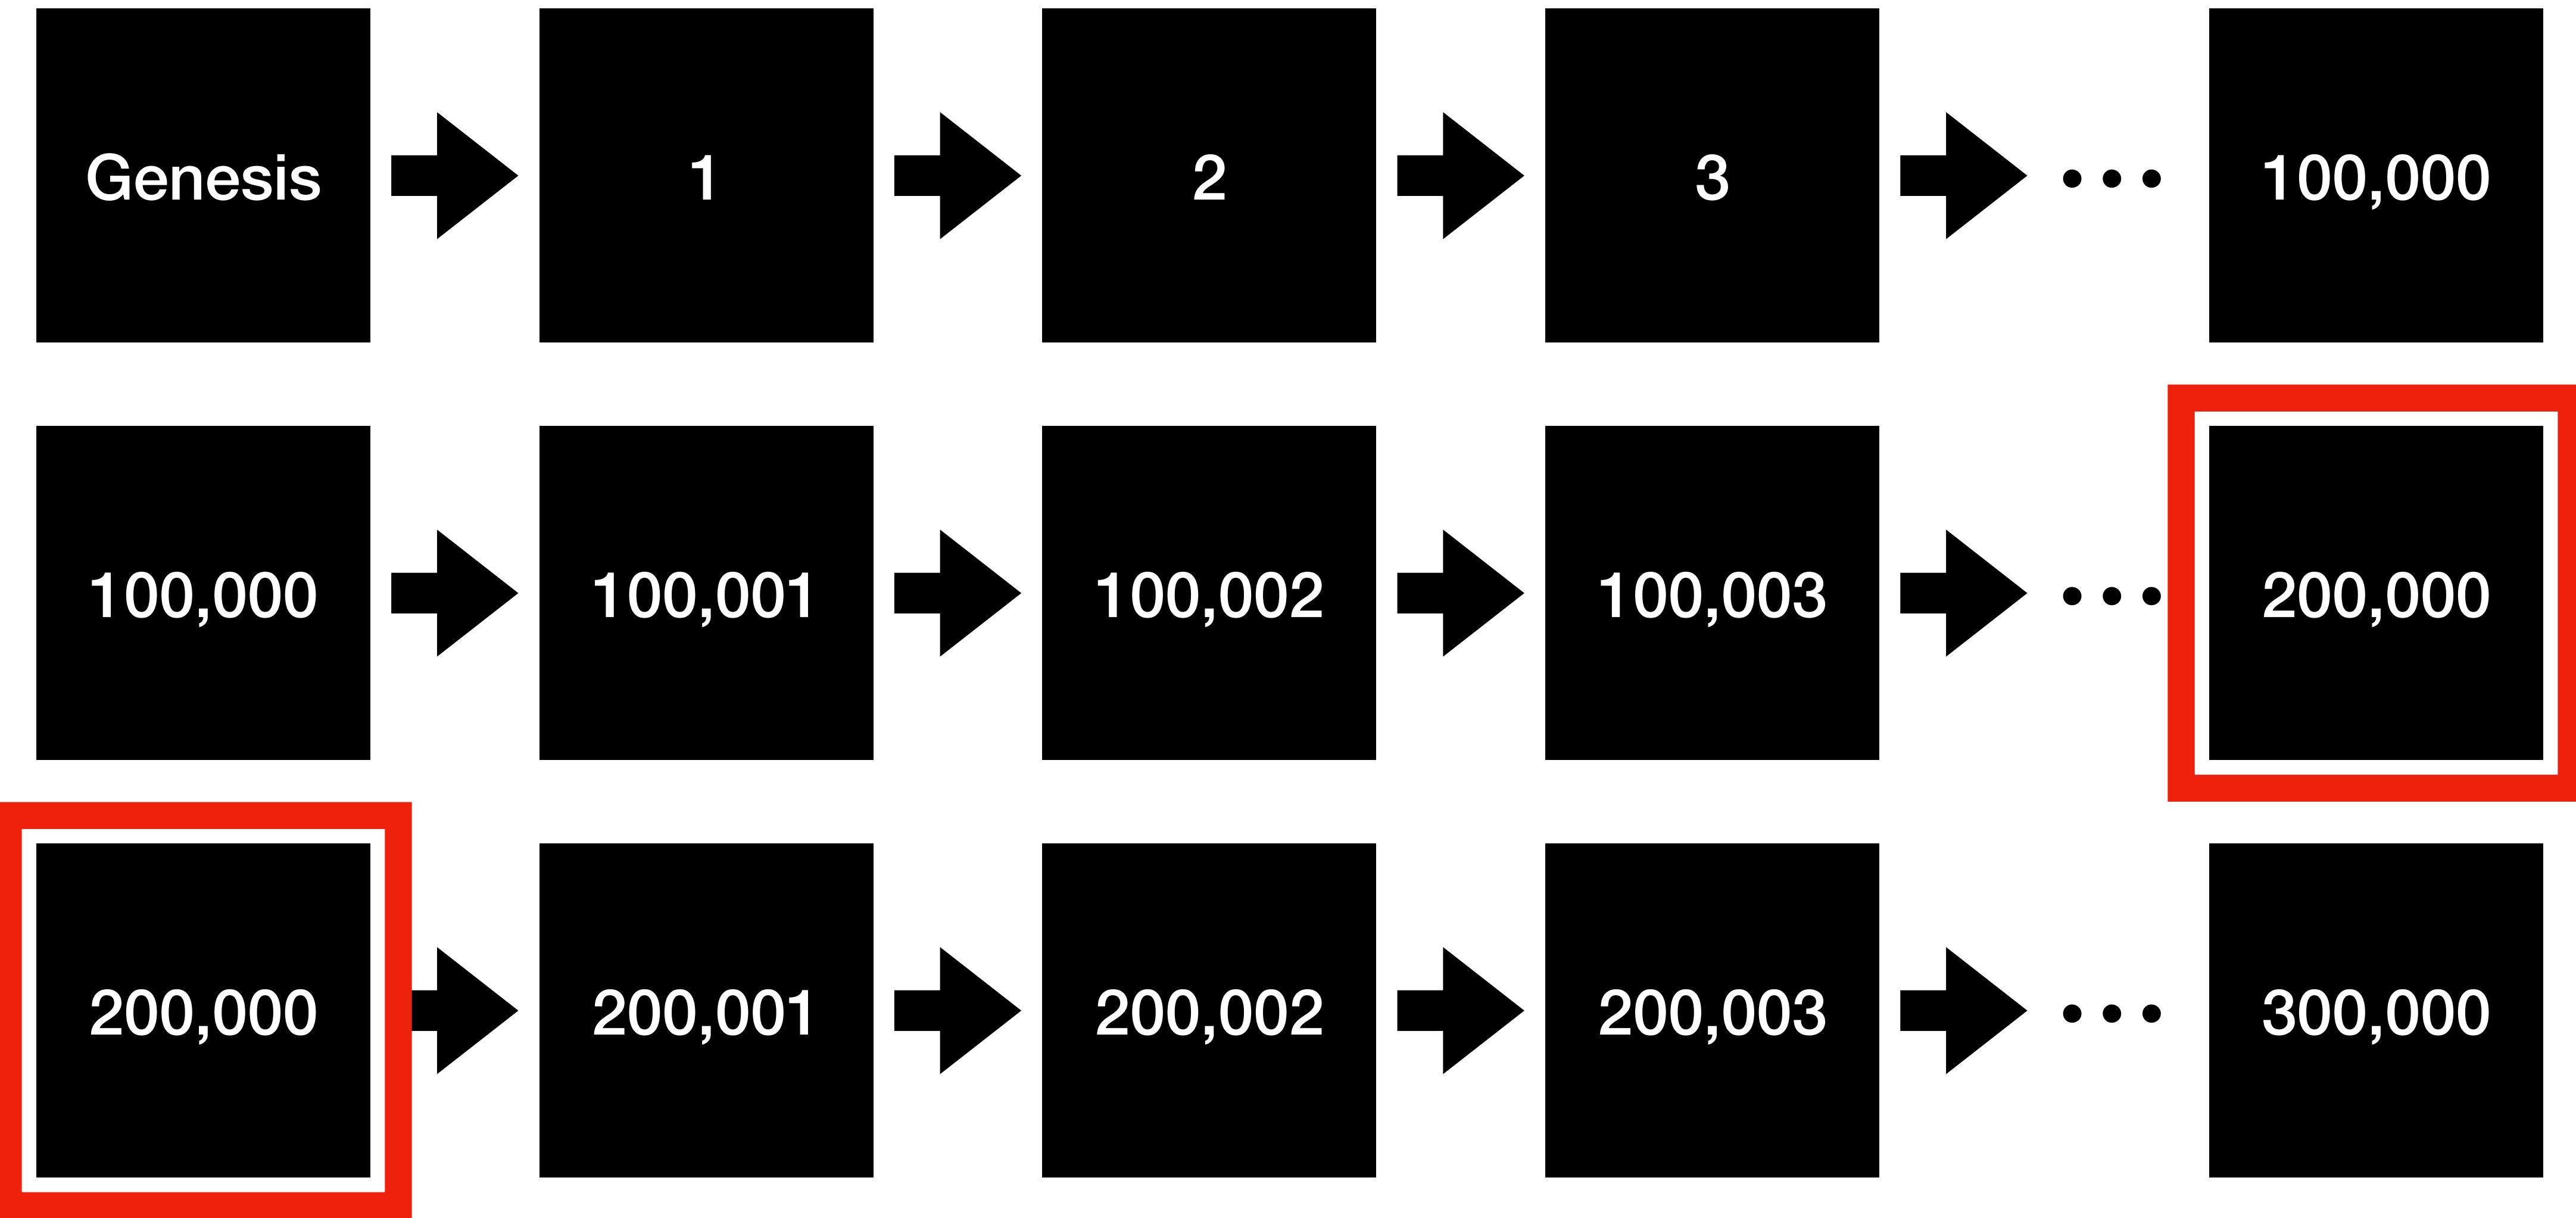
With Utreexo

Efficient parallel validation



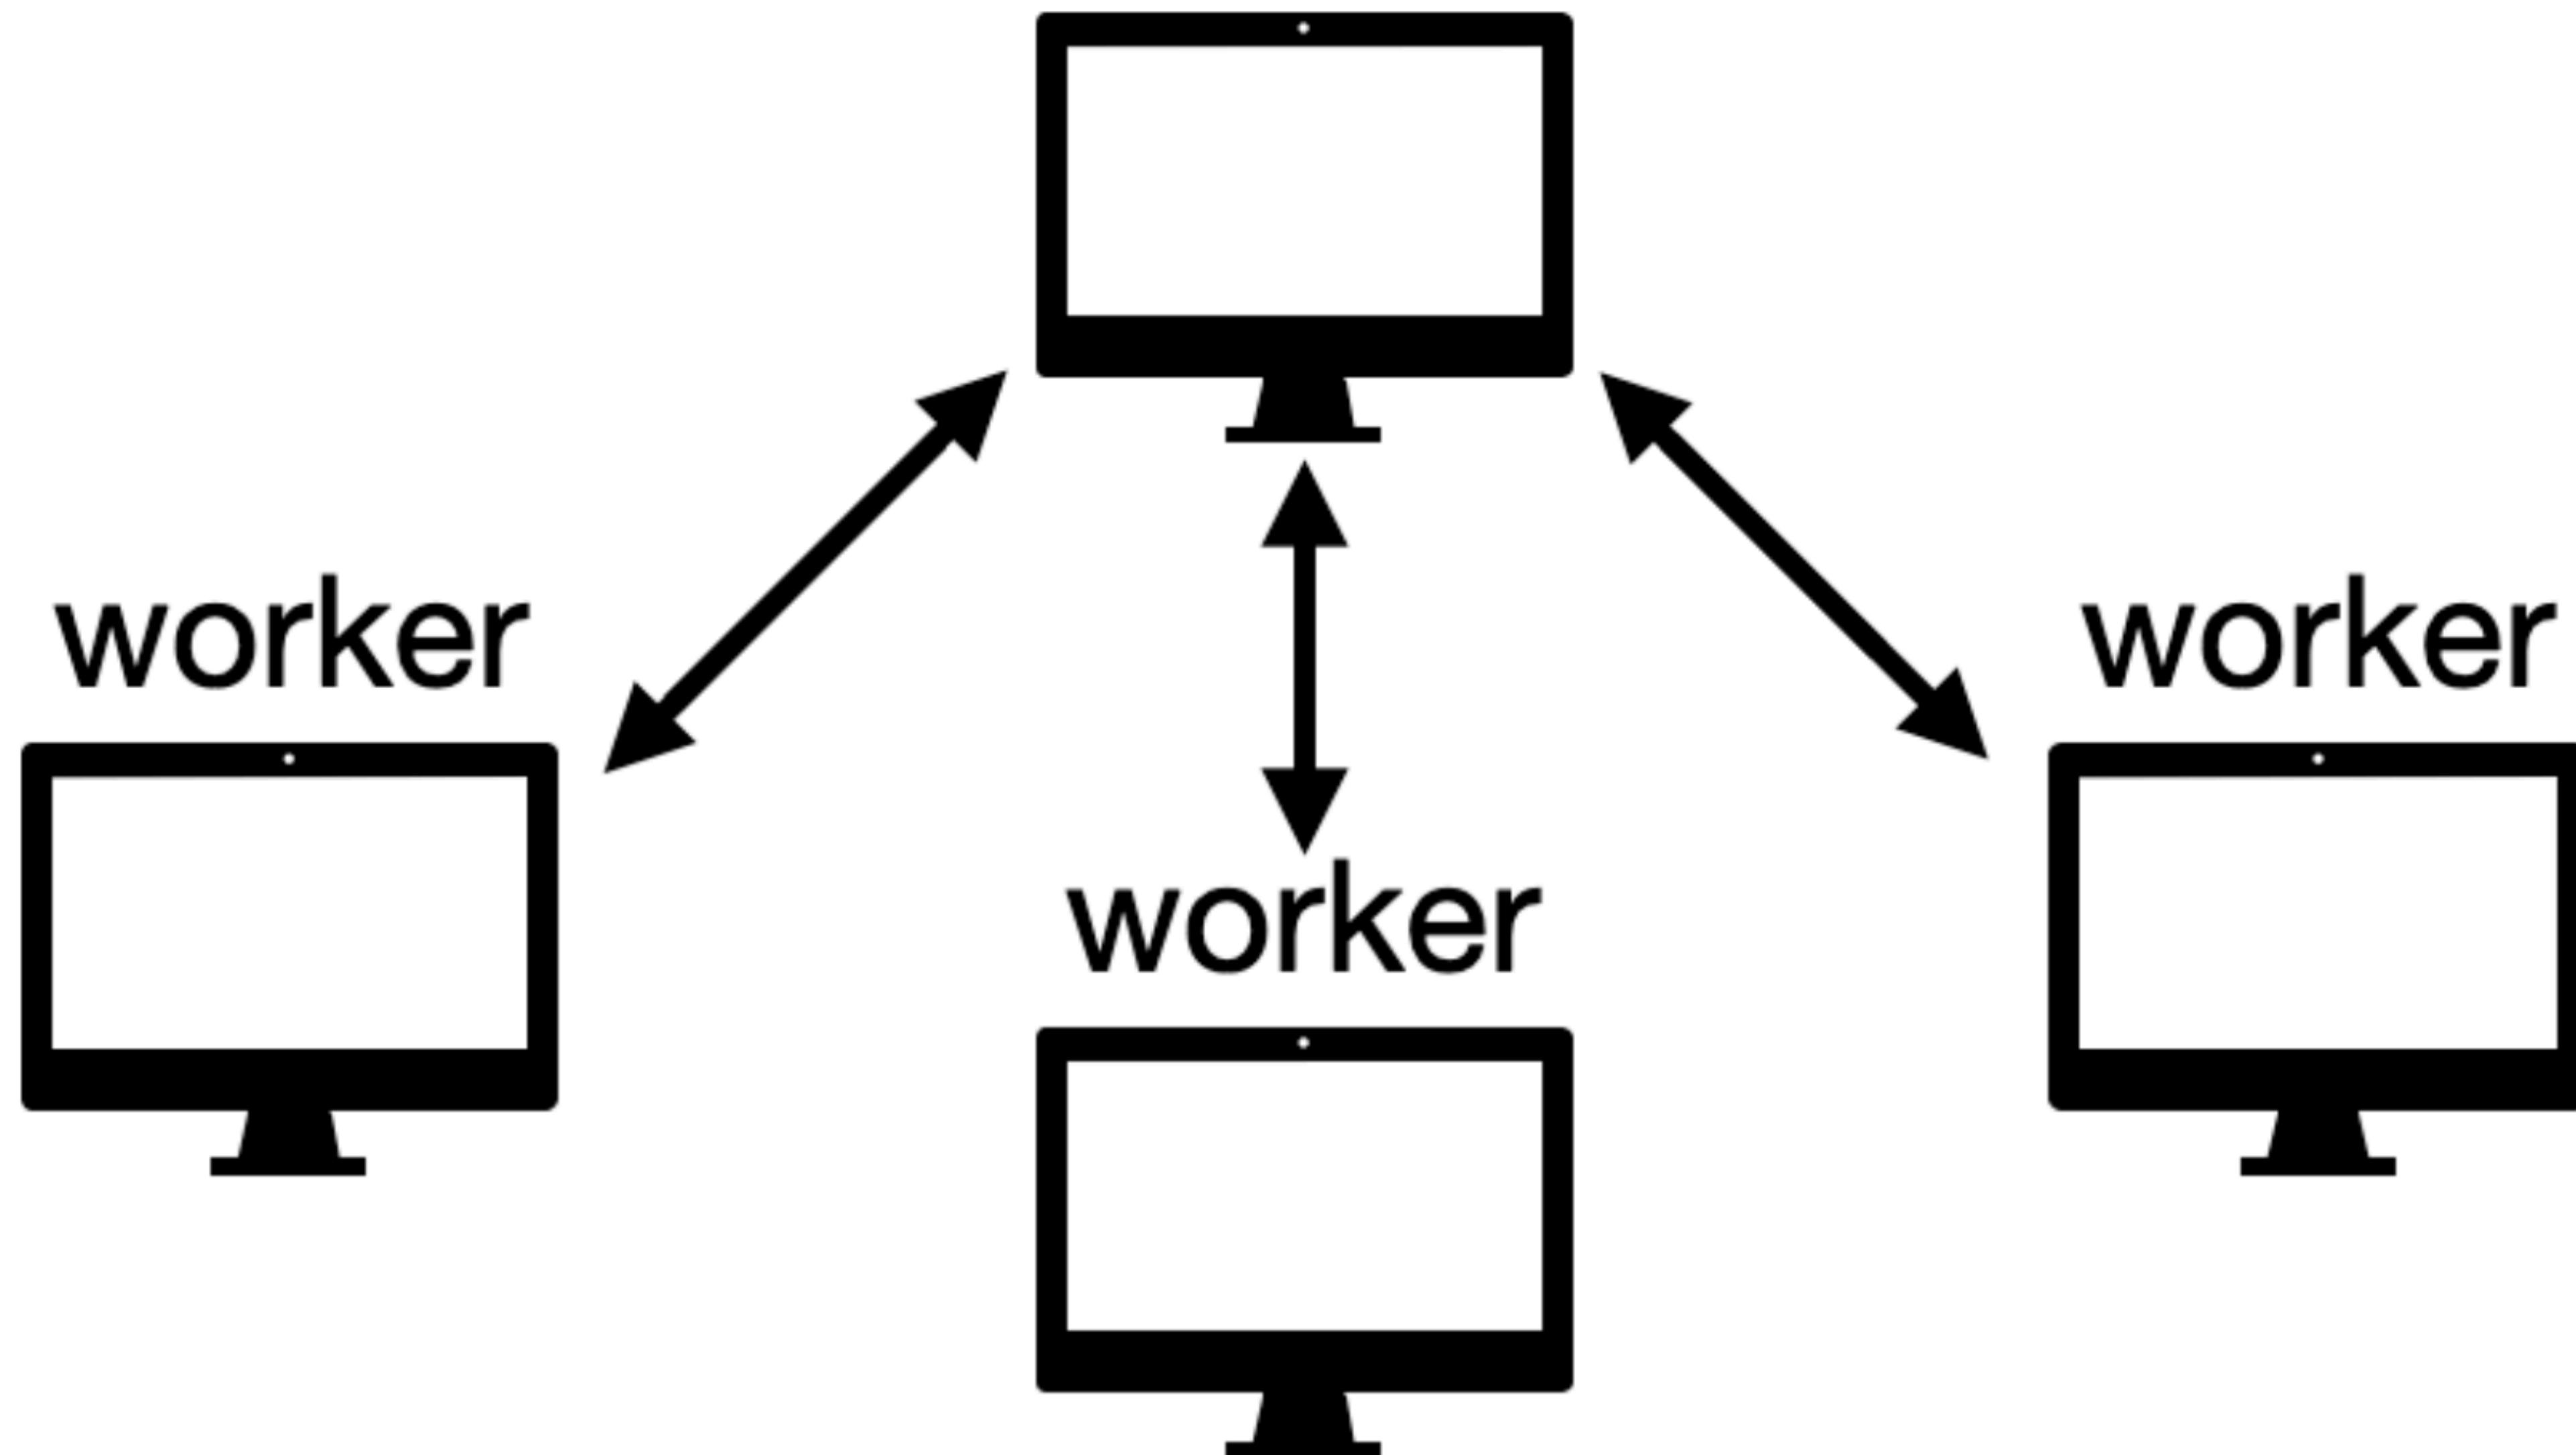
With Utreexo

Efficient parallel validation



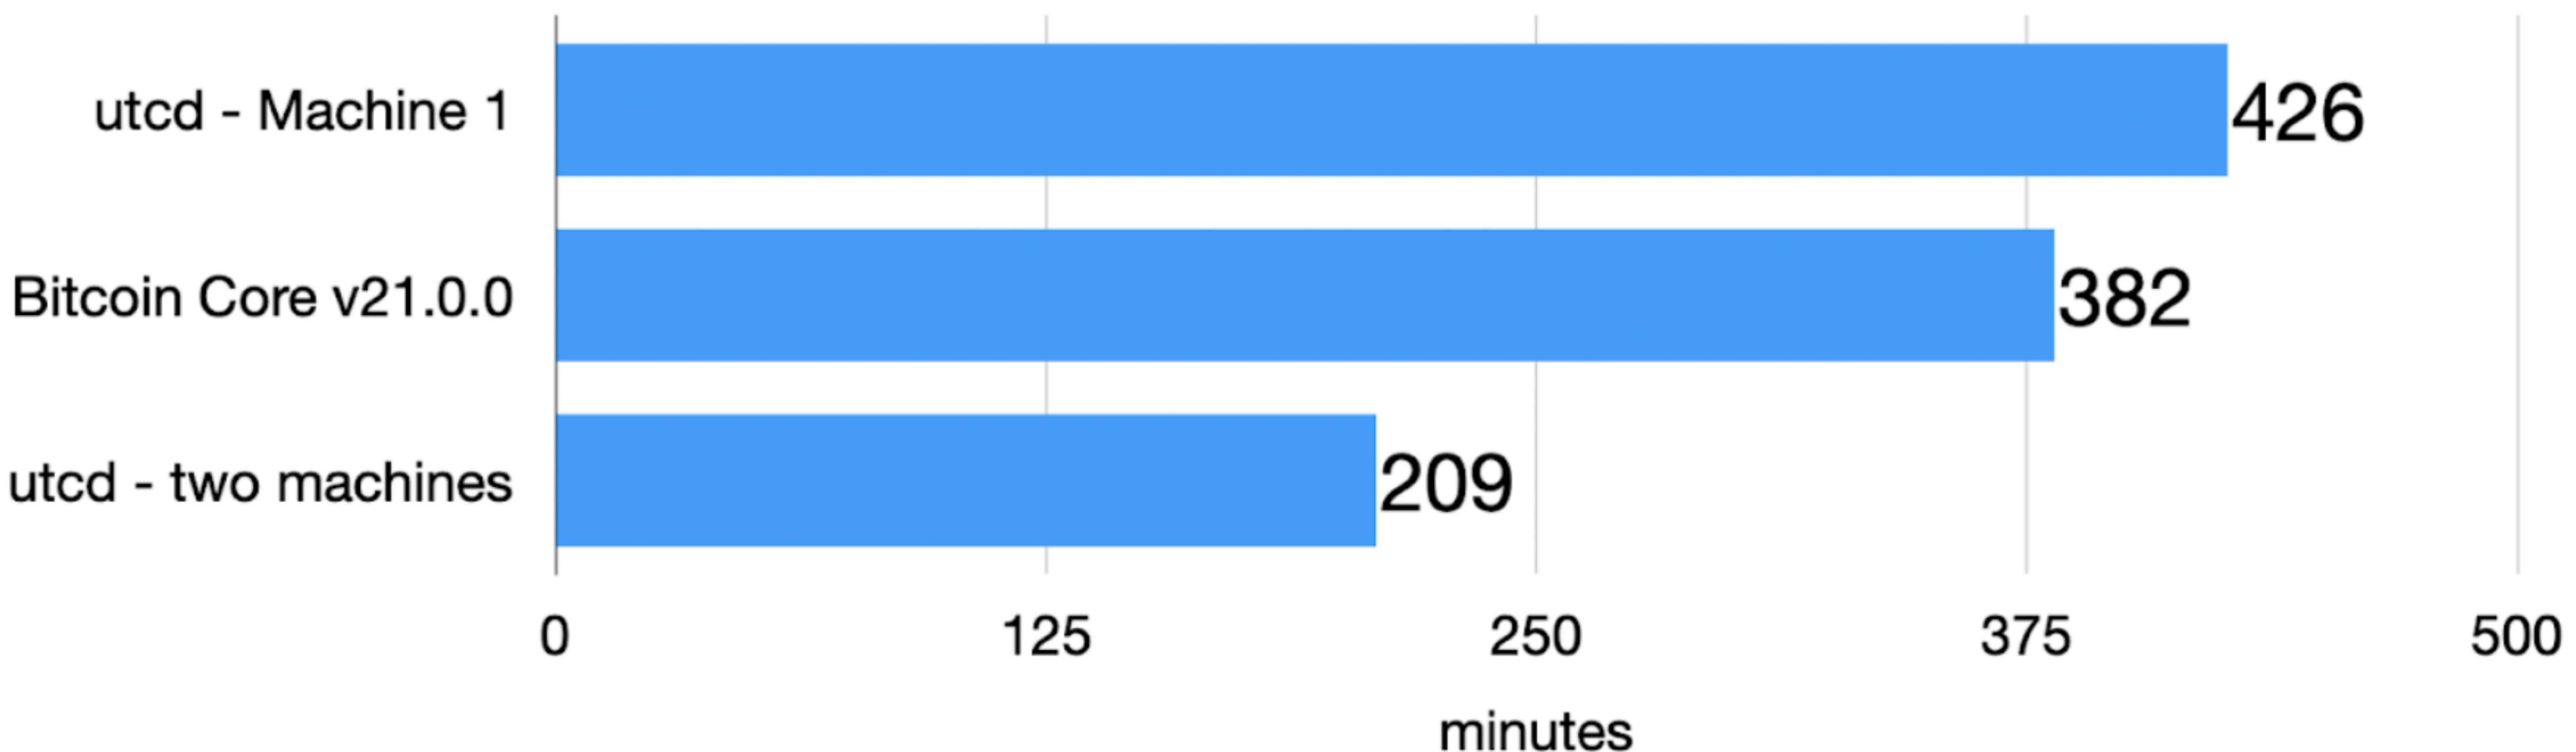
~62% faster
IBD

Coordinator/worker

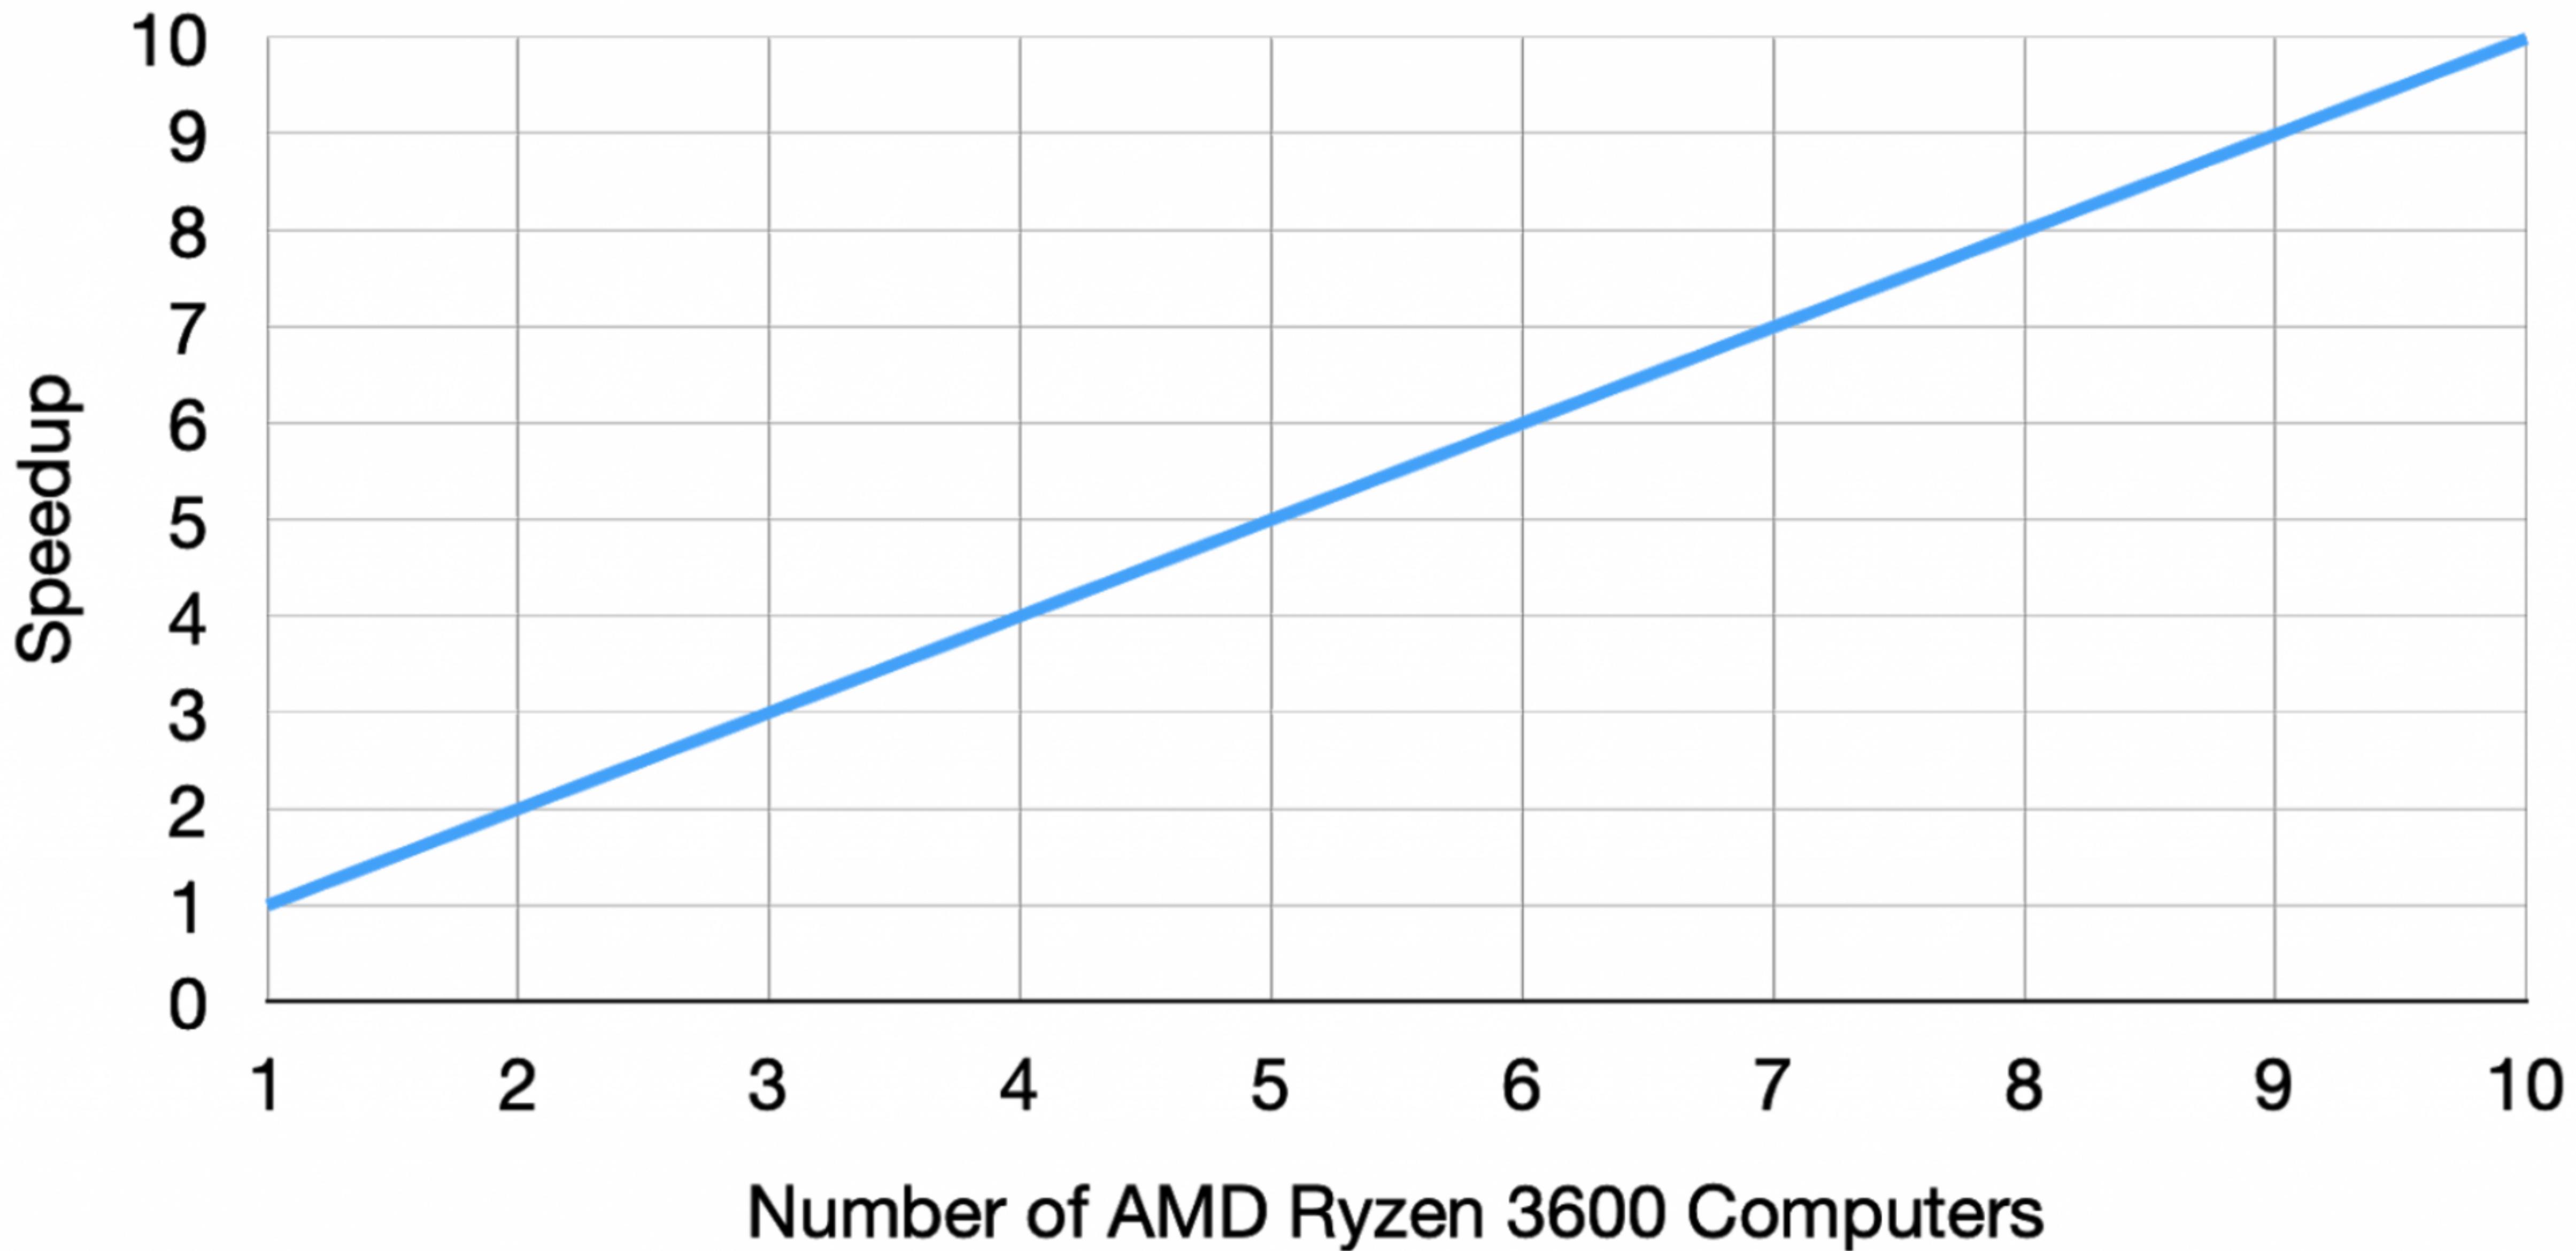


2x faster IBD
When using 2 computers for IBD

Initial Block Download Speeds to block 671,000 (Full signature check)



Utreexo Parallel Node Speedups (based on Amdahl's Law)



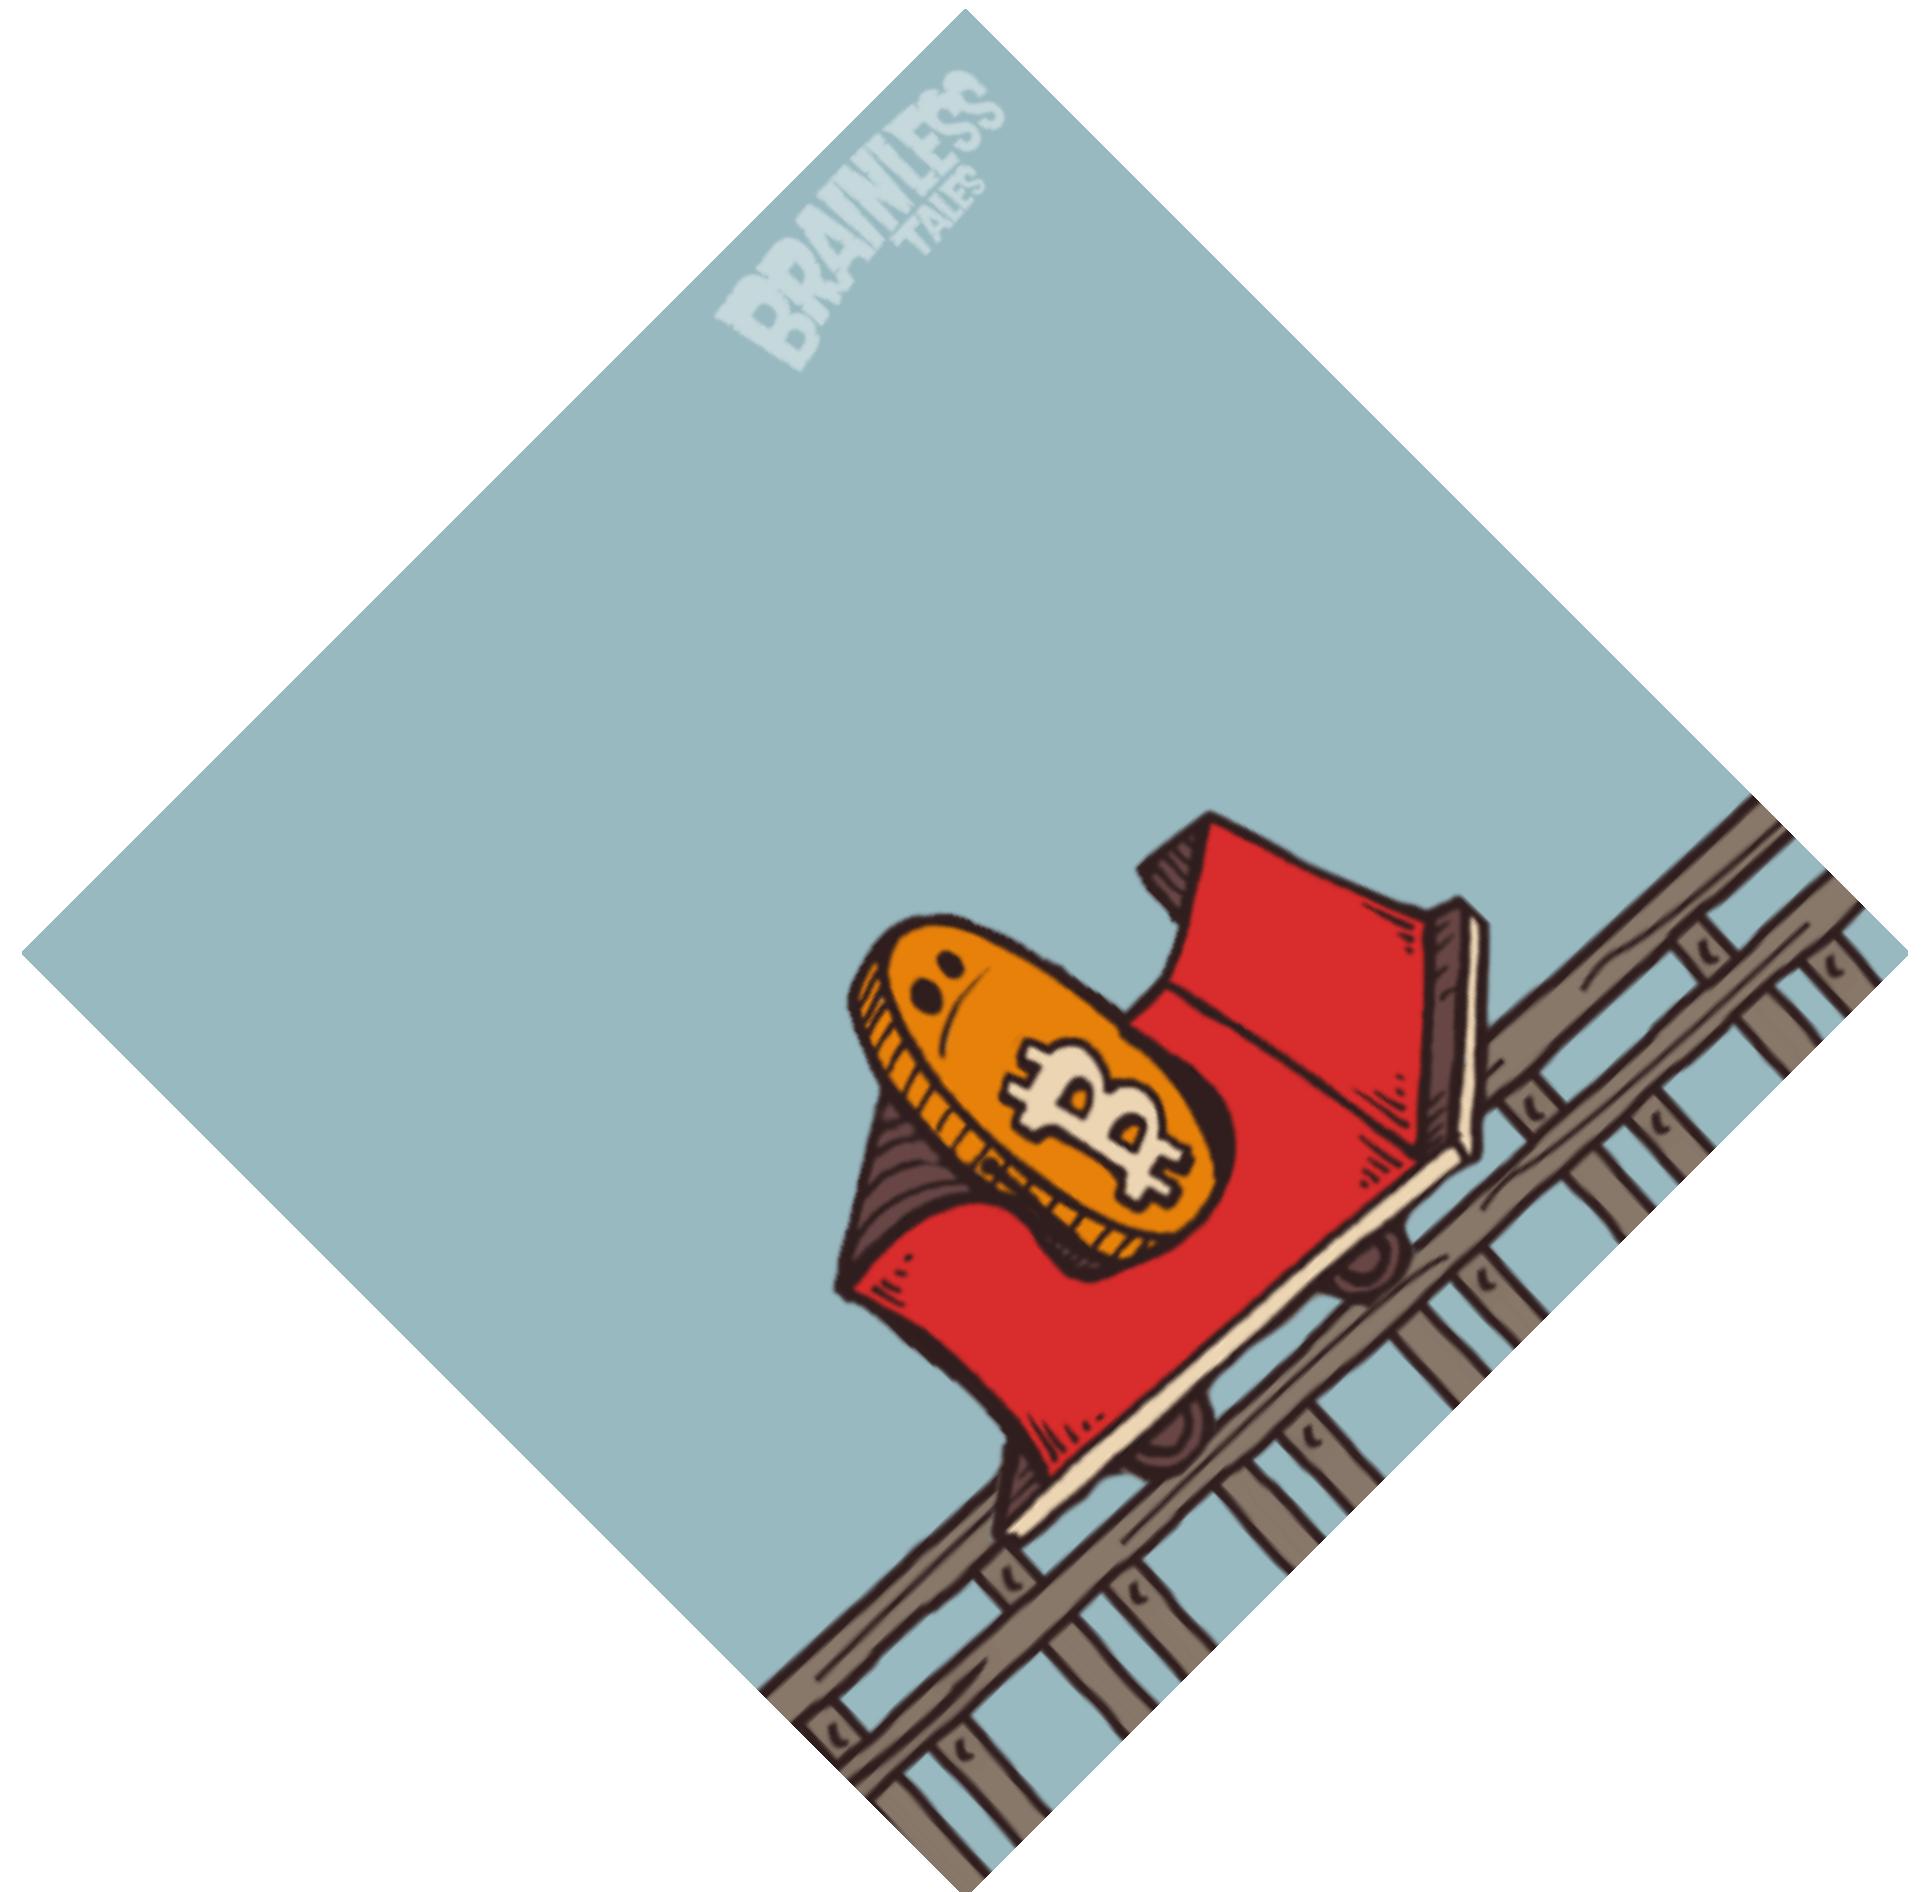
What's the catch?

More data to download

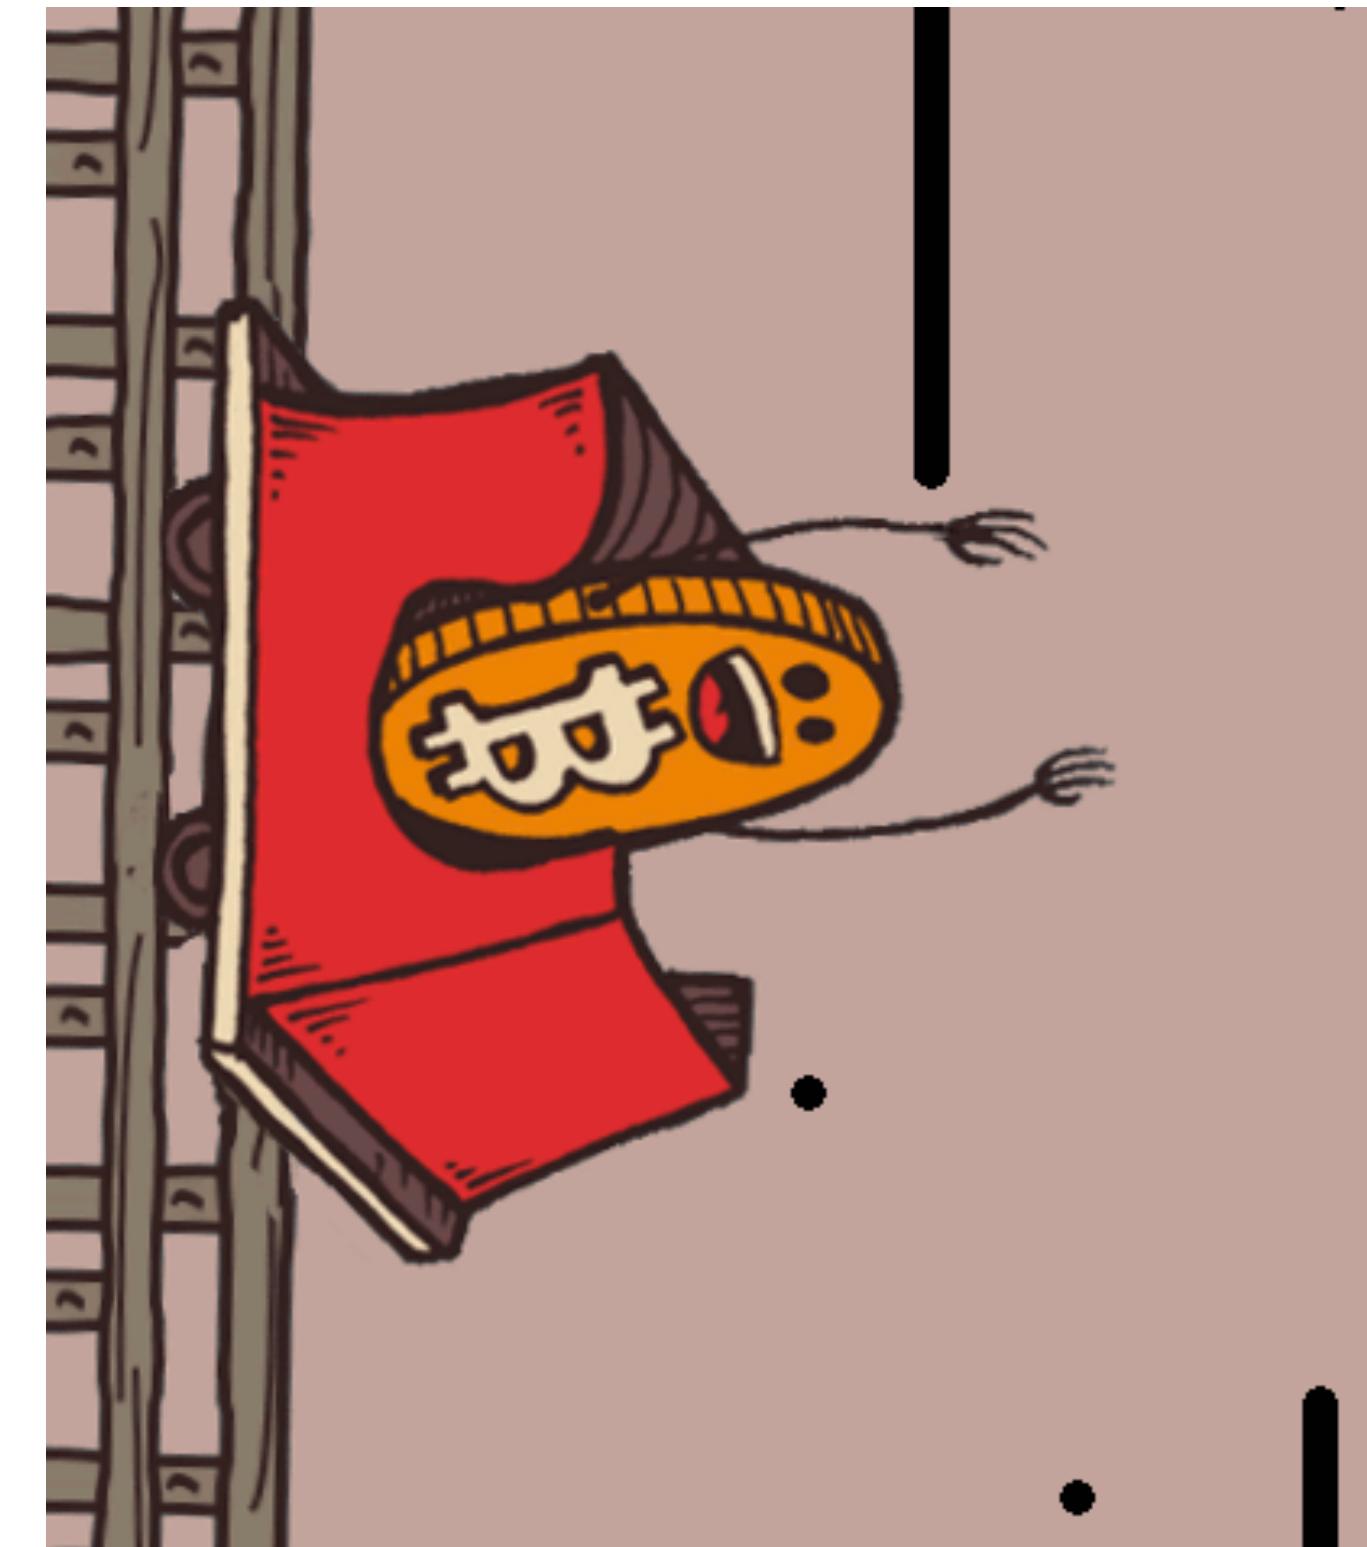
Extra data that a utreexo node will download

- 364GB of proof data (as of block 710,000)
- Optimizations being explored

Bandwidth CPU



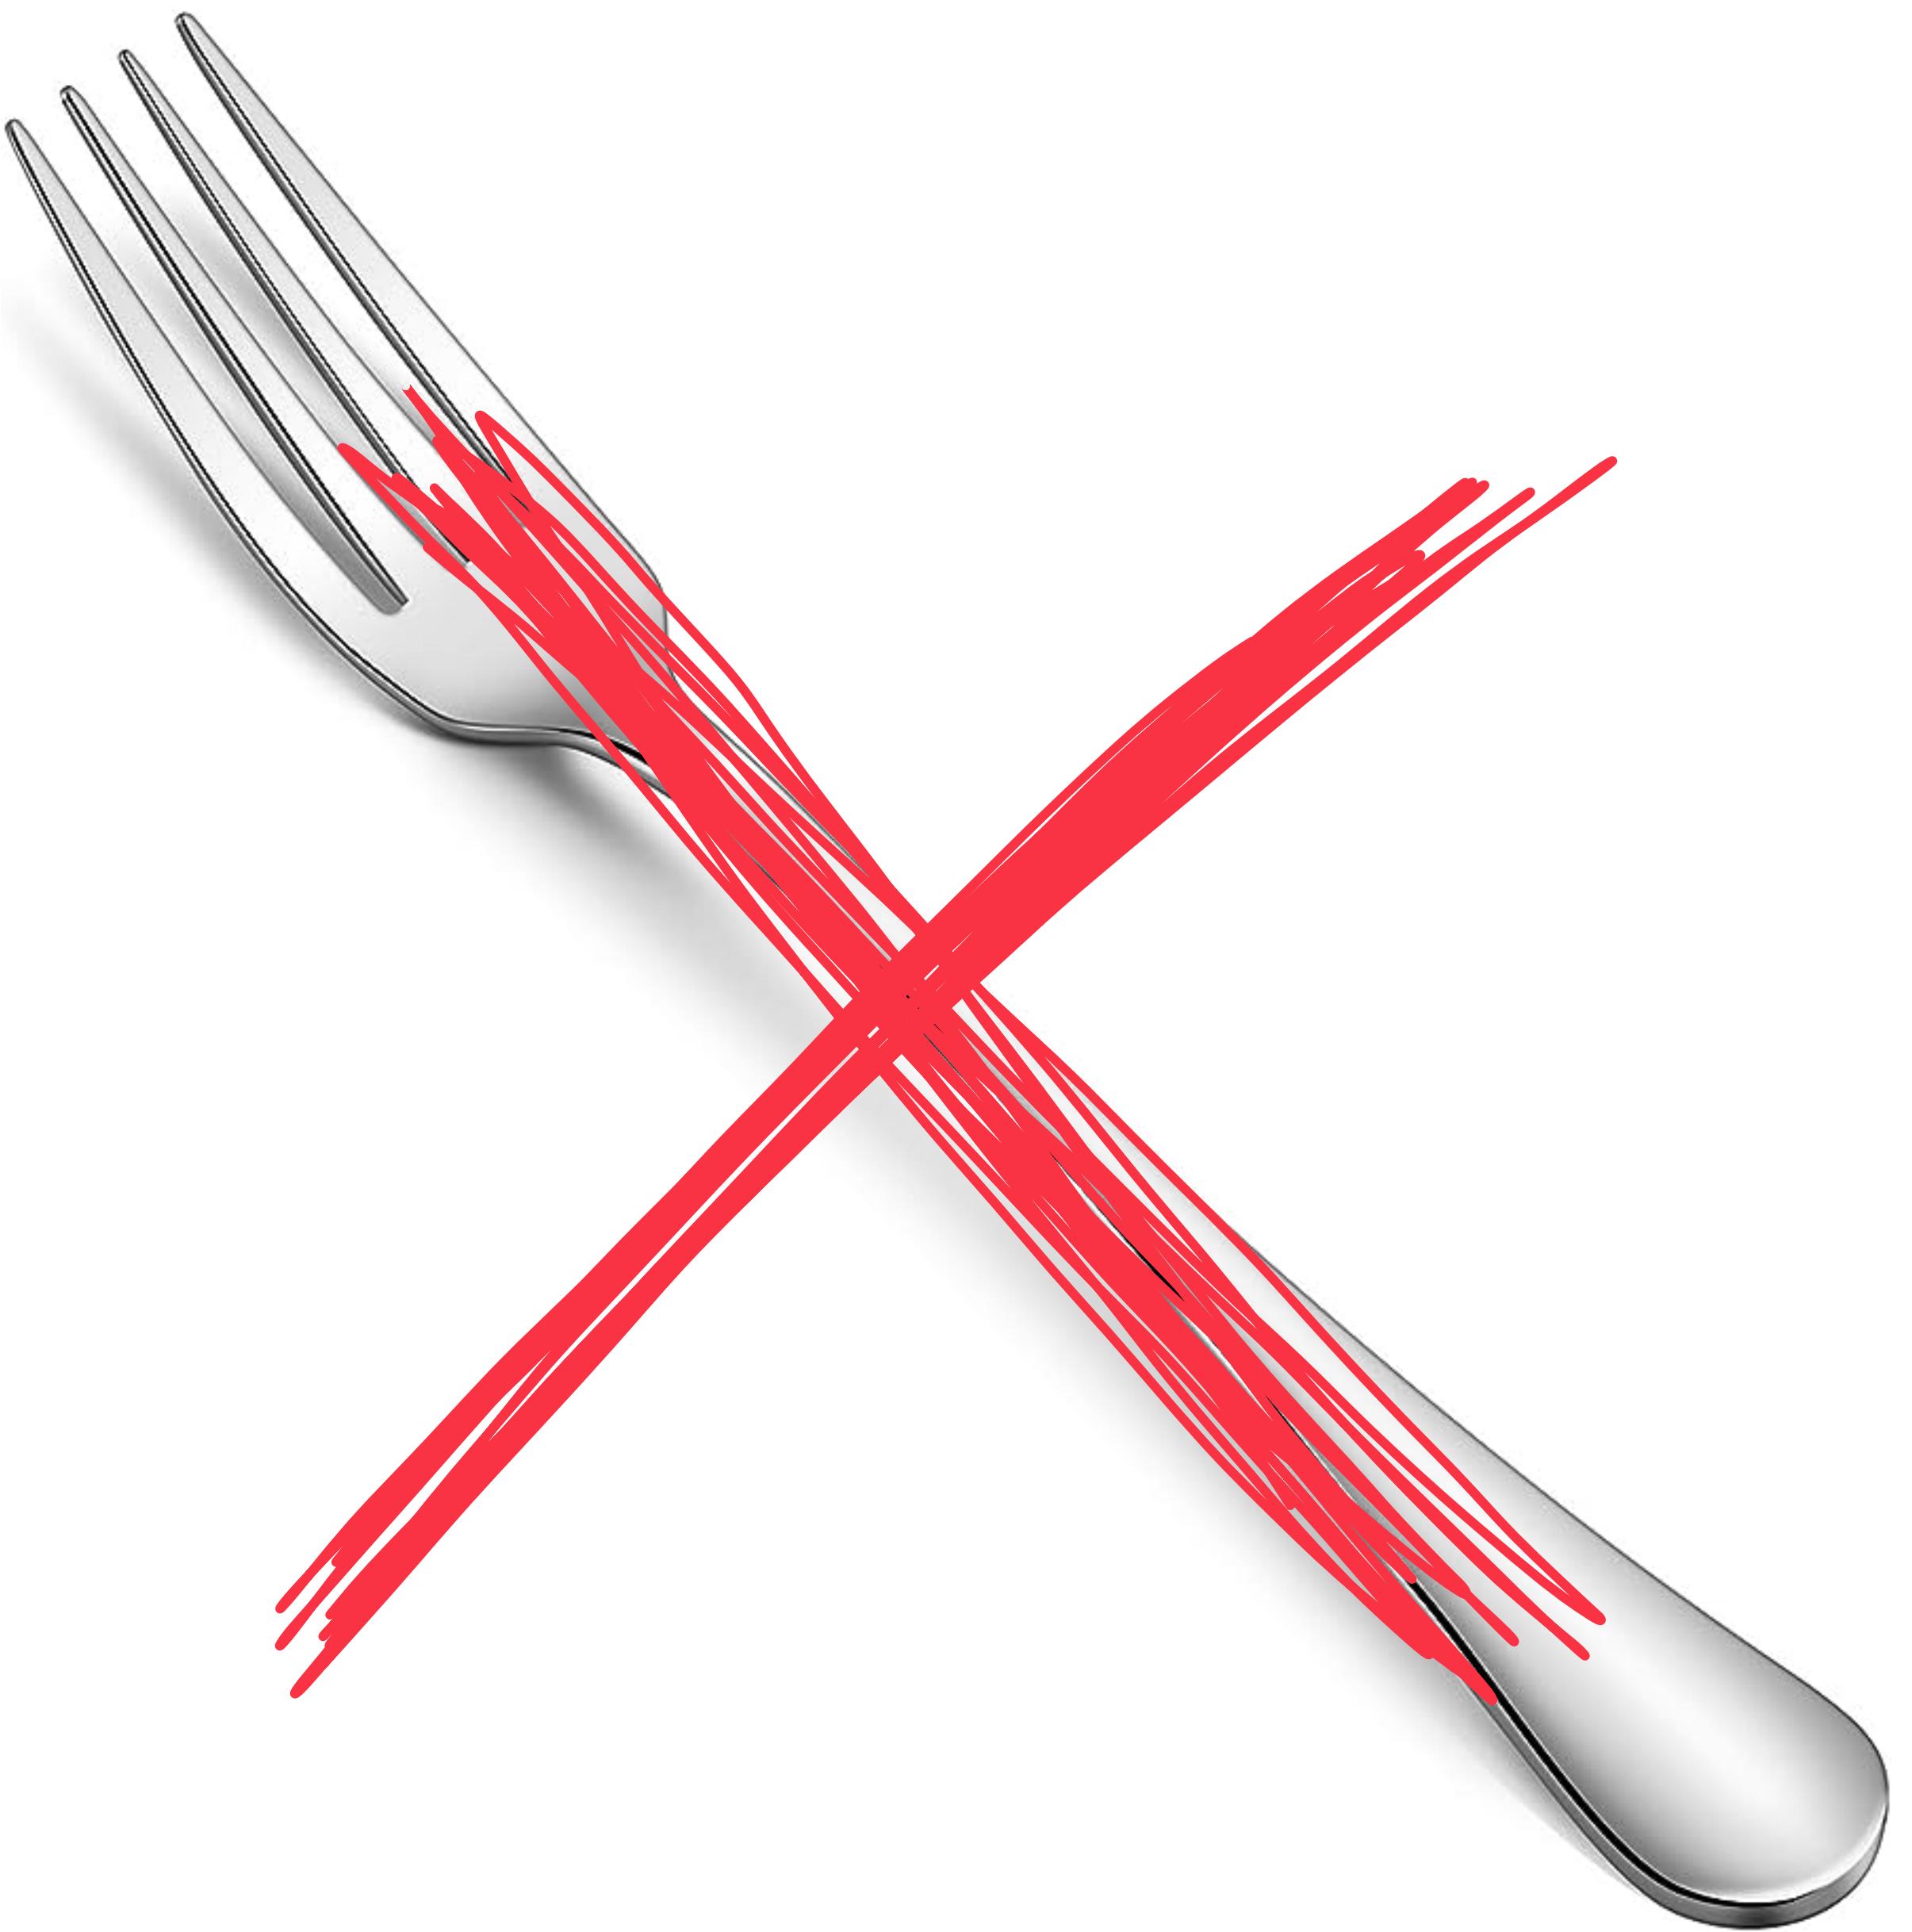
disk i/o RAM usage



Deployment

100% optional

Ignore it if you don't want it



Need nodes to generate proofs

"bridge nodes"

- Need to generate proofs for new blocks/txs
- These "bridge nodes" keep the entire utxo set and the entire merkle forest

Need nodes to generate proofs

"bridge nodes"

- Can be pruned
- Implemented as an indexer in
github.com/utreexo/utreexod

Aligning incentives

**No burden of
UTXO storage**

coinbase

coinbase

The Binance logo icon is a yellow diamond shape composed of eight smaller yellow triangles pointing inward toward the center.

BINANCE

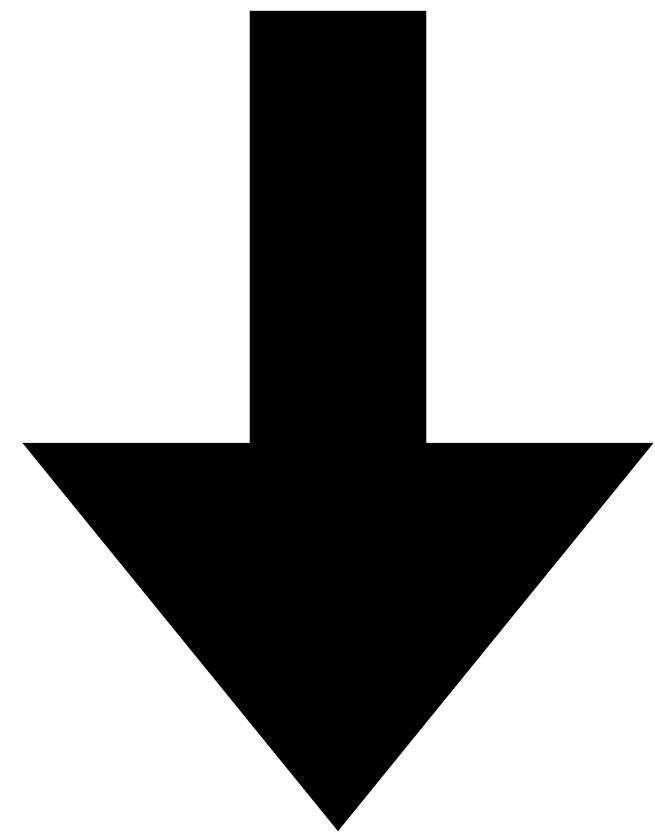
coinbase



BITCOIN STAMPS

Unprunable UTXO Art, Because Sats Don't Exist.

Everyone



Each store their own

How does it work?

Algorithms

Consensus critical stuff

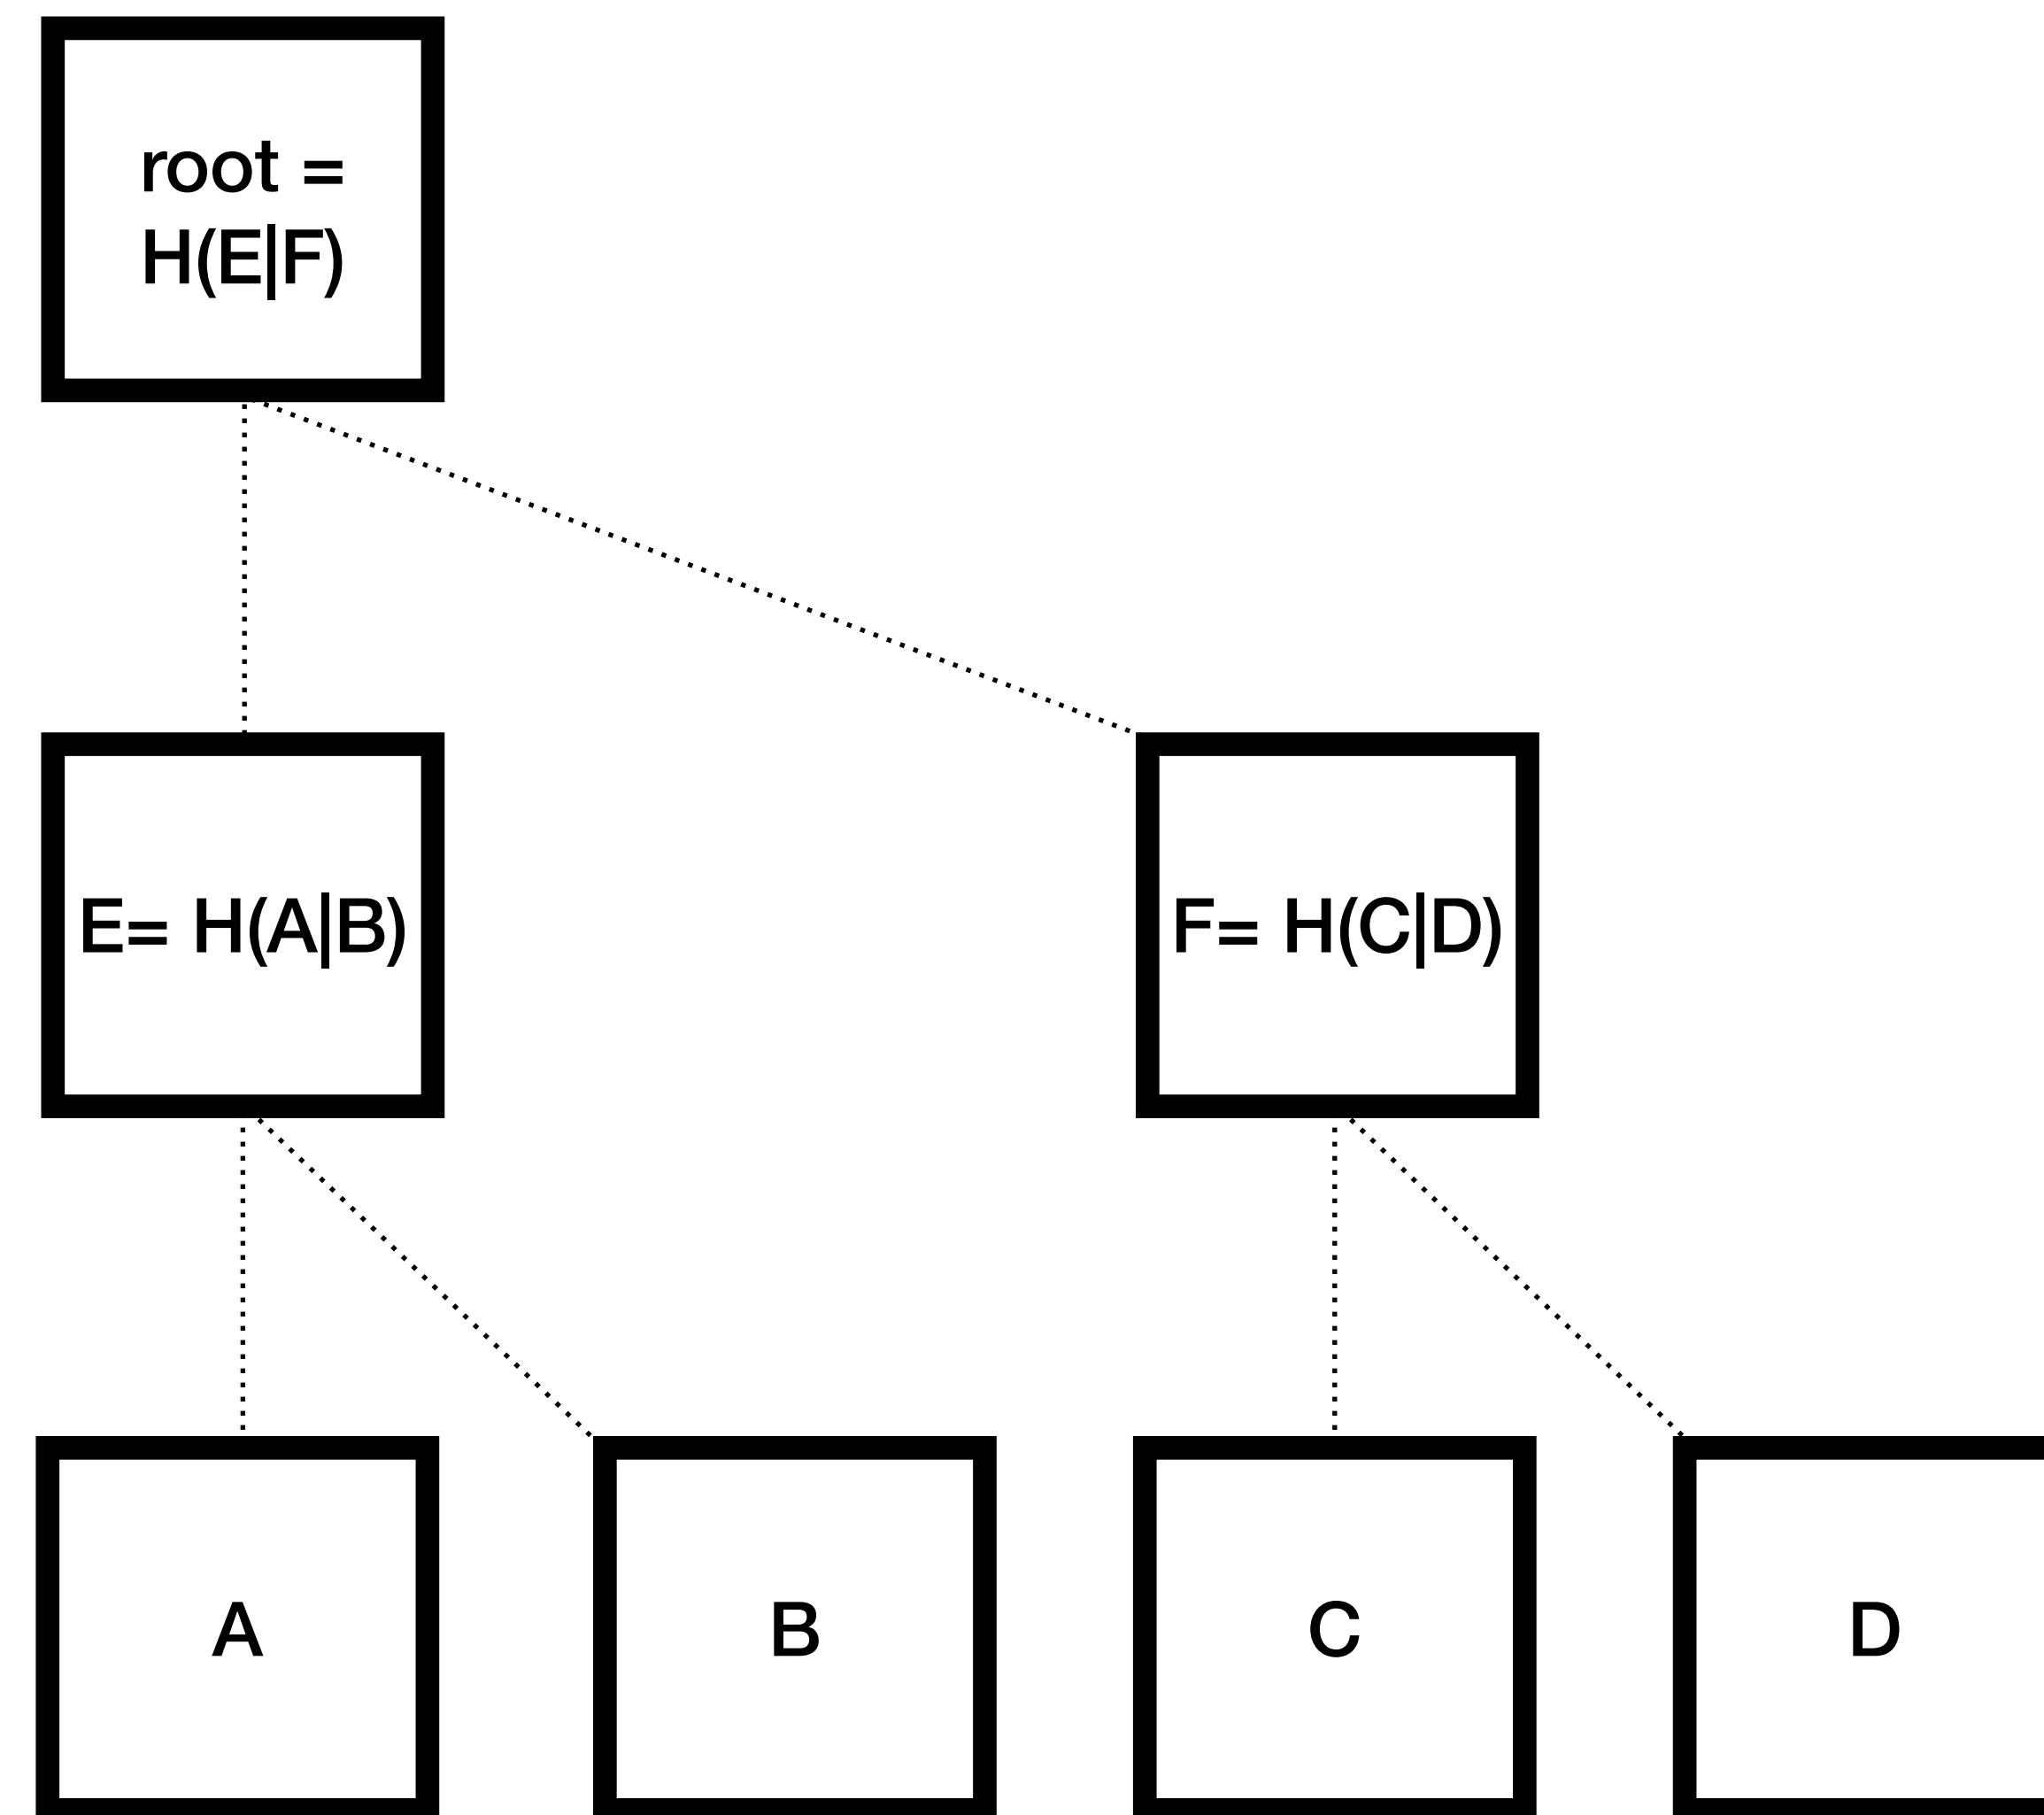
1. Add a UTXO
2. Delete a UTXO
3. Verify the existence

Algorithms

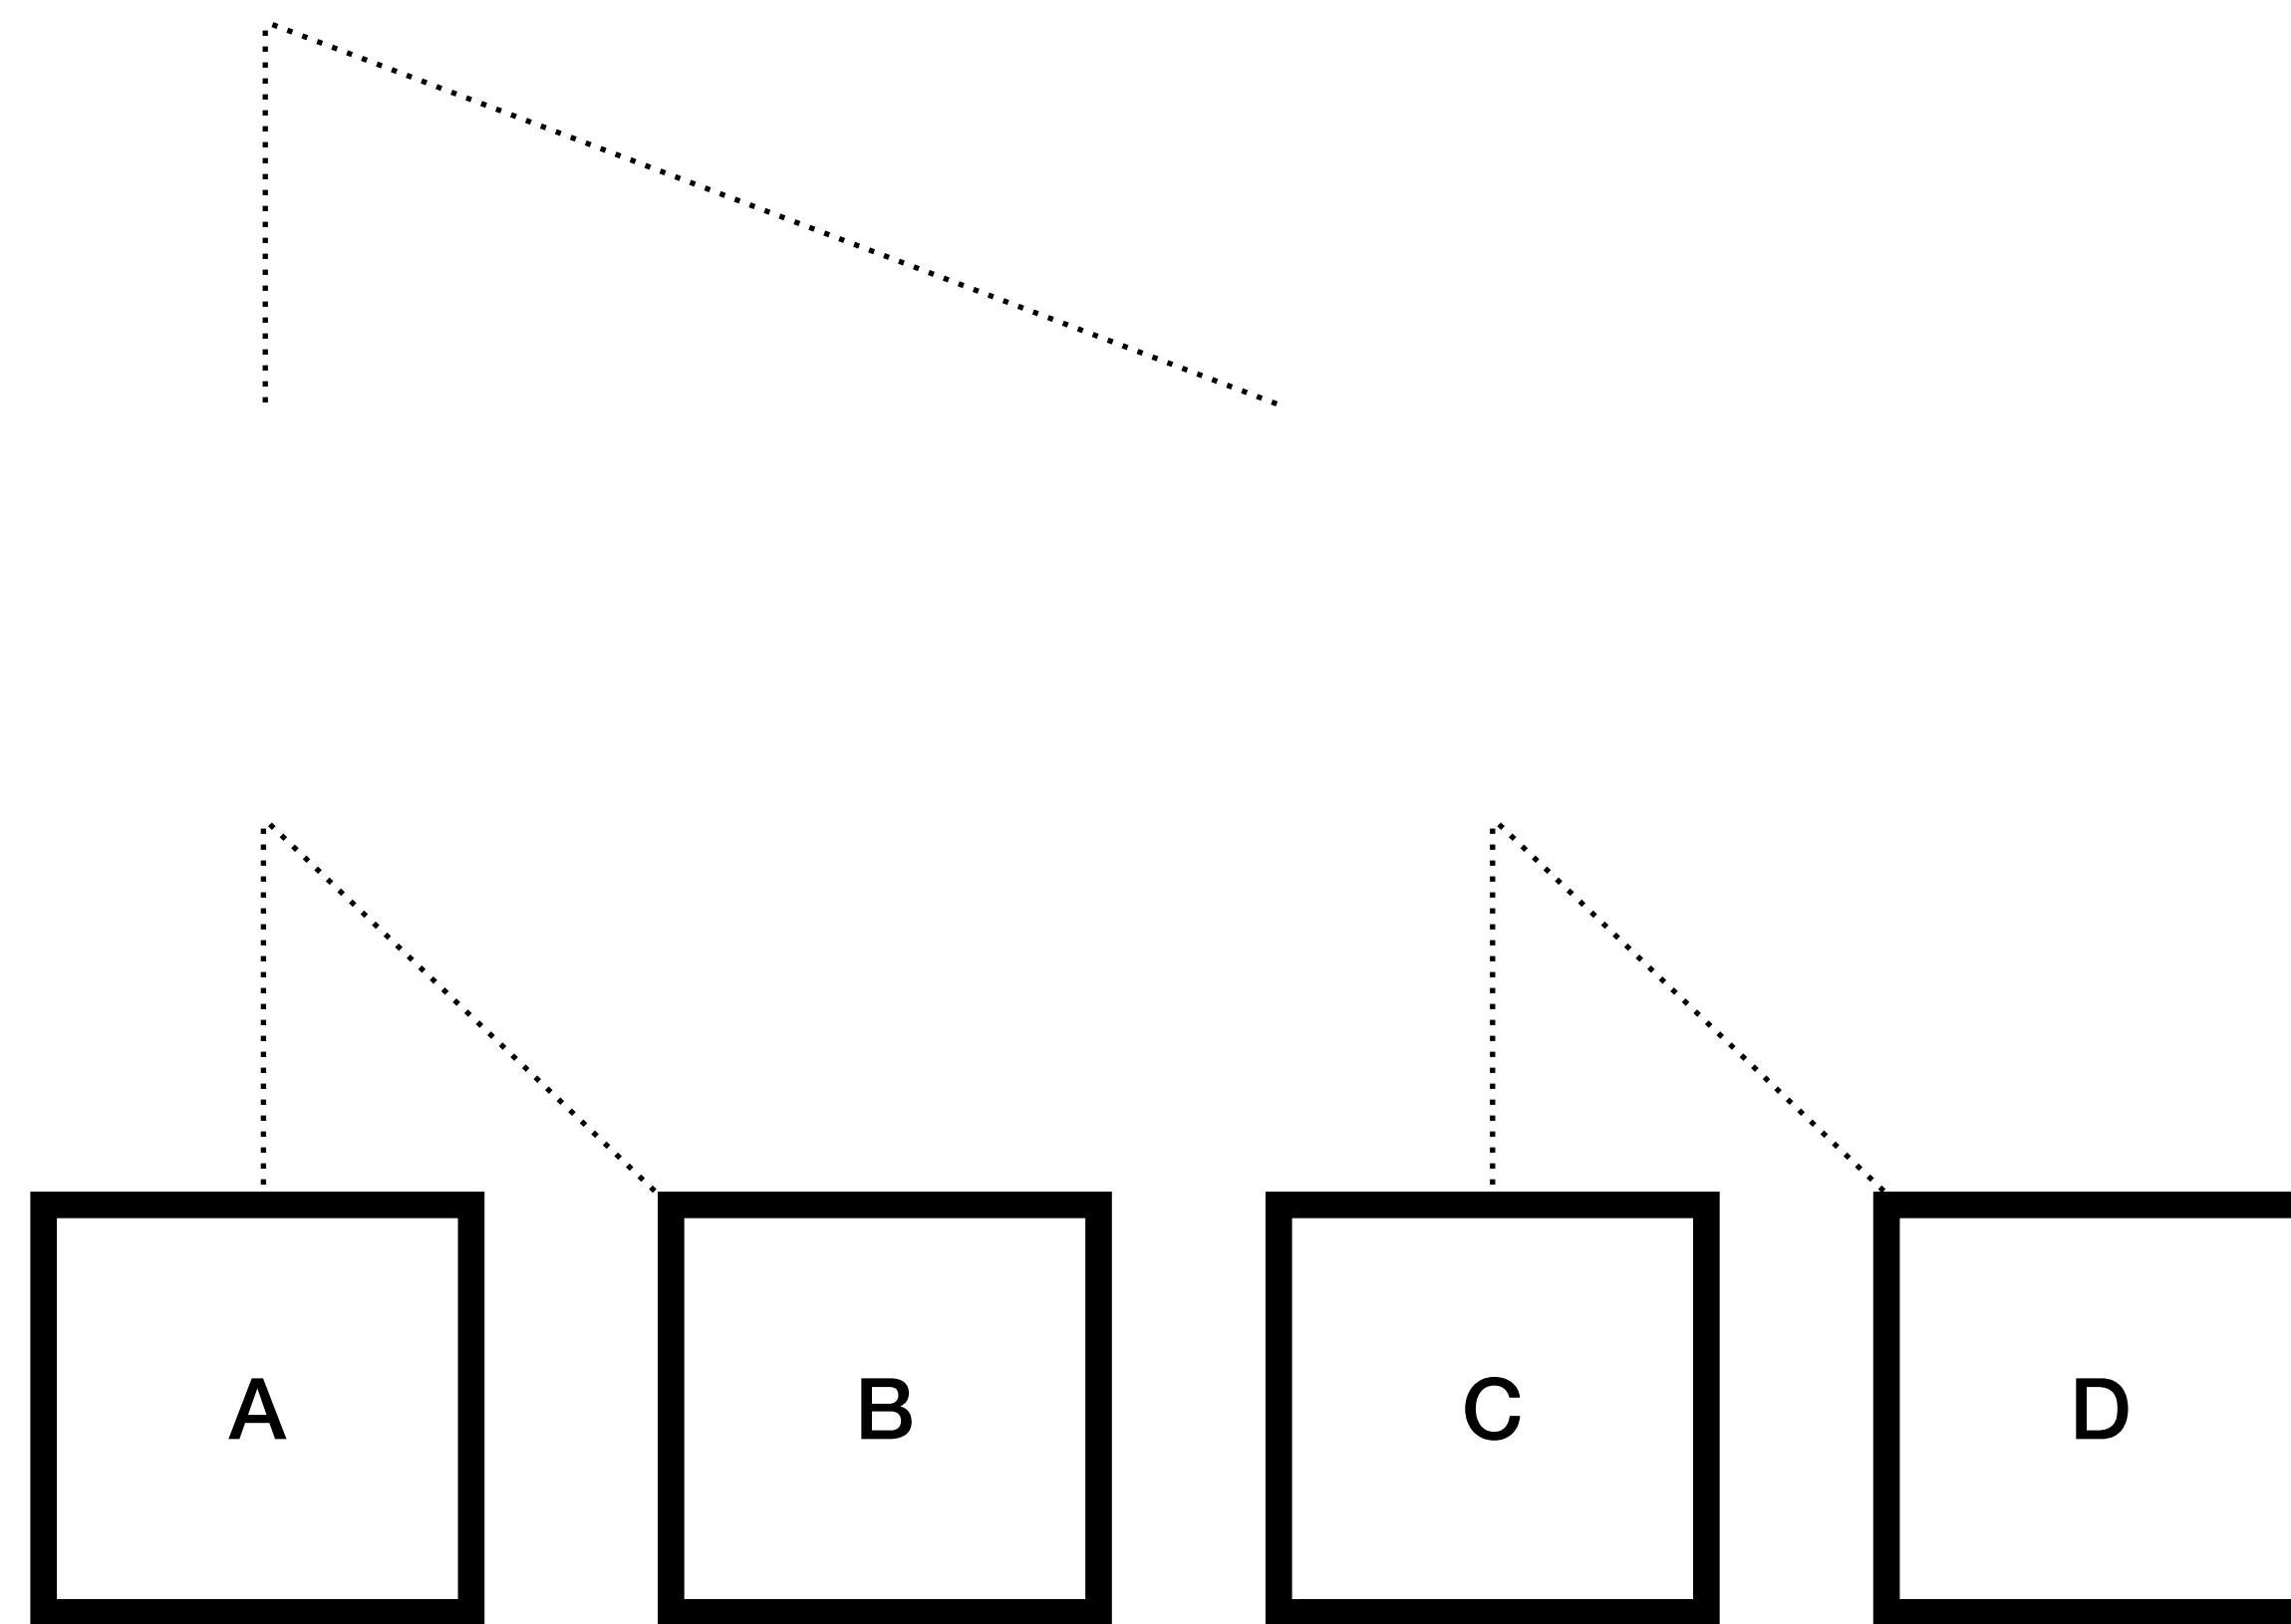
Consensus critical stuff

- 1. Add a UTXO**
- 2. Delete a UTXO**
- 3. Verify the existence**

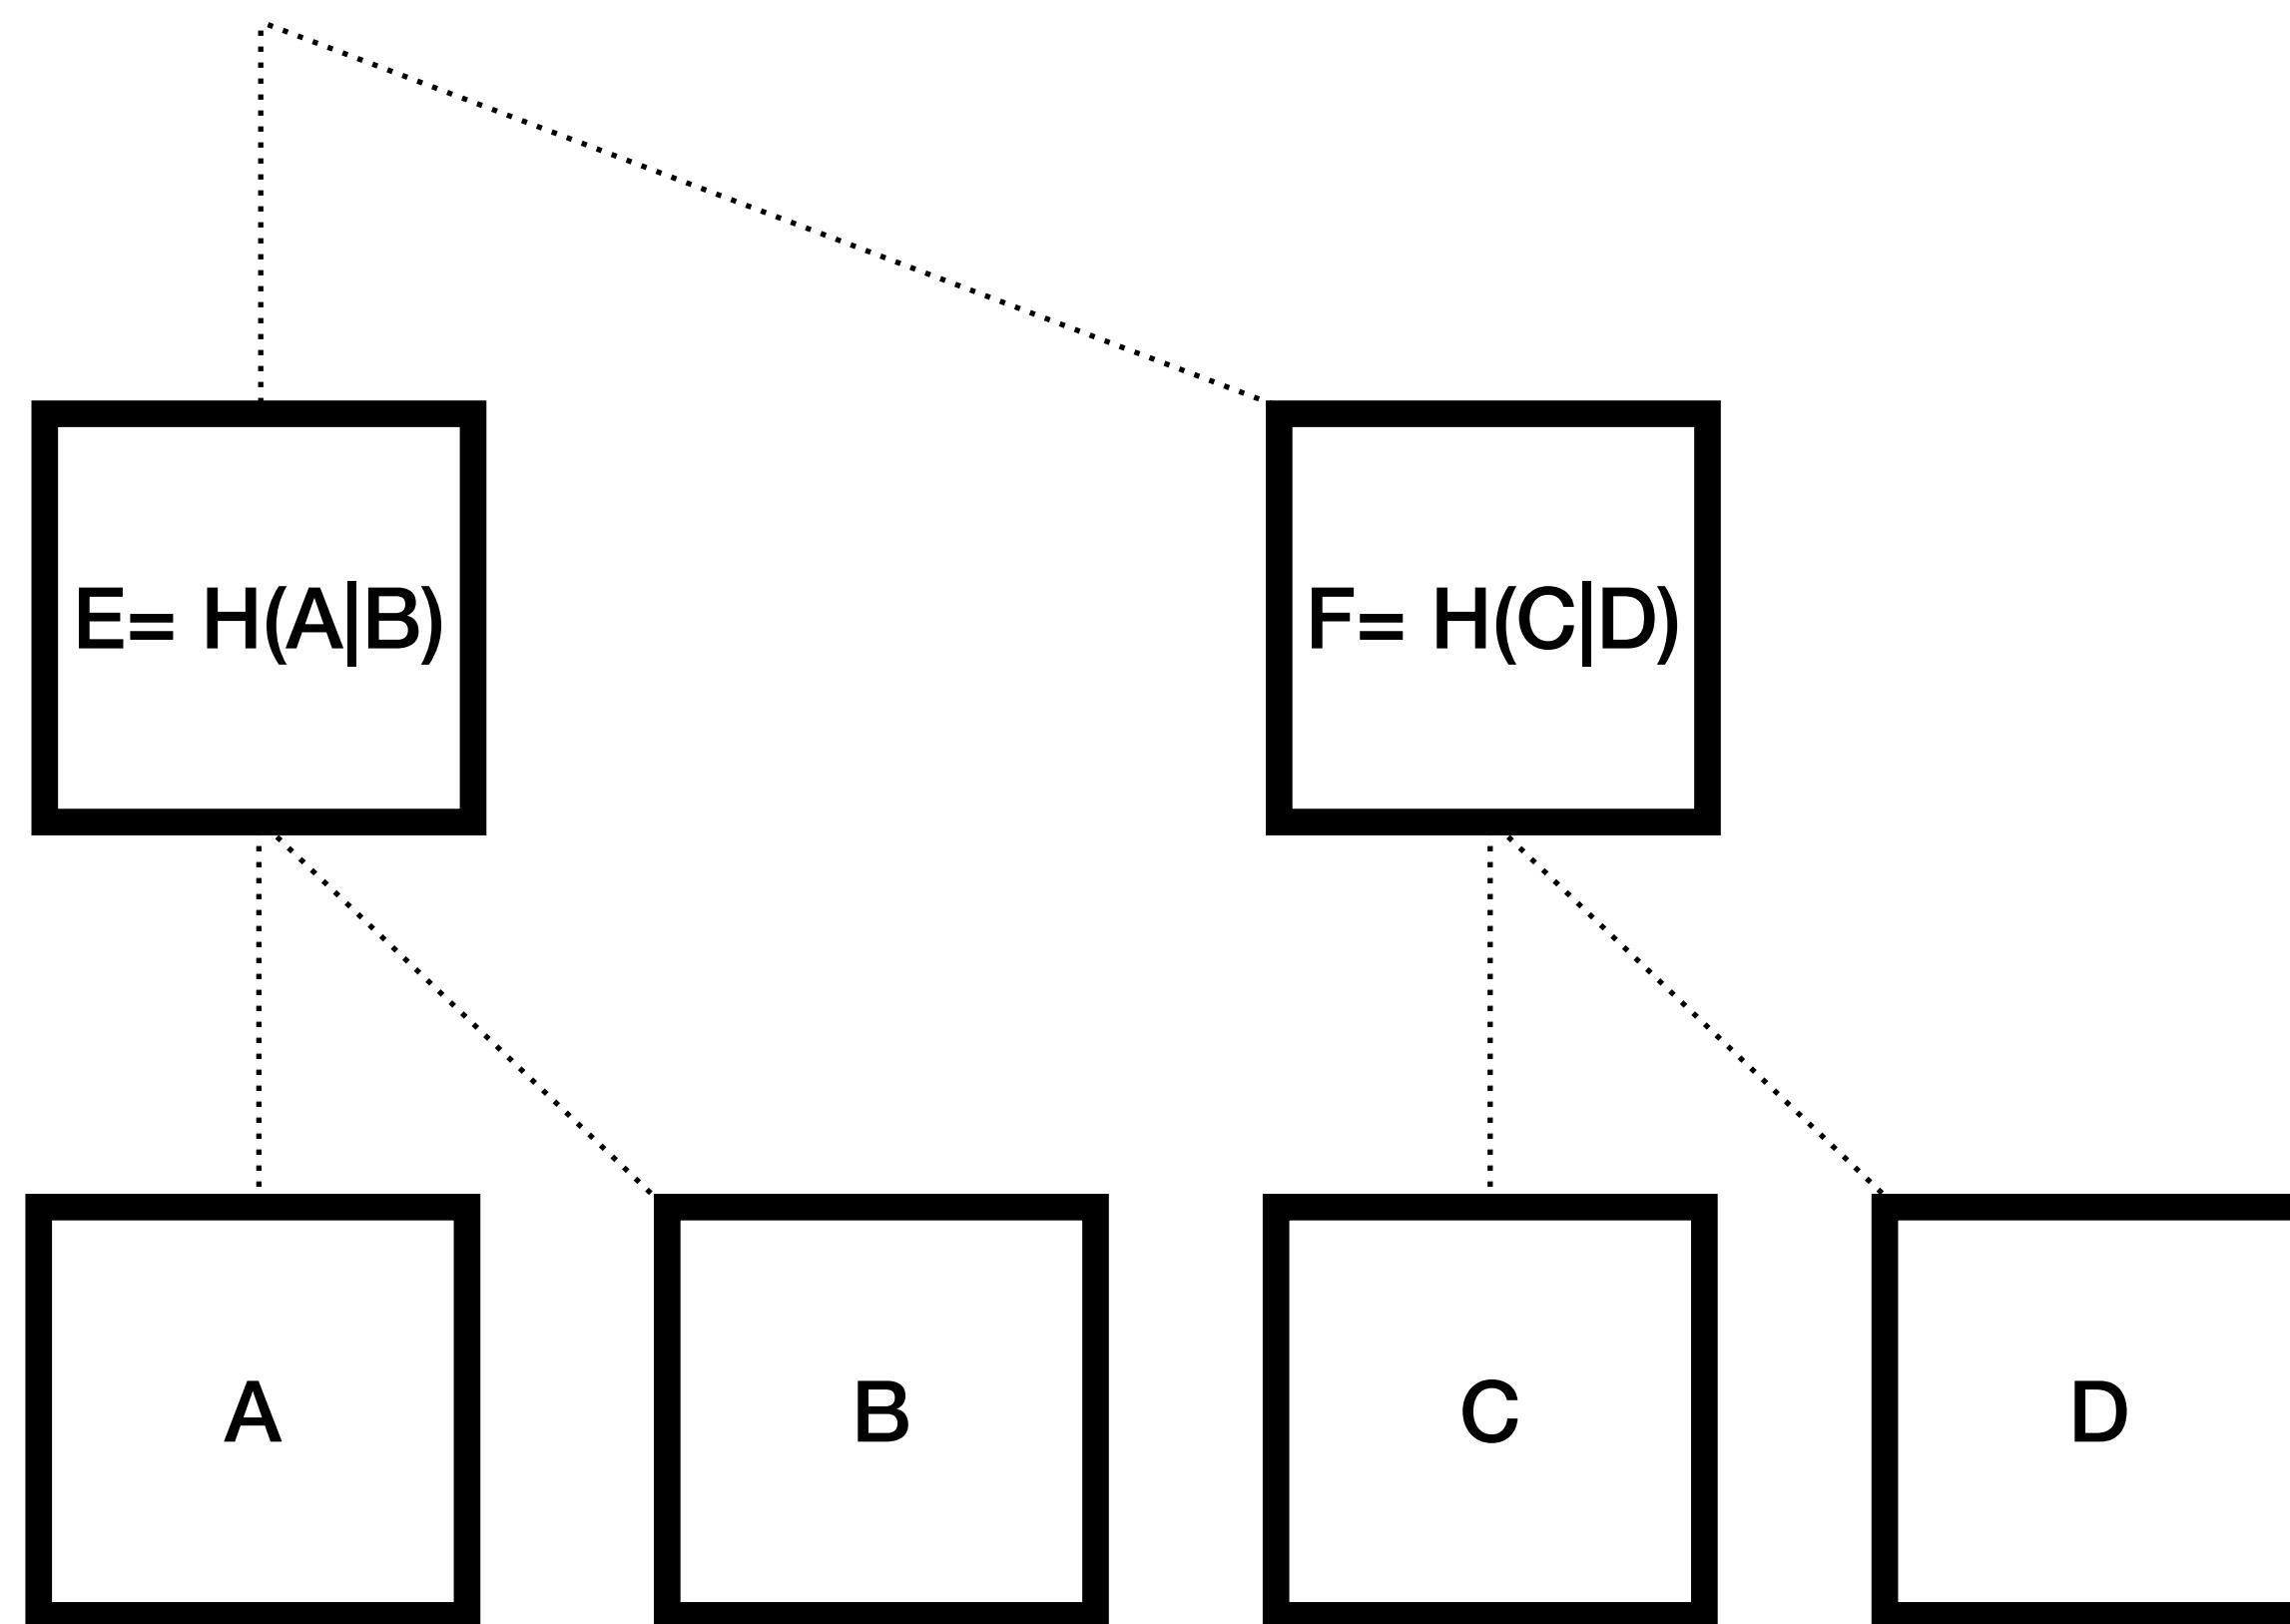
Quick review of merkle tree creation



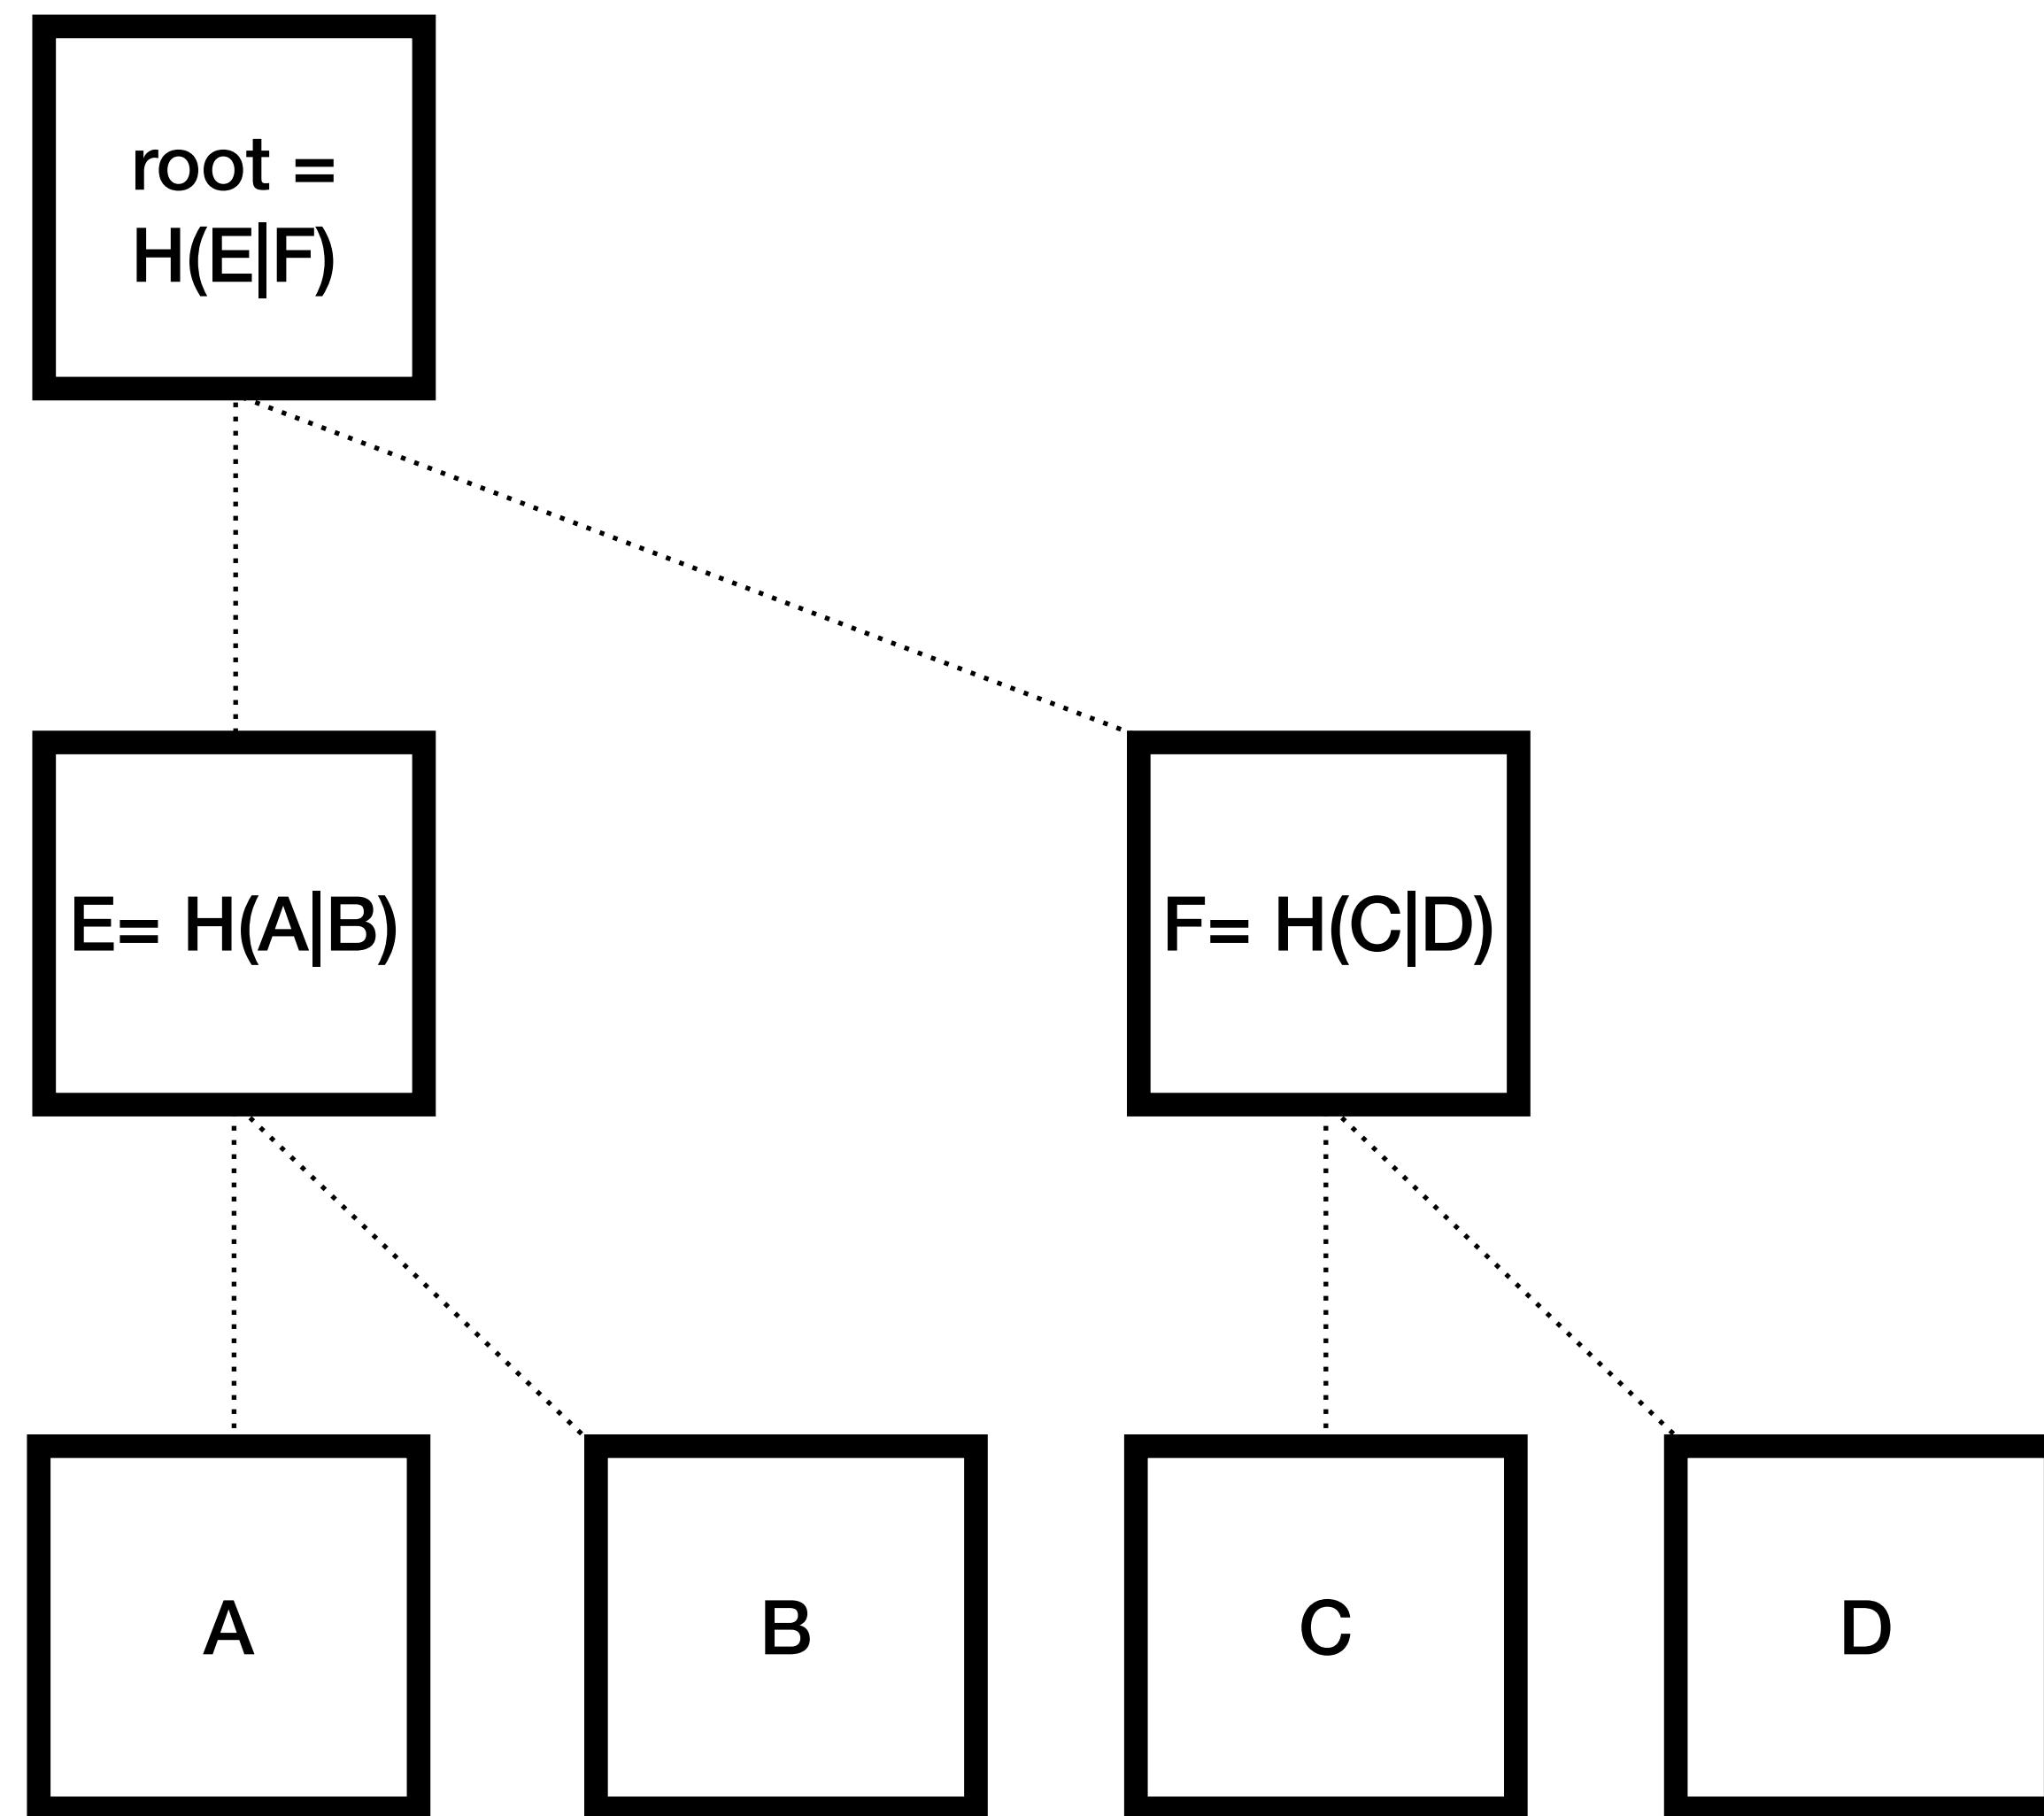
Start off with the hash of the elements



Hash each pairs

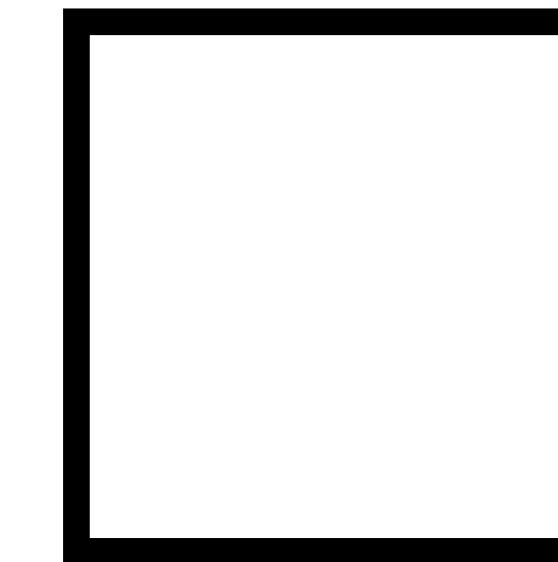


Hash last pair

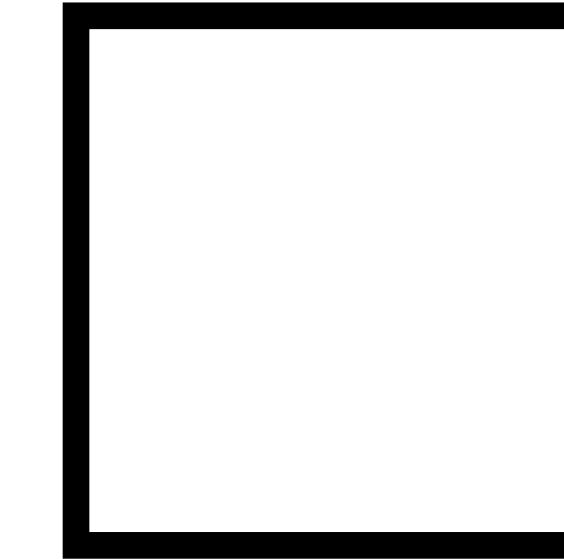
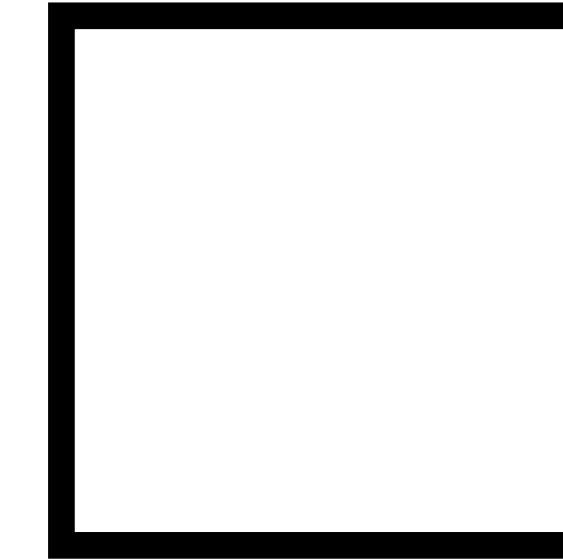


1 UTXO

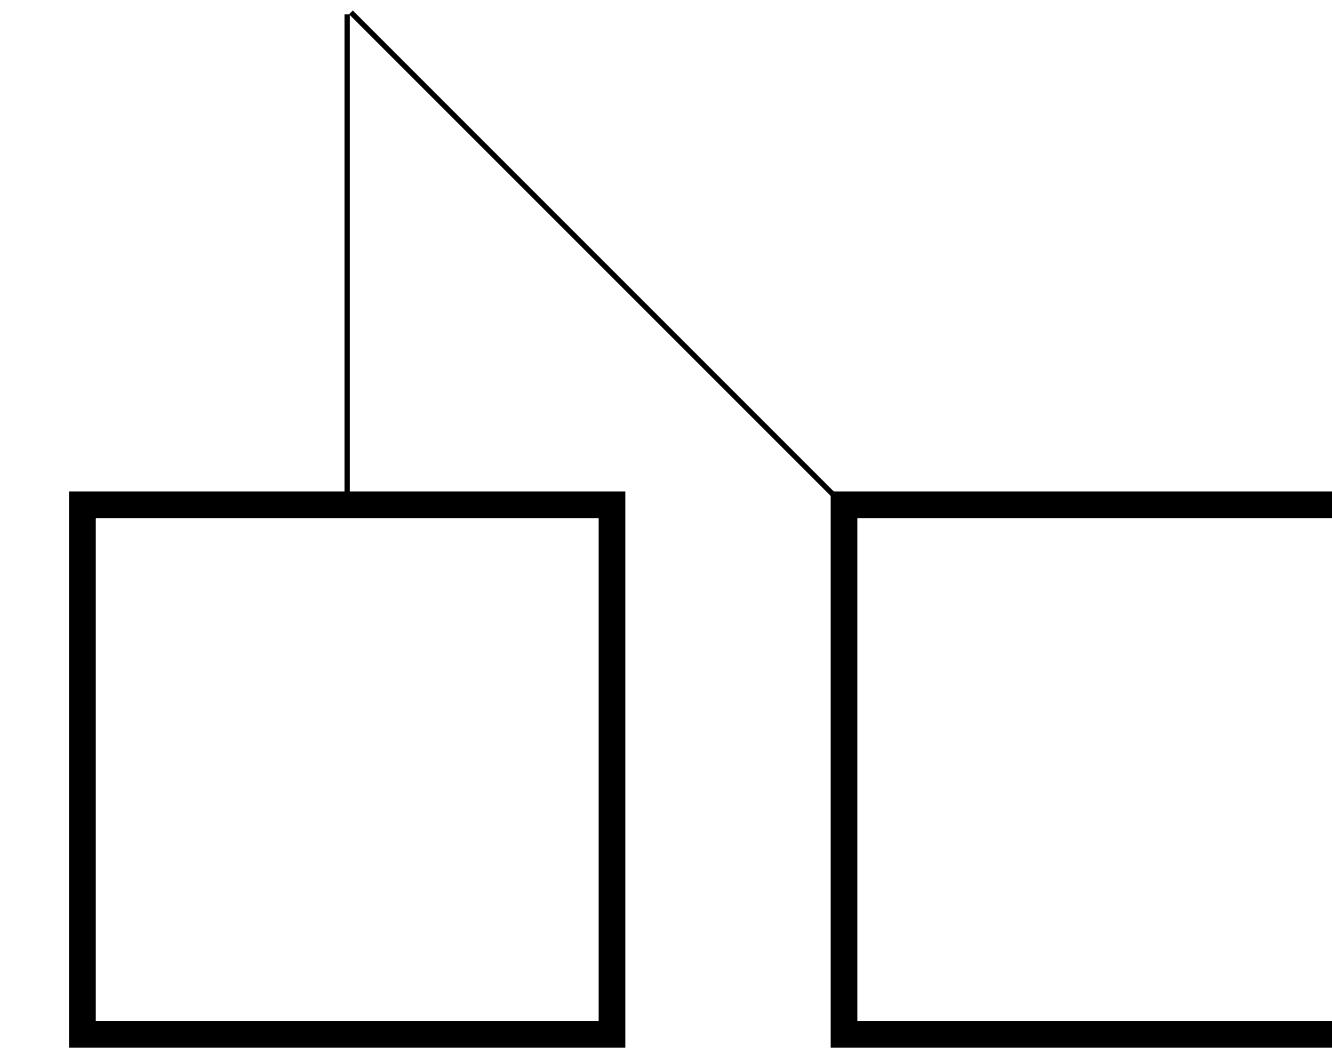
Single root



2 UTXOs

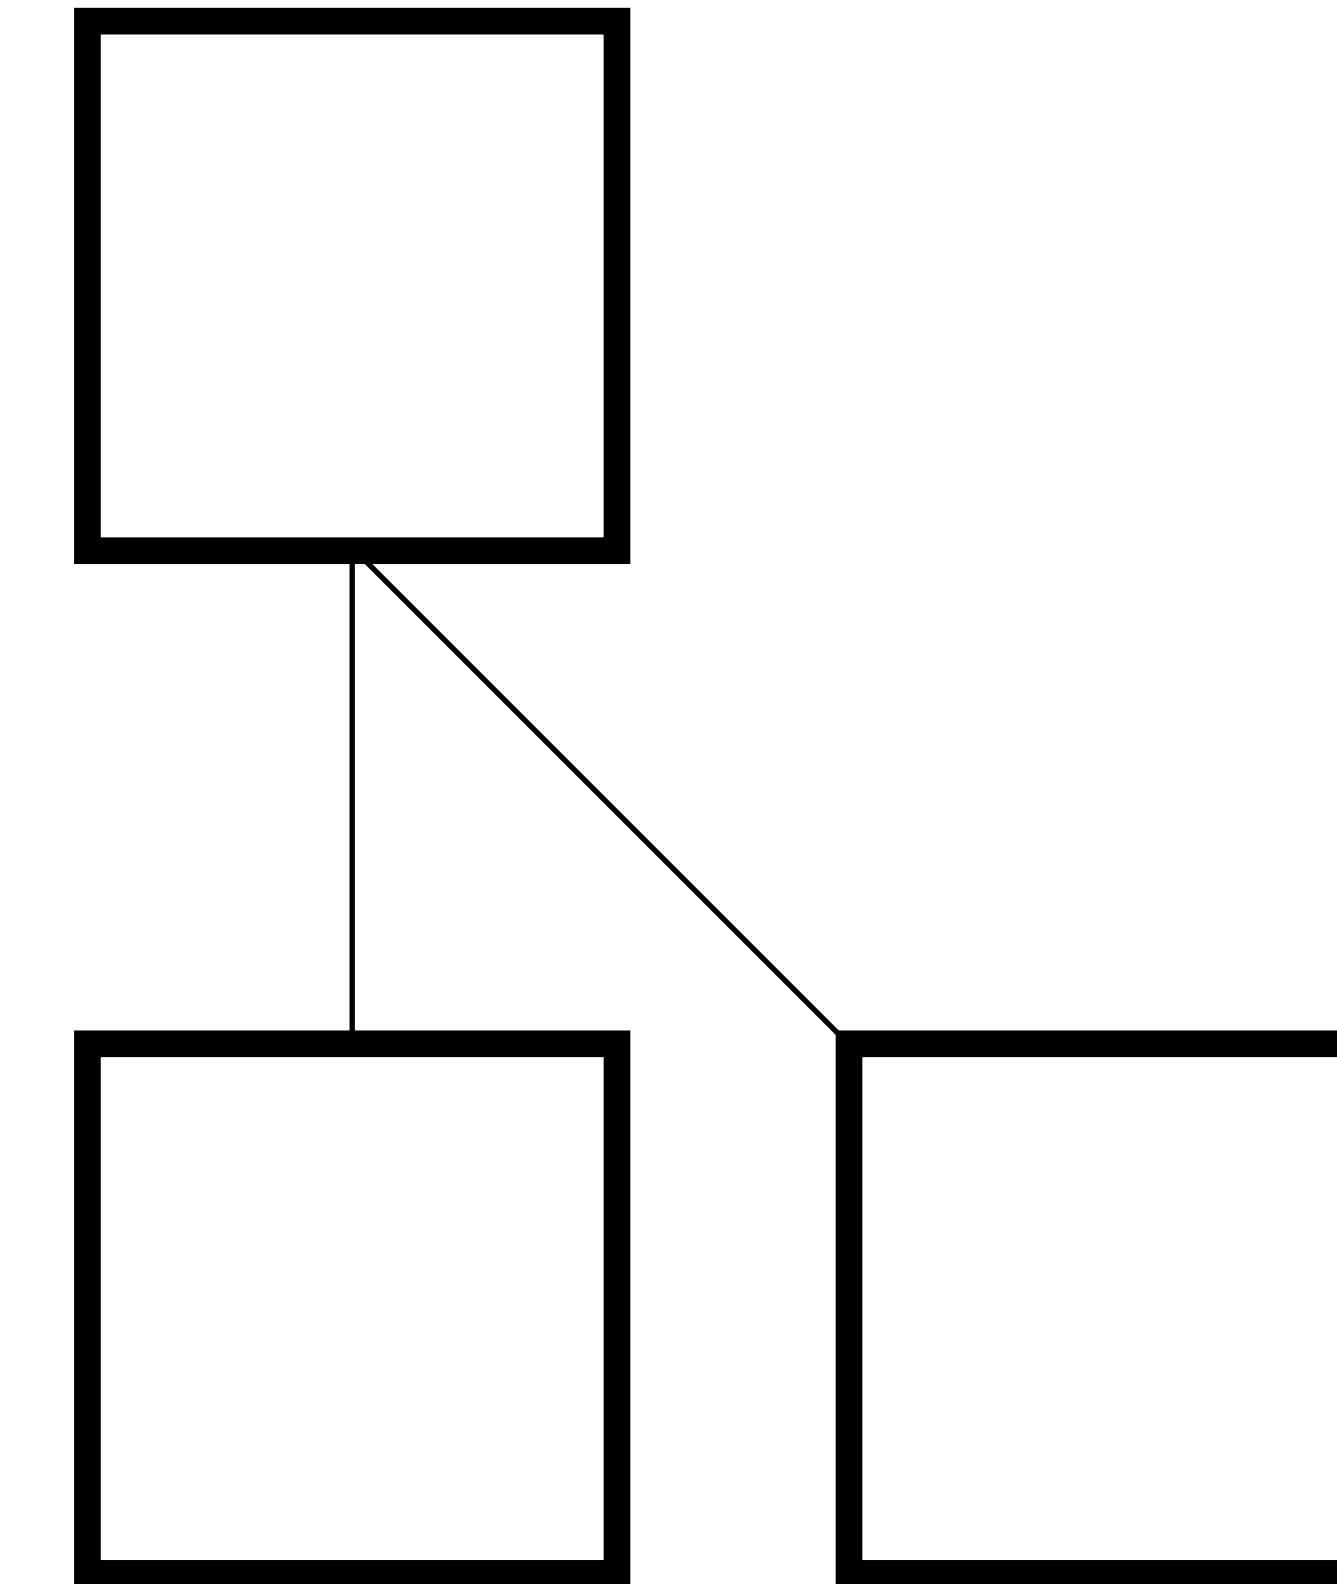


2 UTXOs Concatenate



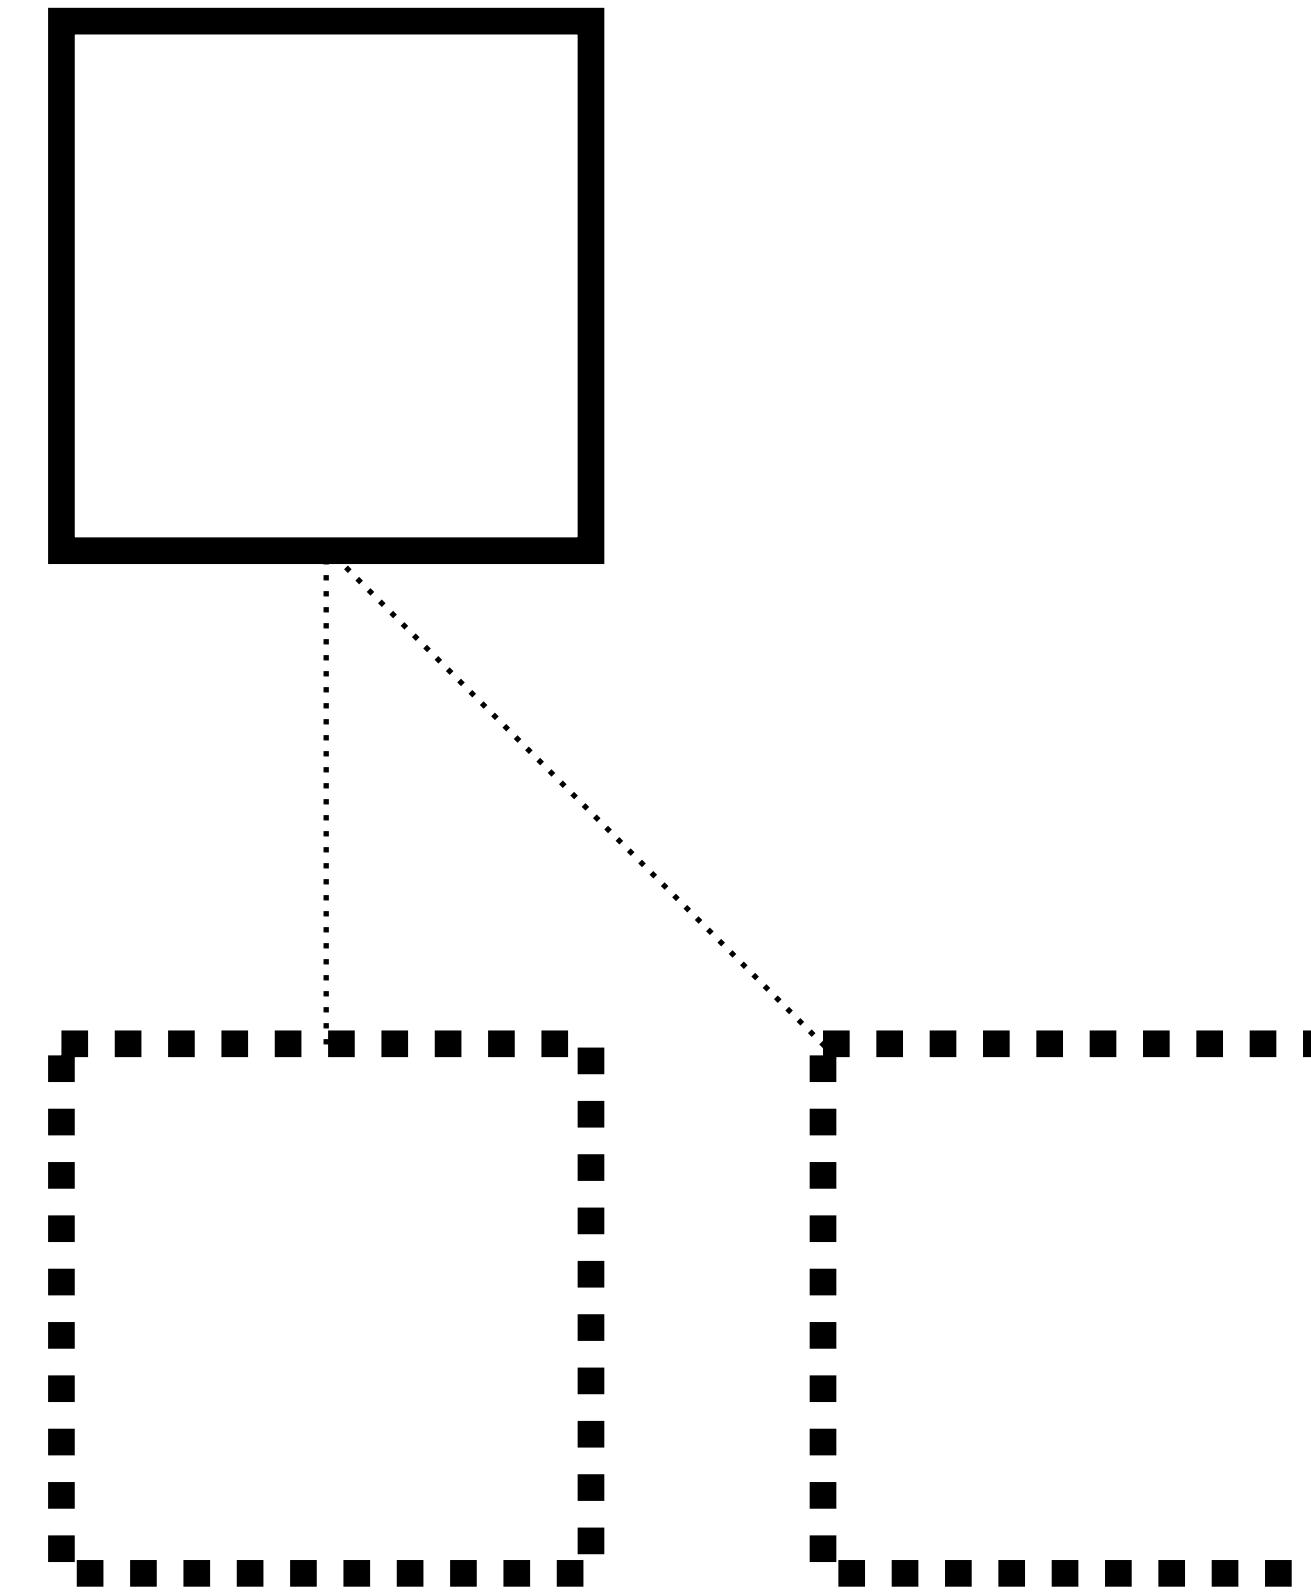
2 UTXOs

Hash to create a new root



2 UTXOs

Can now delete the leaves



2 UTXOs



<https://gist.github.com/kcalvinalvin/a790d524832e1b7f96a70c642315fffc>

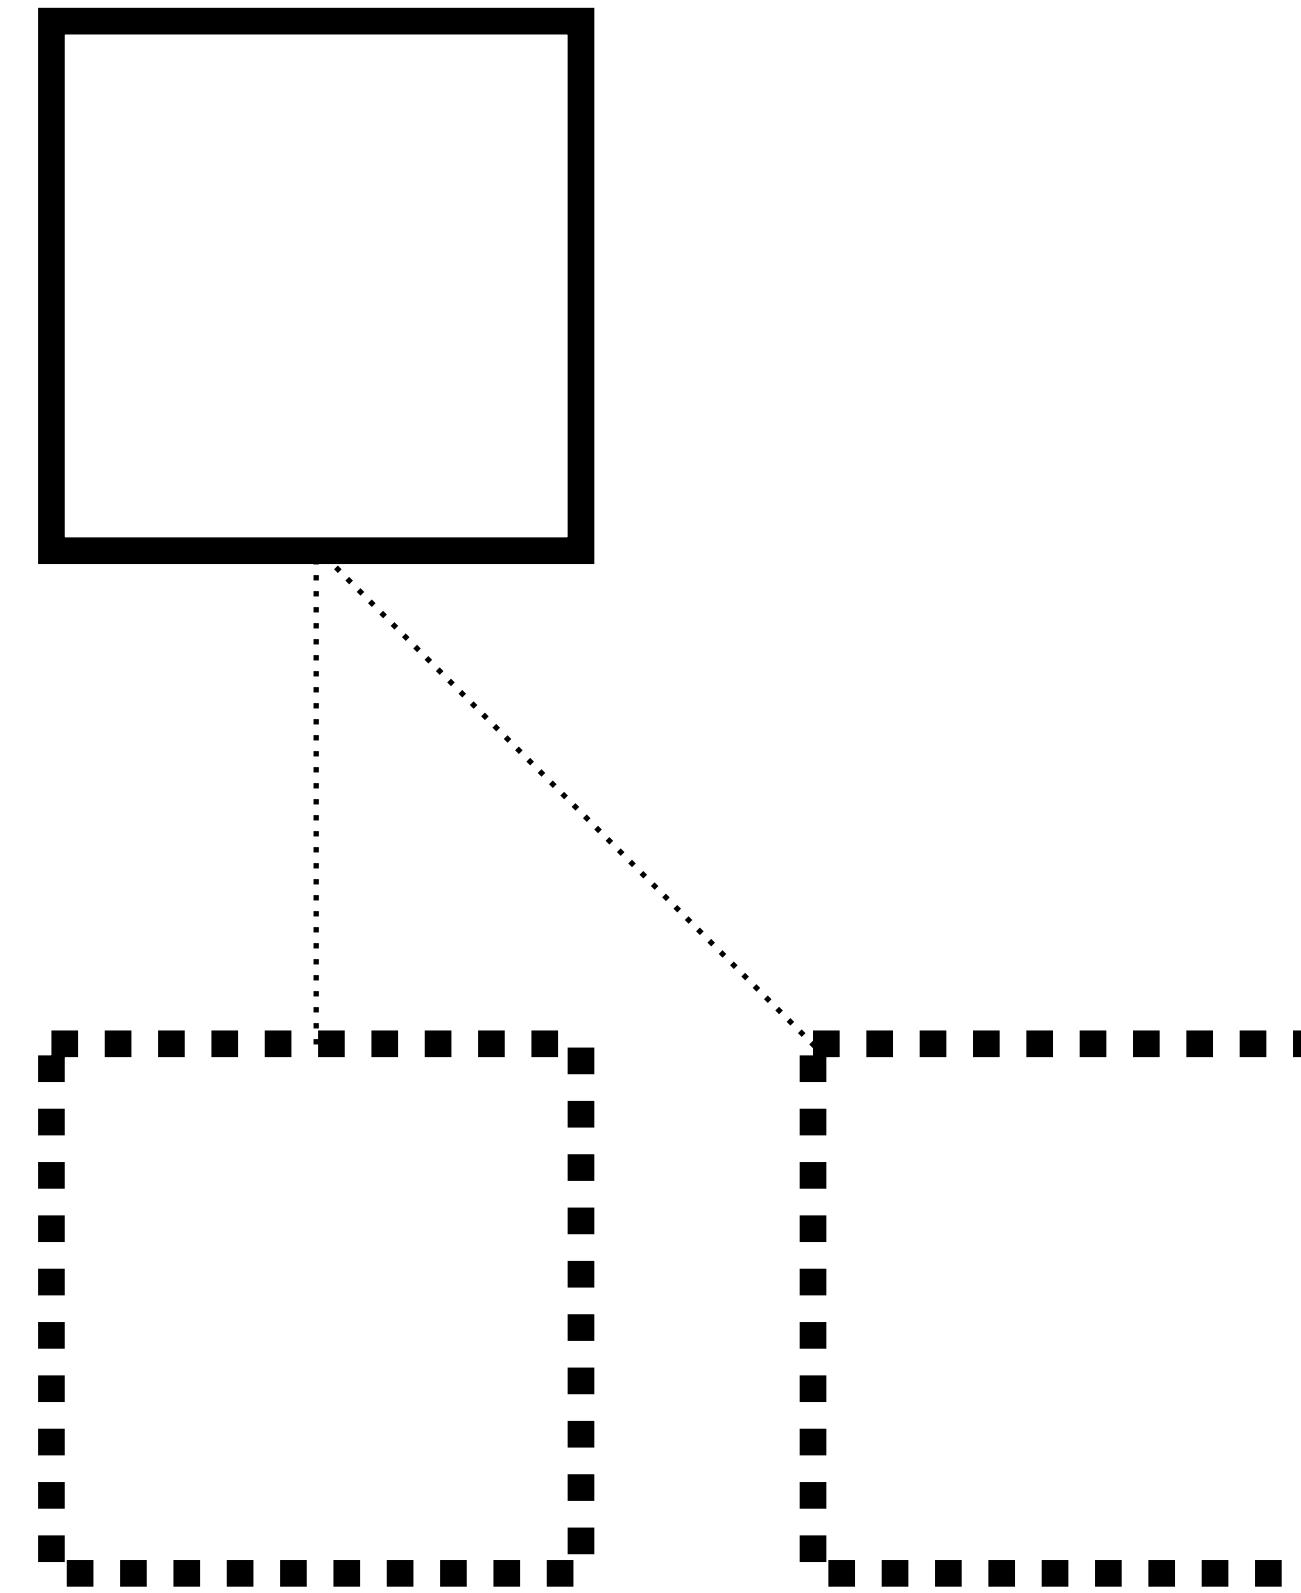
3 UTXOs

Add 1 more



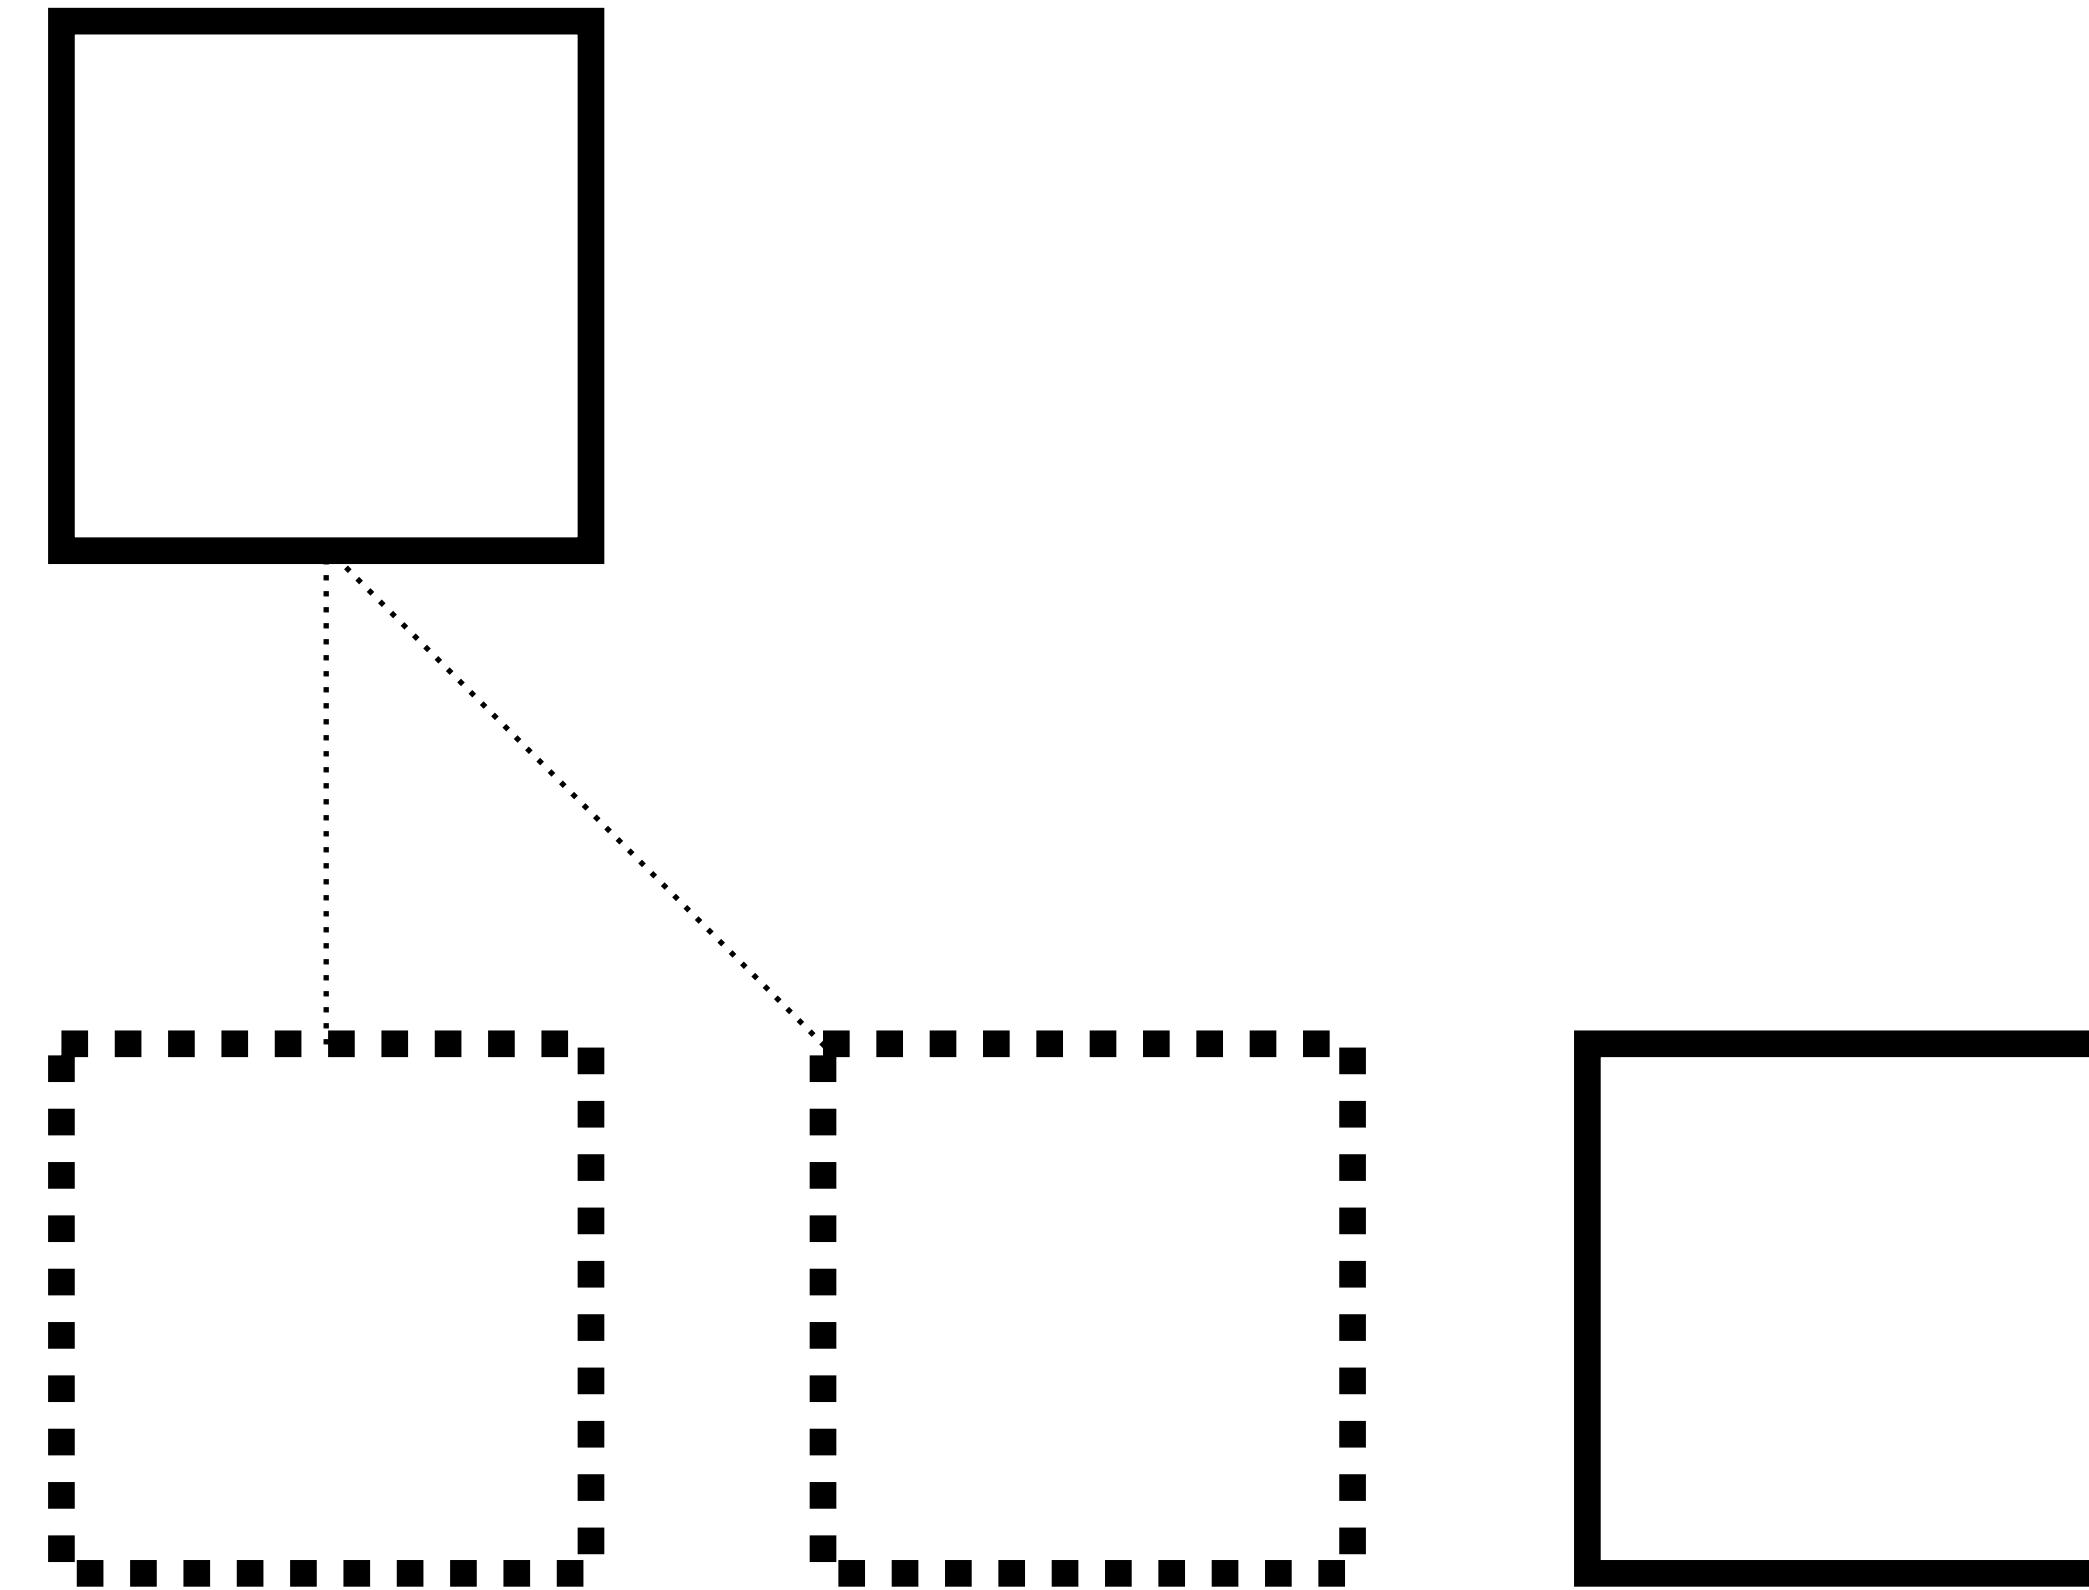
3 UTXOs

Add 1 more



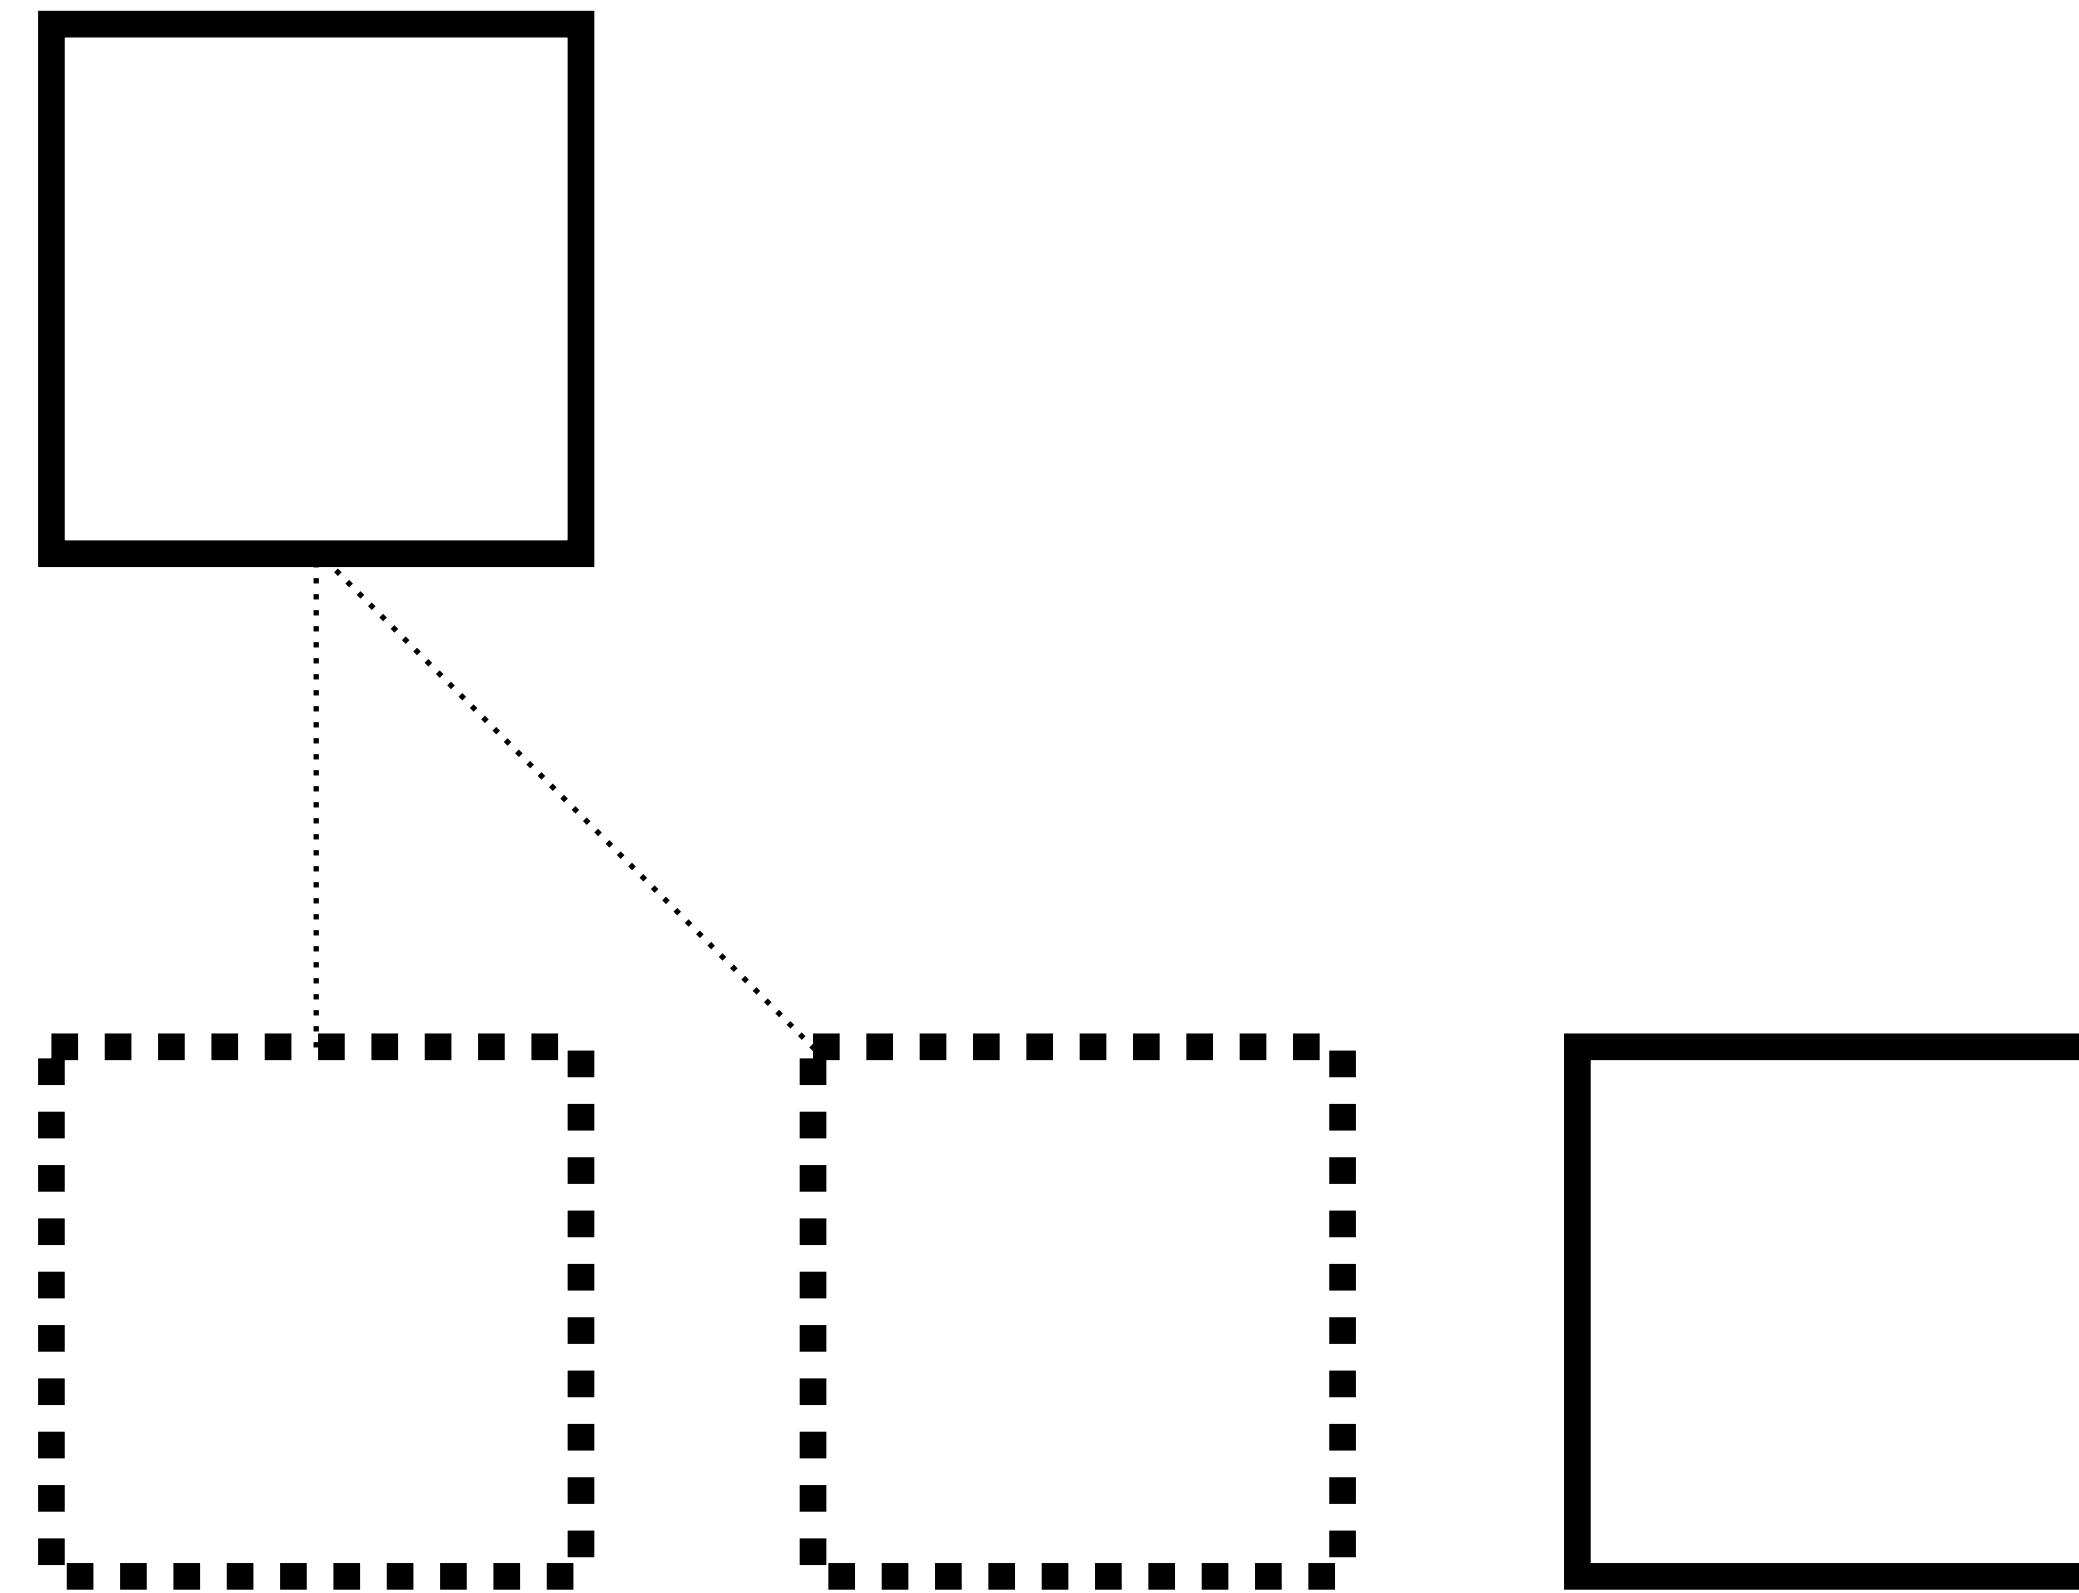
3 UTXOs

Add 1 more



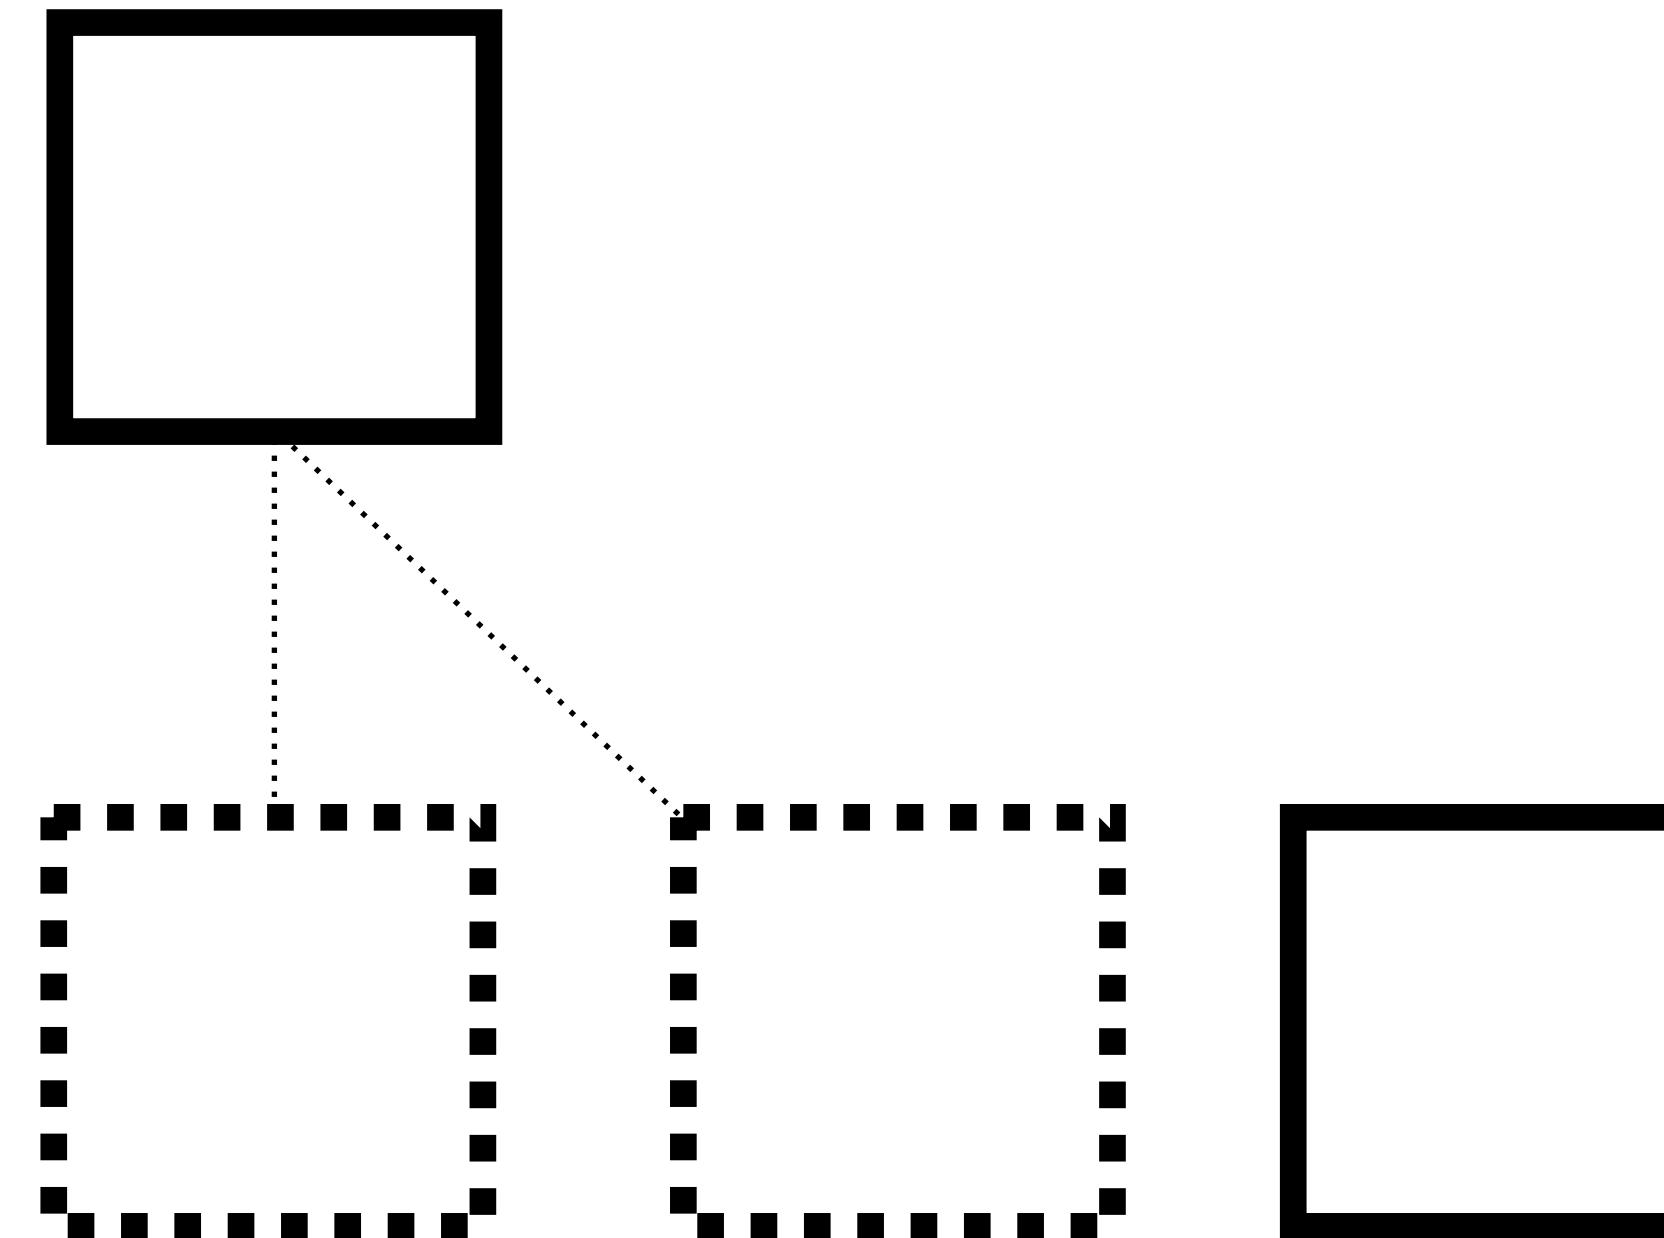
3 UTXOs

Nothing to hash with. Becomes a root



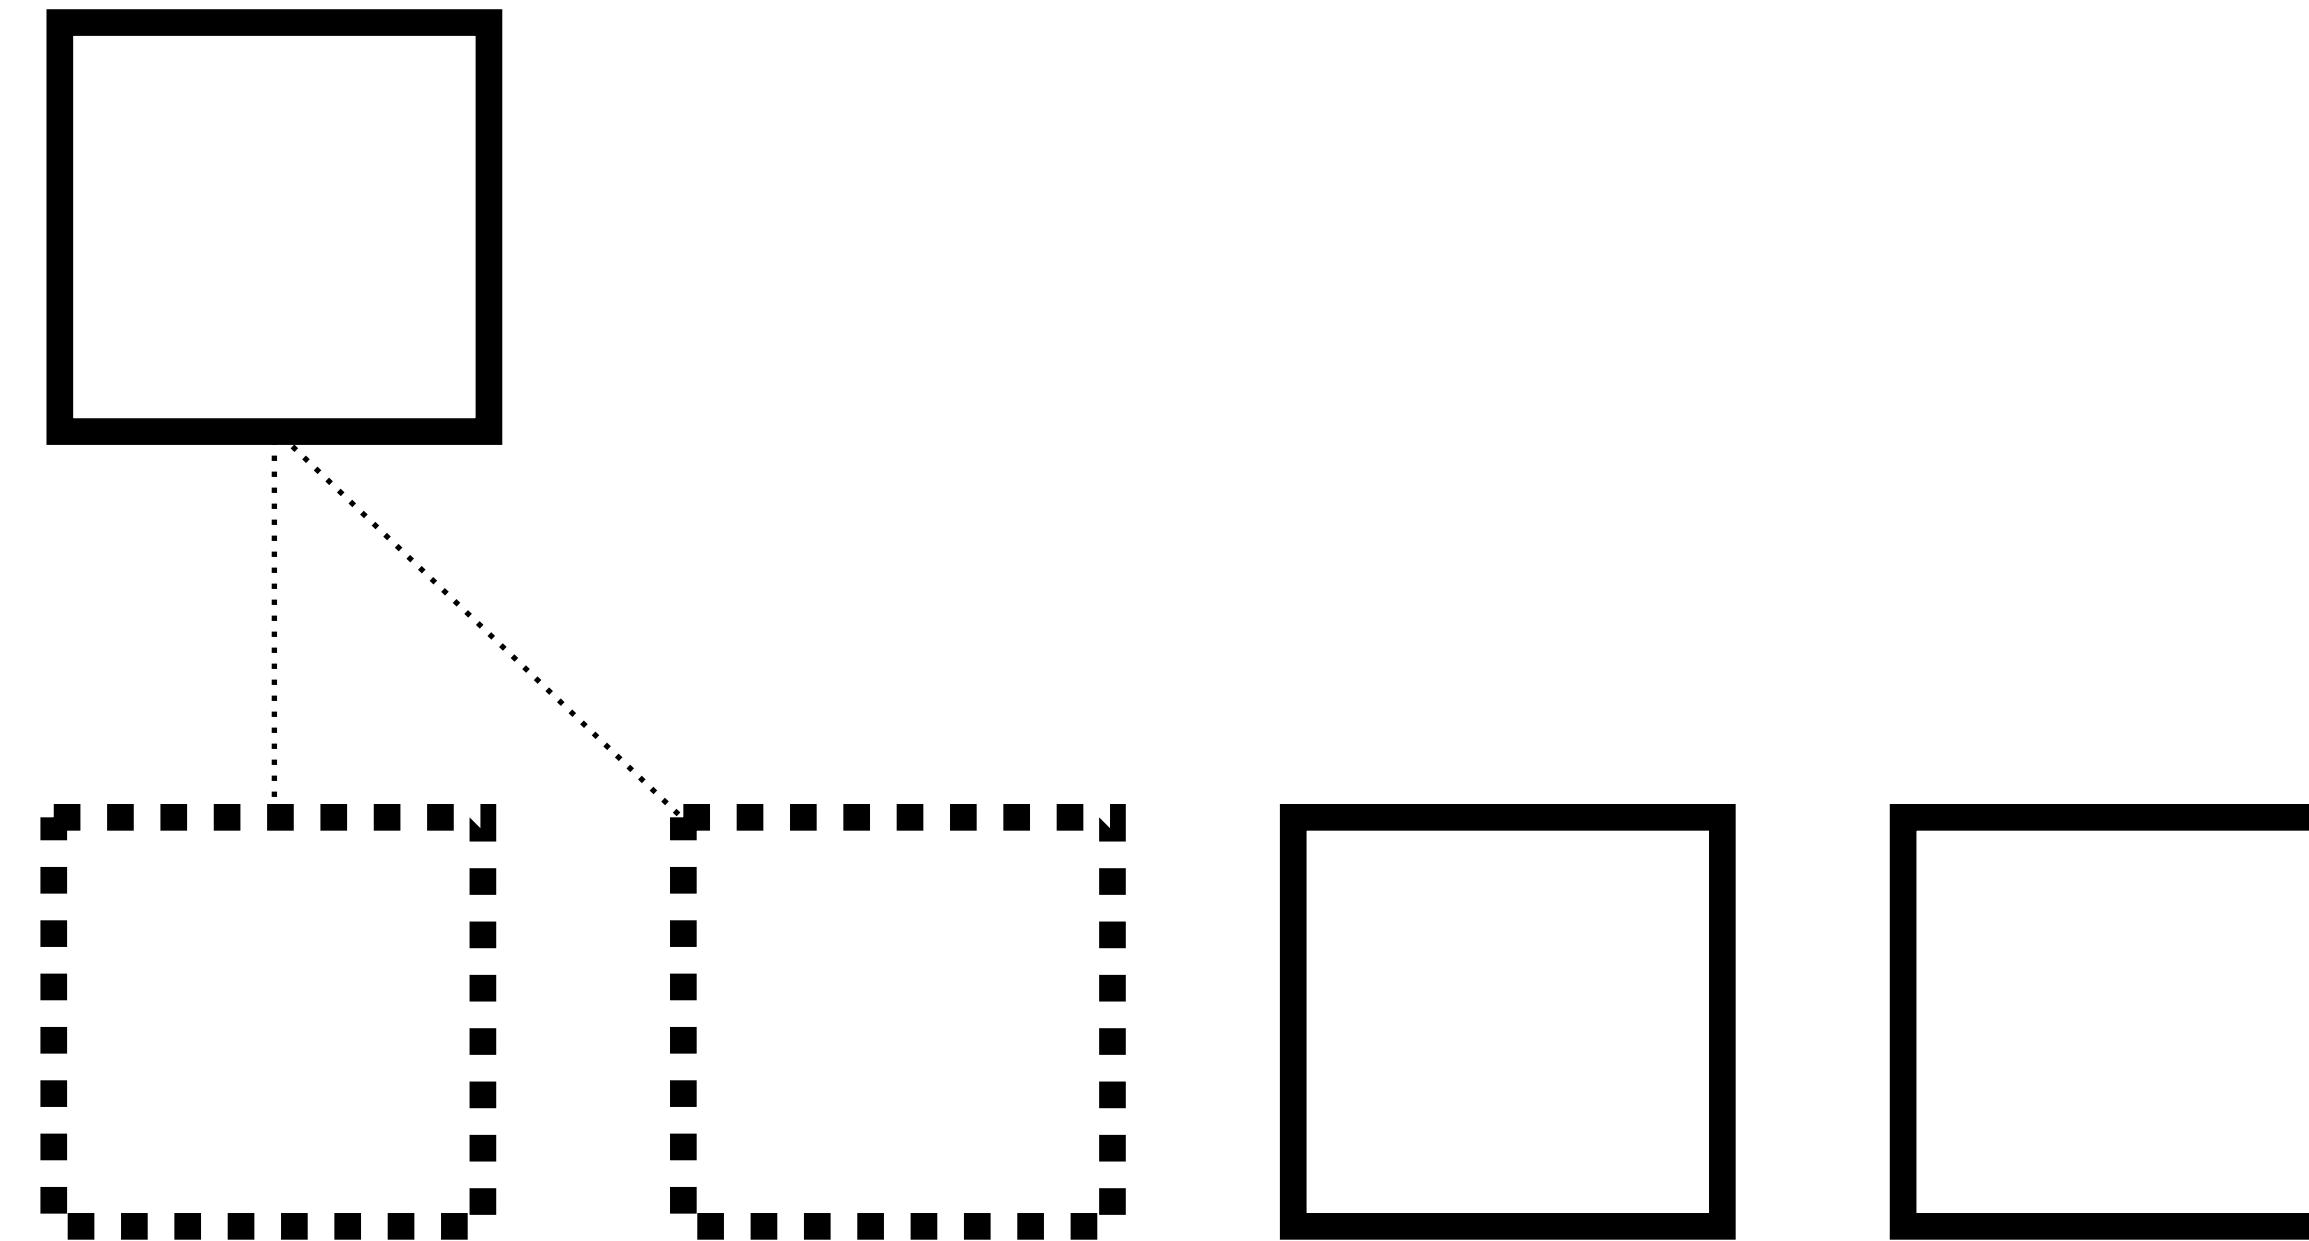
4 UTXOs

Another one

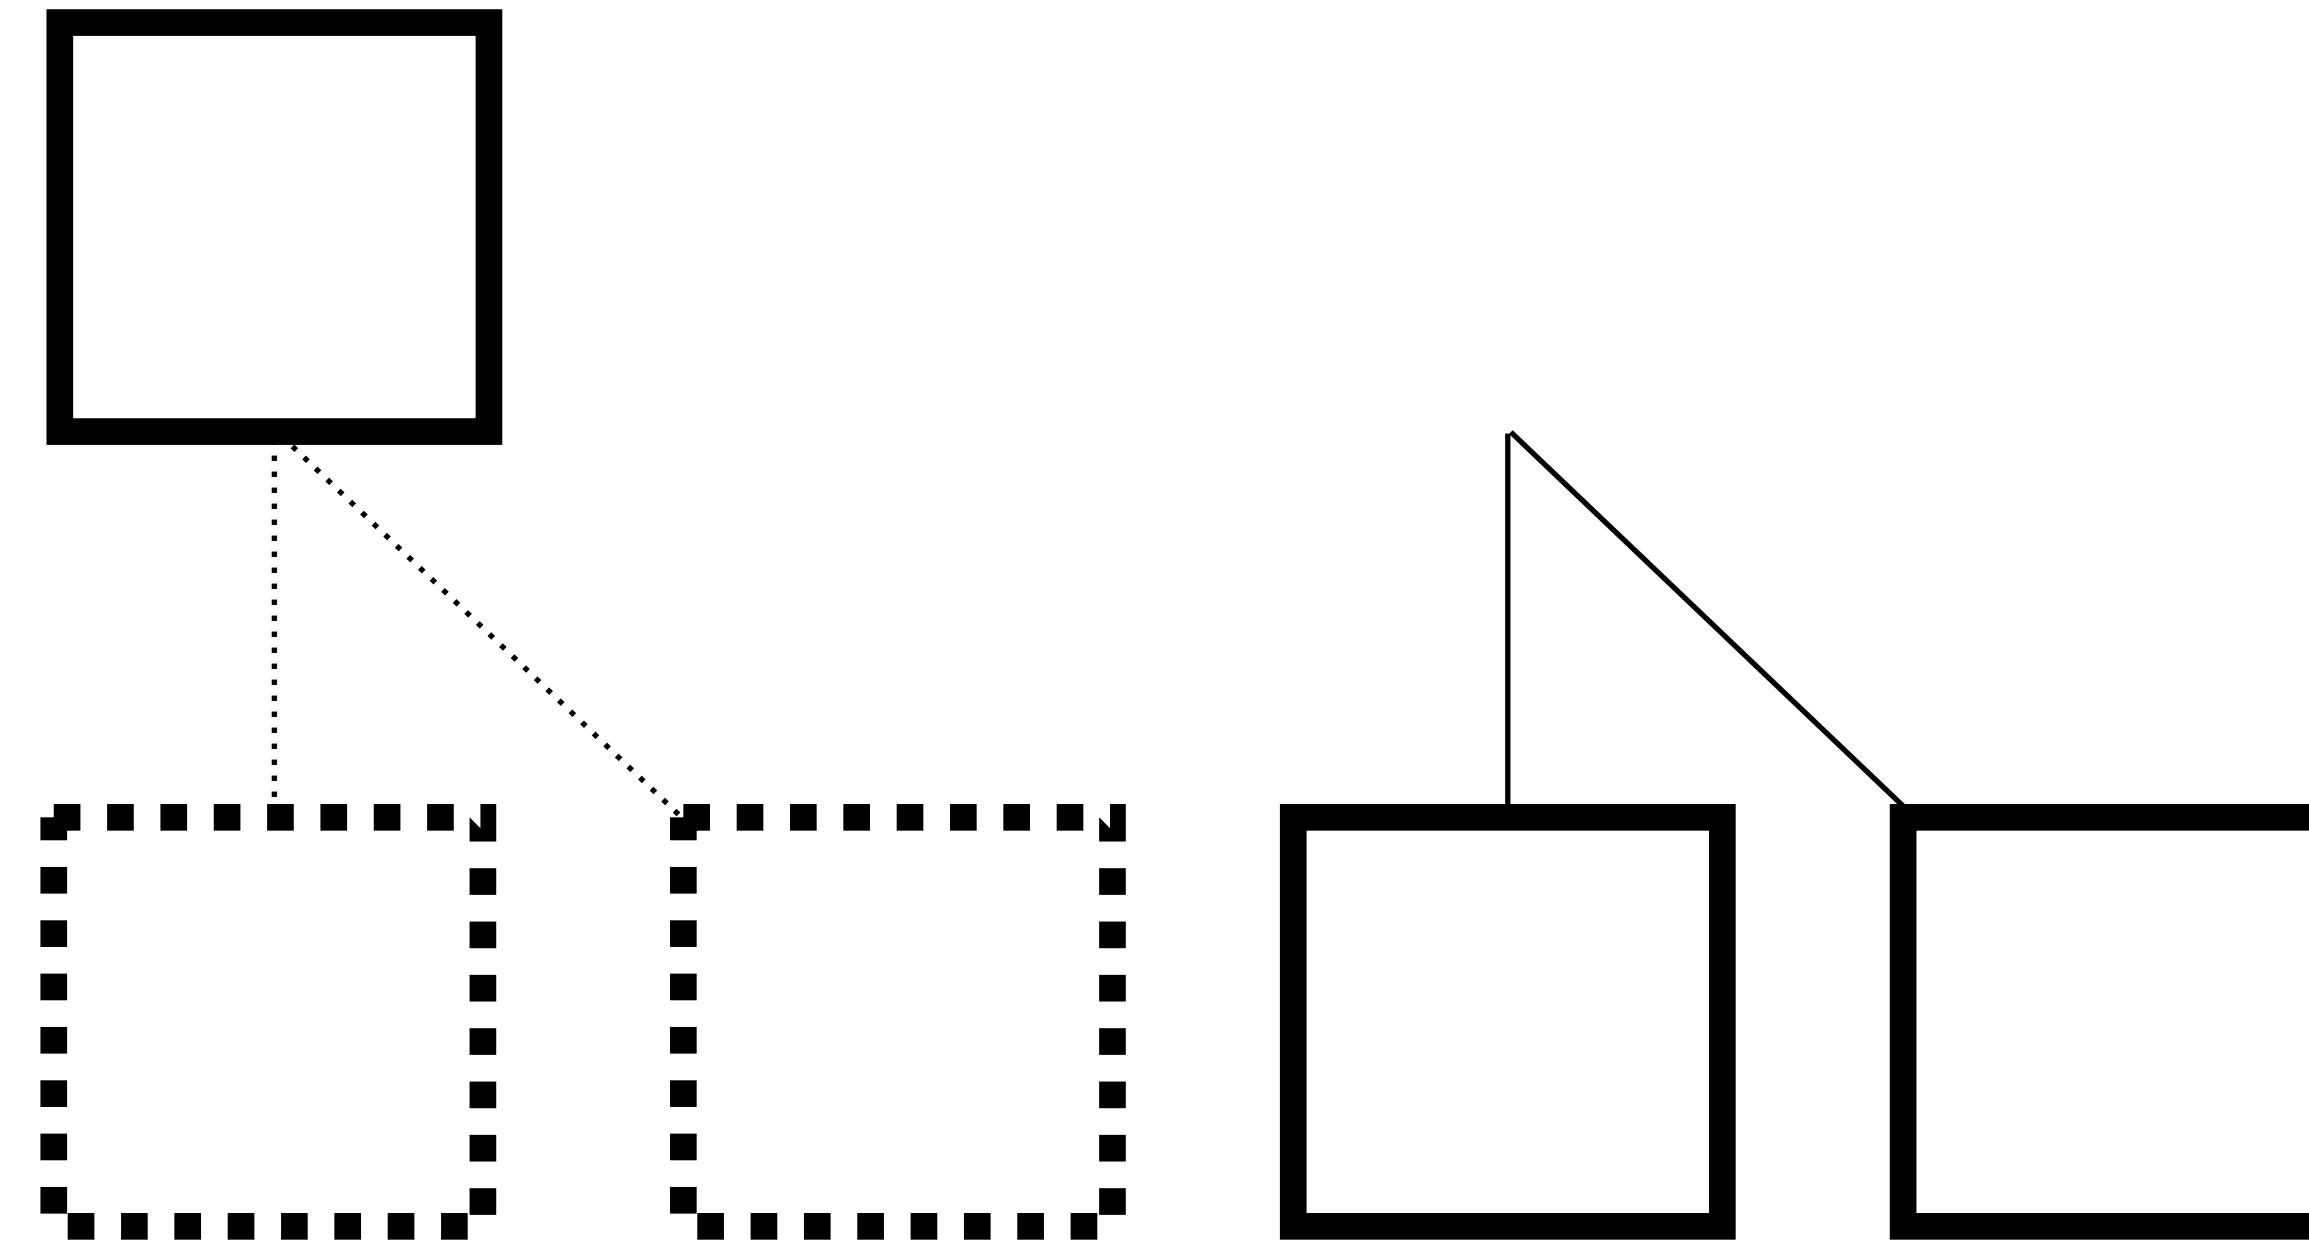


4 UTXOs

Another one



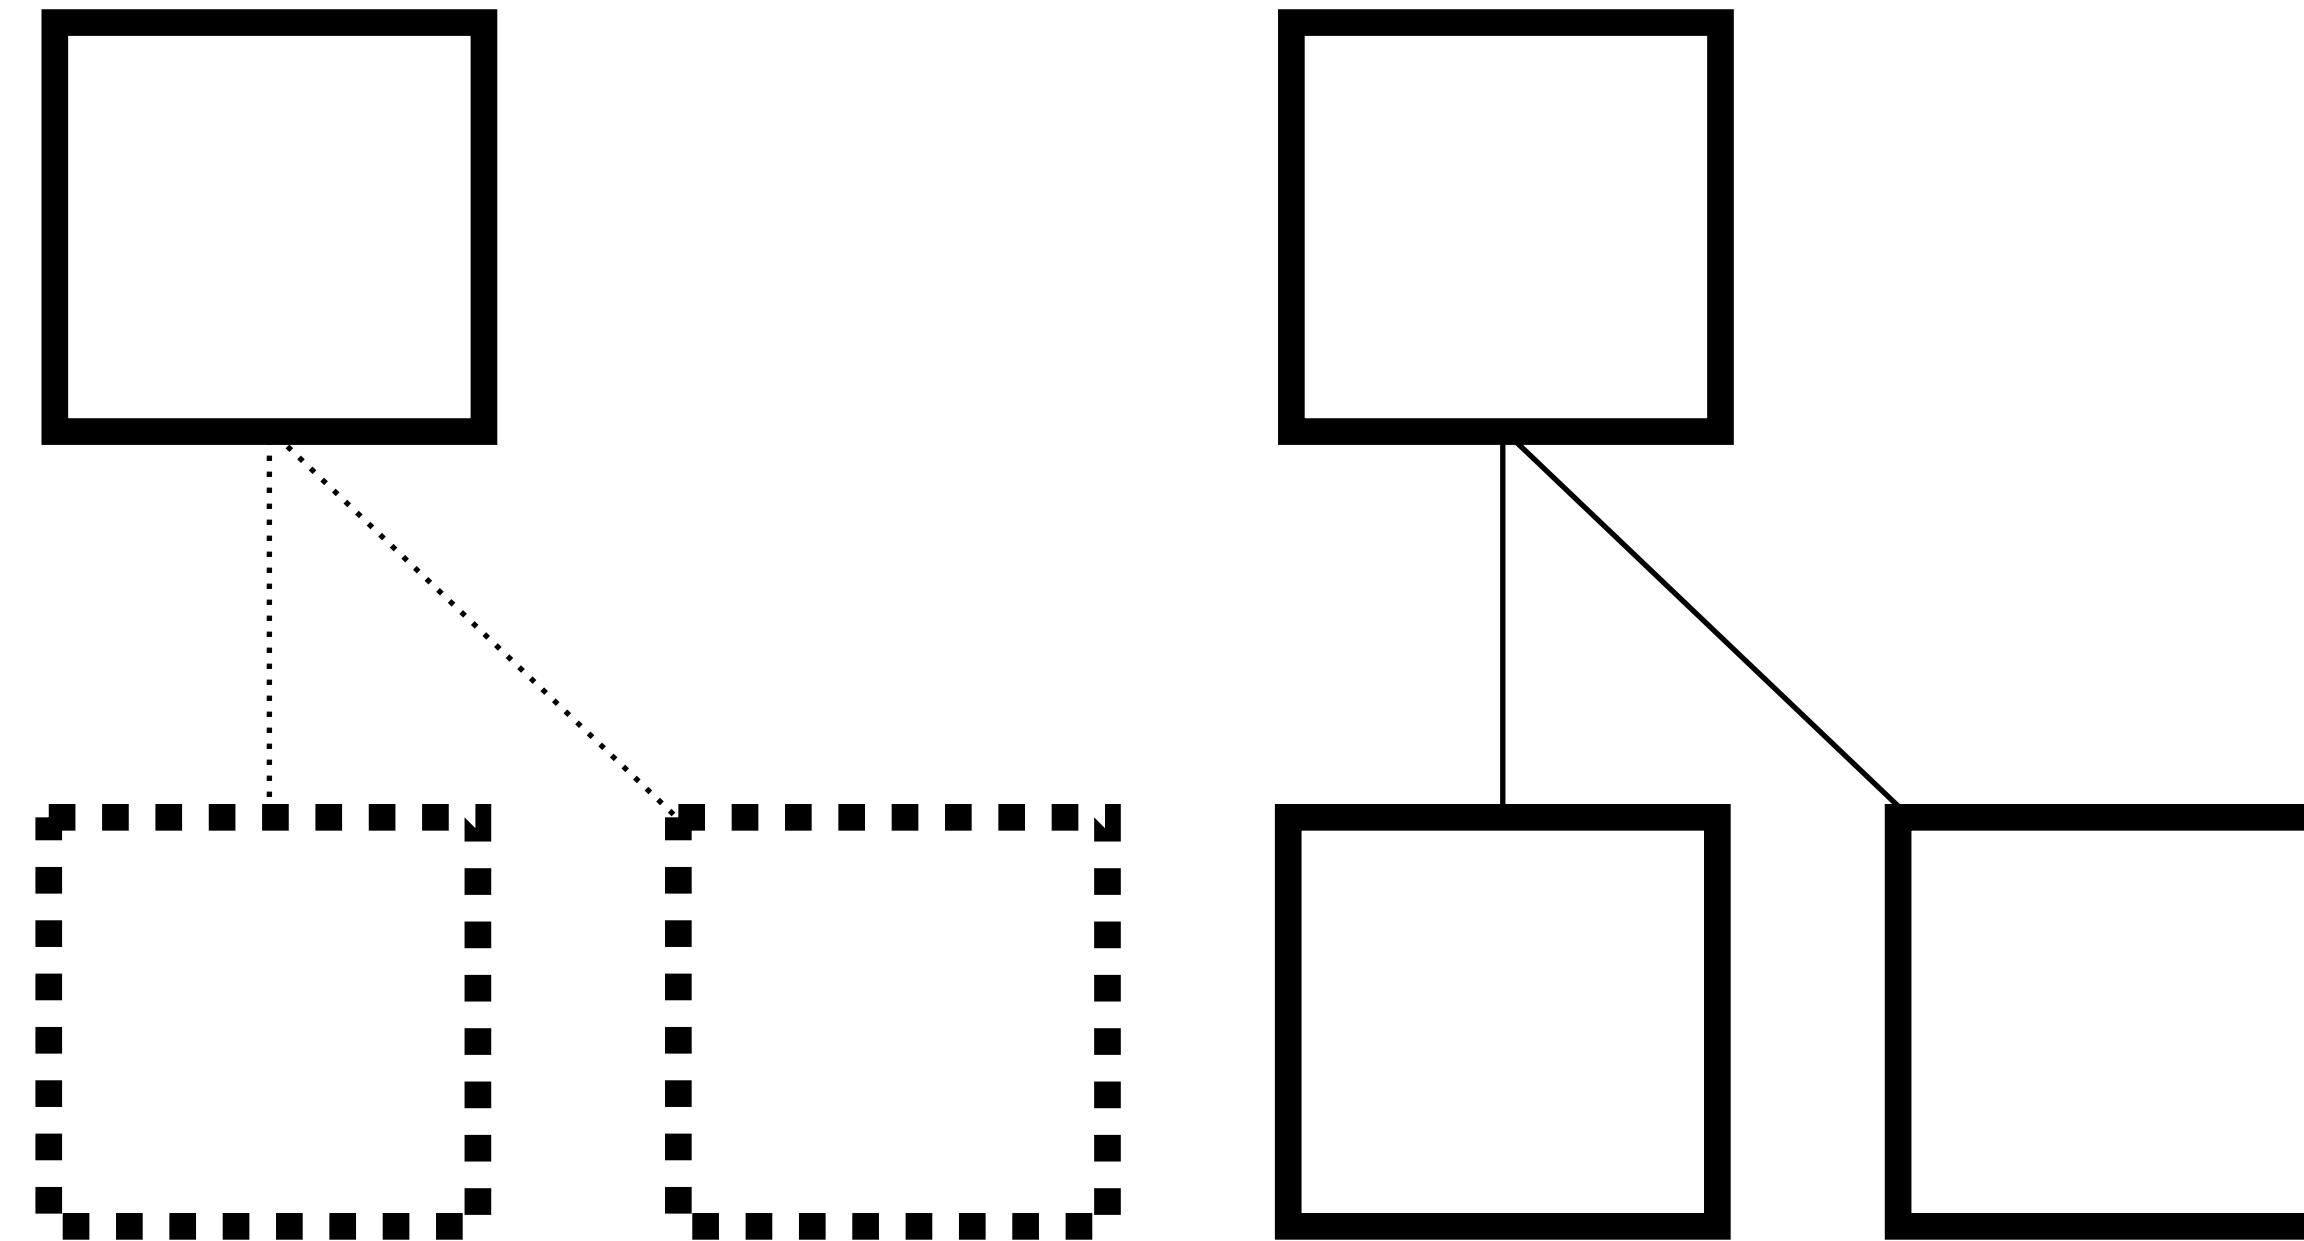
4 UTXOs Concatenate



<https://gist.github.com/kcalvinalvin/a790d524832e1b7f96a70c642315fffc>

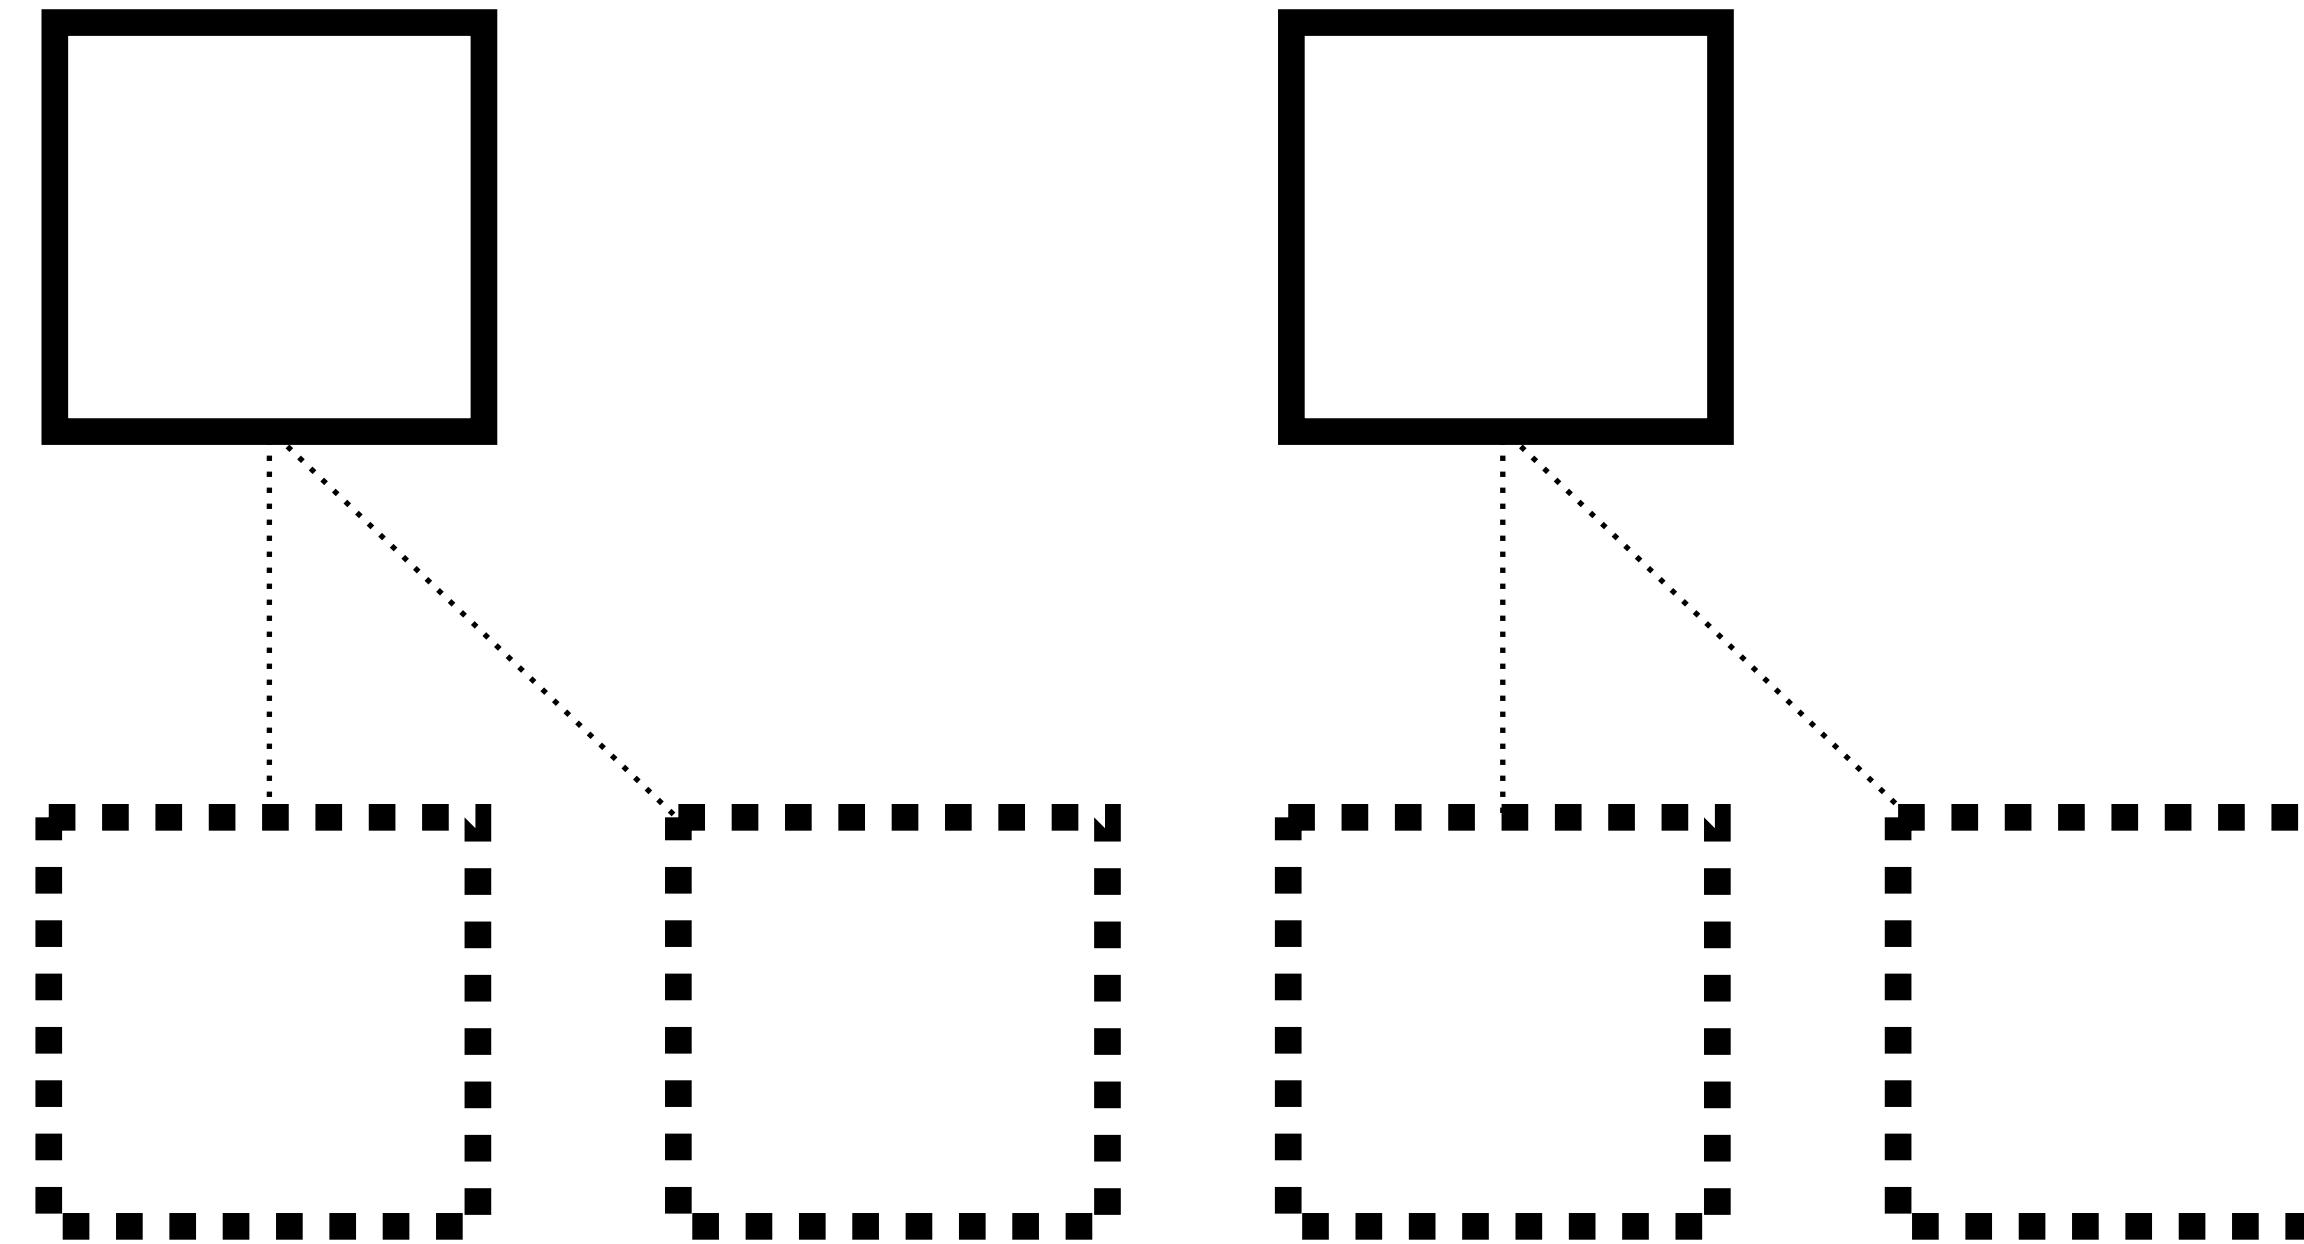
4 UTXOs

Hash

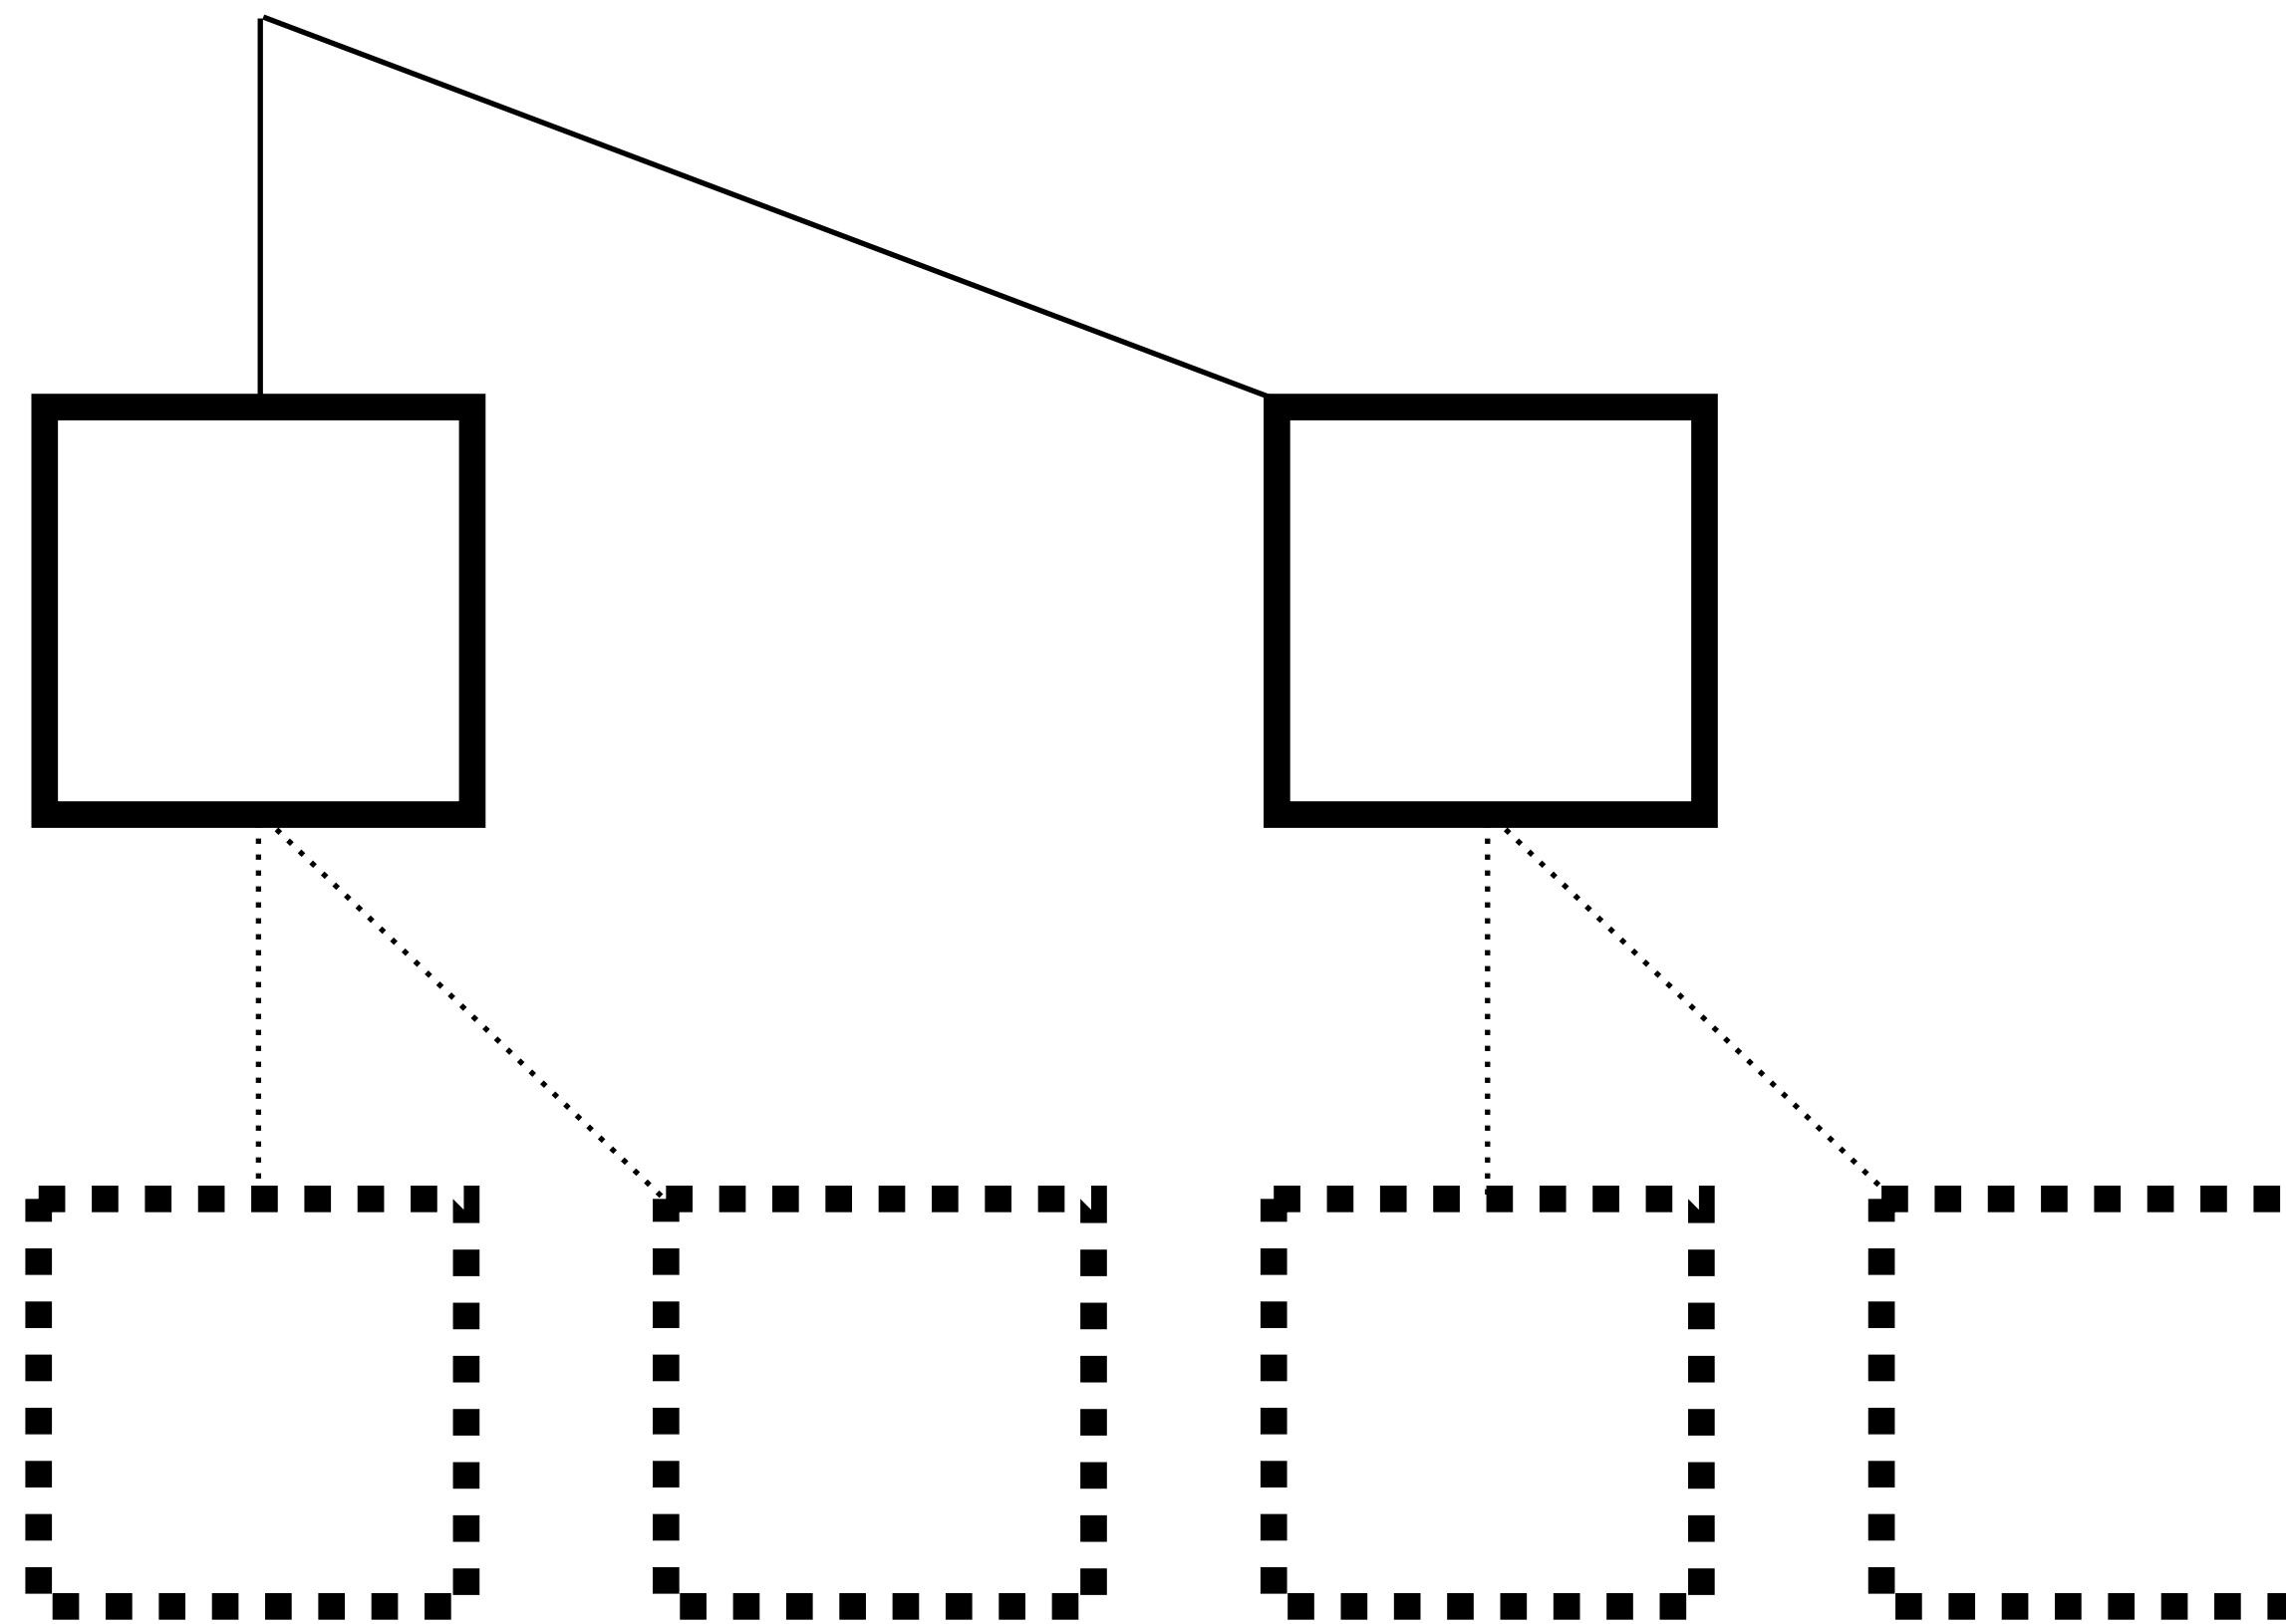


4 UTXOs

Can now throw away the leaves

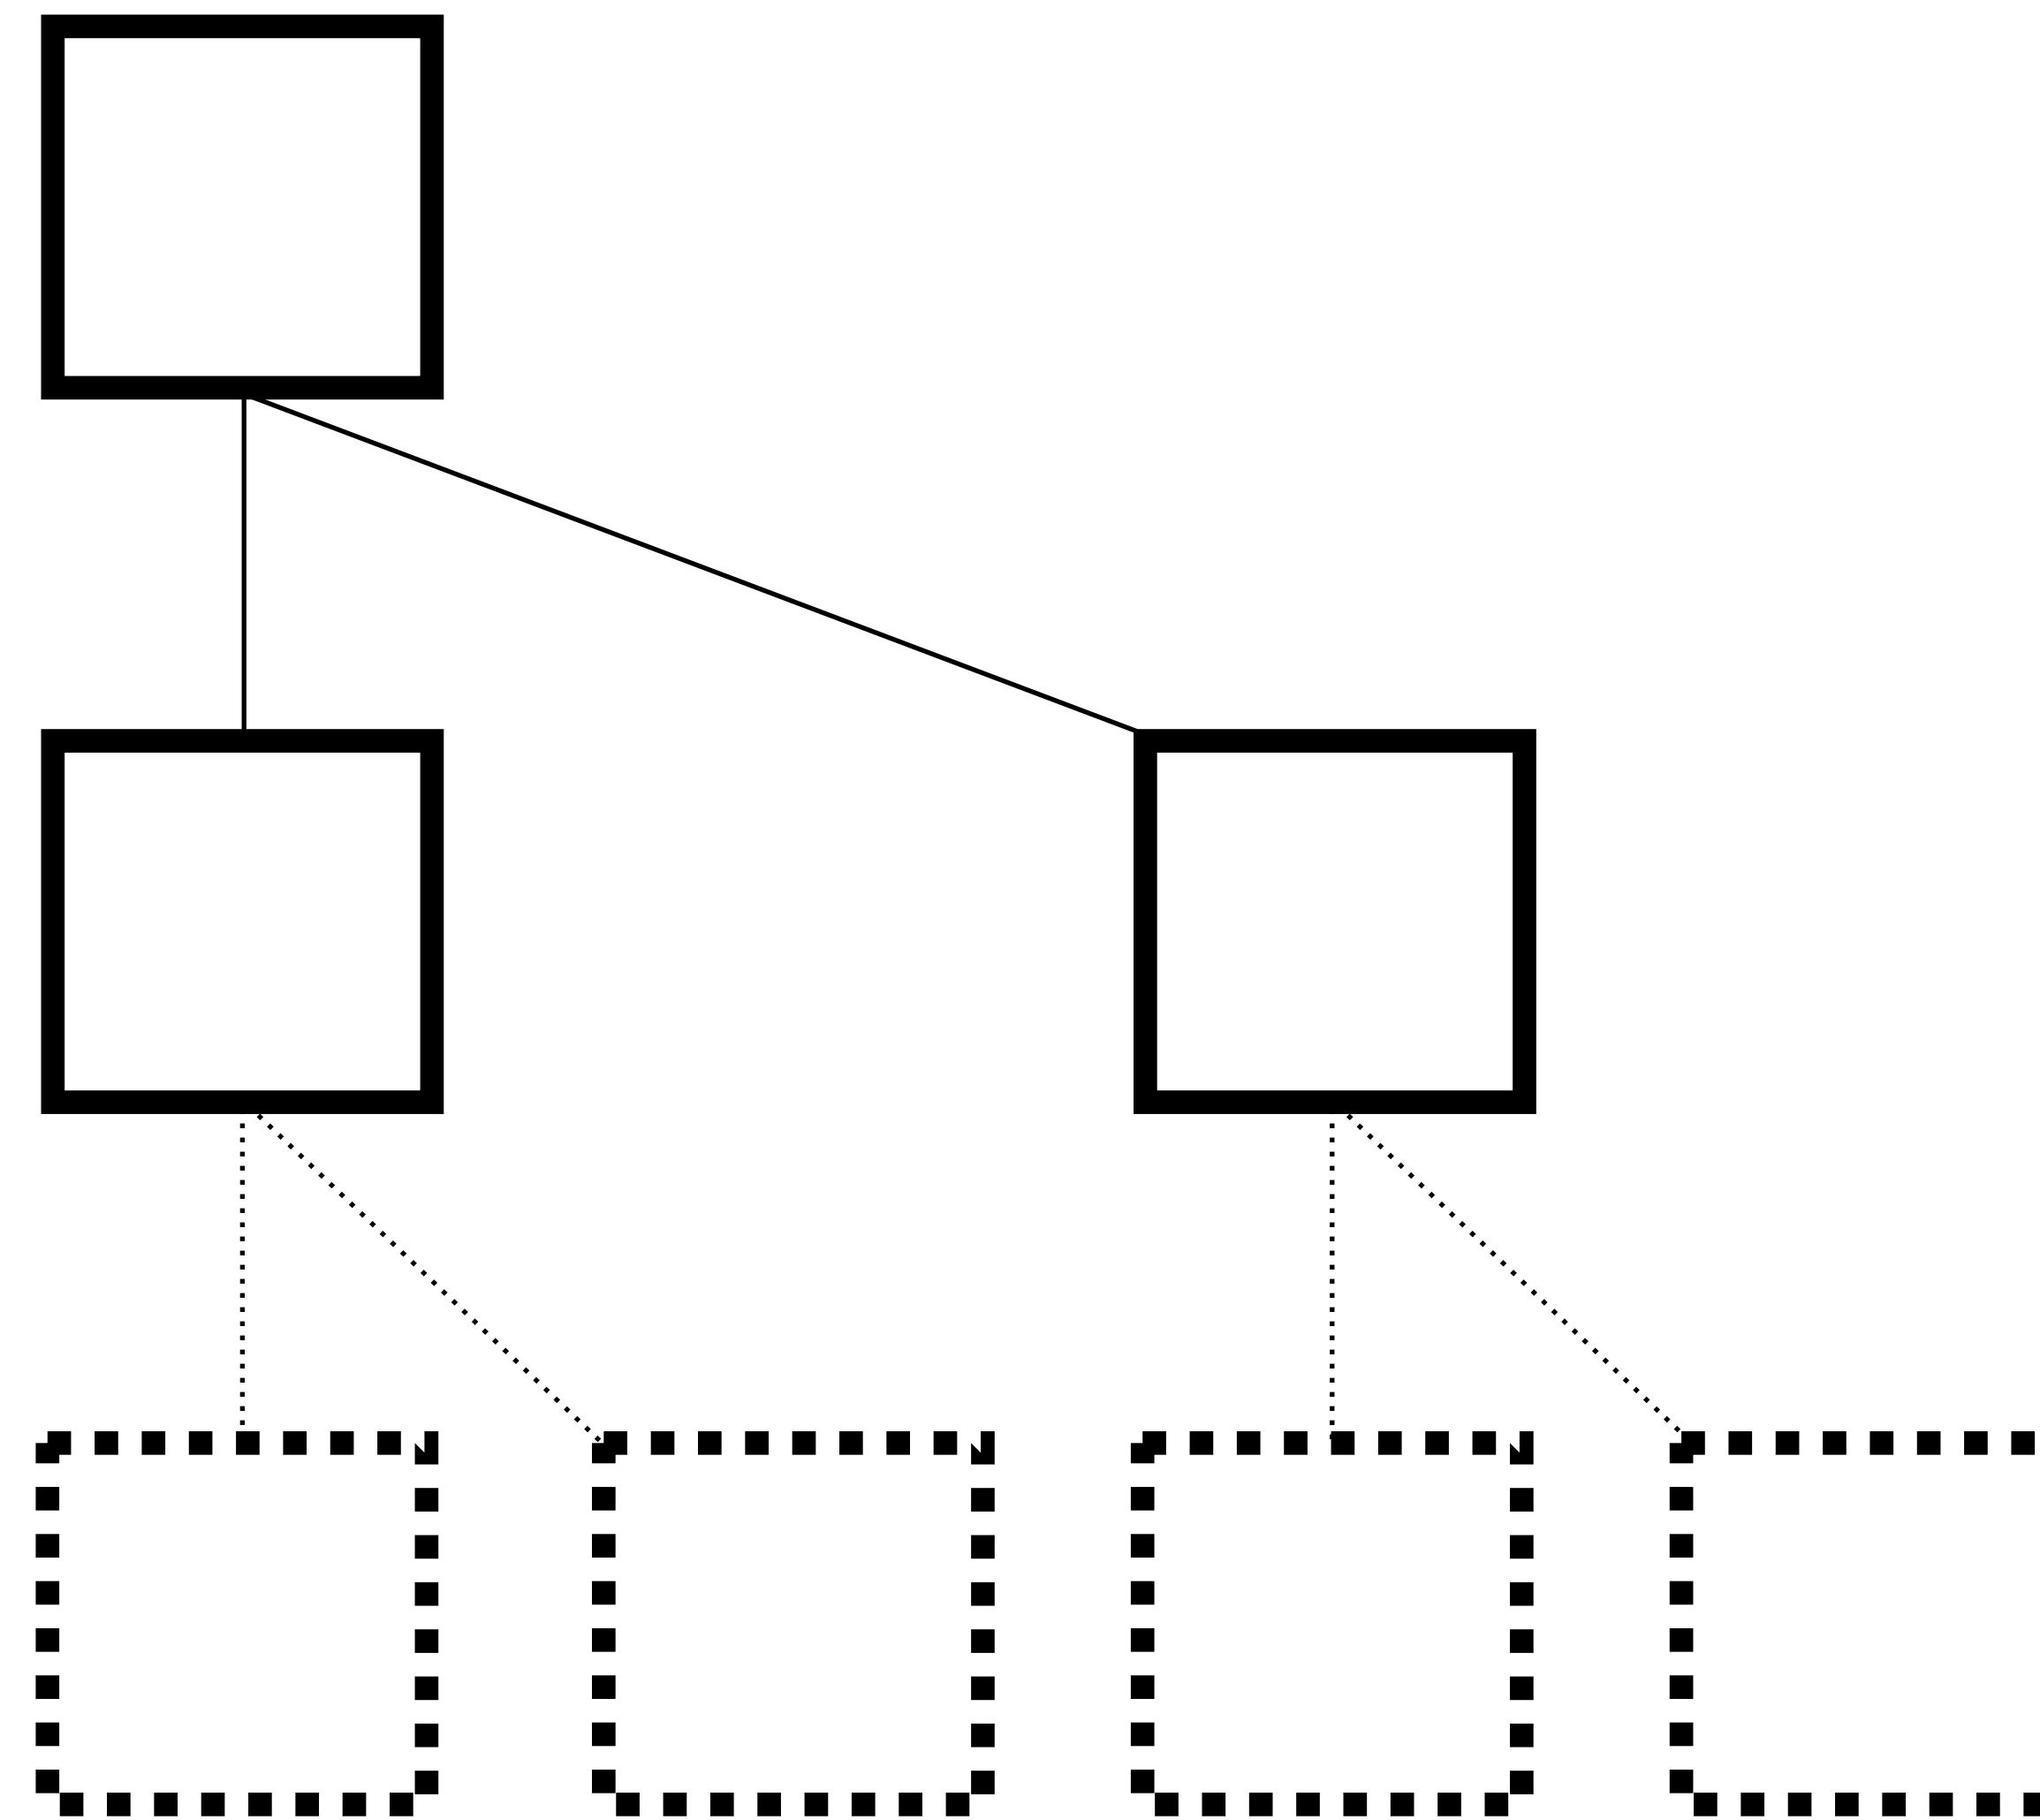


4 UTXOs Concatenate



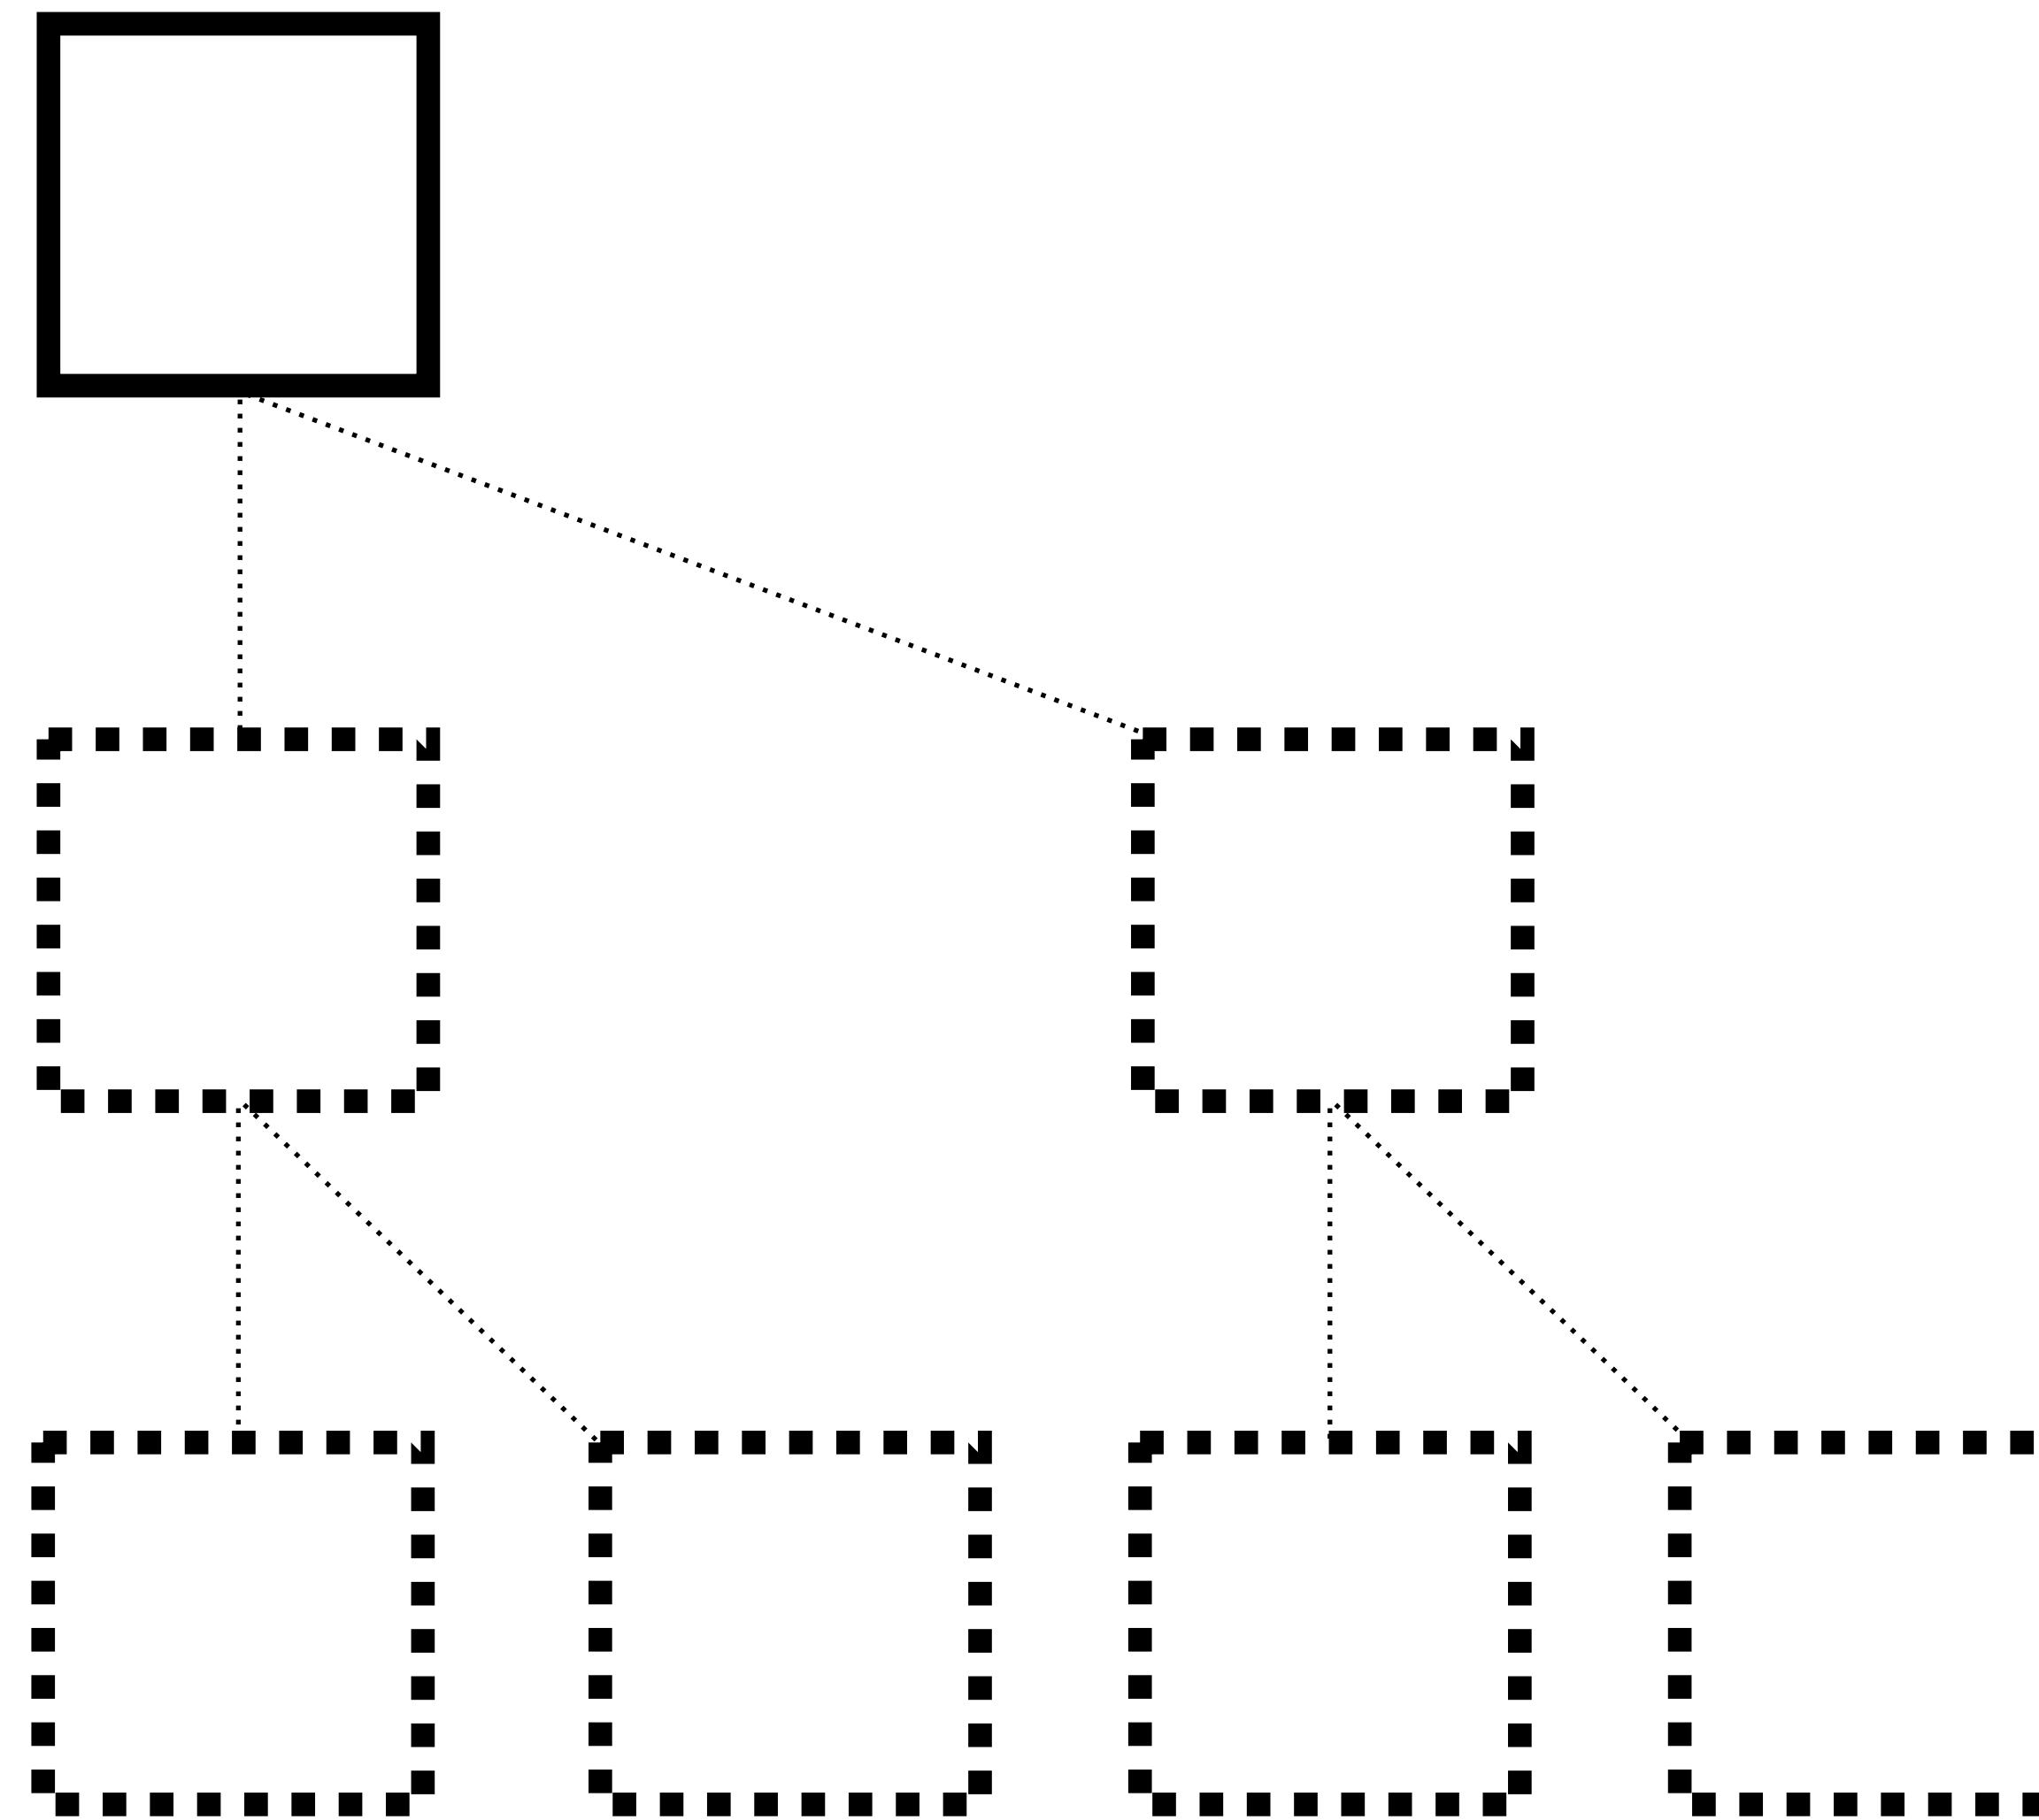
4 UTXOs

Hash



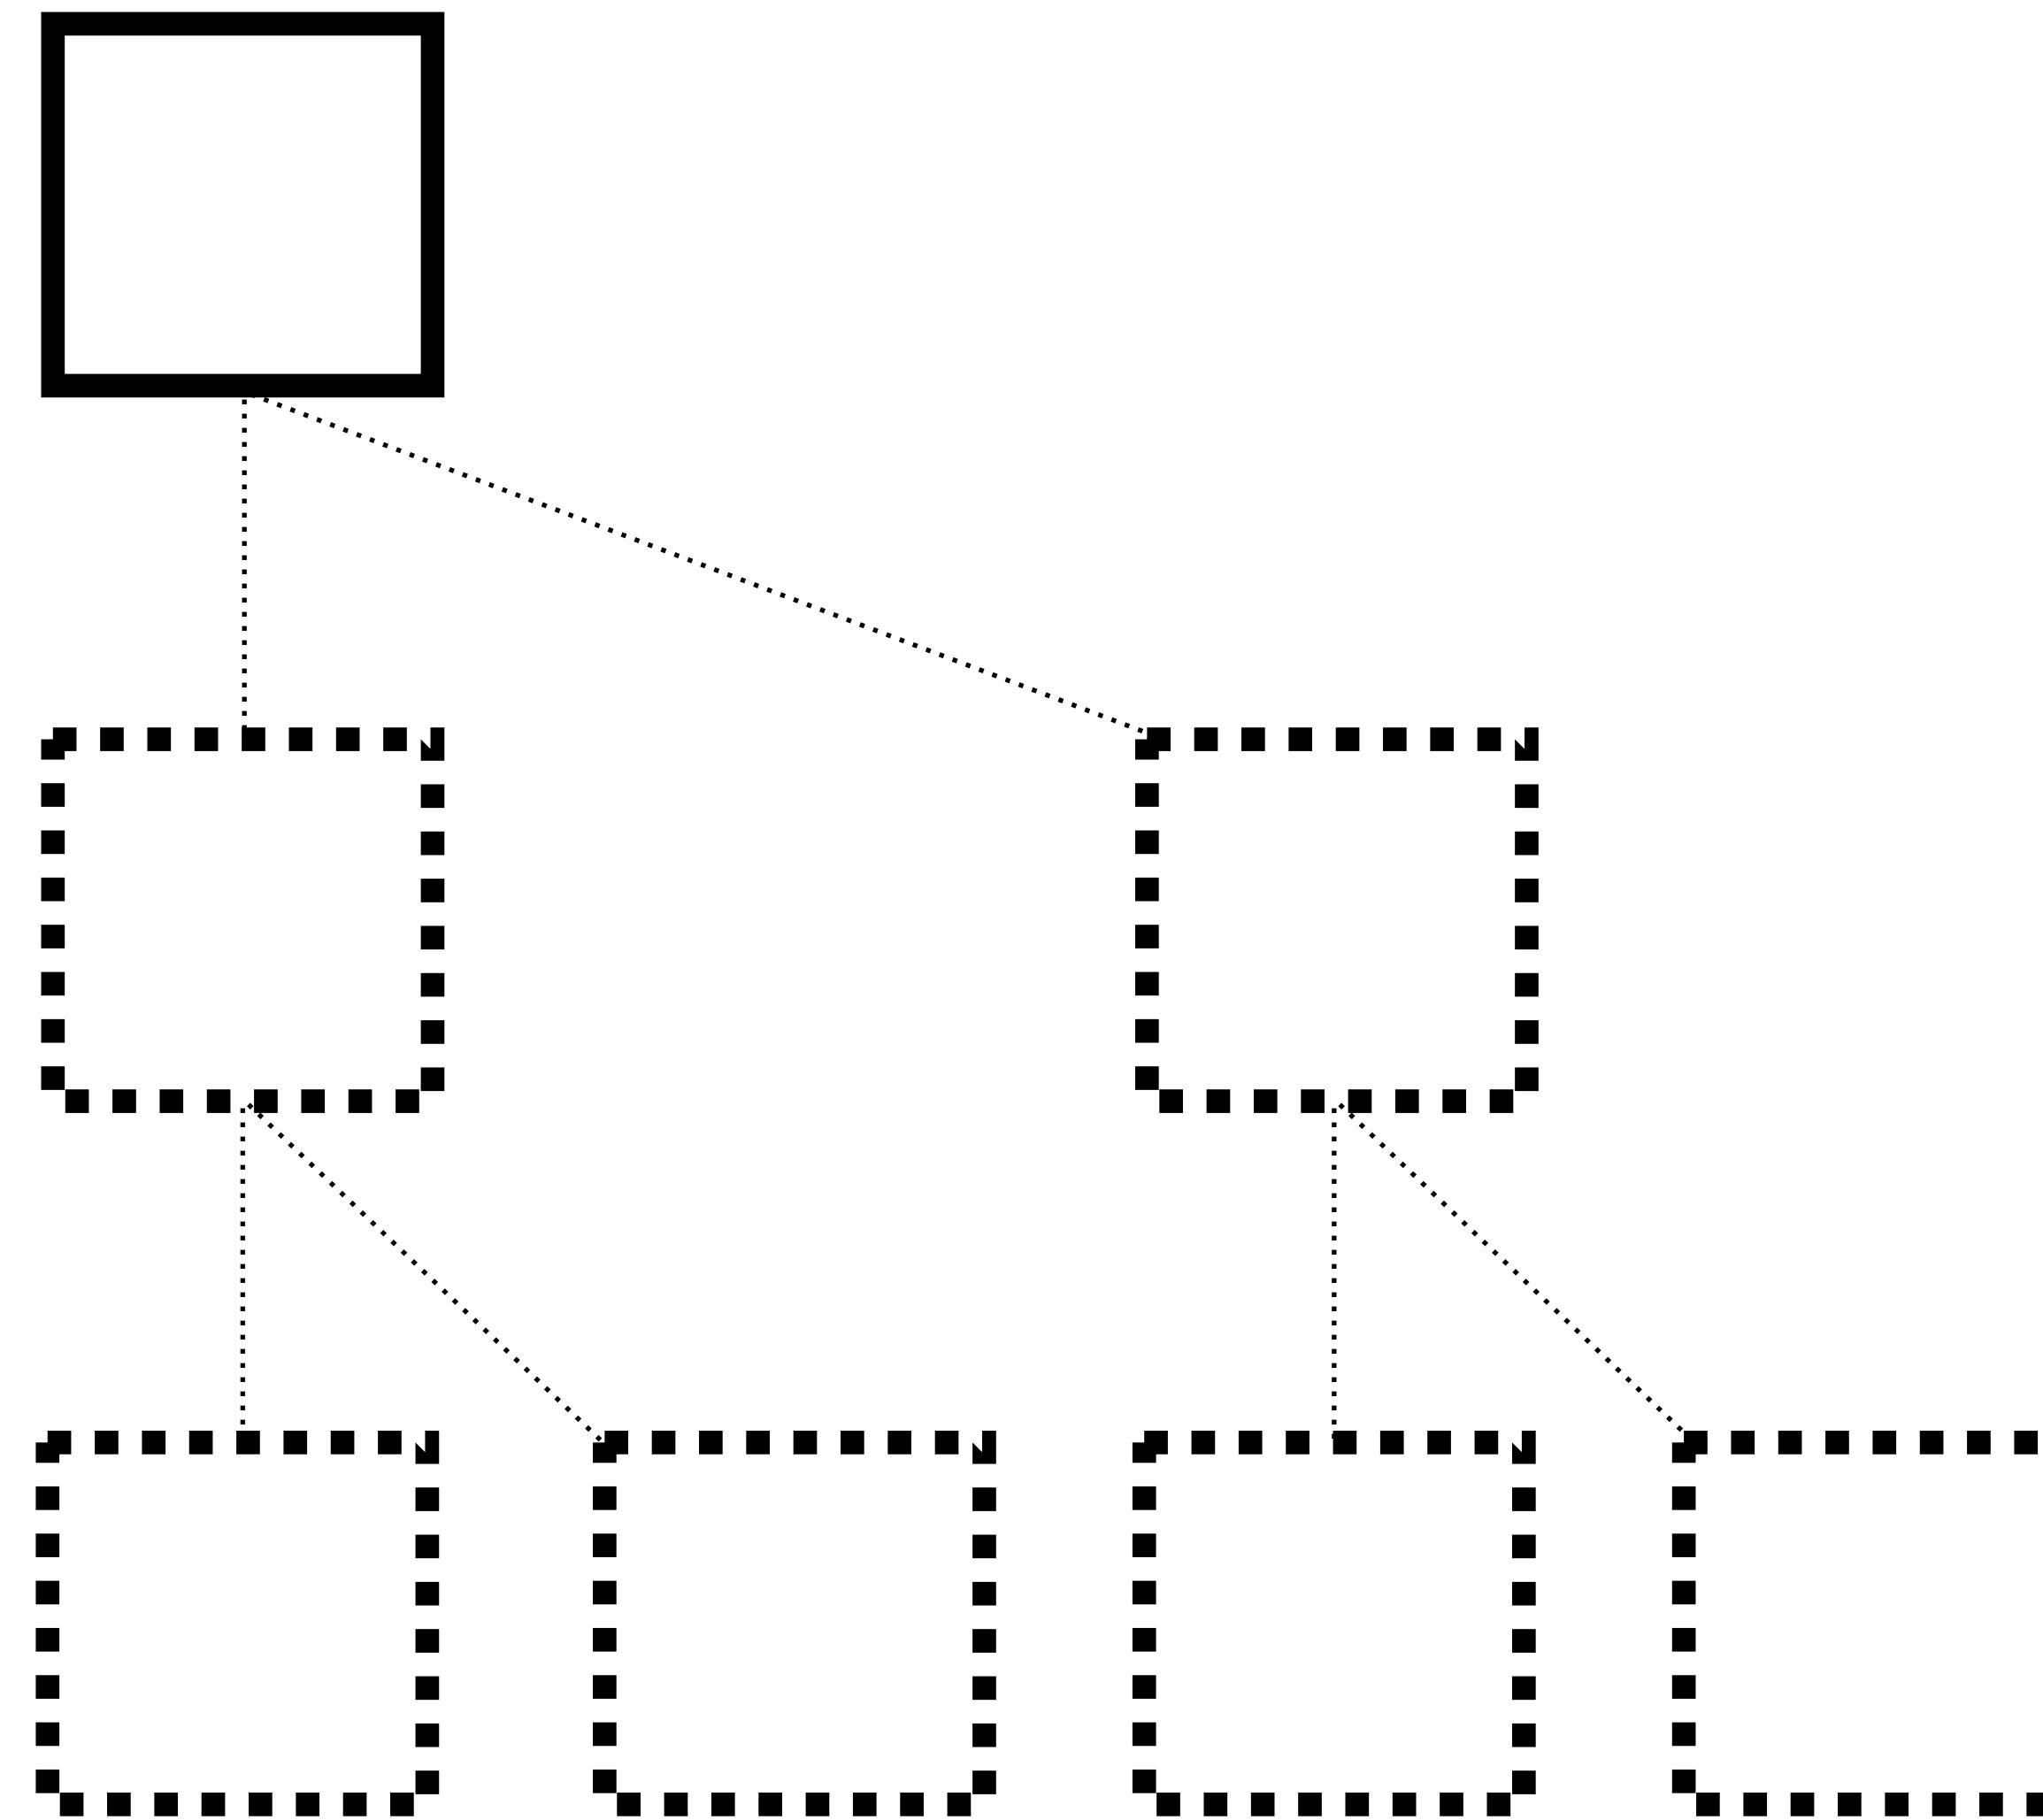
4 UTXOs

Throw away



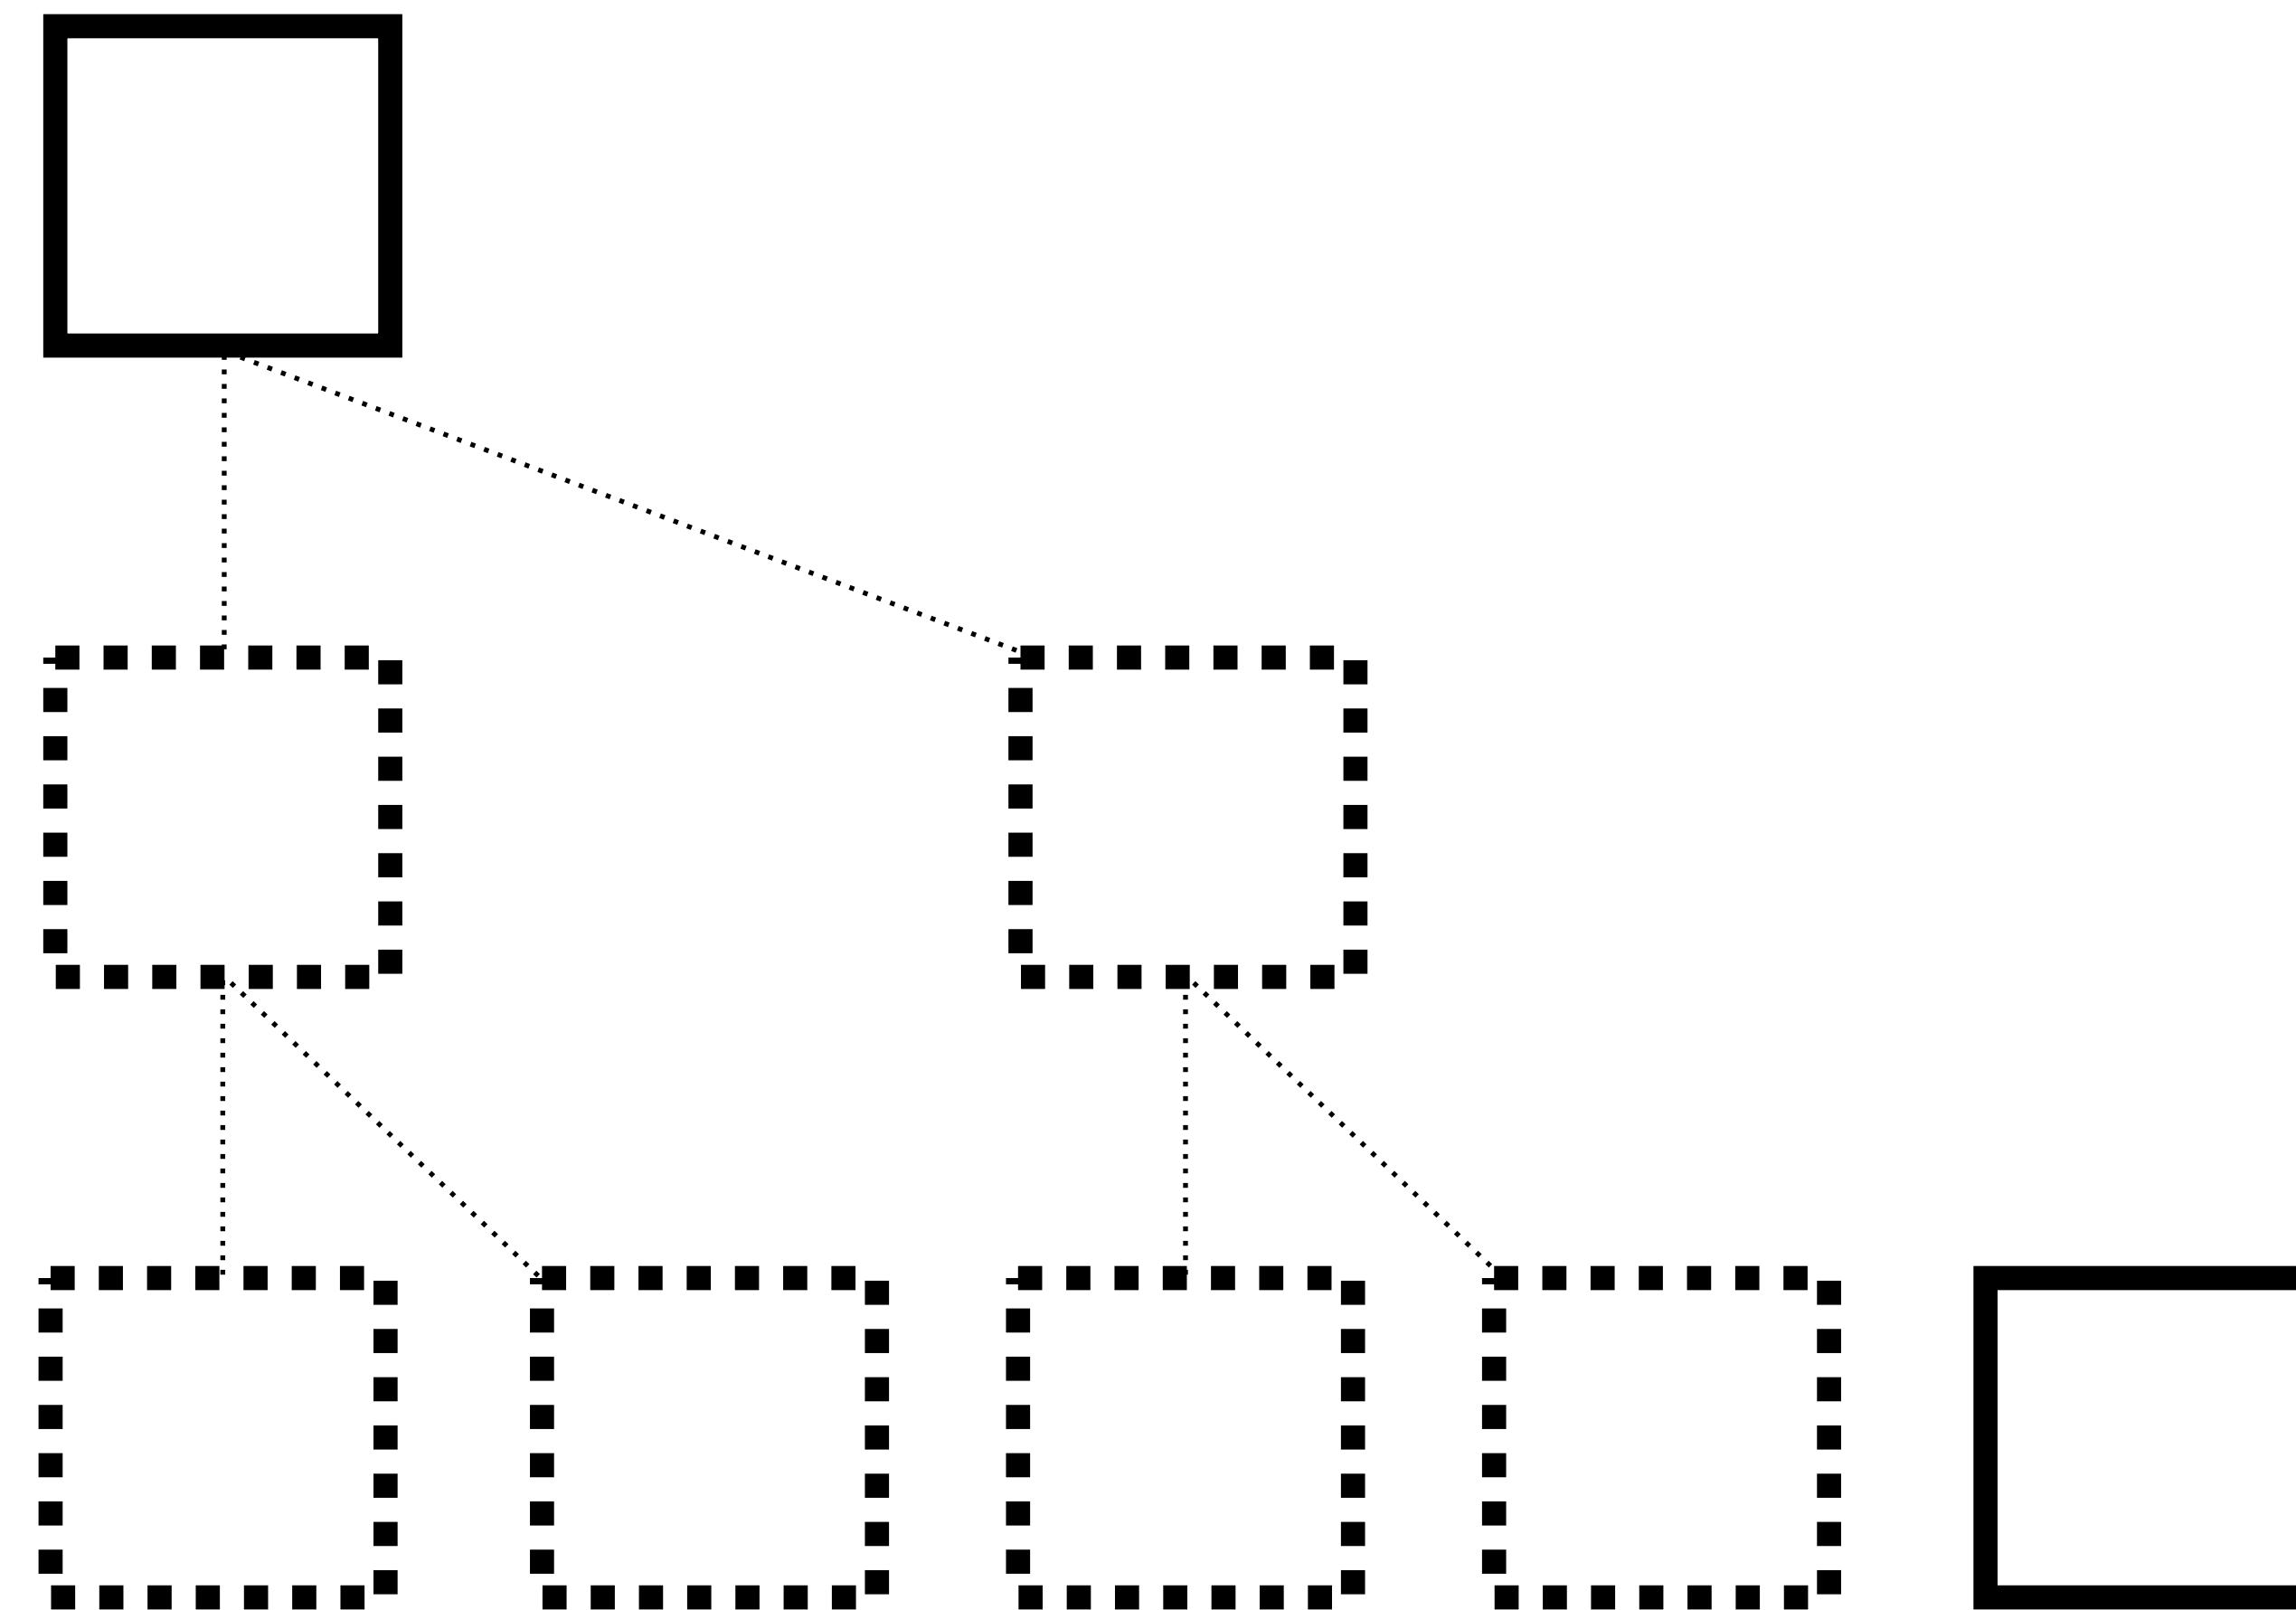
4 UTXOs

Finished tree



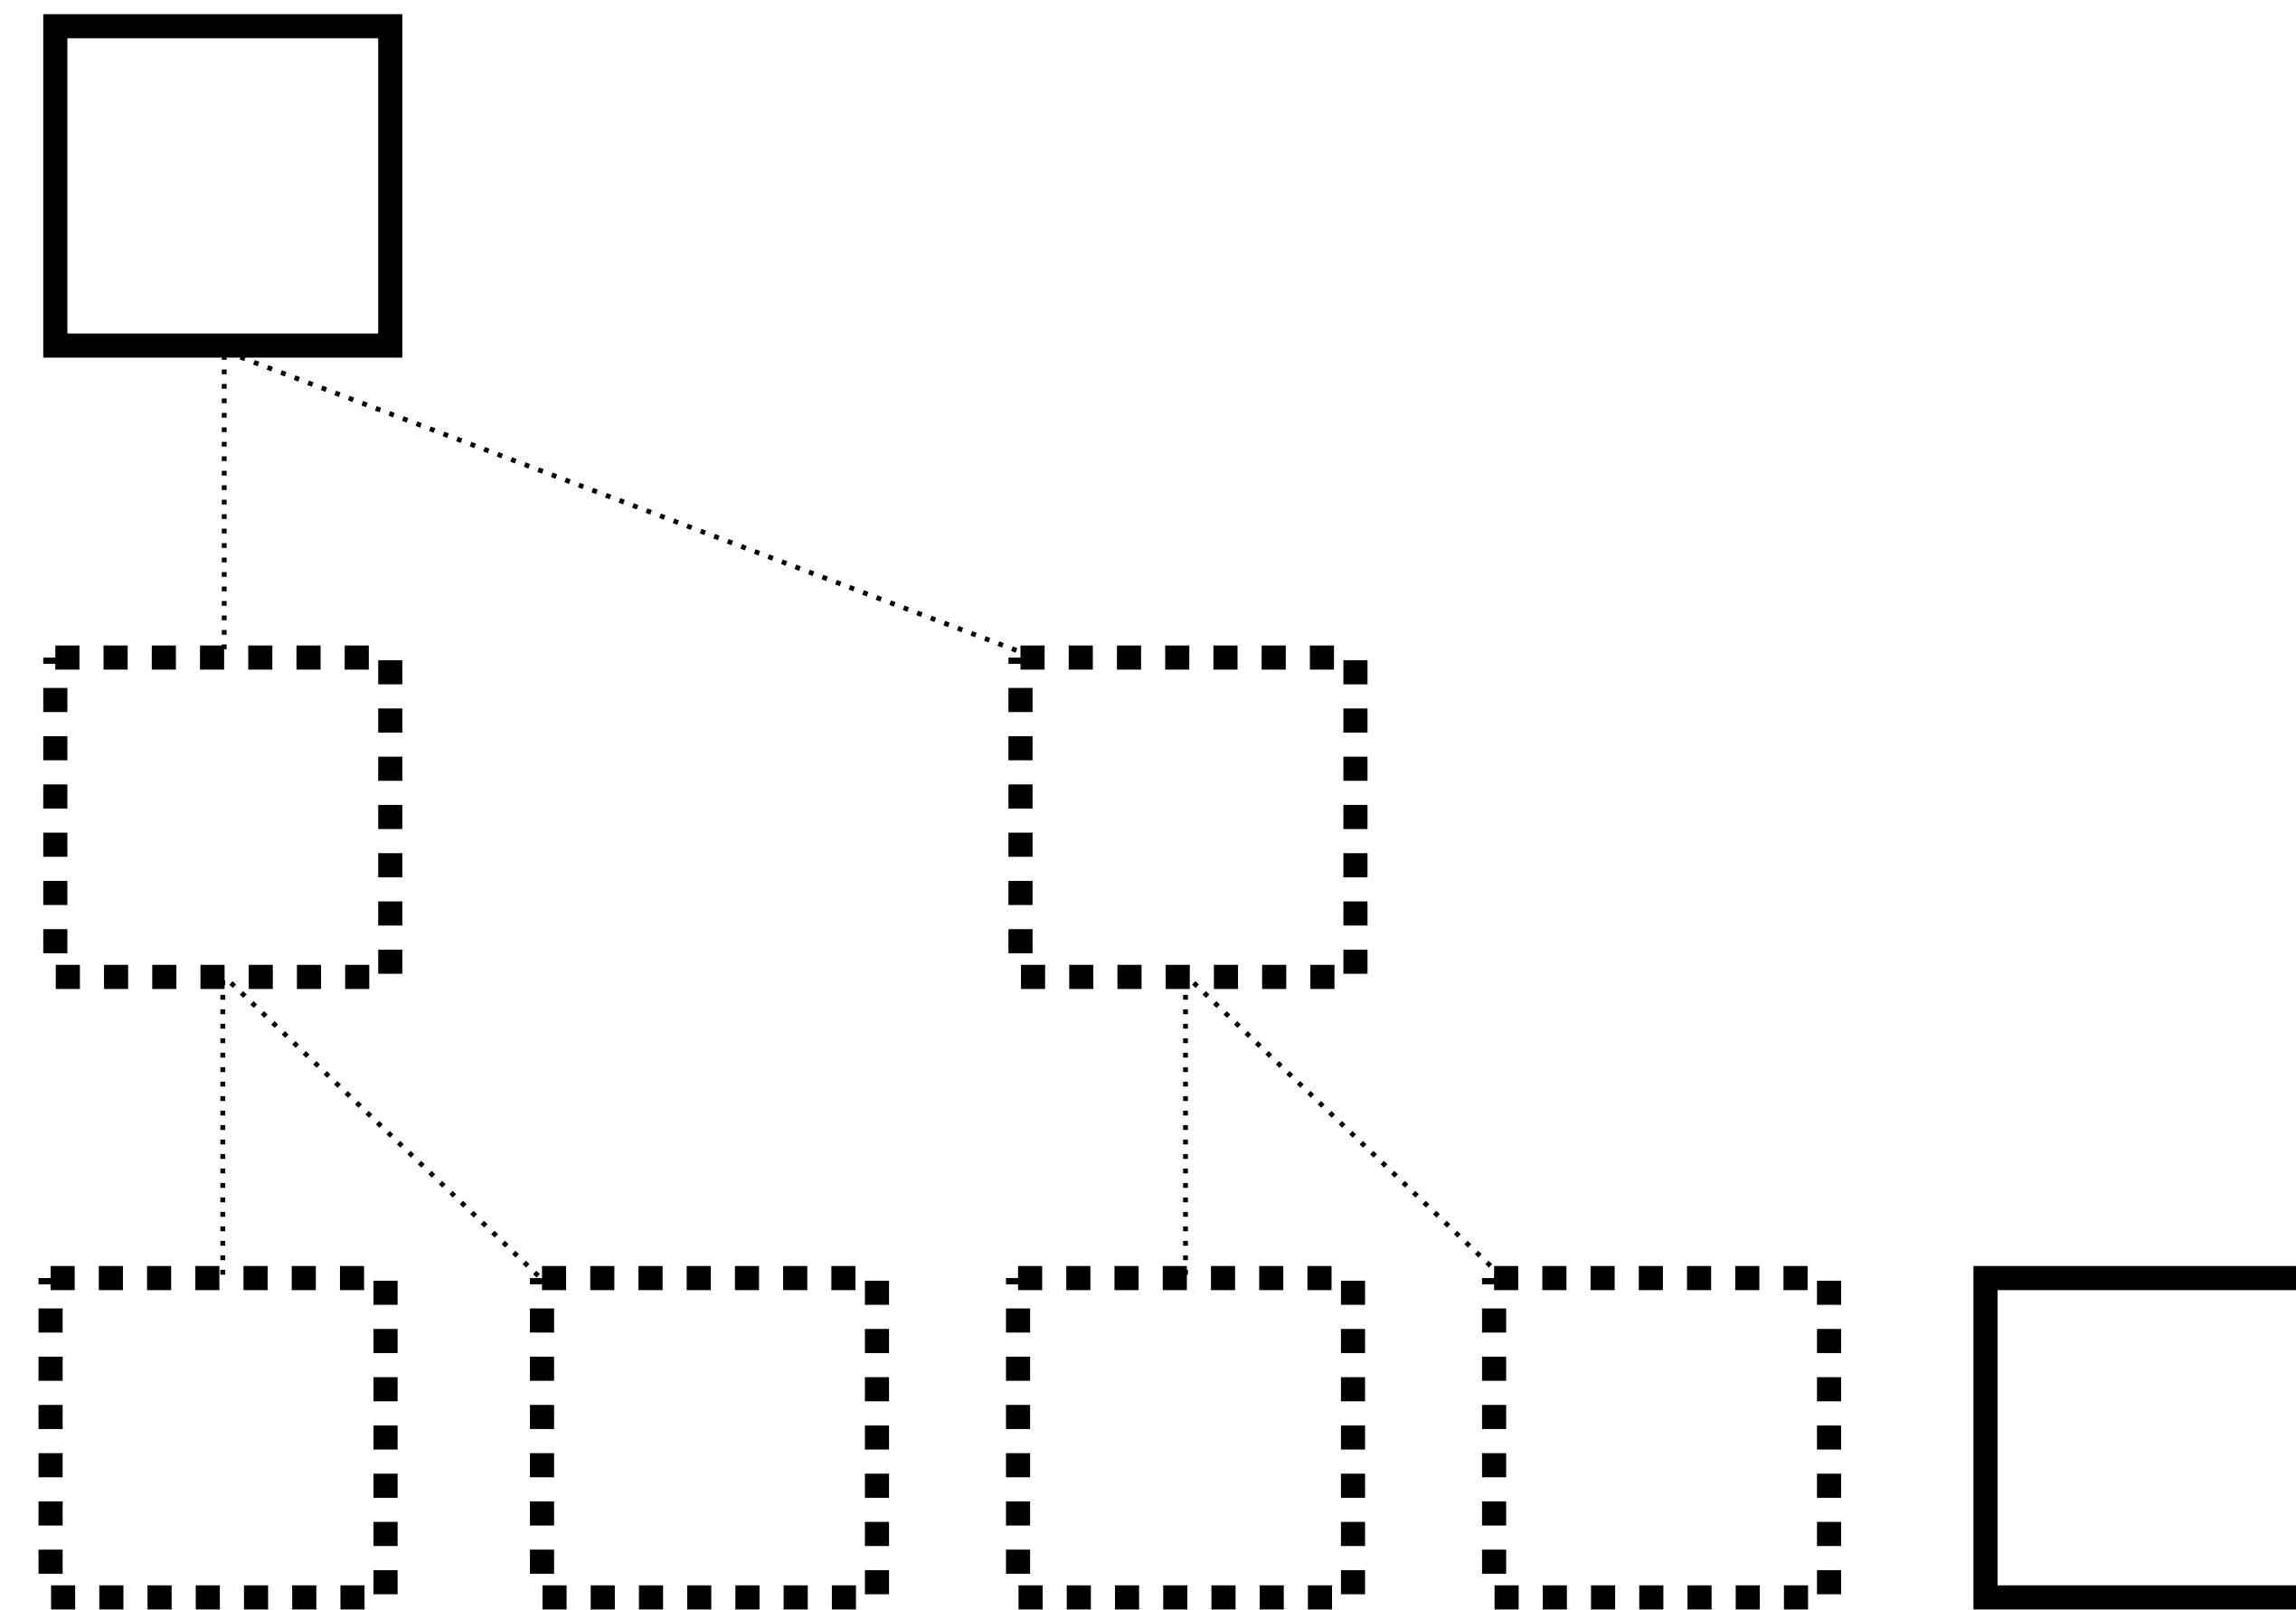
5 UTXOs

Add one more



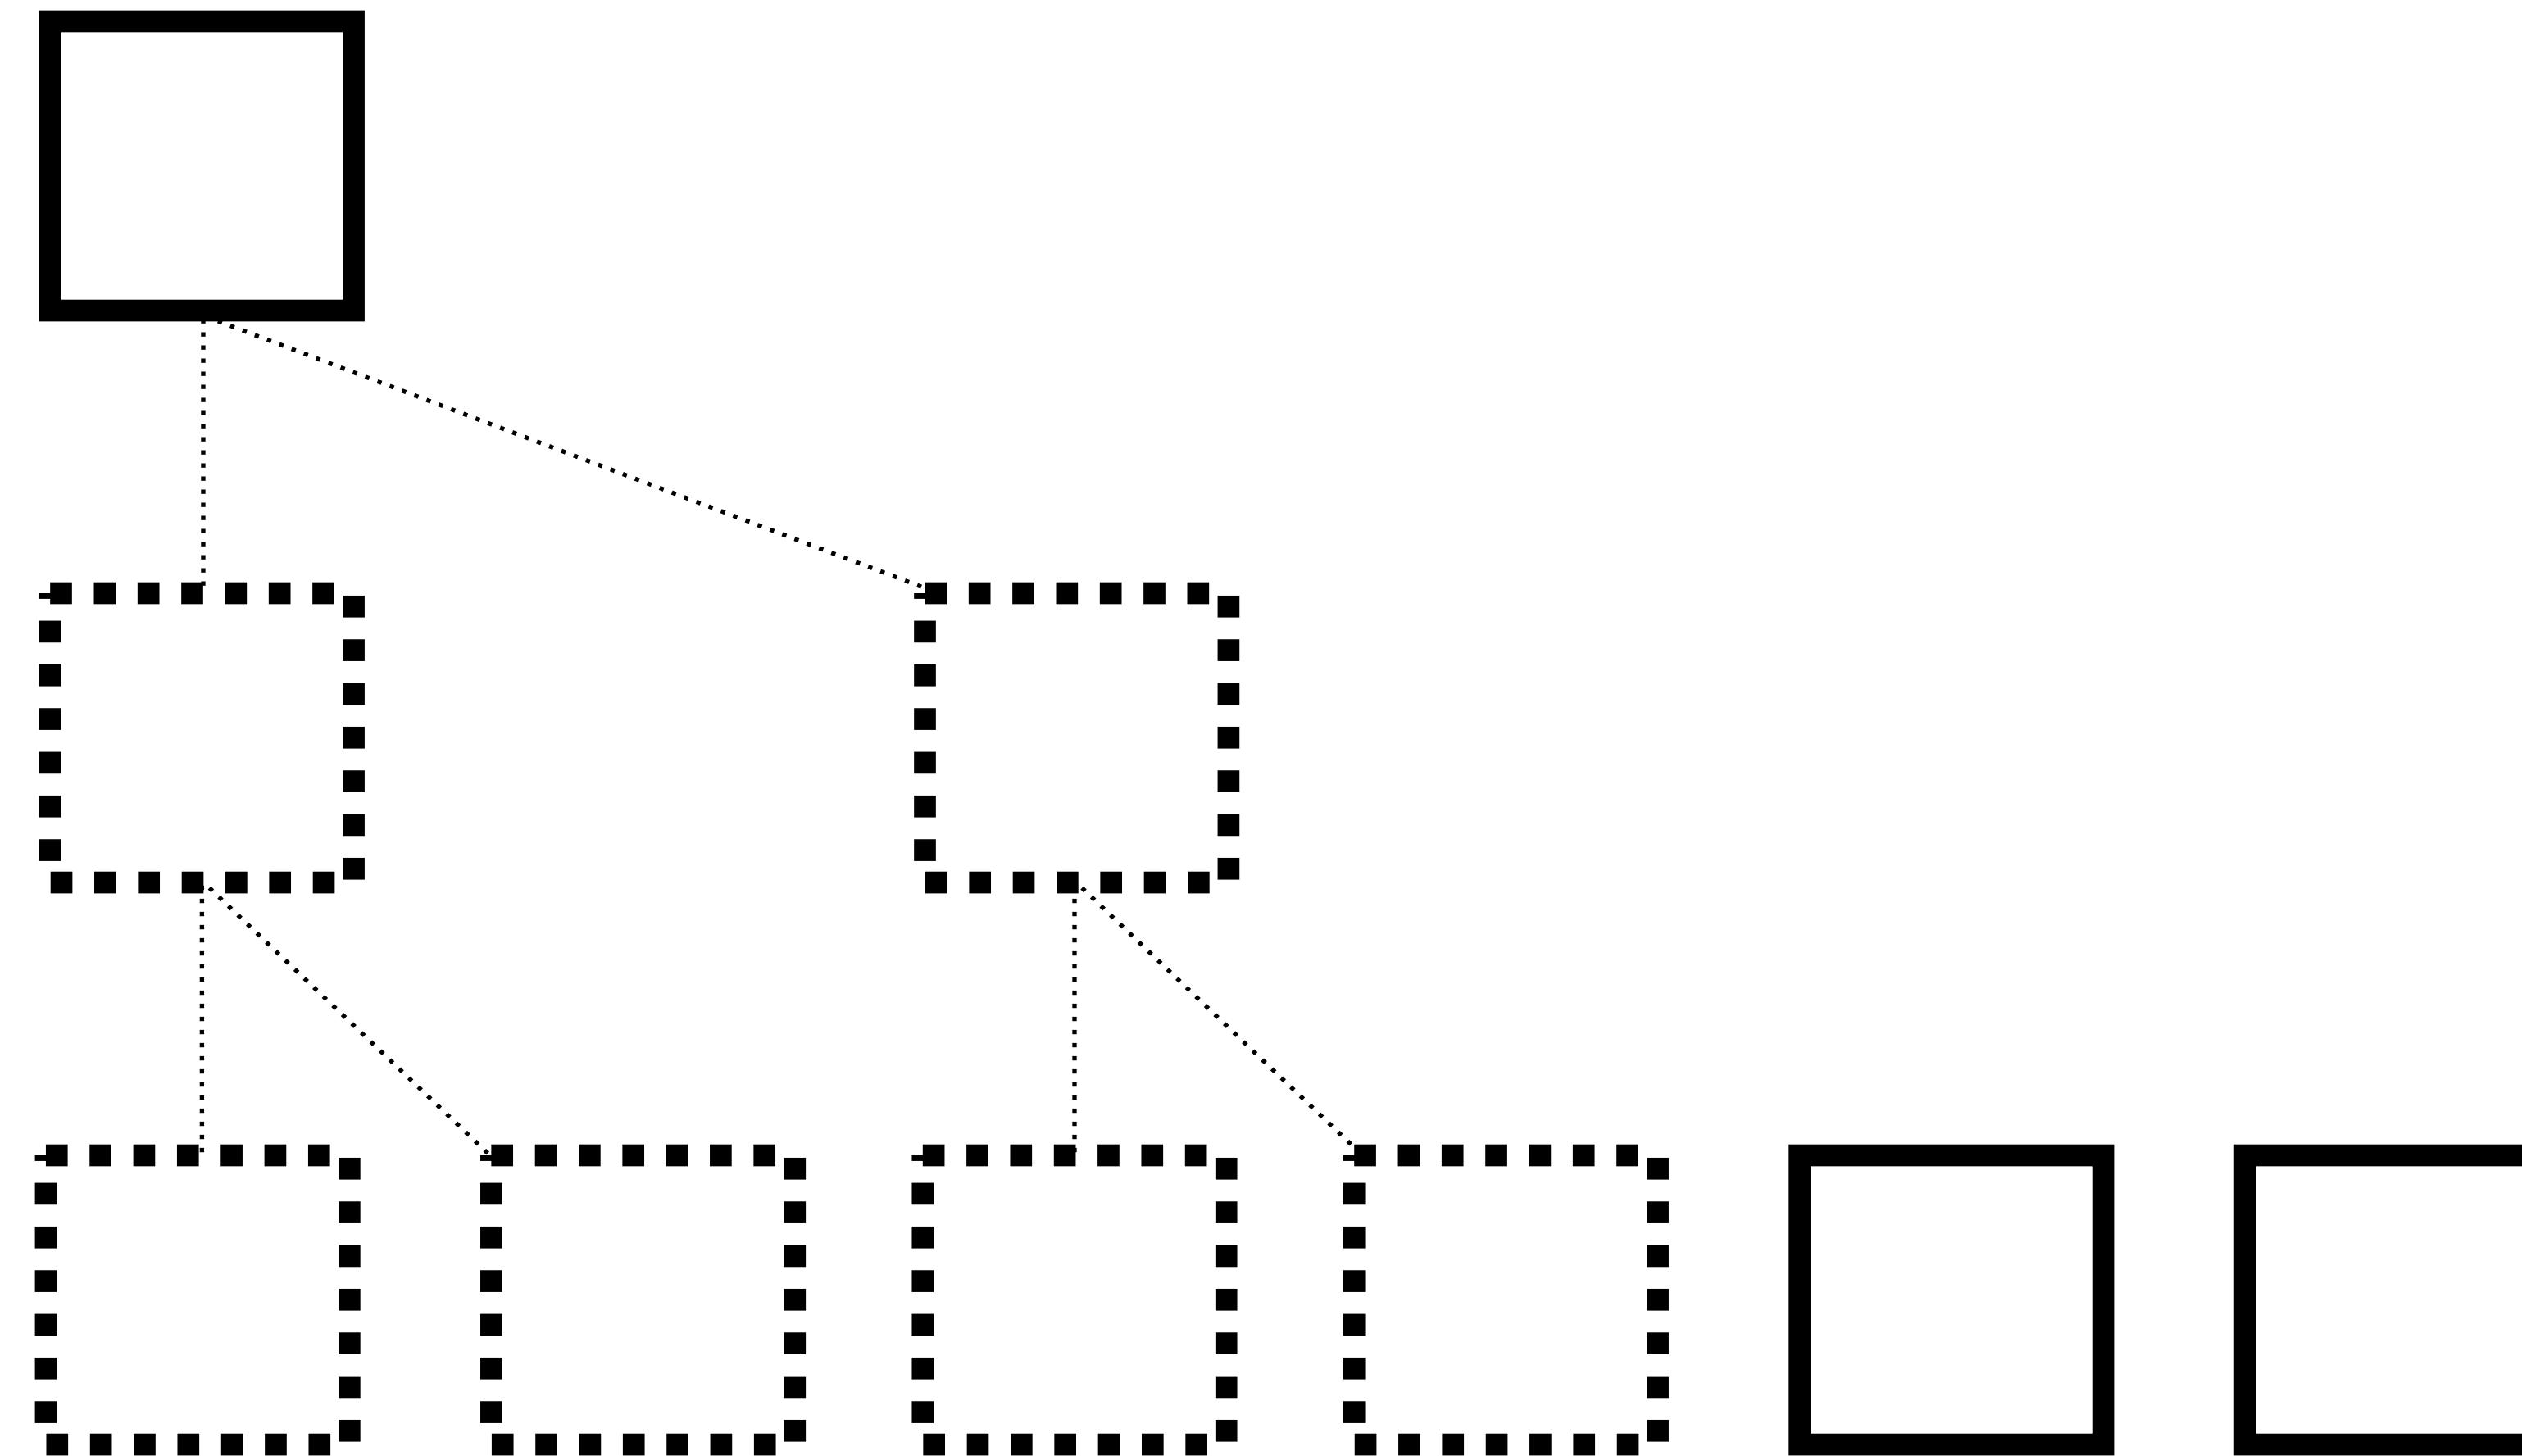
5 UTXOs

Nothing to hash with so it becomes a root

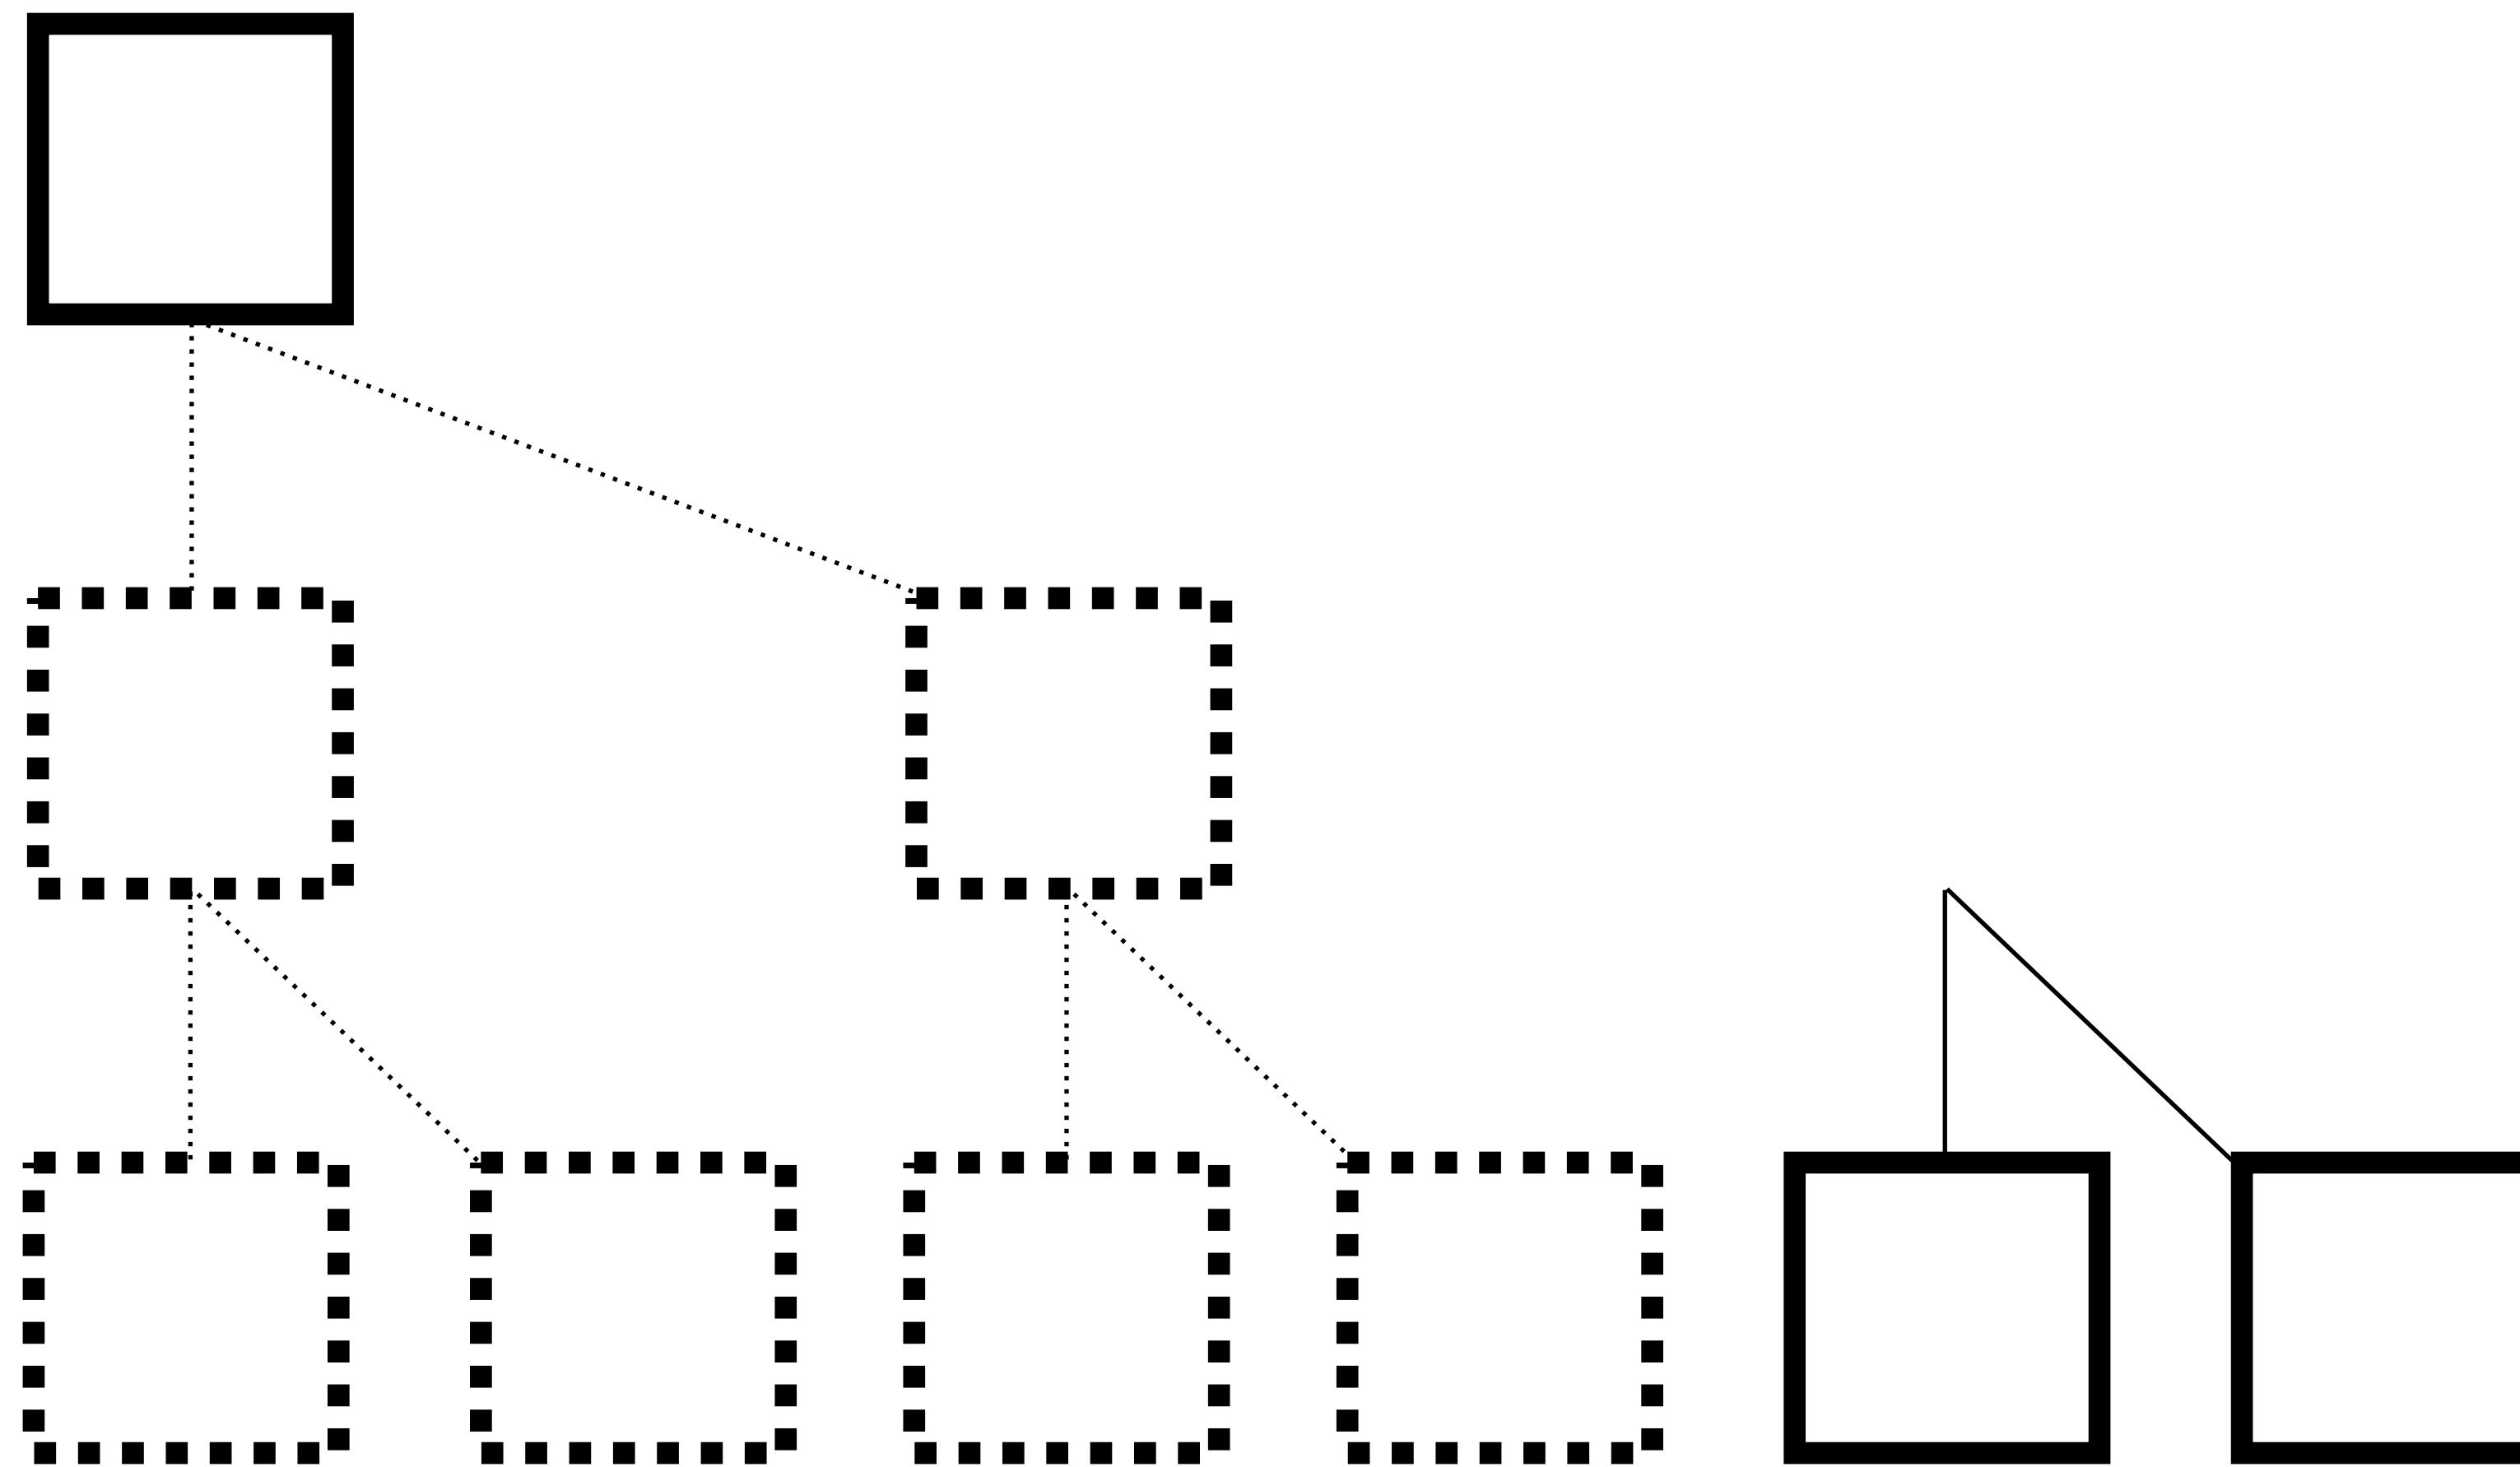


6 UTXOs

1 more

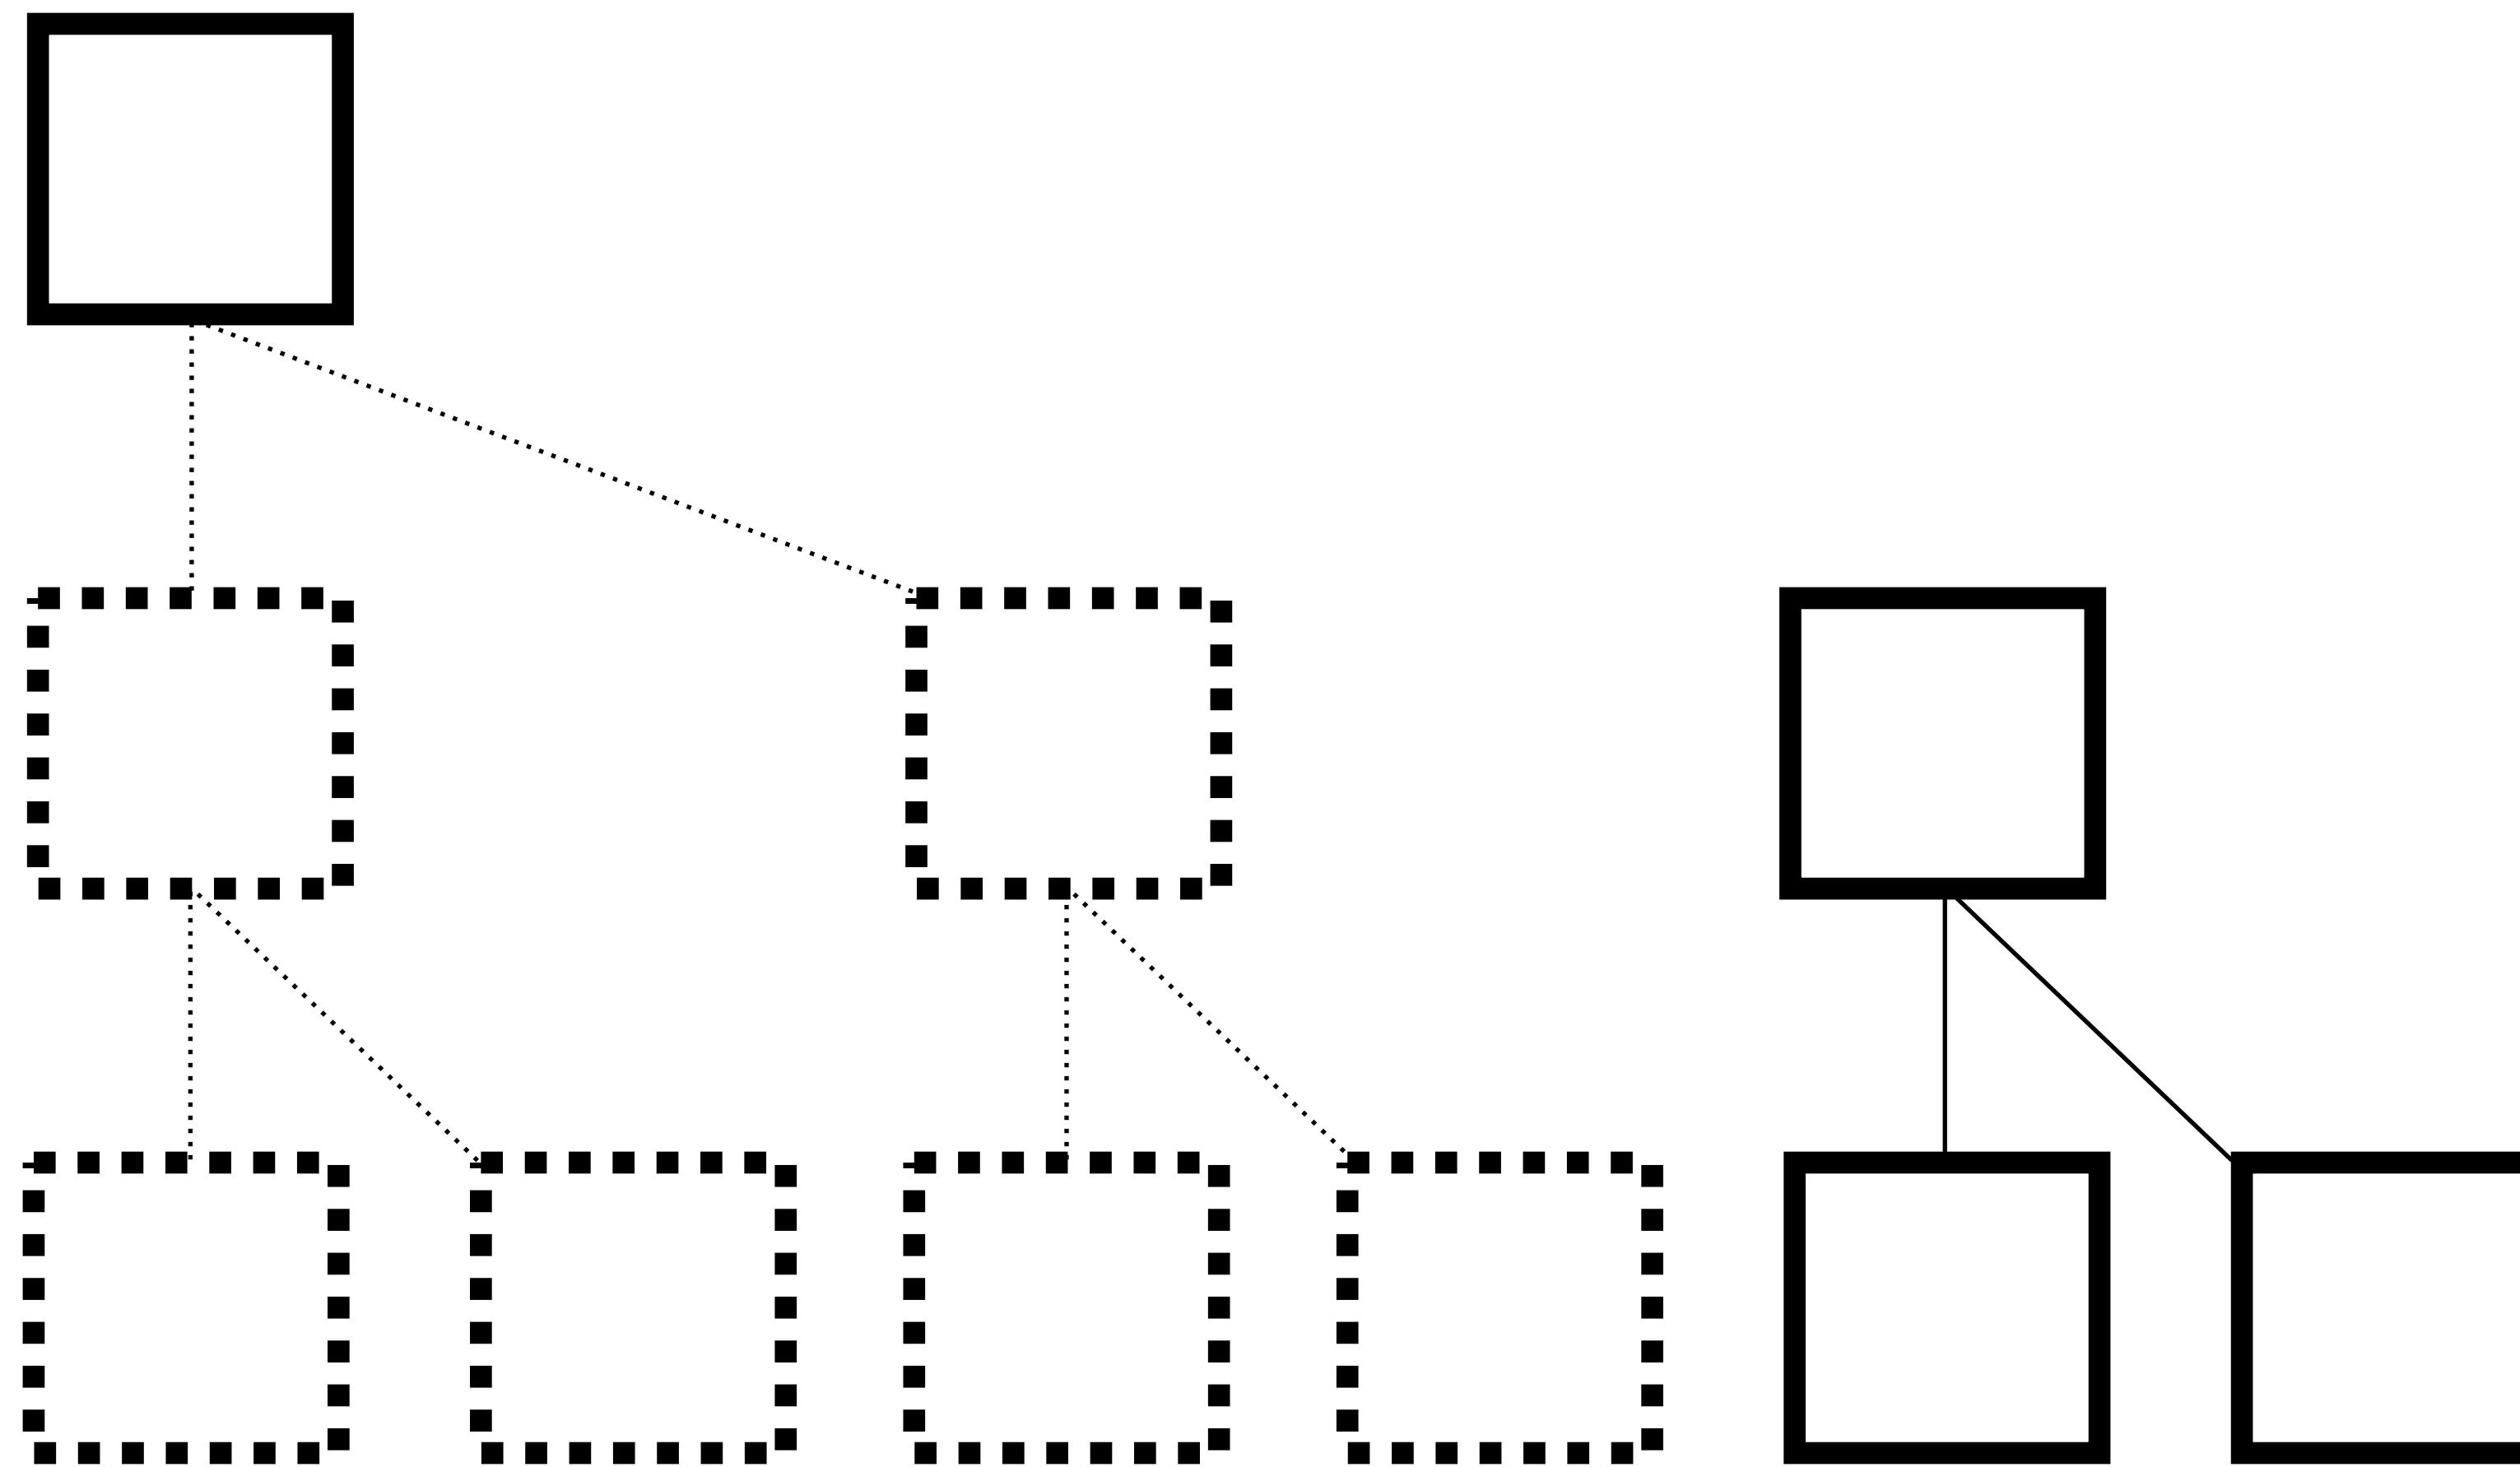


6 UTXOs Concatenate



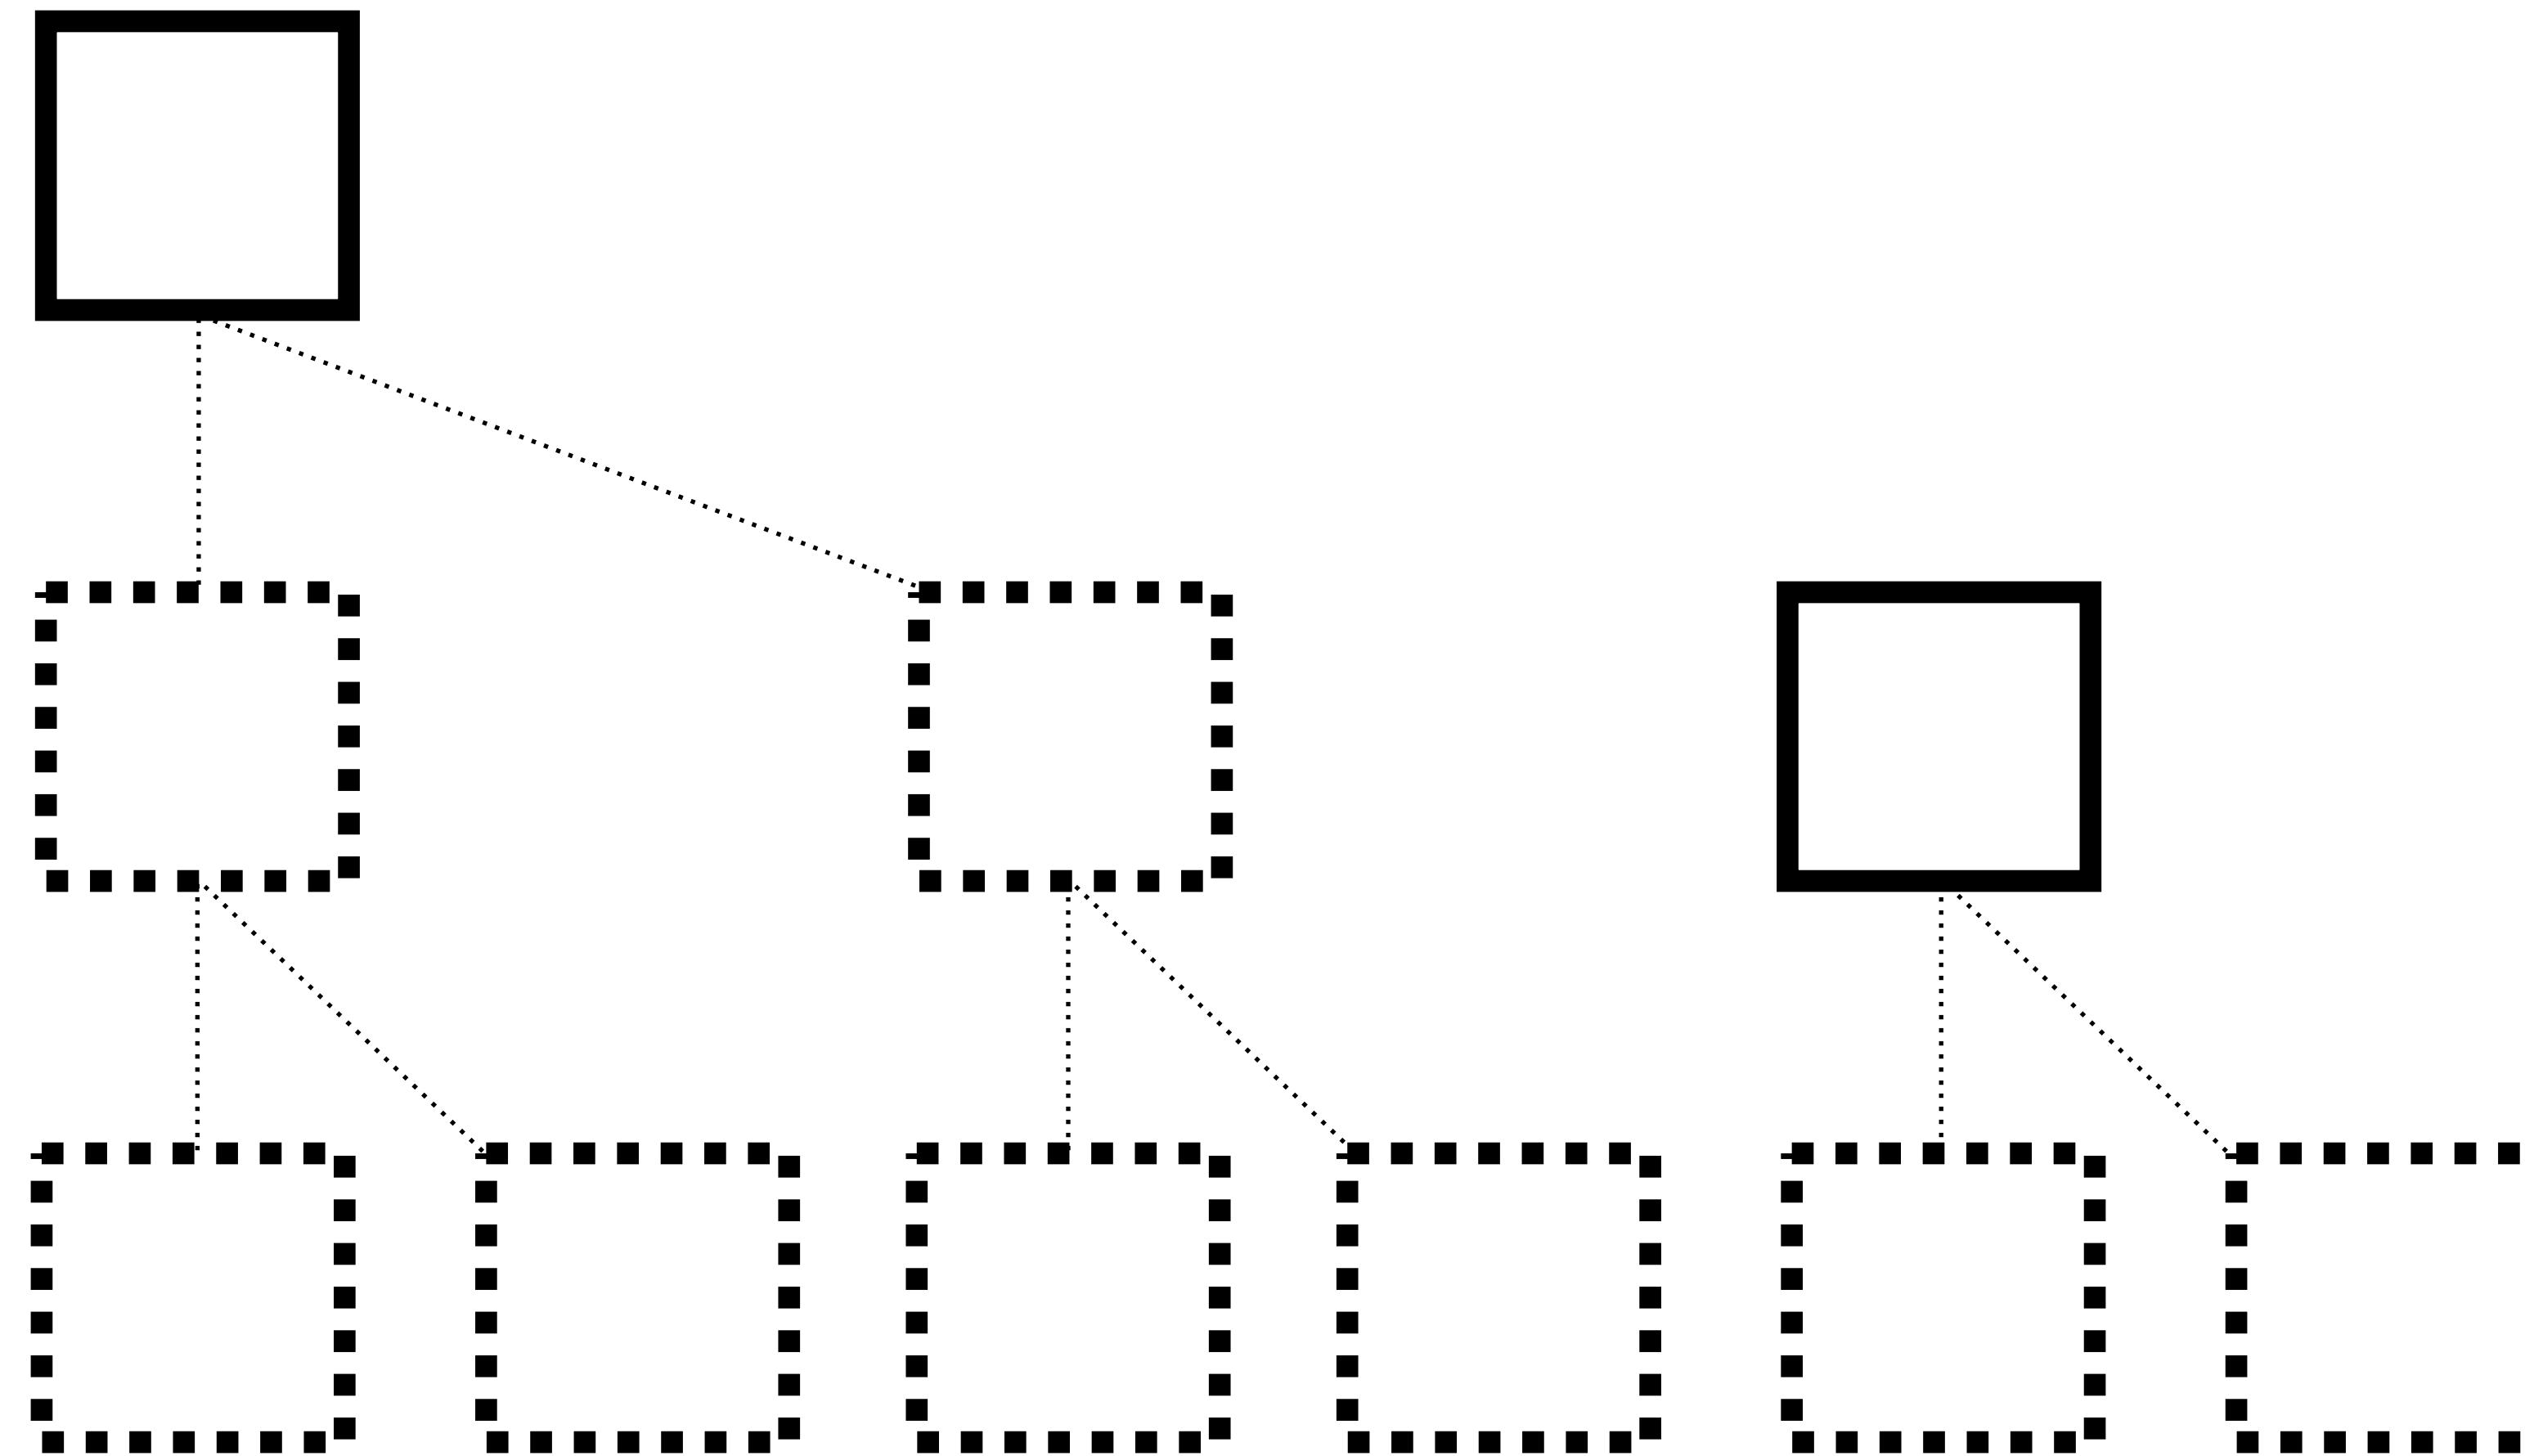
6 UTXOs

Hash



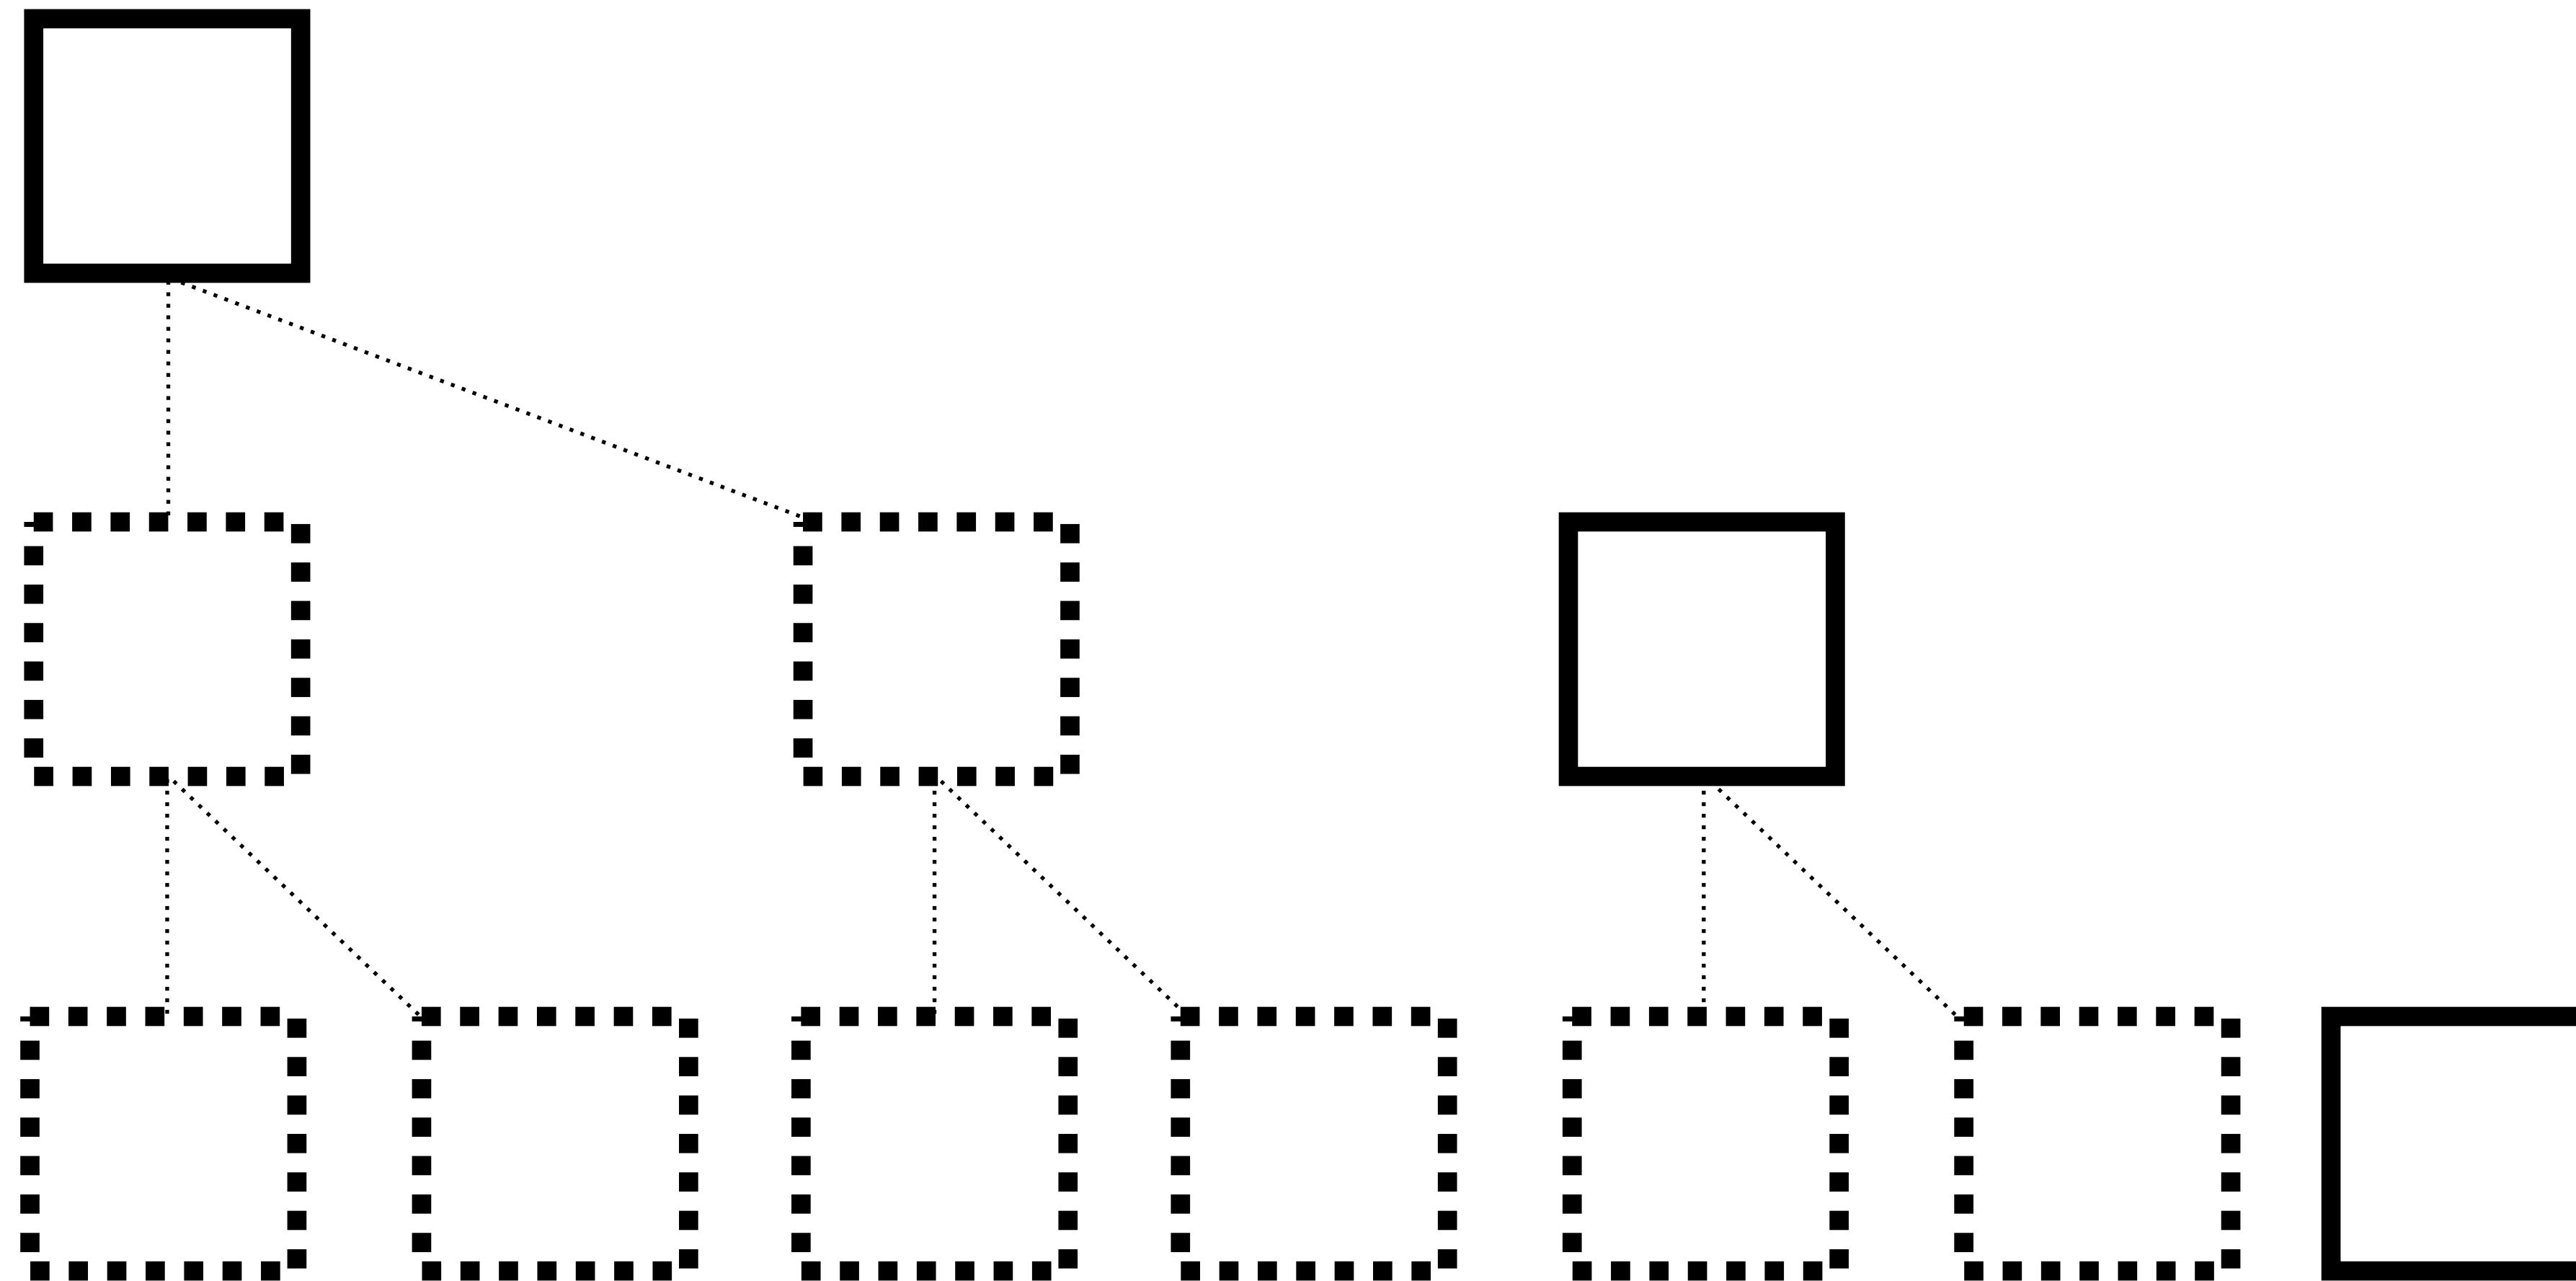
6 UTXOs

Throw away



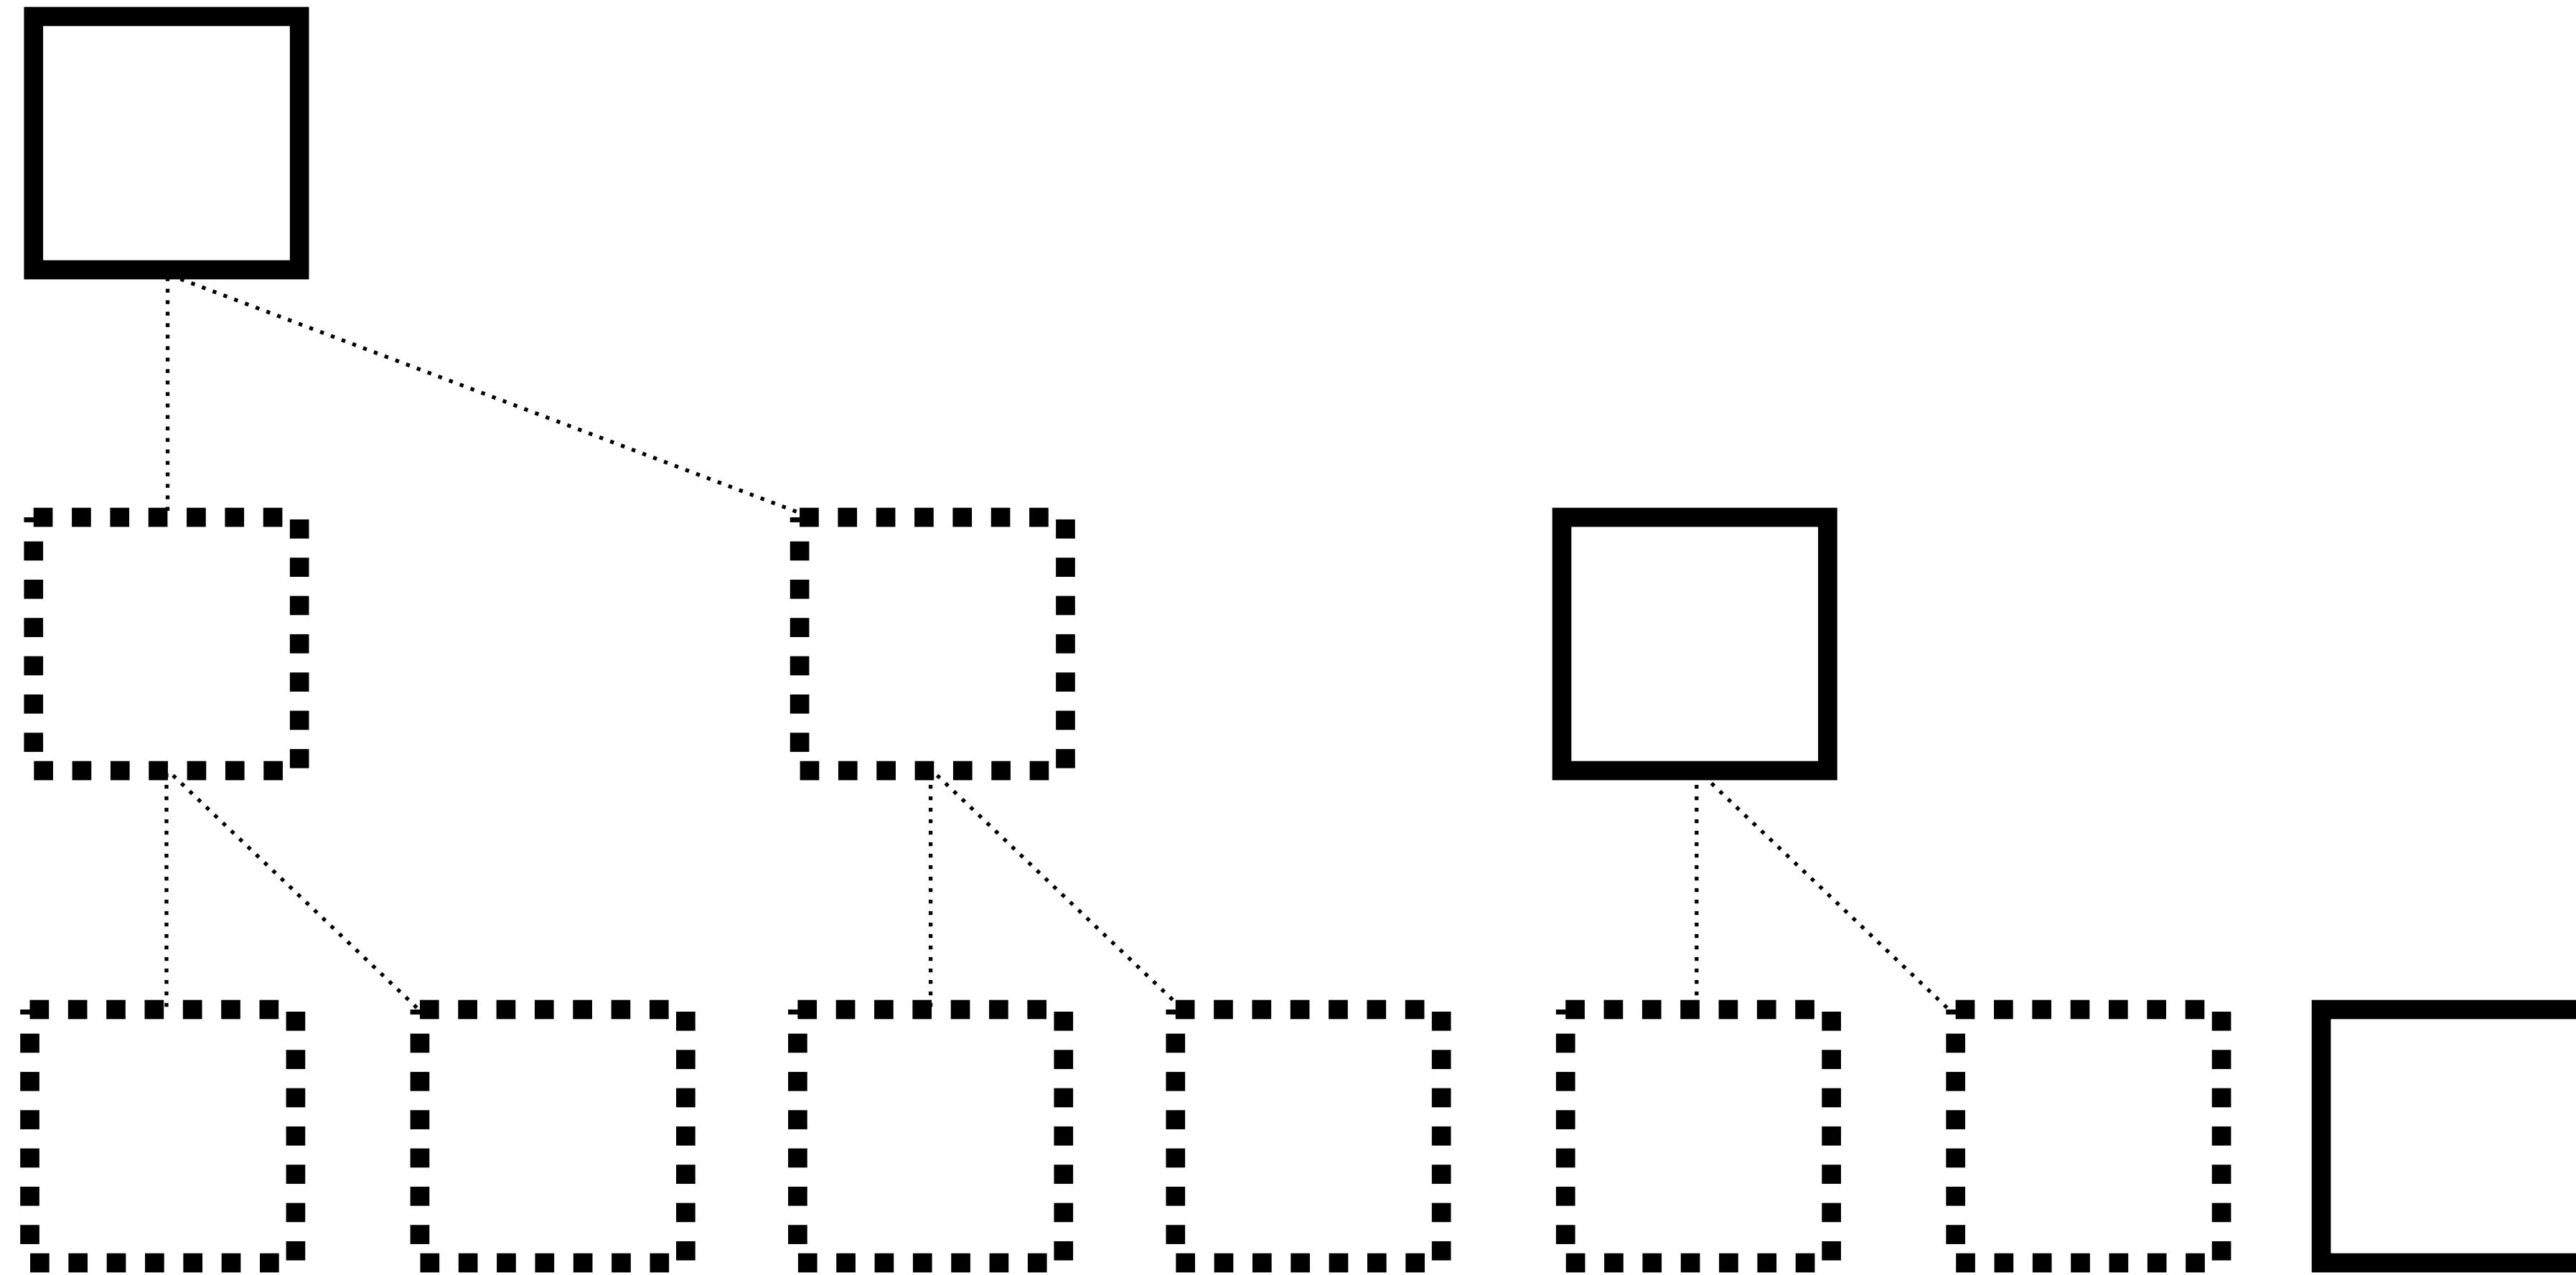
7 UTXOs

Another one



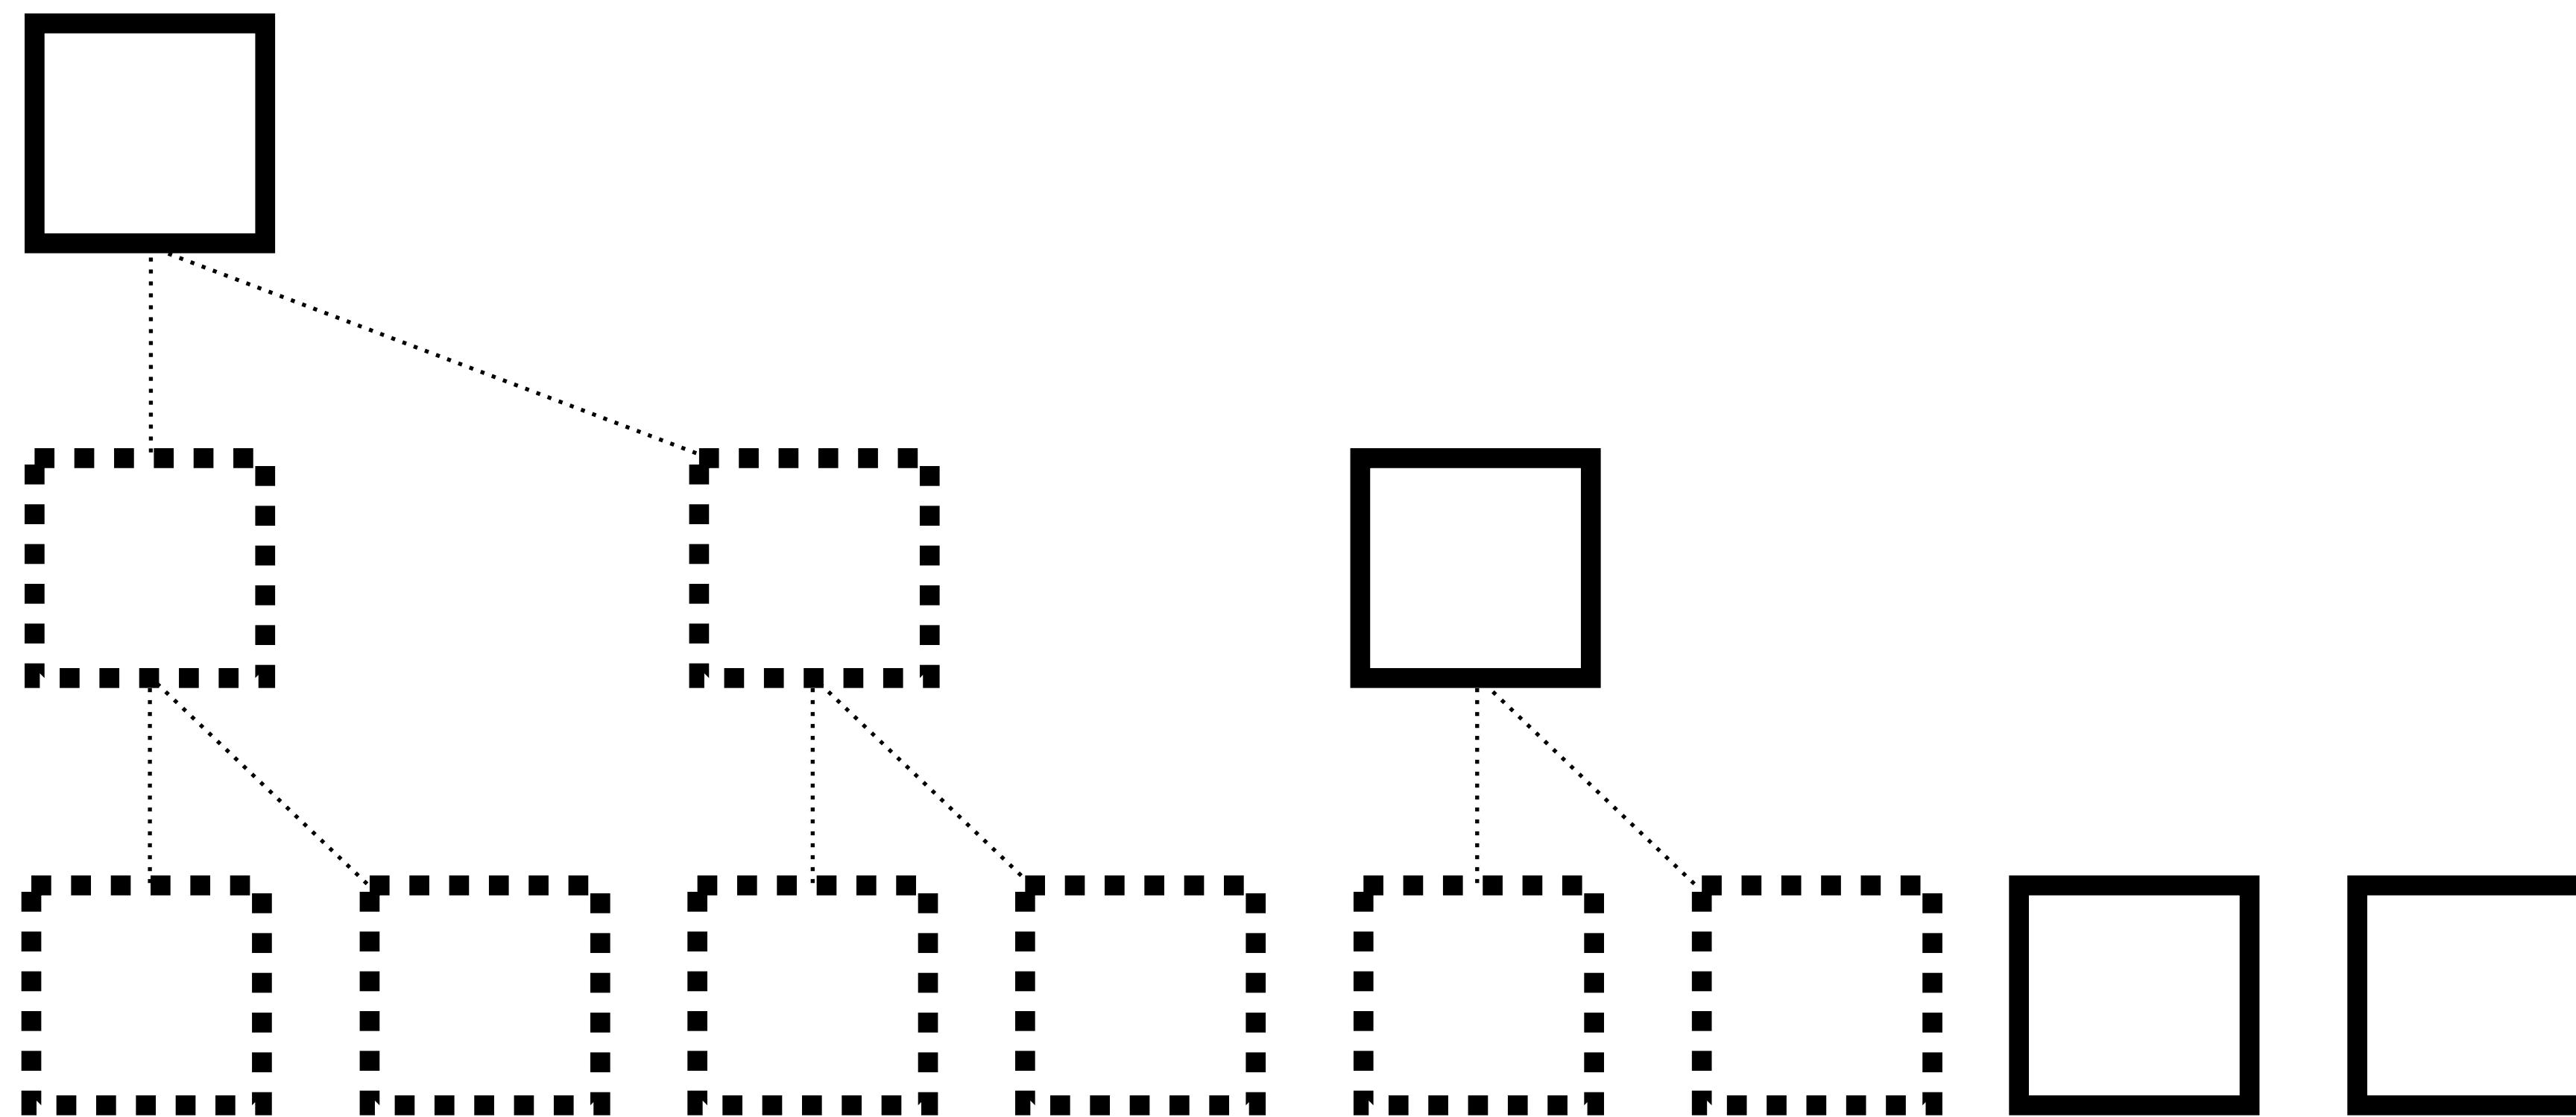
7 UTXOs

Nothing to hash with so it becomes a root



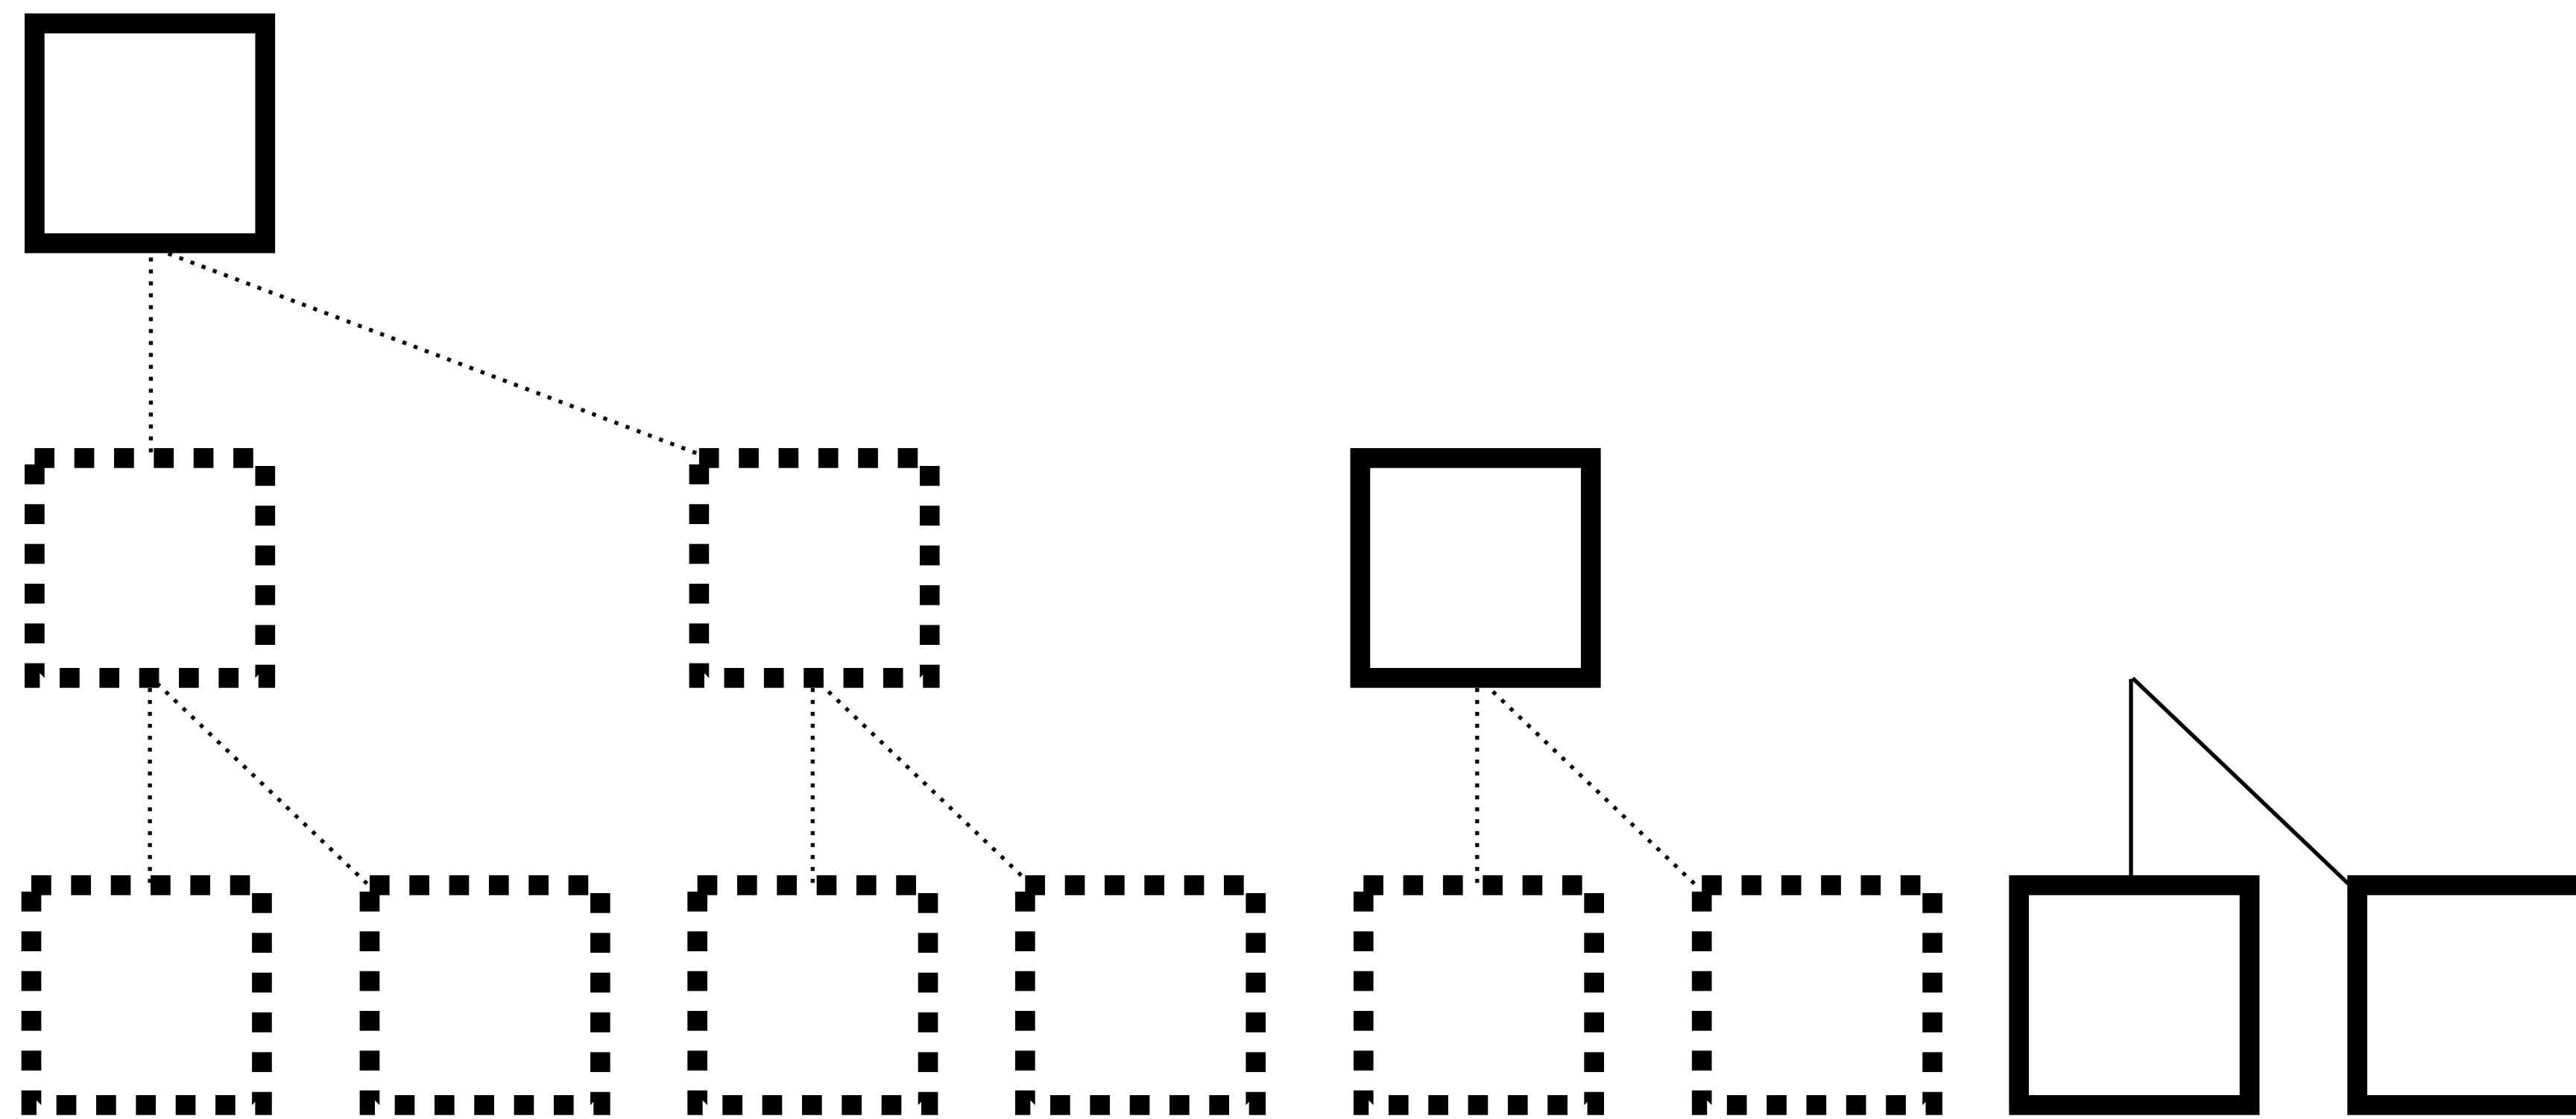
8 UTXOs

Another one



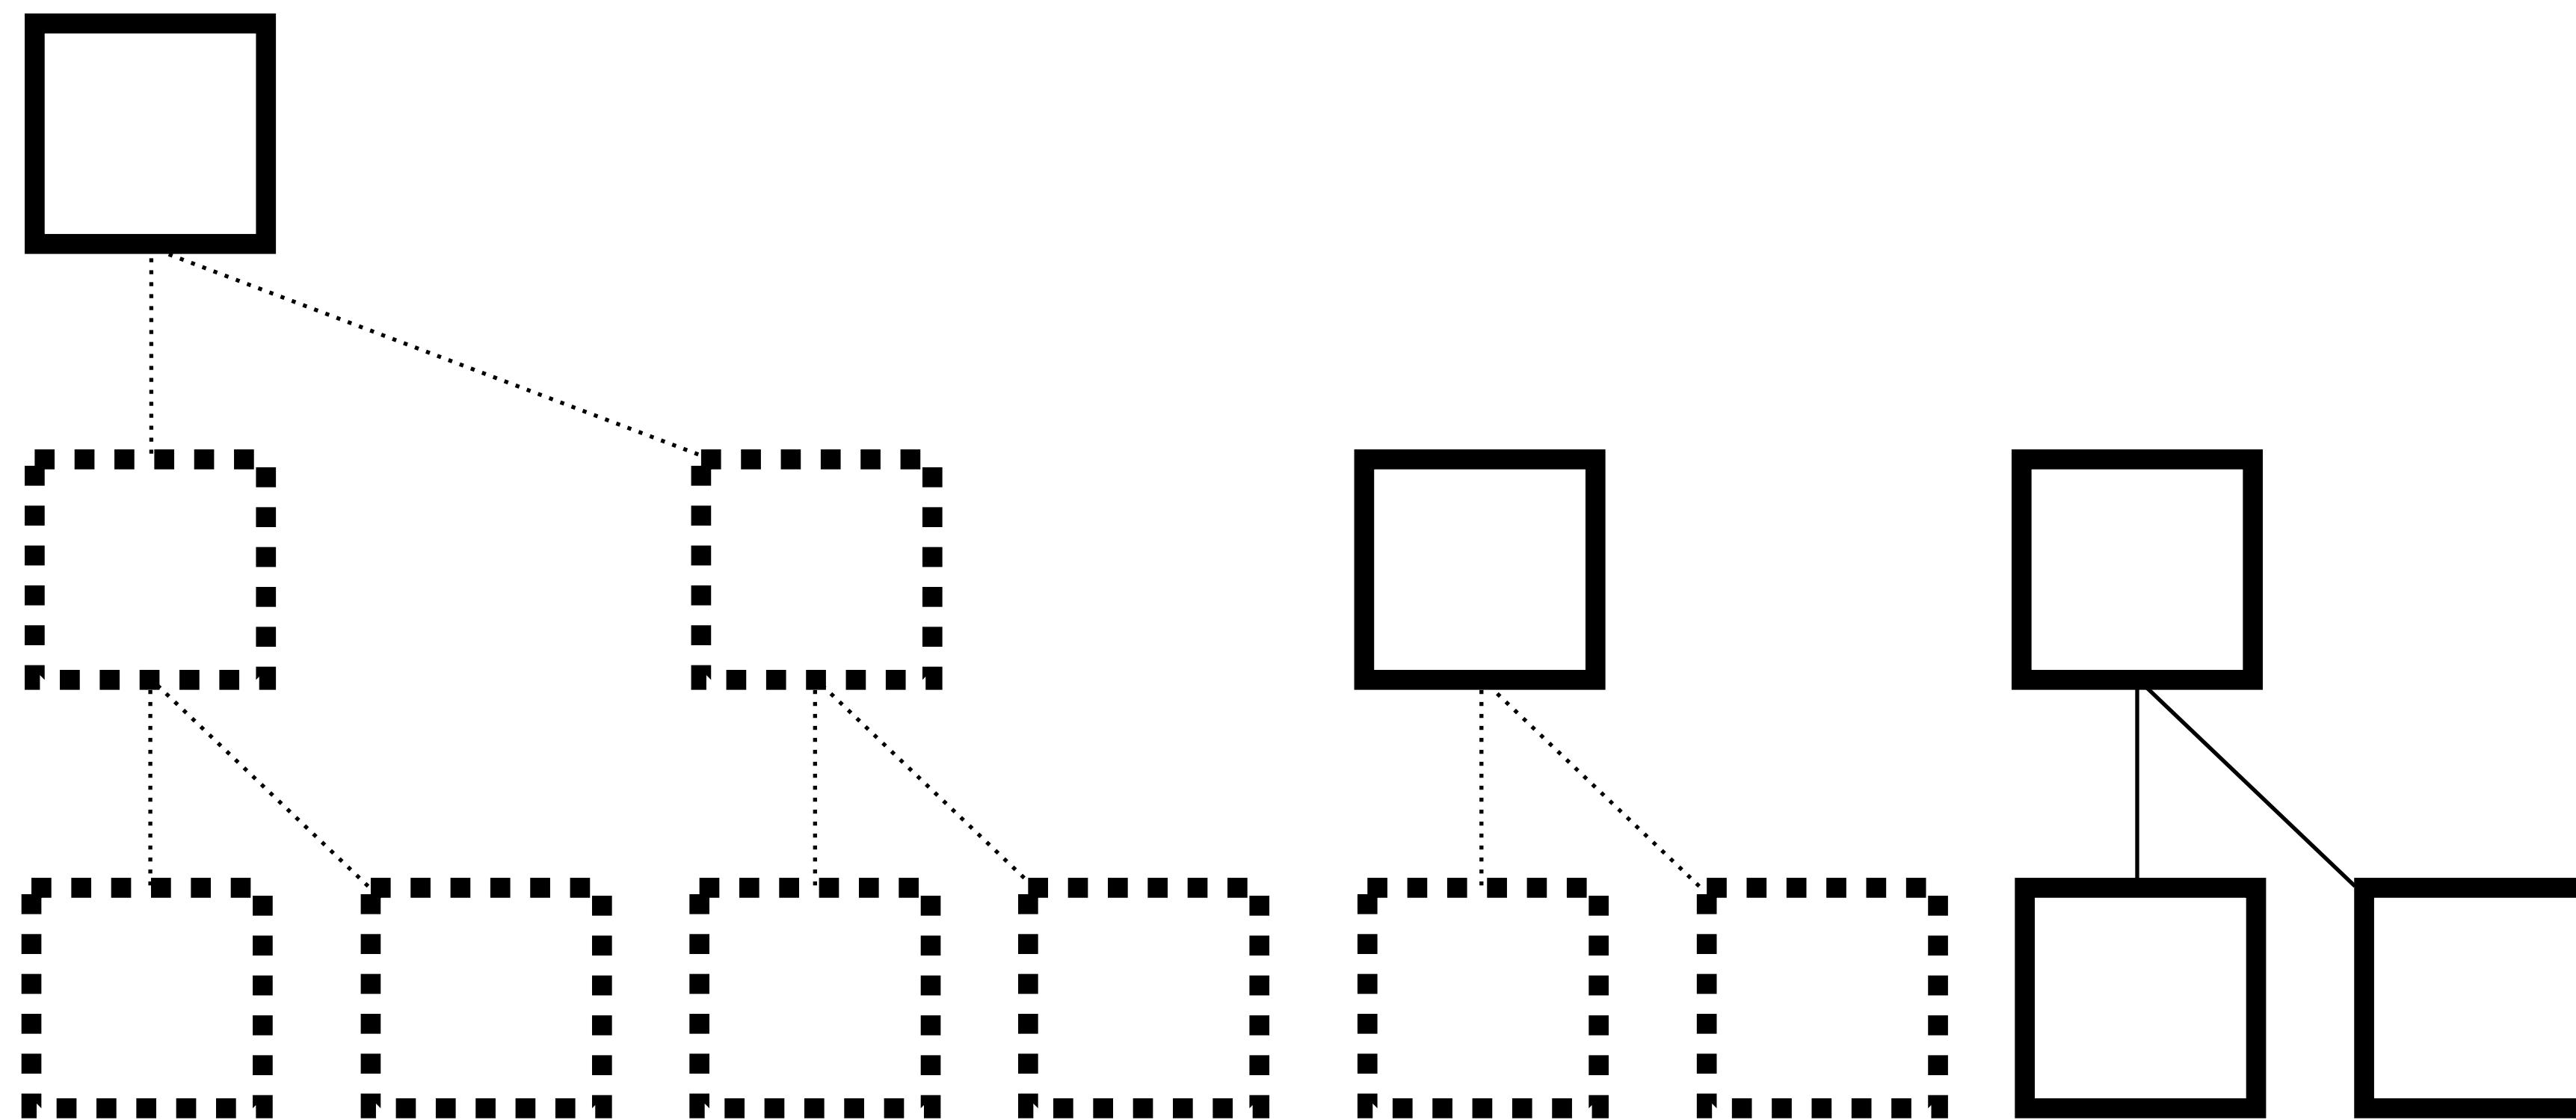
8 UTXOs

Concatenate



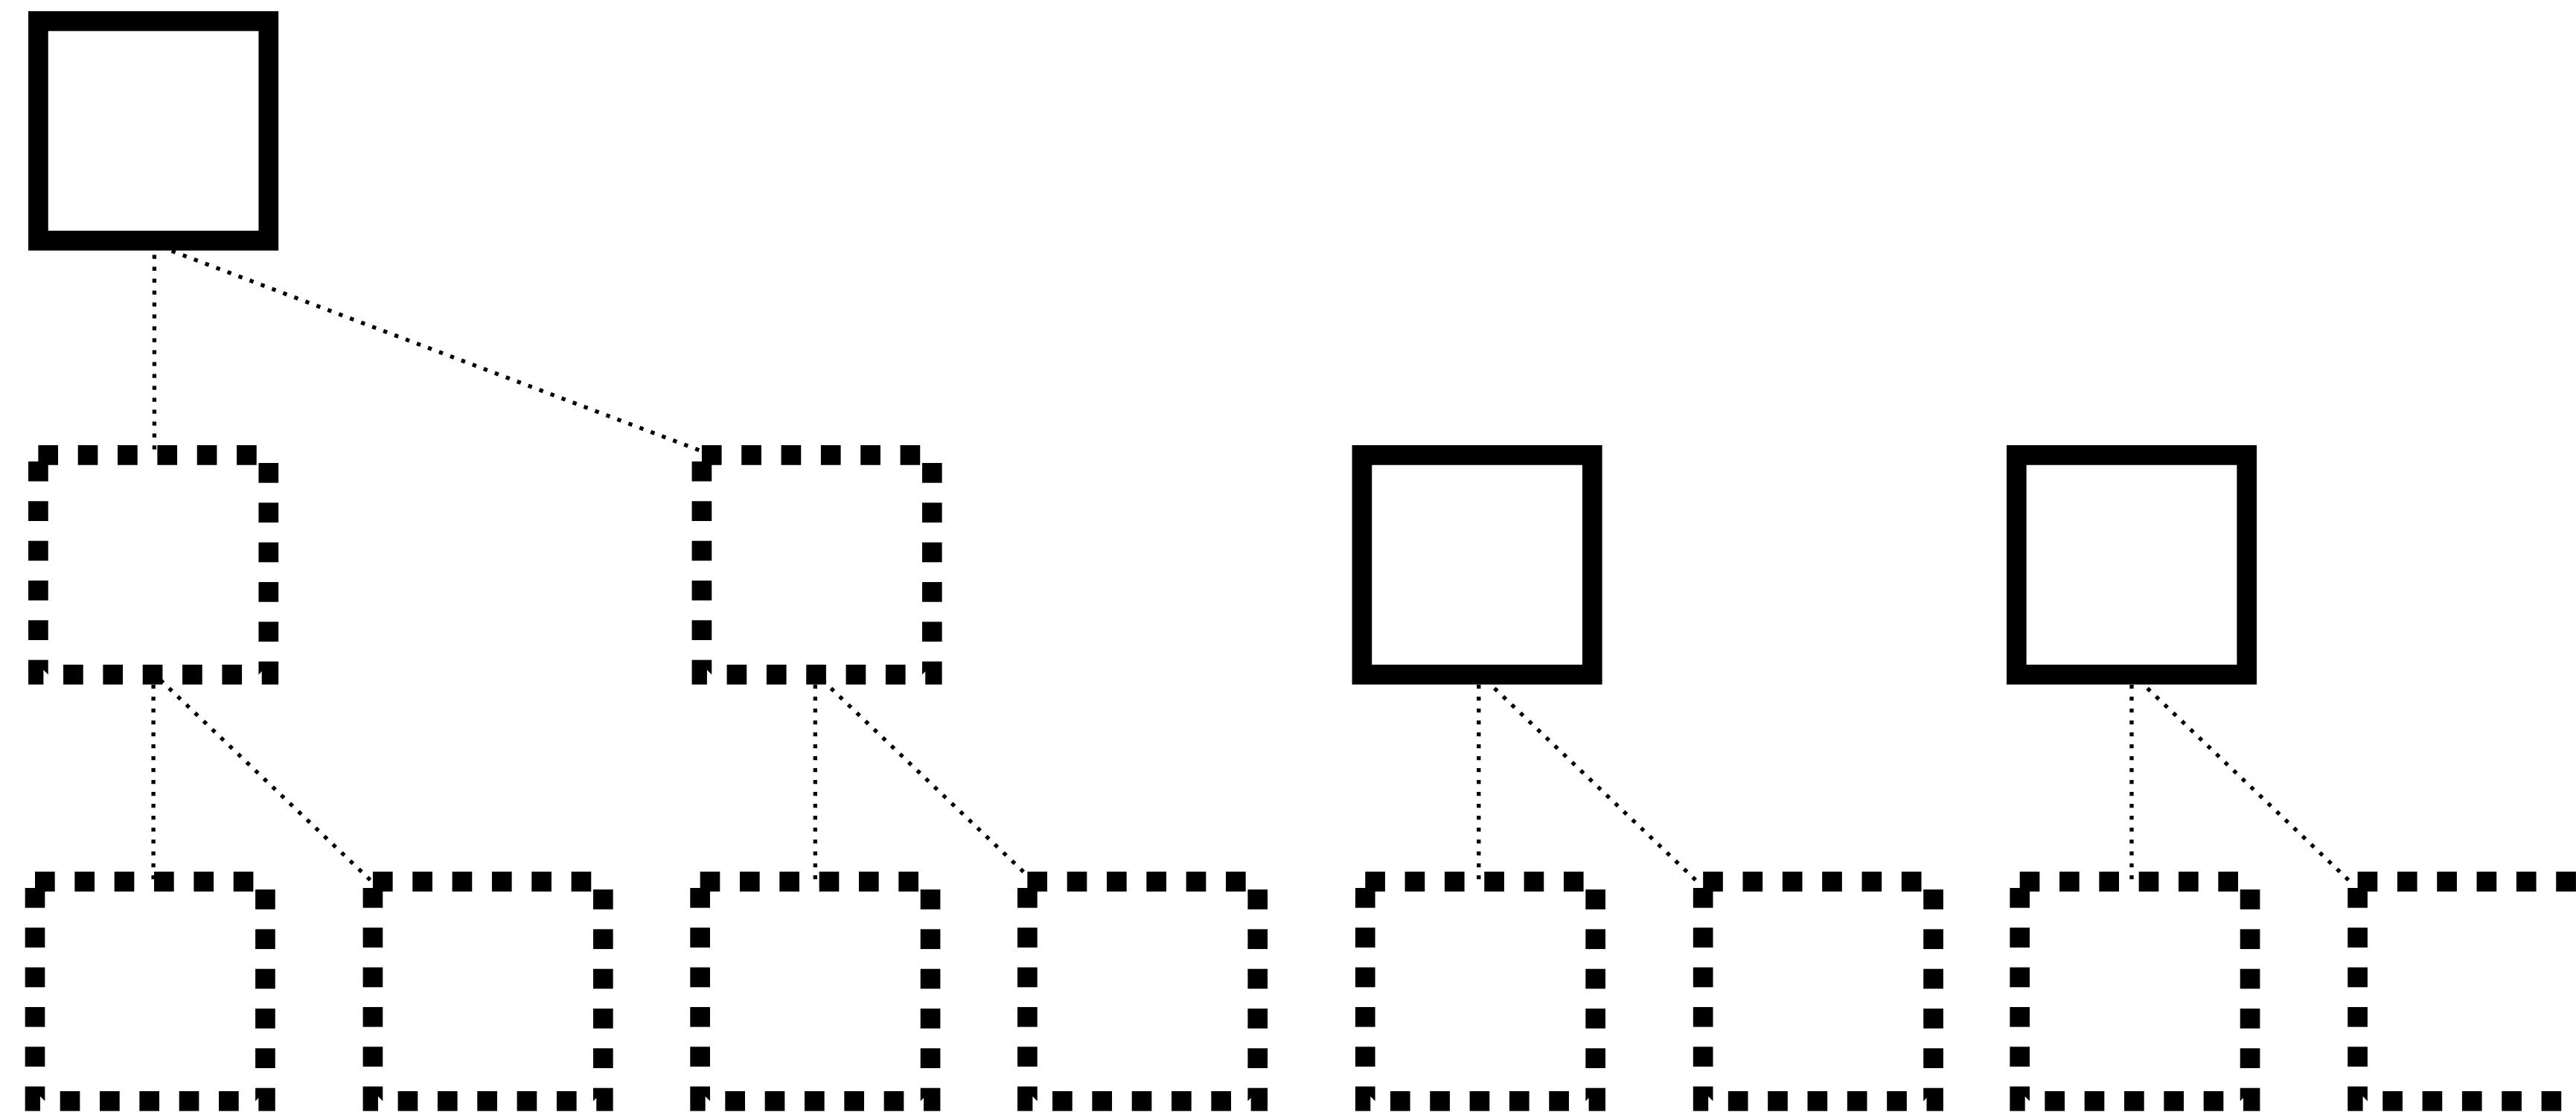
8 UTXOs

Hash



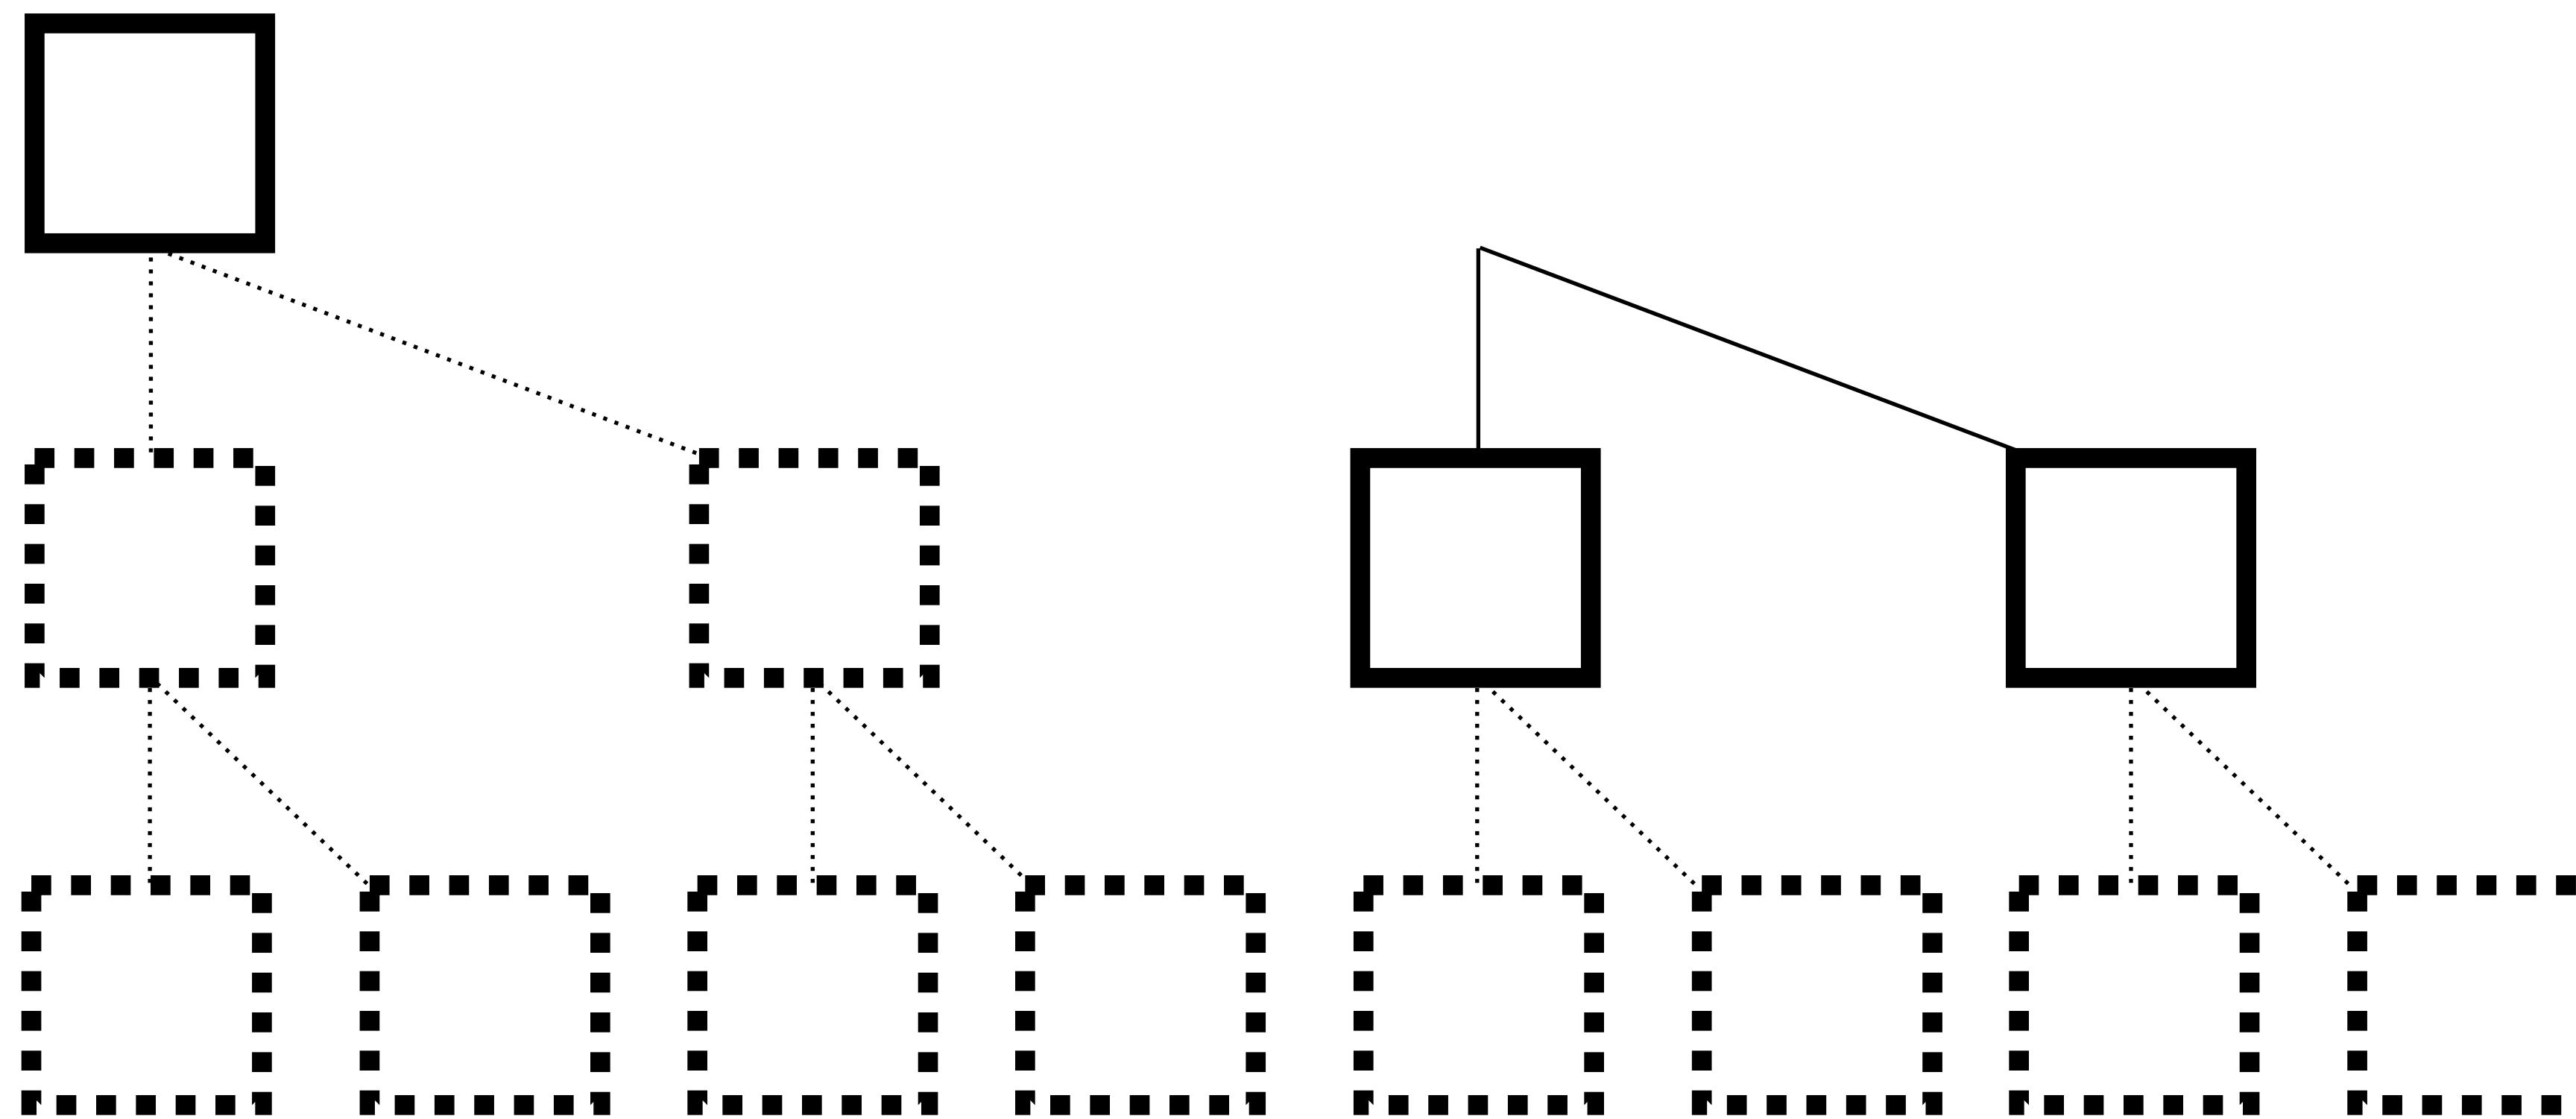
8 UTXOs

Throw away



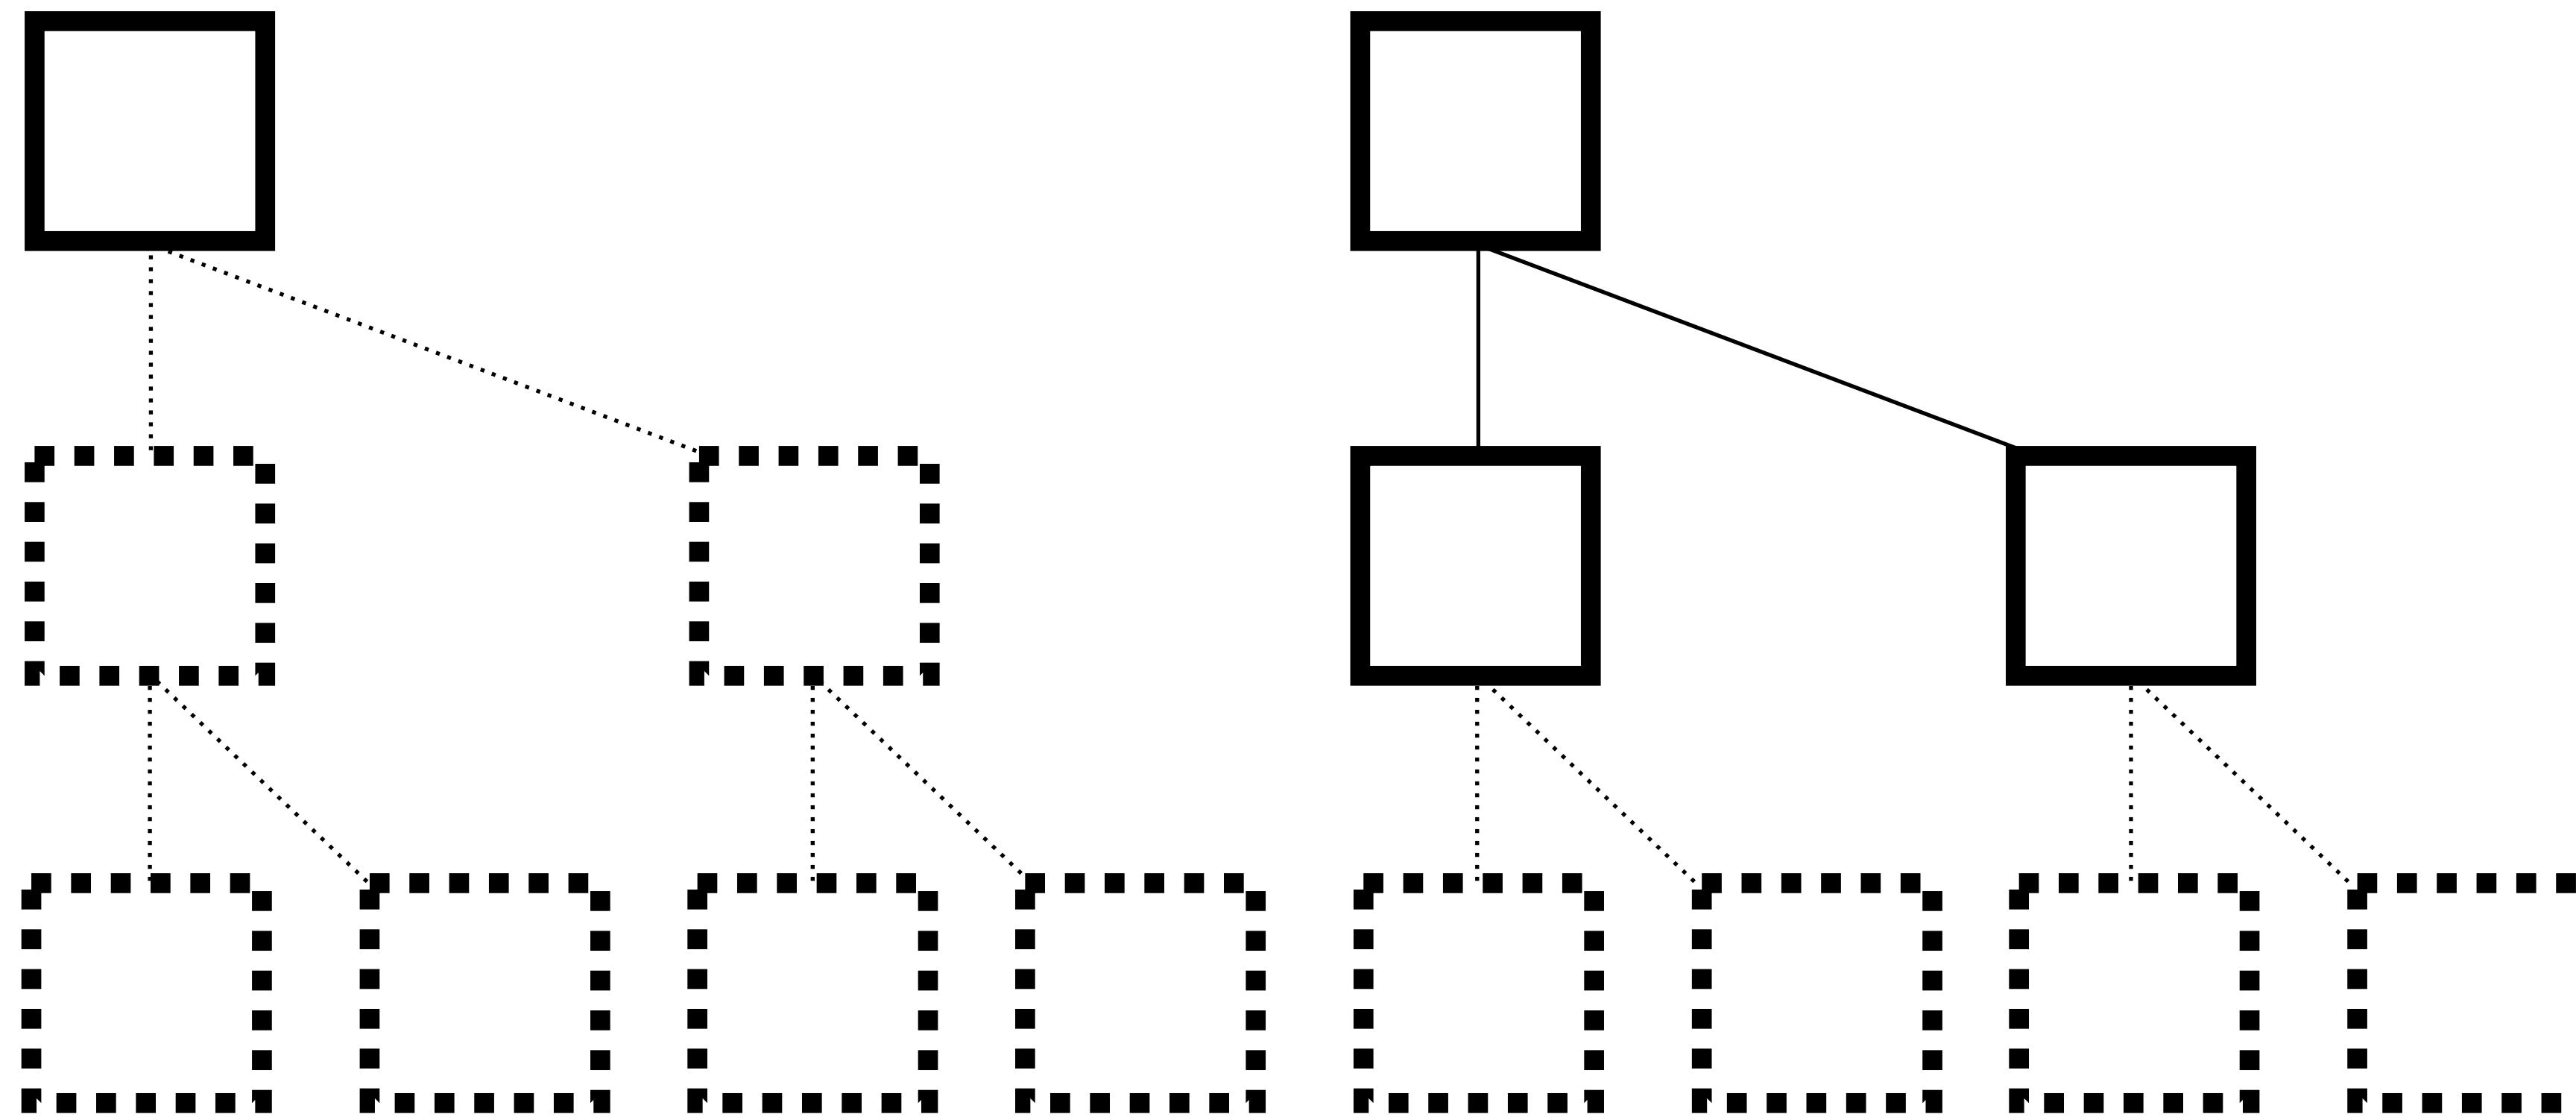
8 UTXOs

Concatenate



8 UTXOs

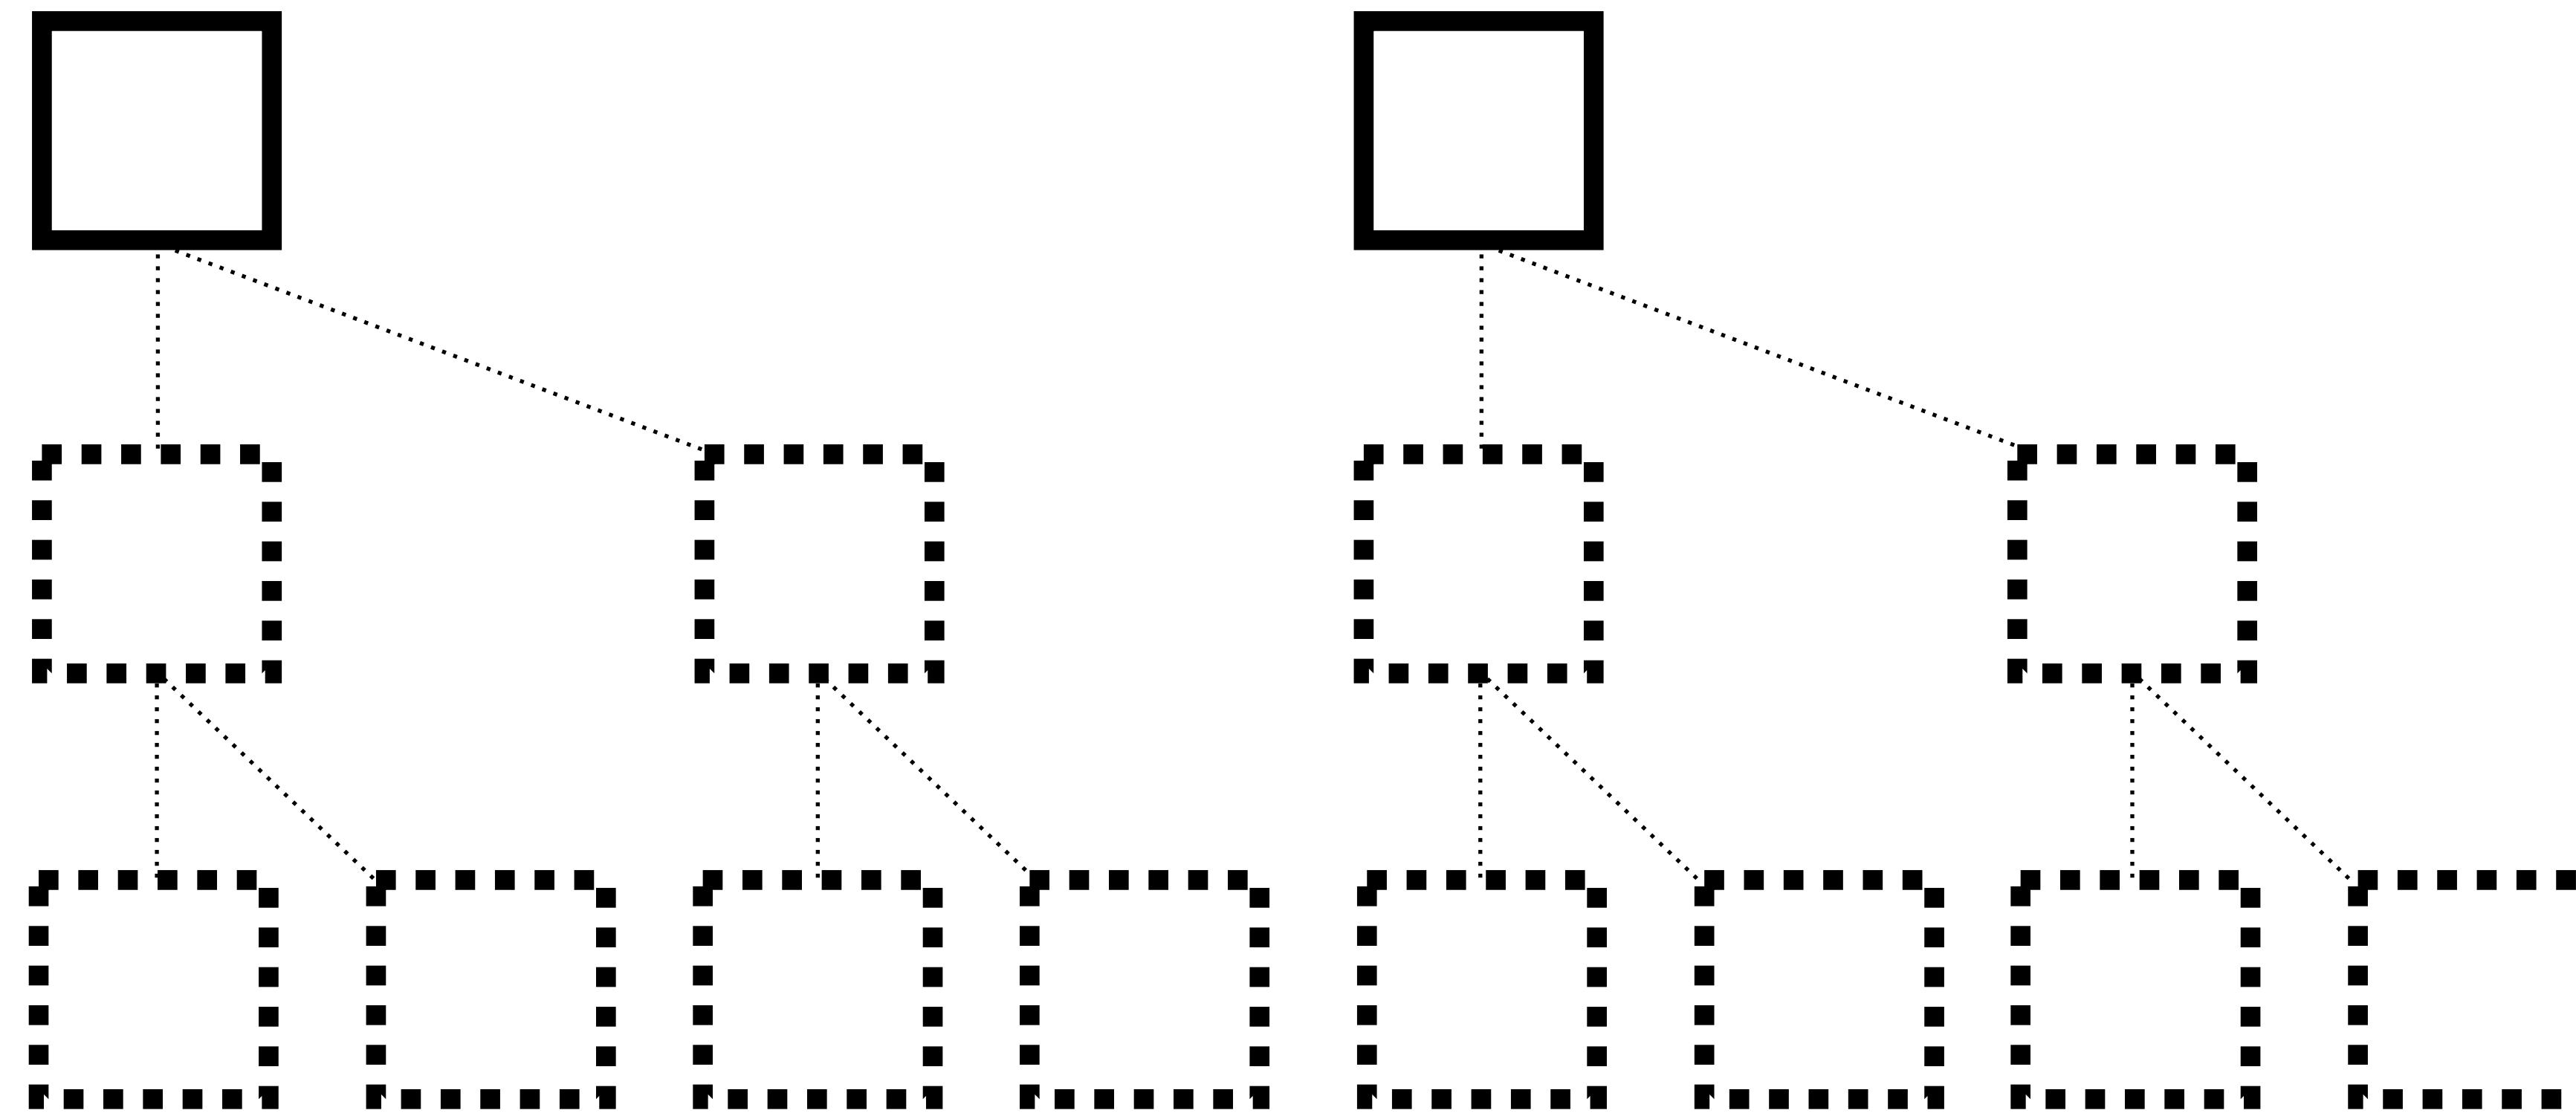
Hash



<https://gist.github.com/kcalvinalvin/a790d524832e1b7f96a70c642315fffc>

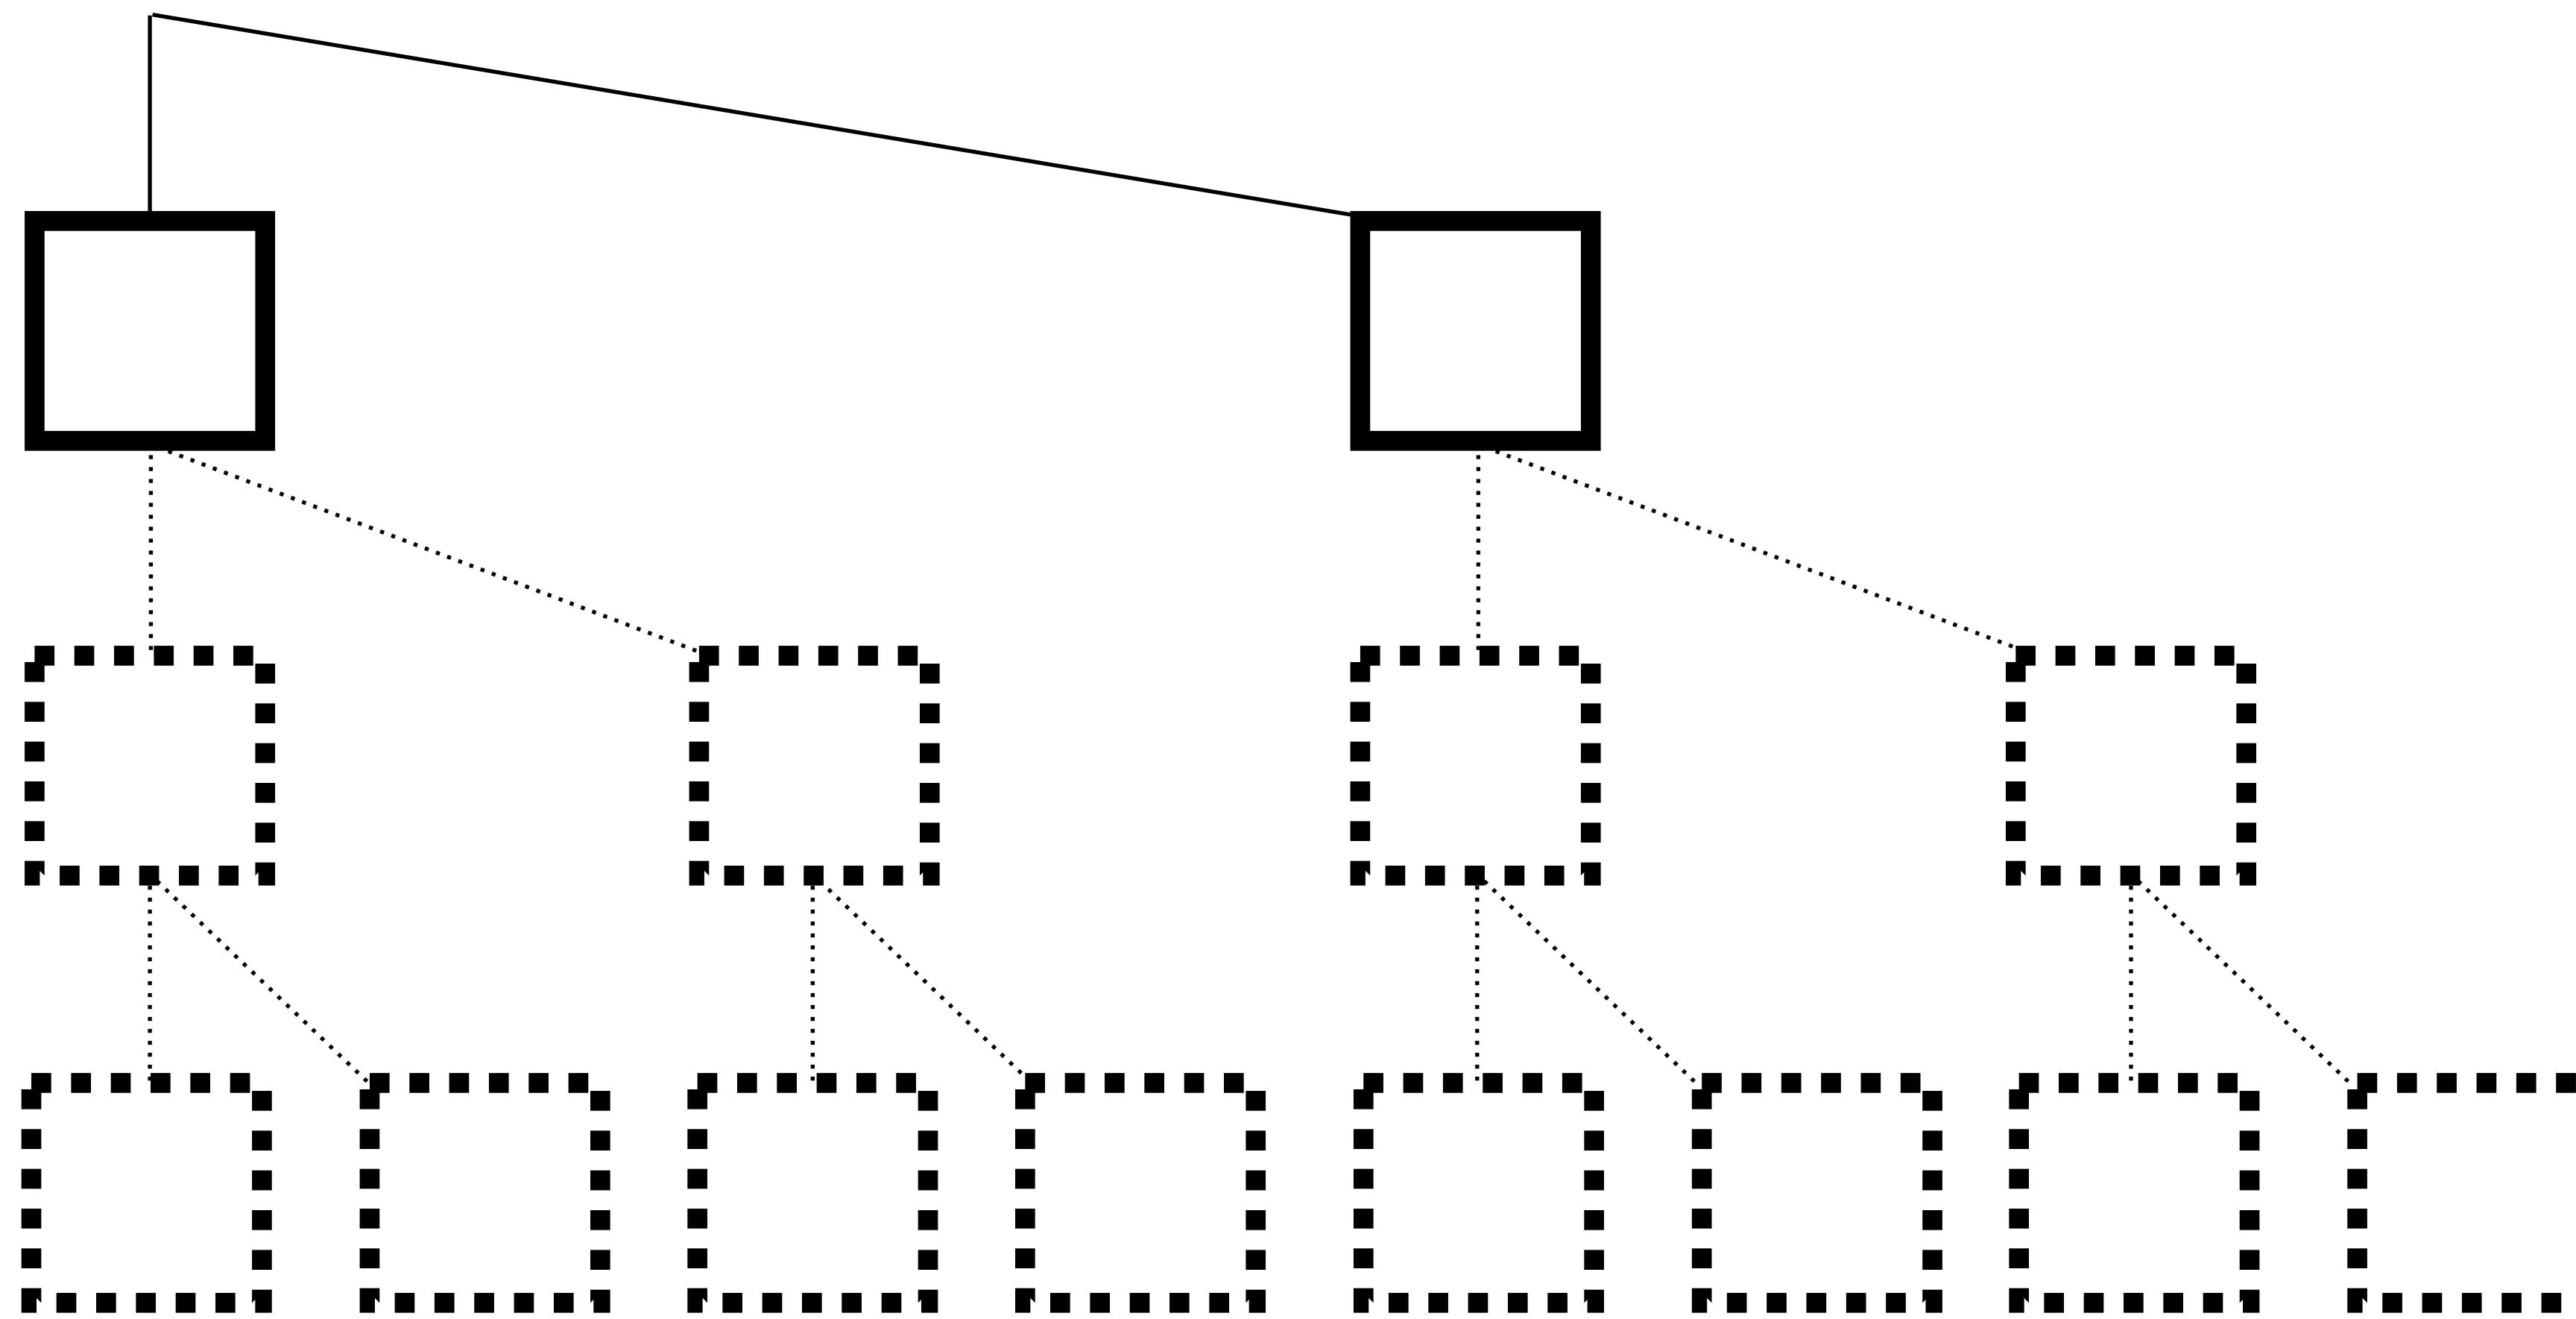
8 UTXOs

Throw away

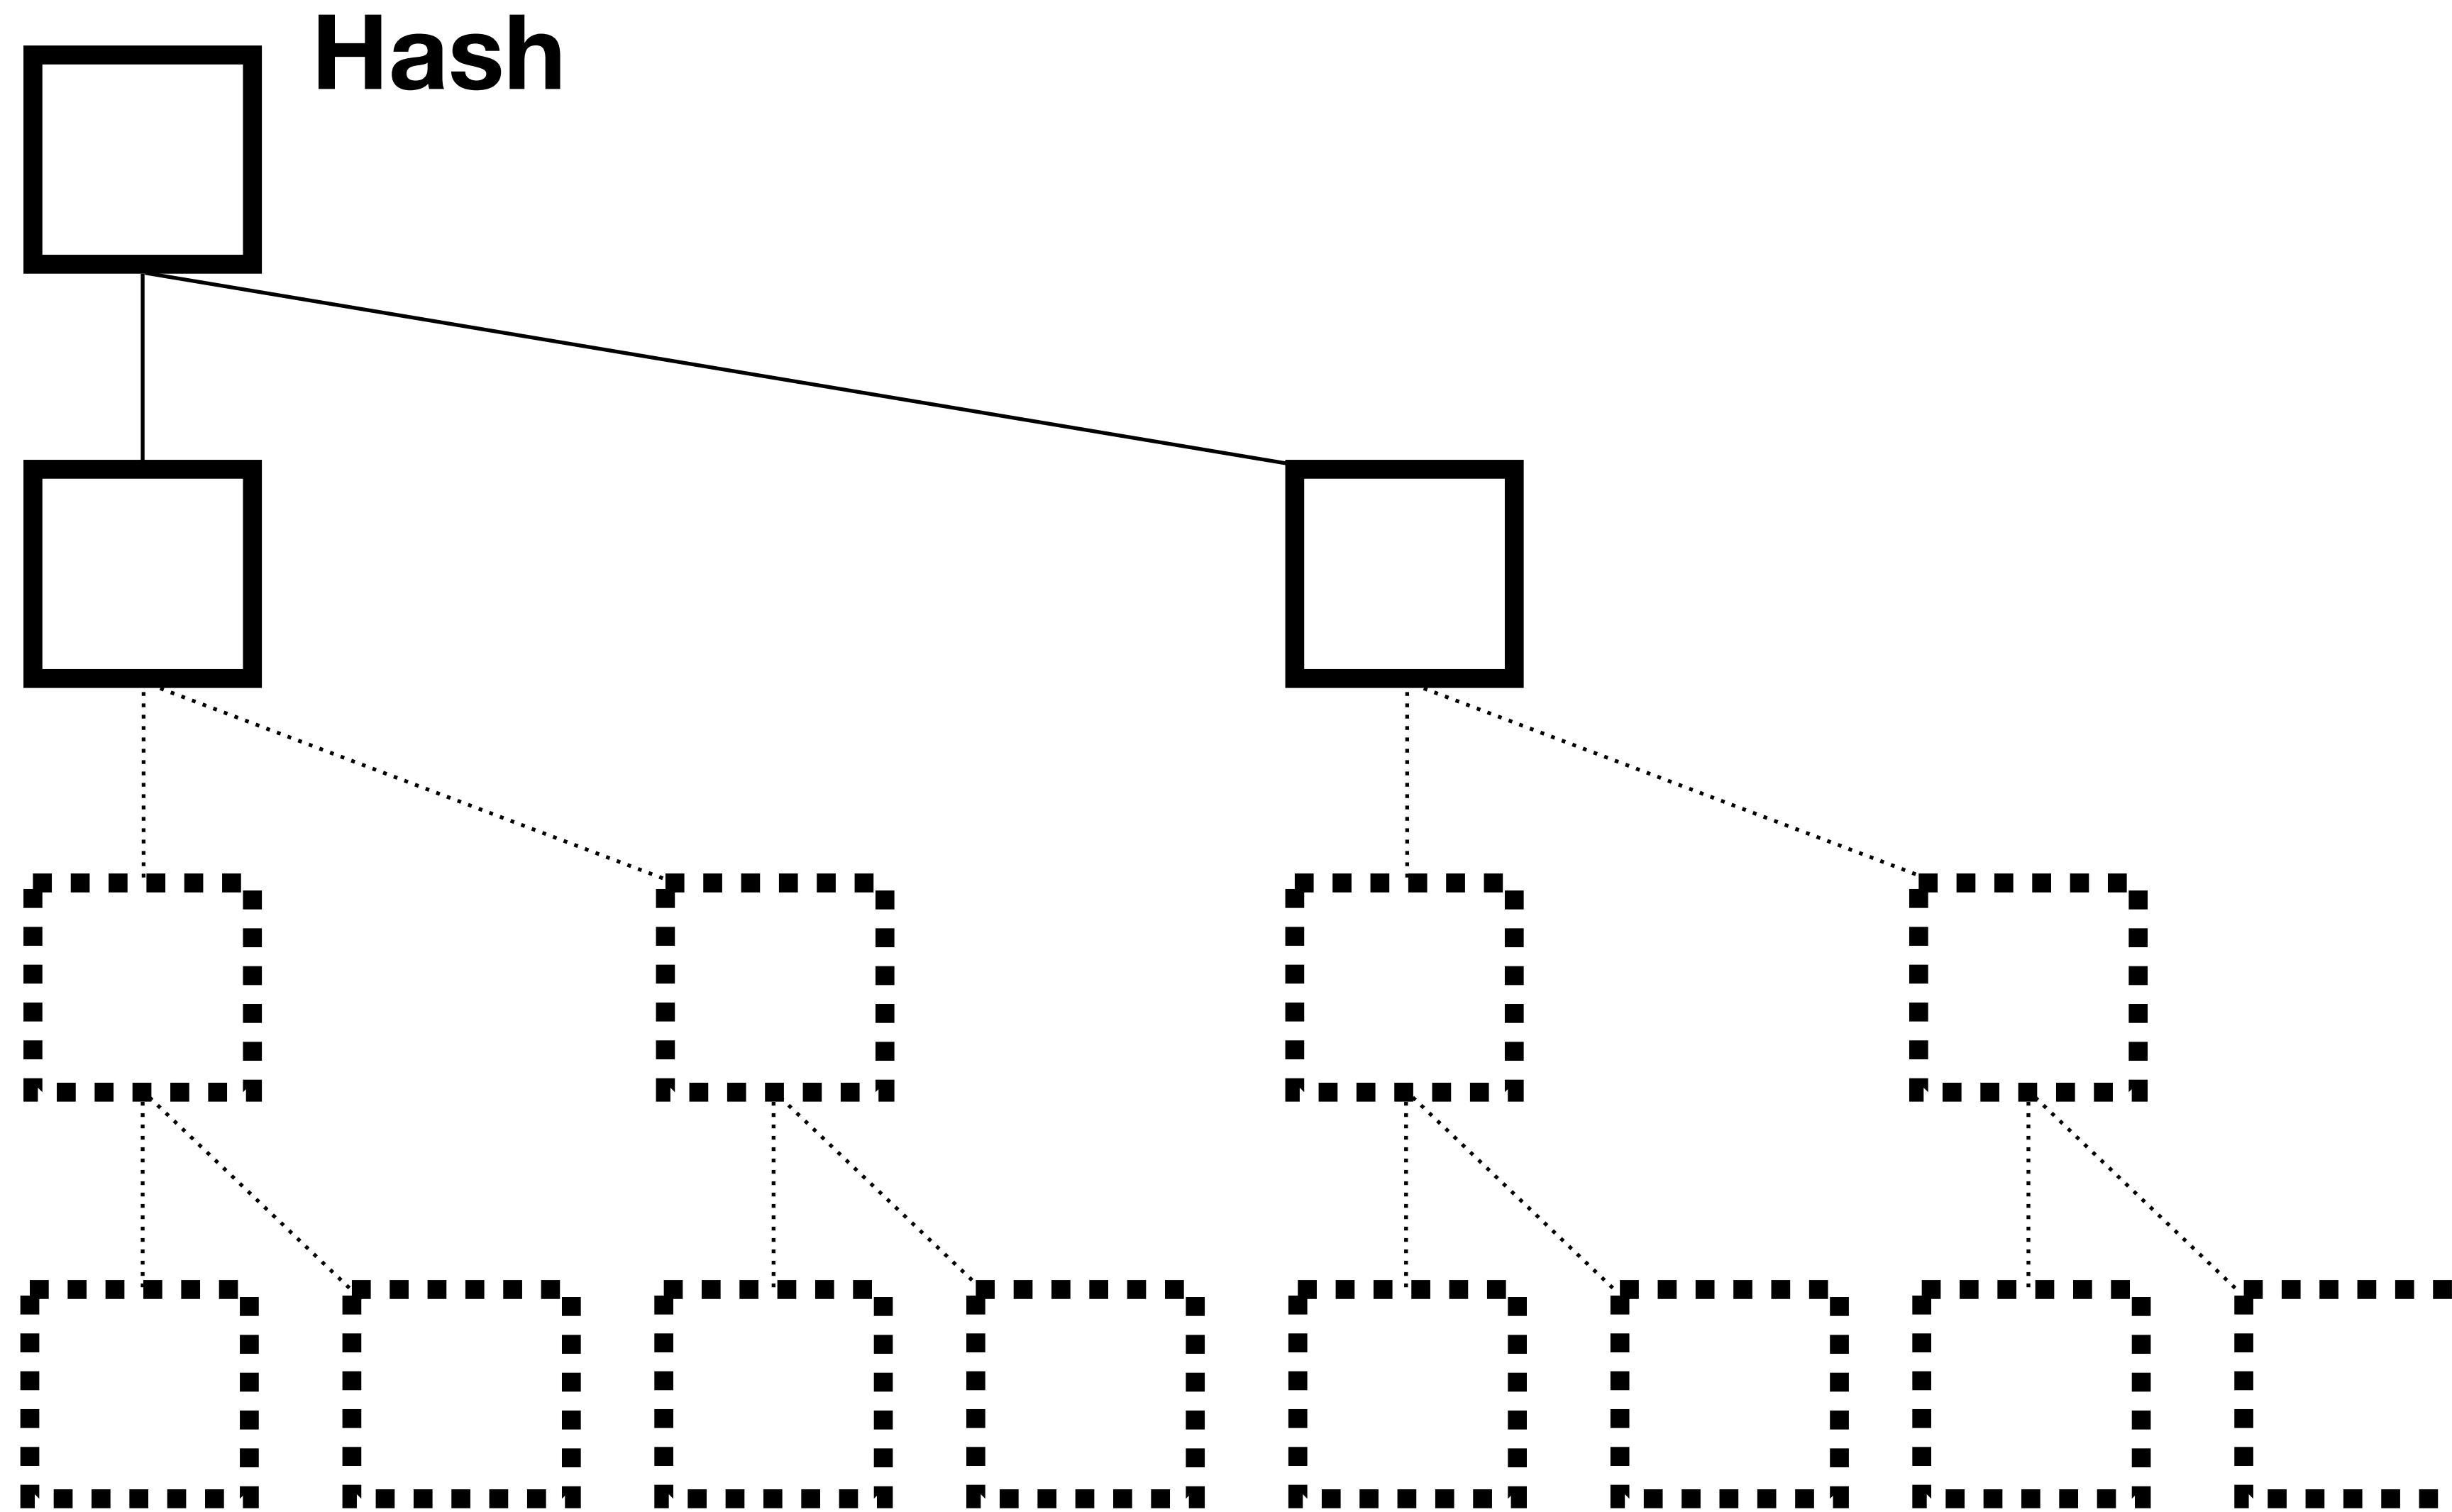


8 UTXOs

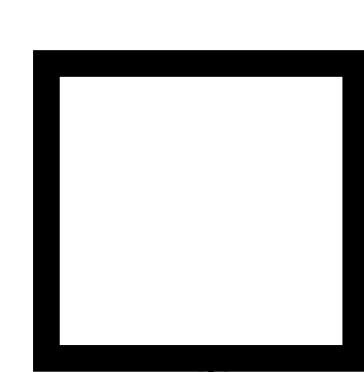
Concatenate



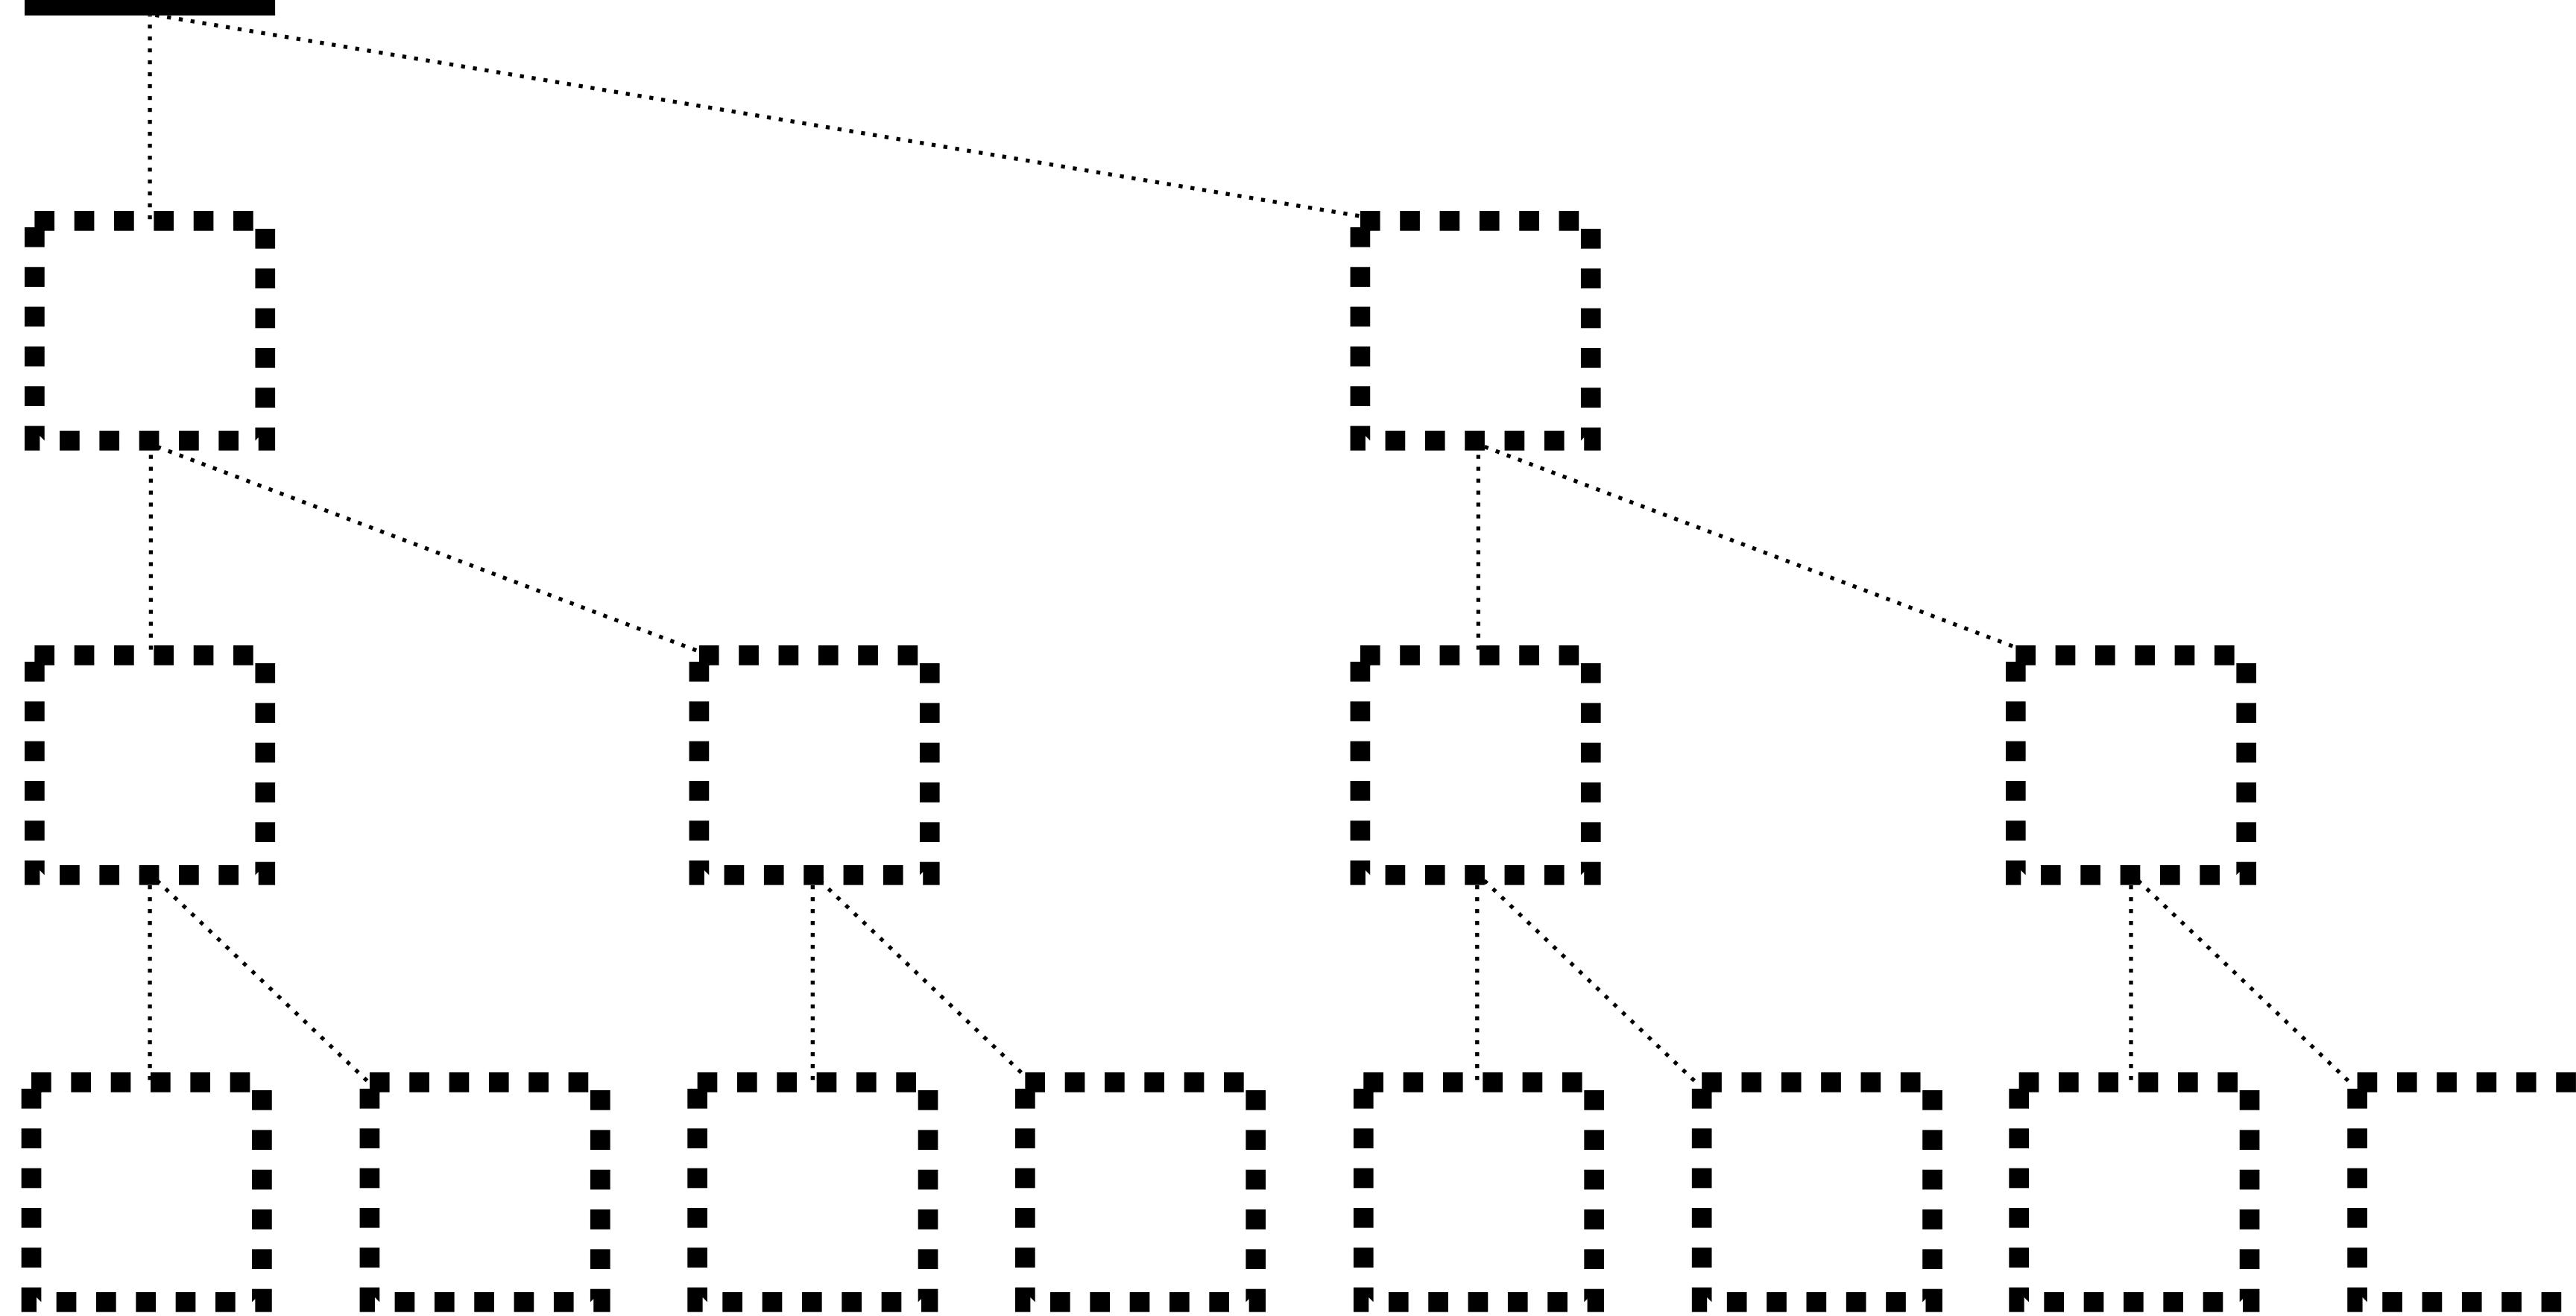
8 UTXOs



8 UTXOs



Throw away



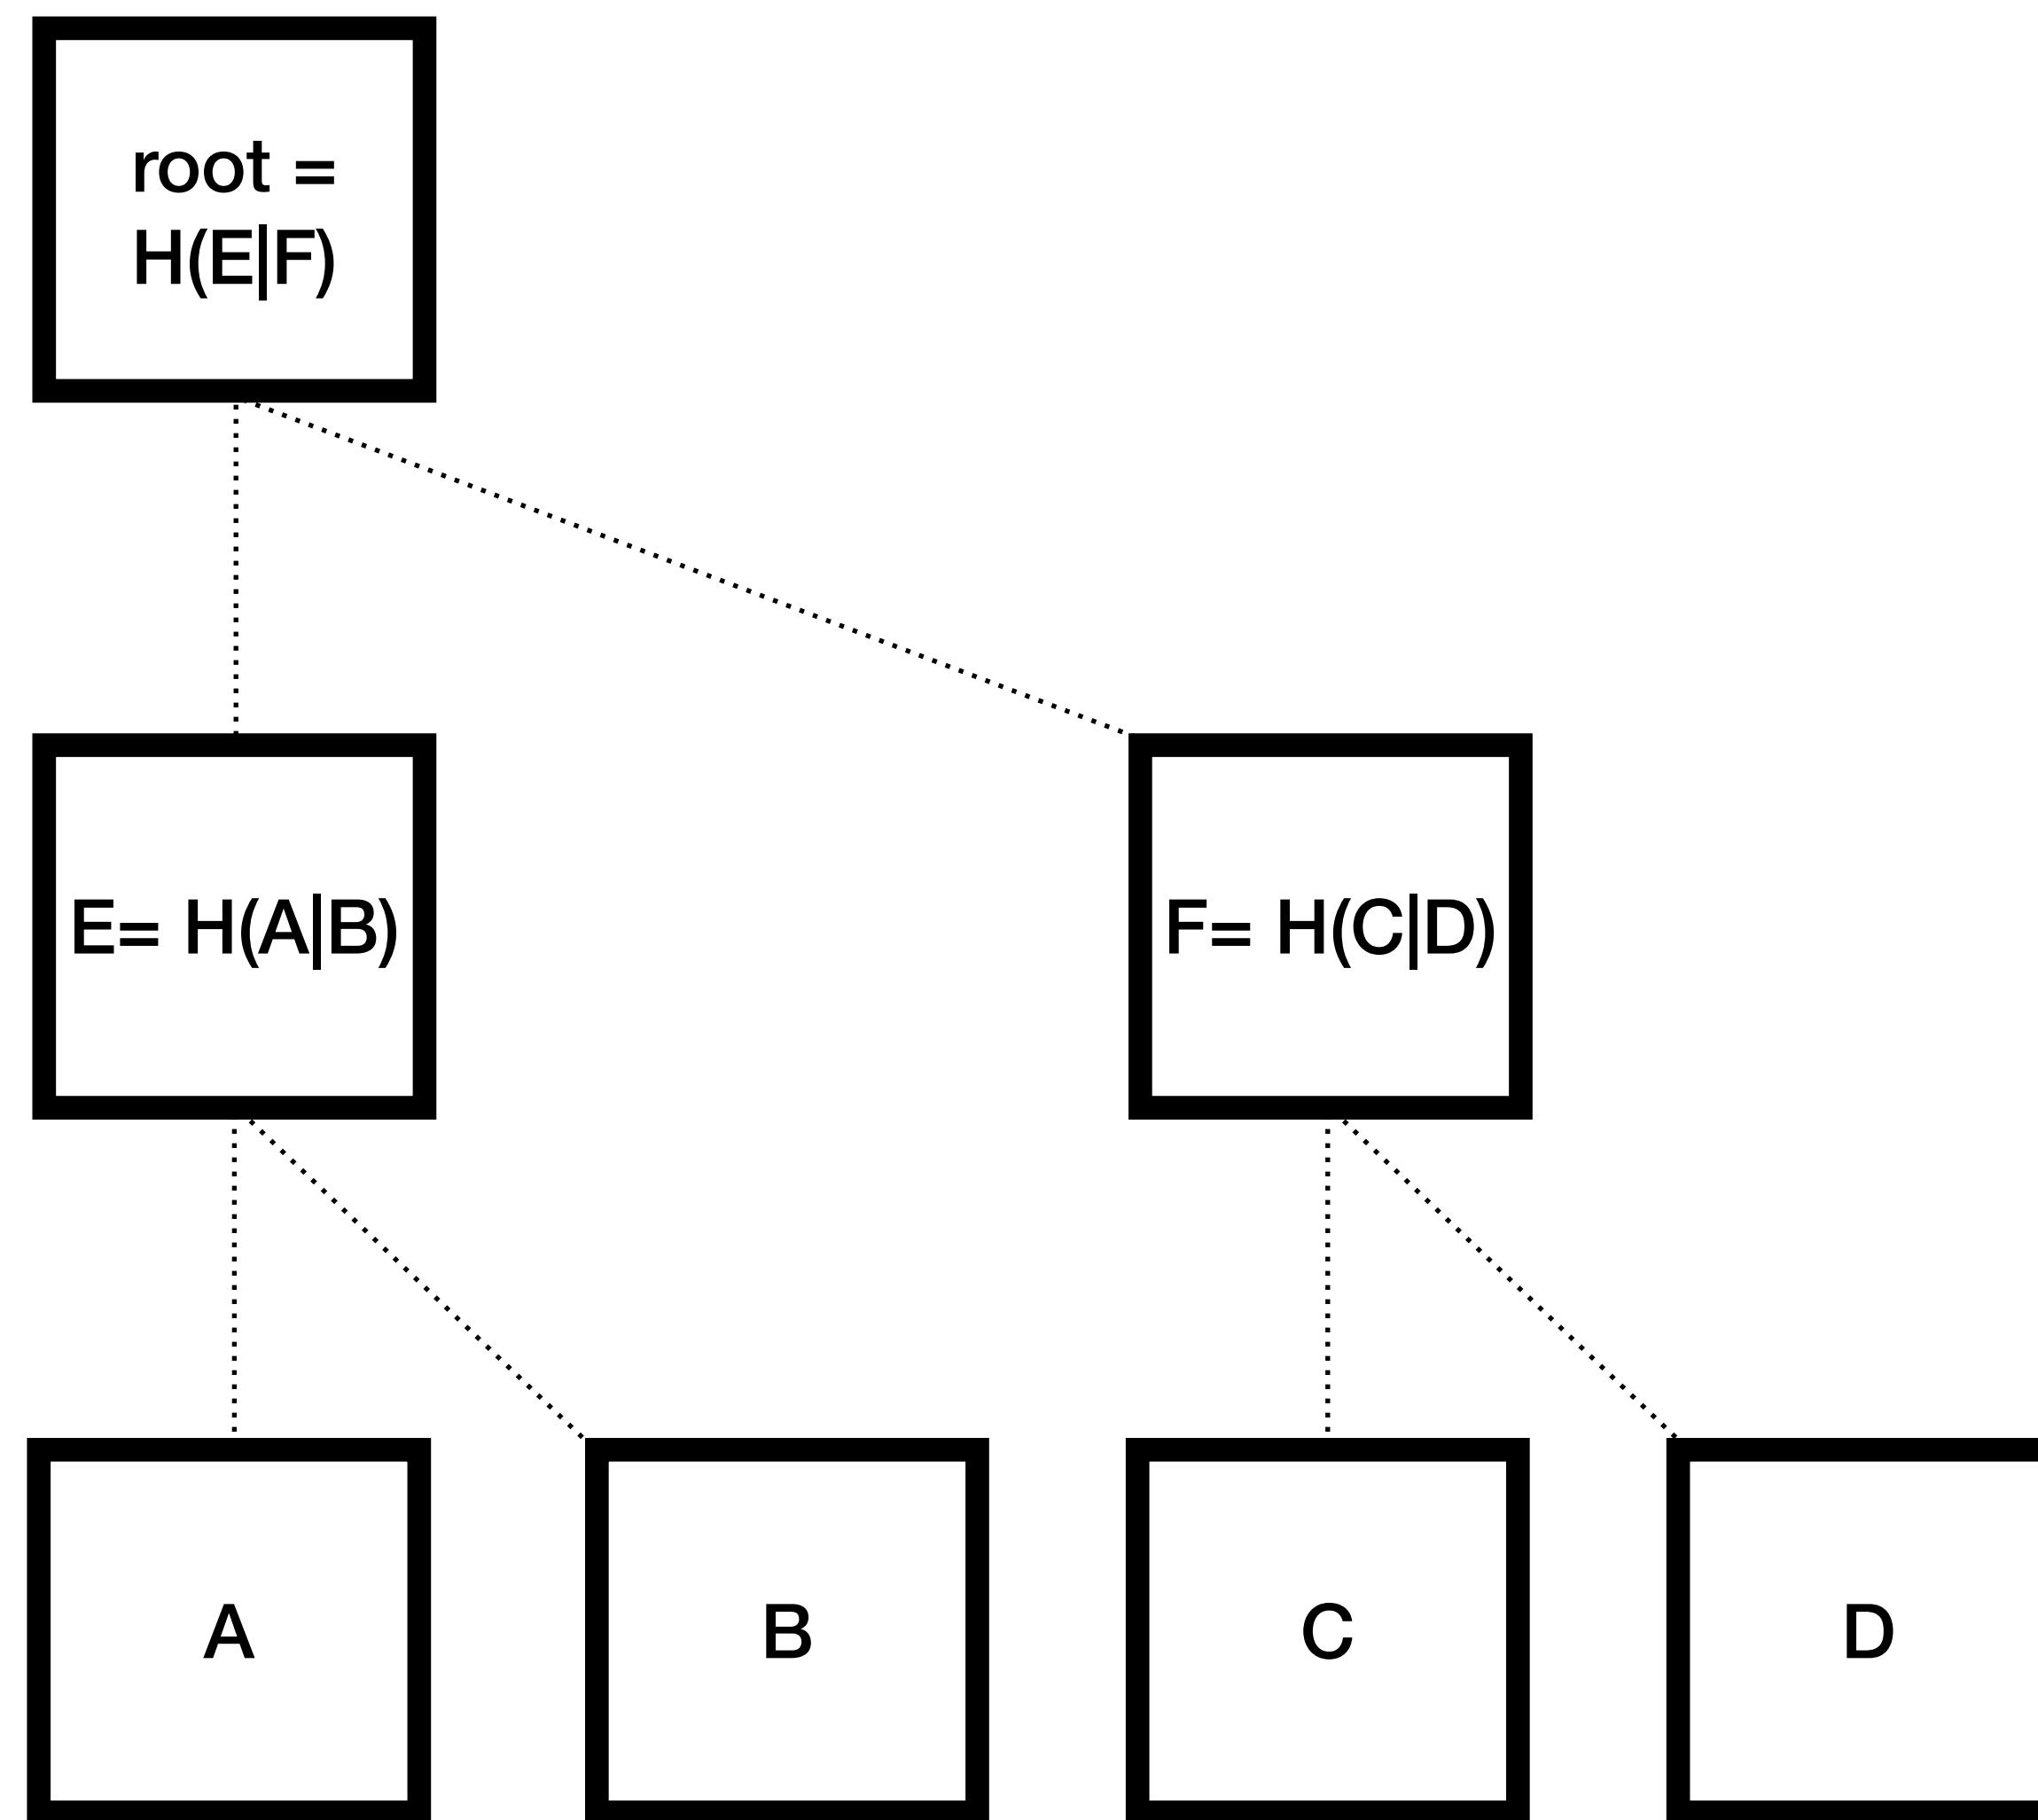
Algorithms

Consensus critical stuff

1. Add a UTXO
2. Delete a UTXO
3. Verify the existence

**Deletion happens with verifying
the existence**

Quick merkle tree prove review

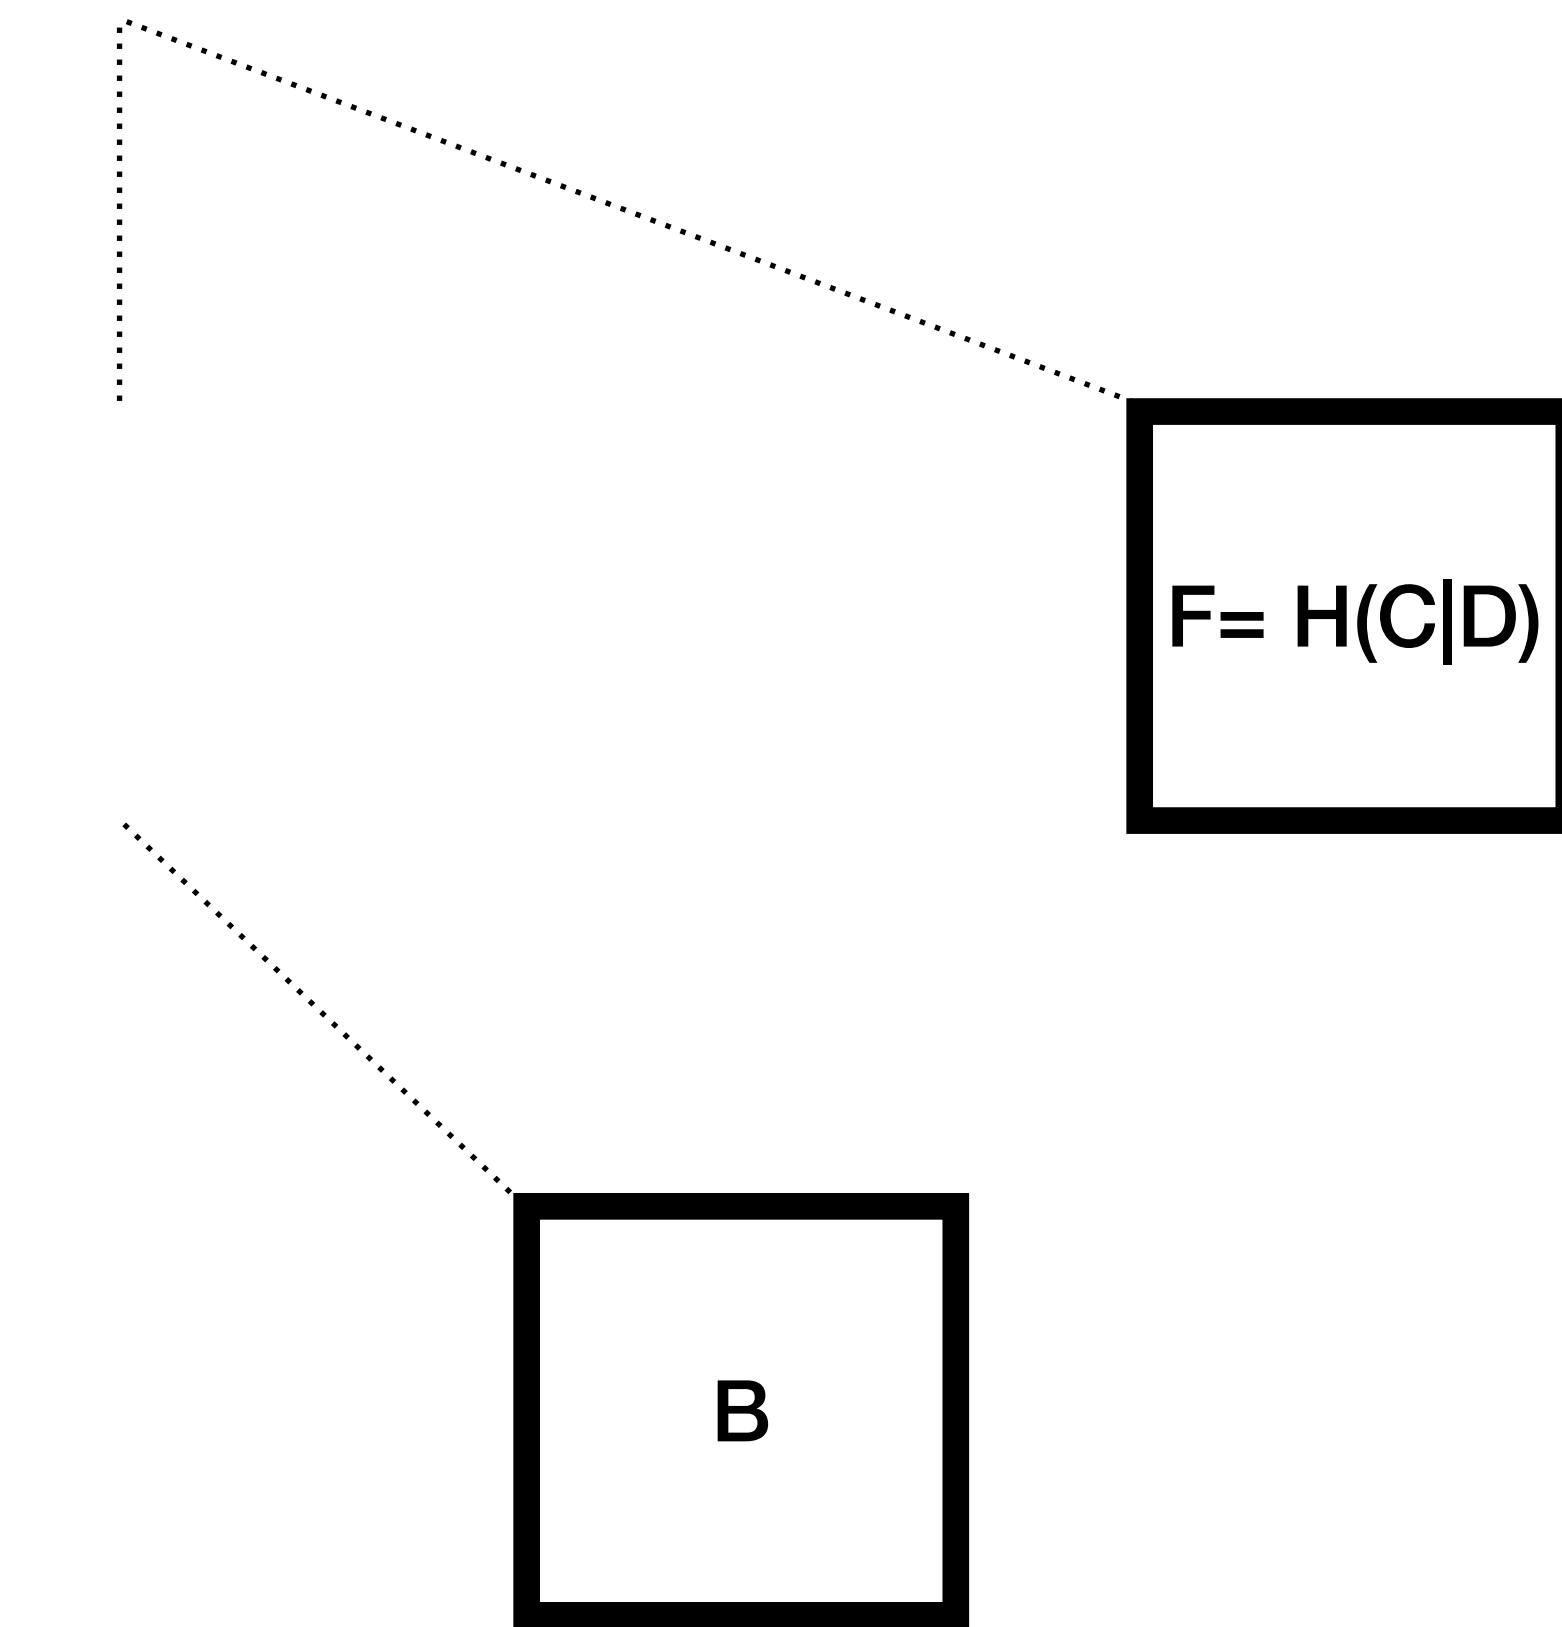


Proving A

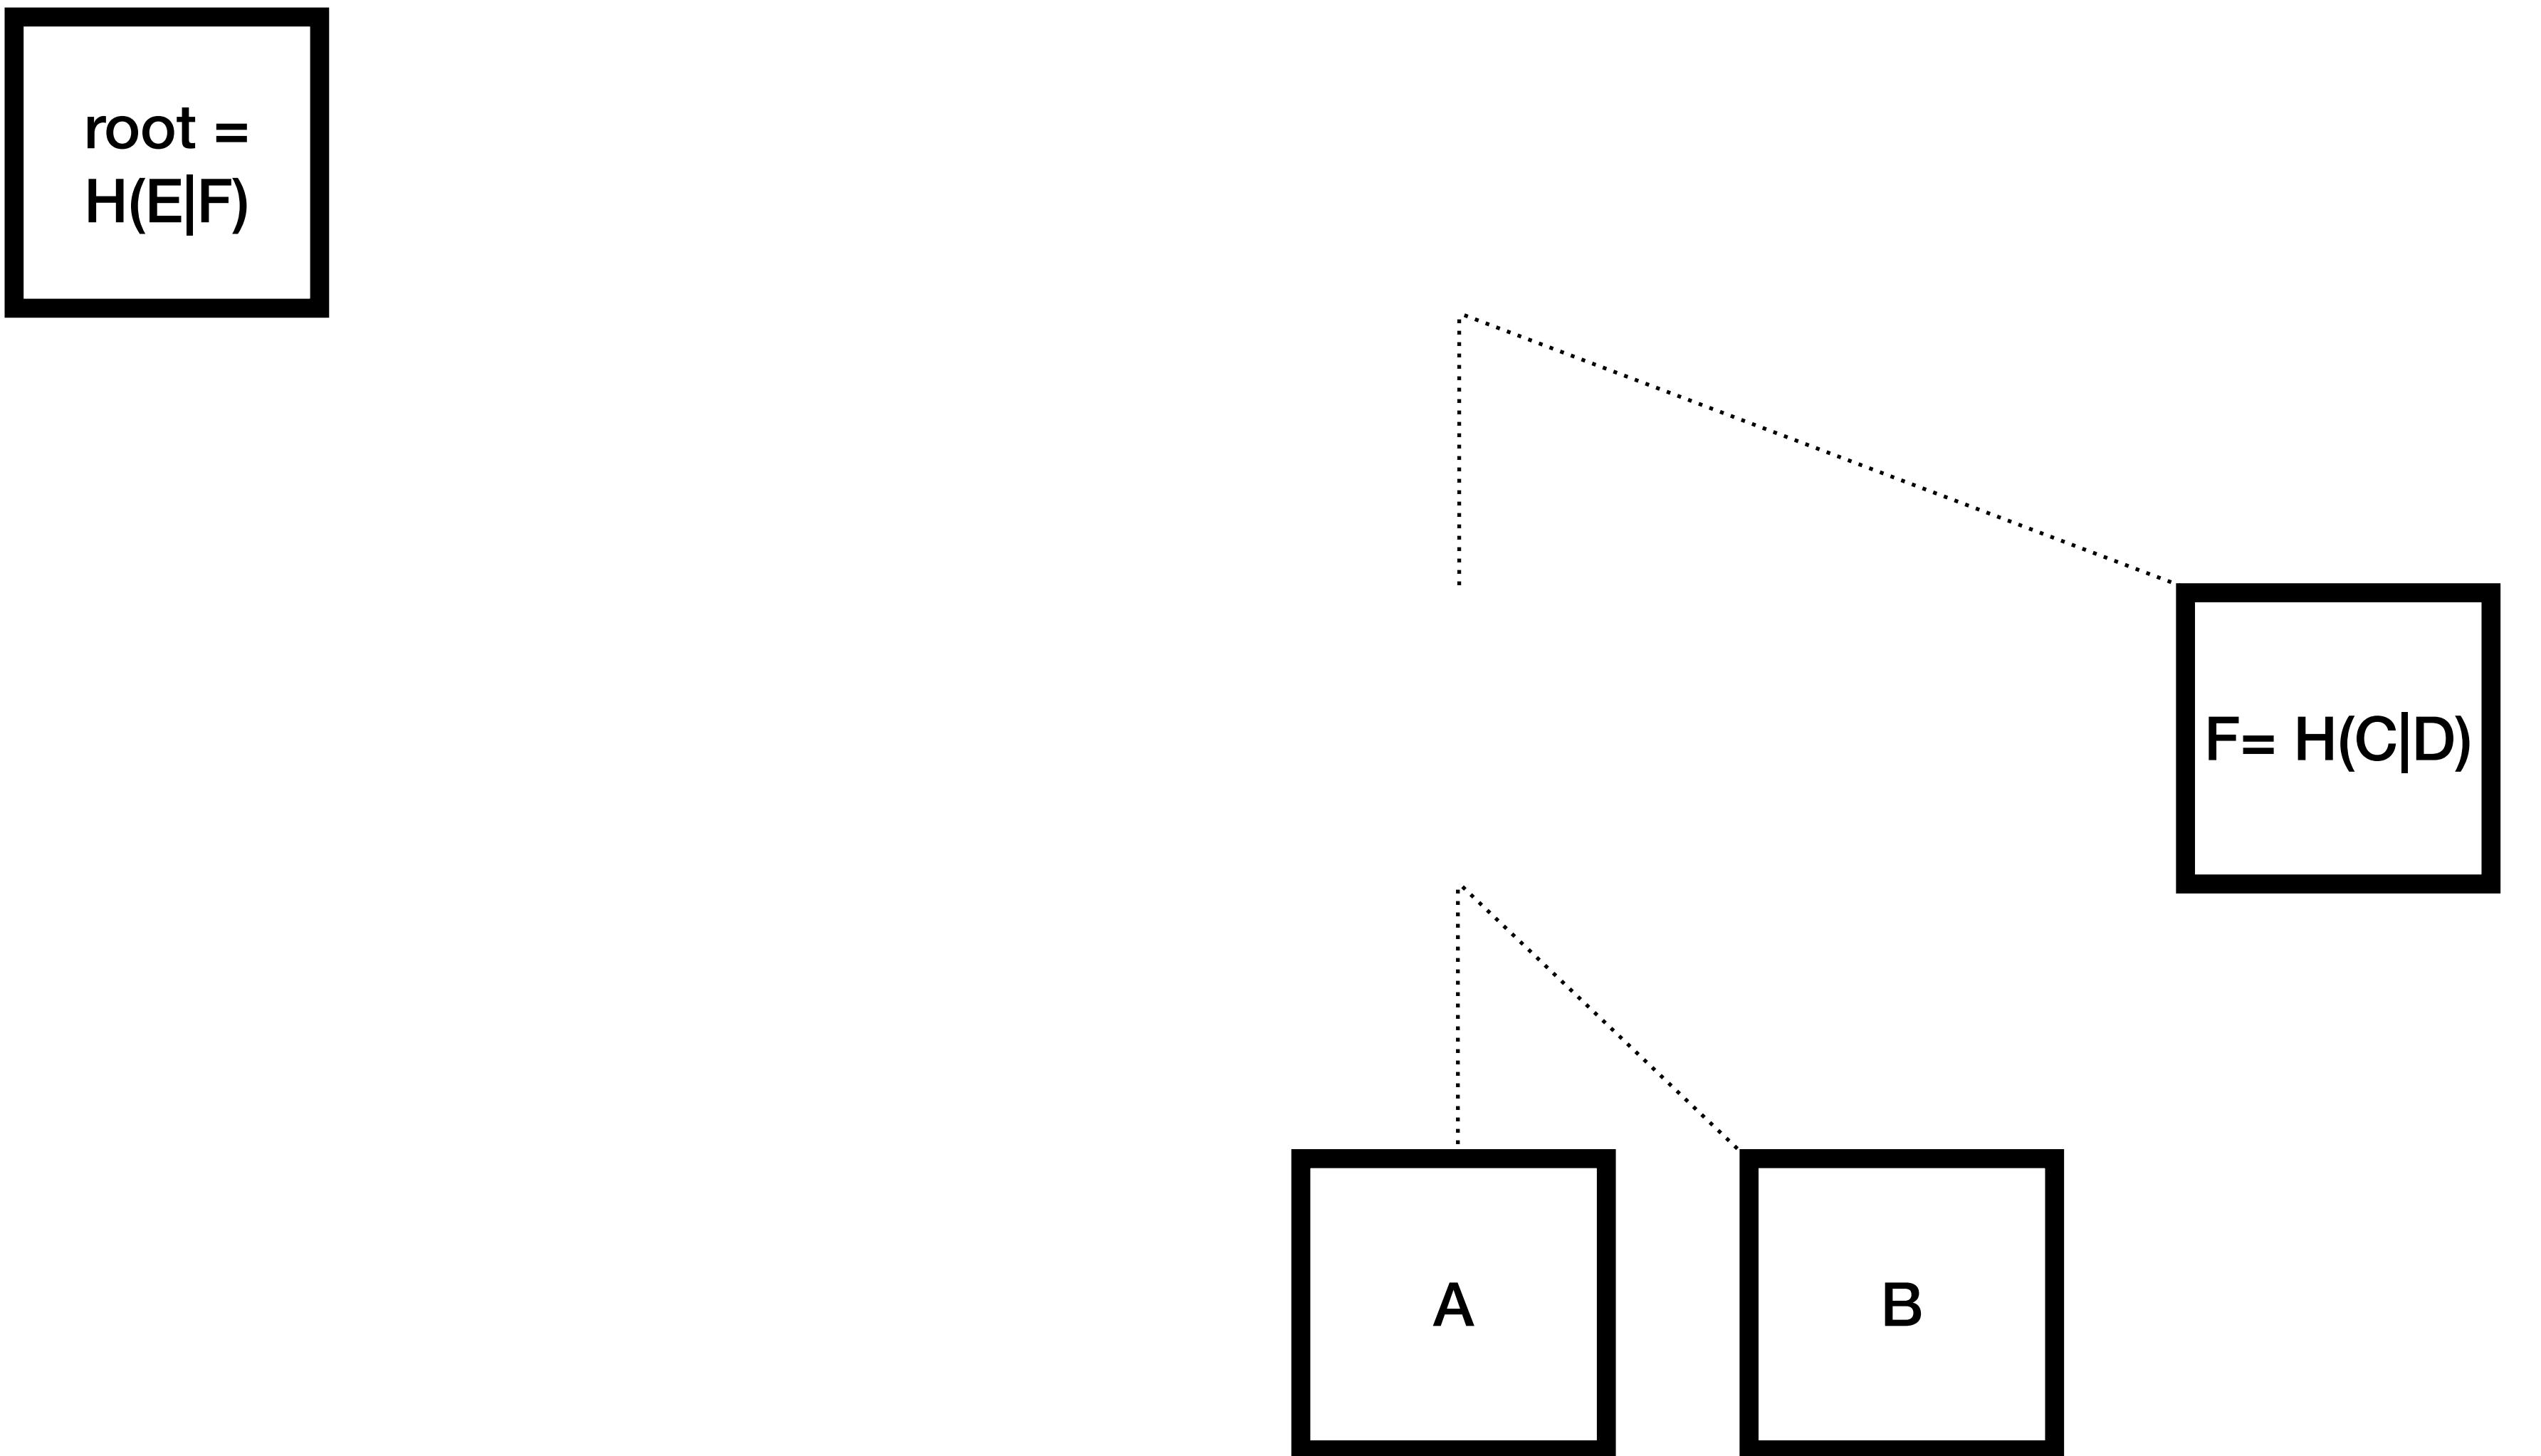
root =
 $H(E|F)$

Receive proof: B, F

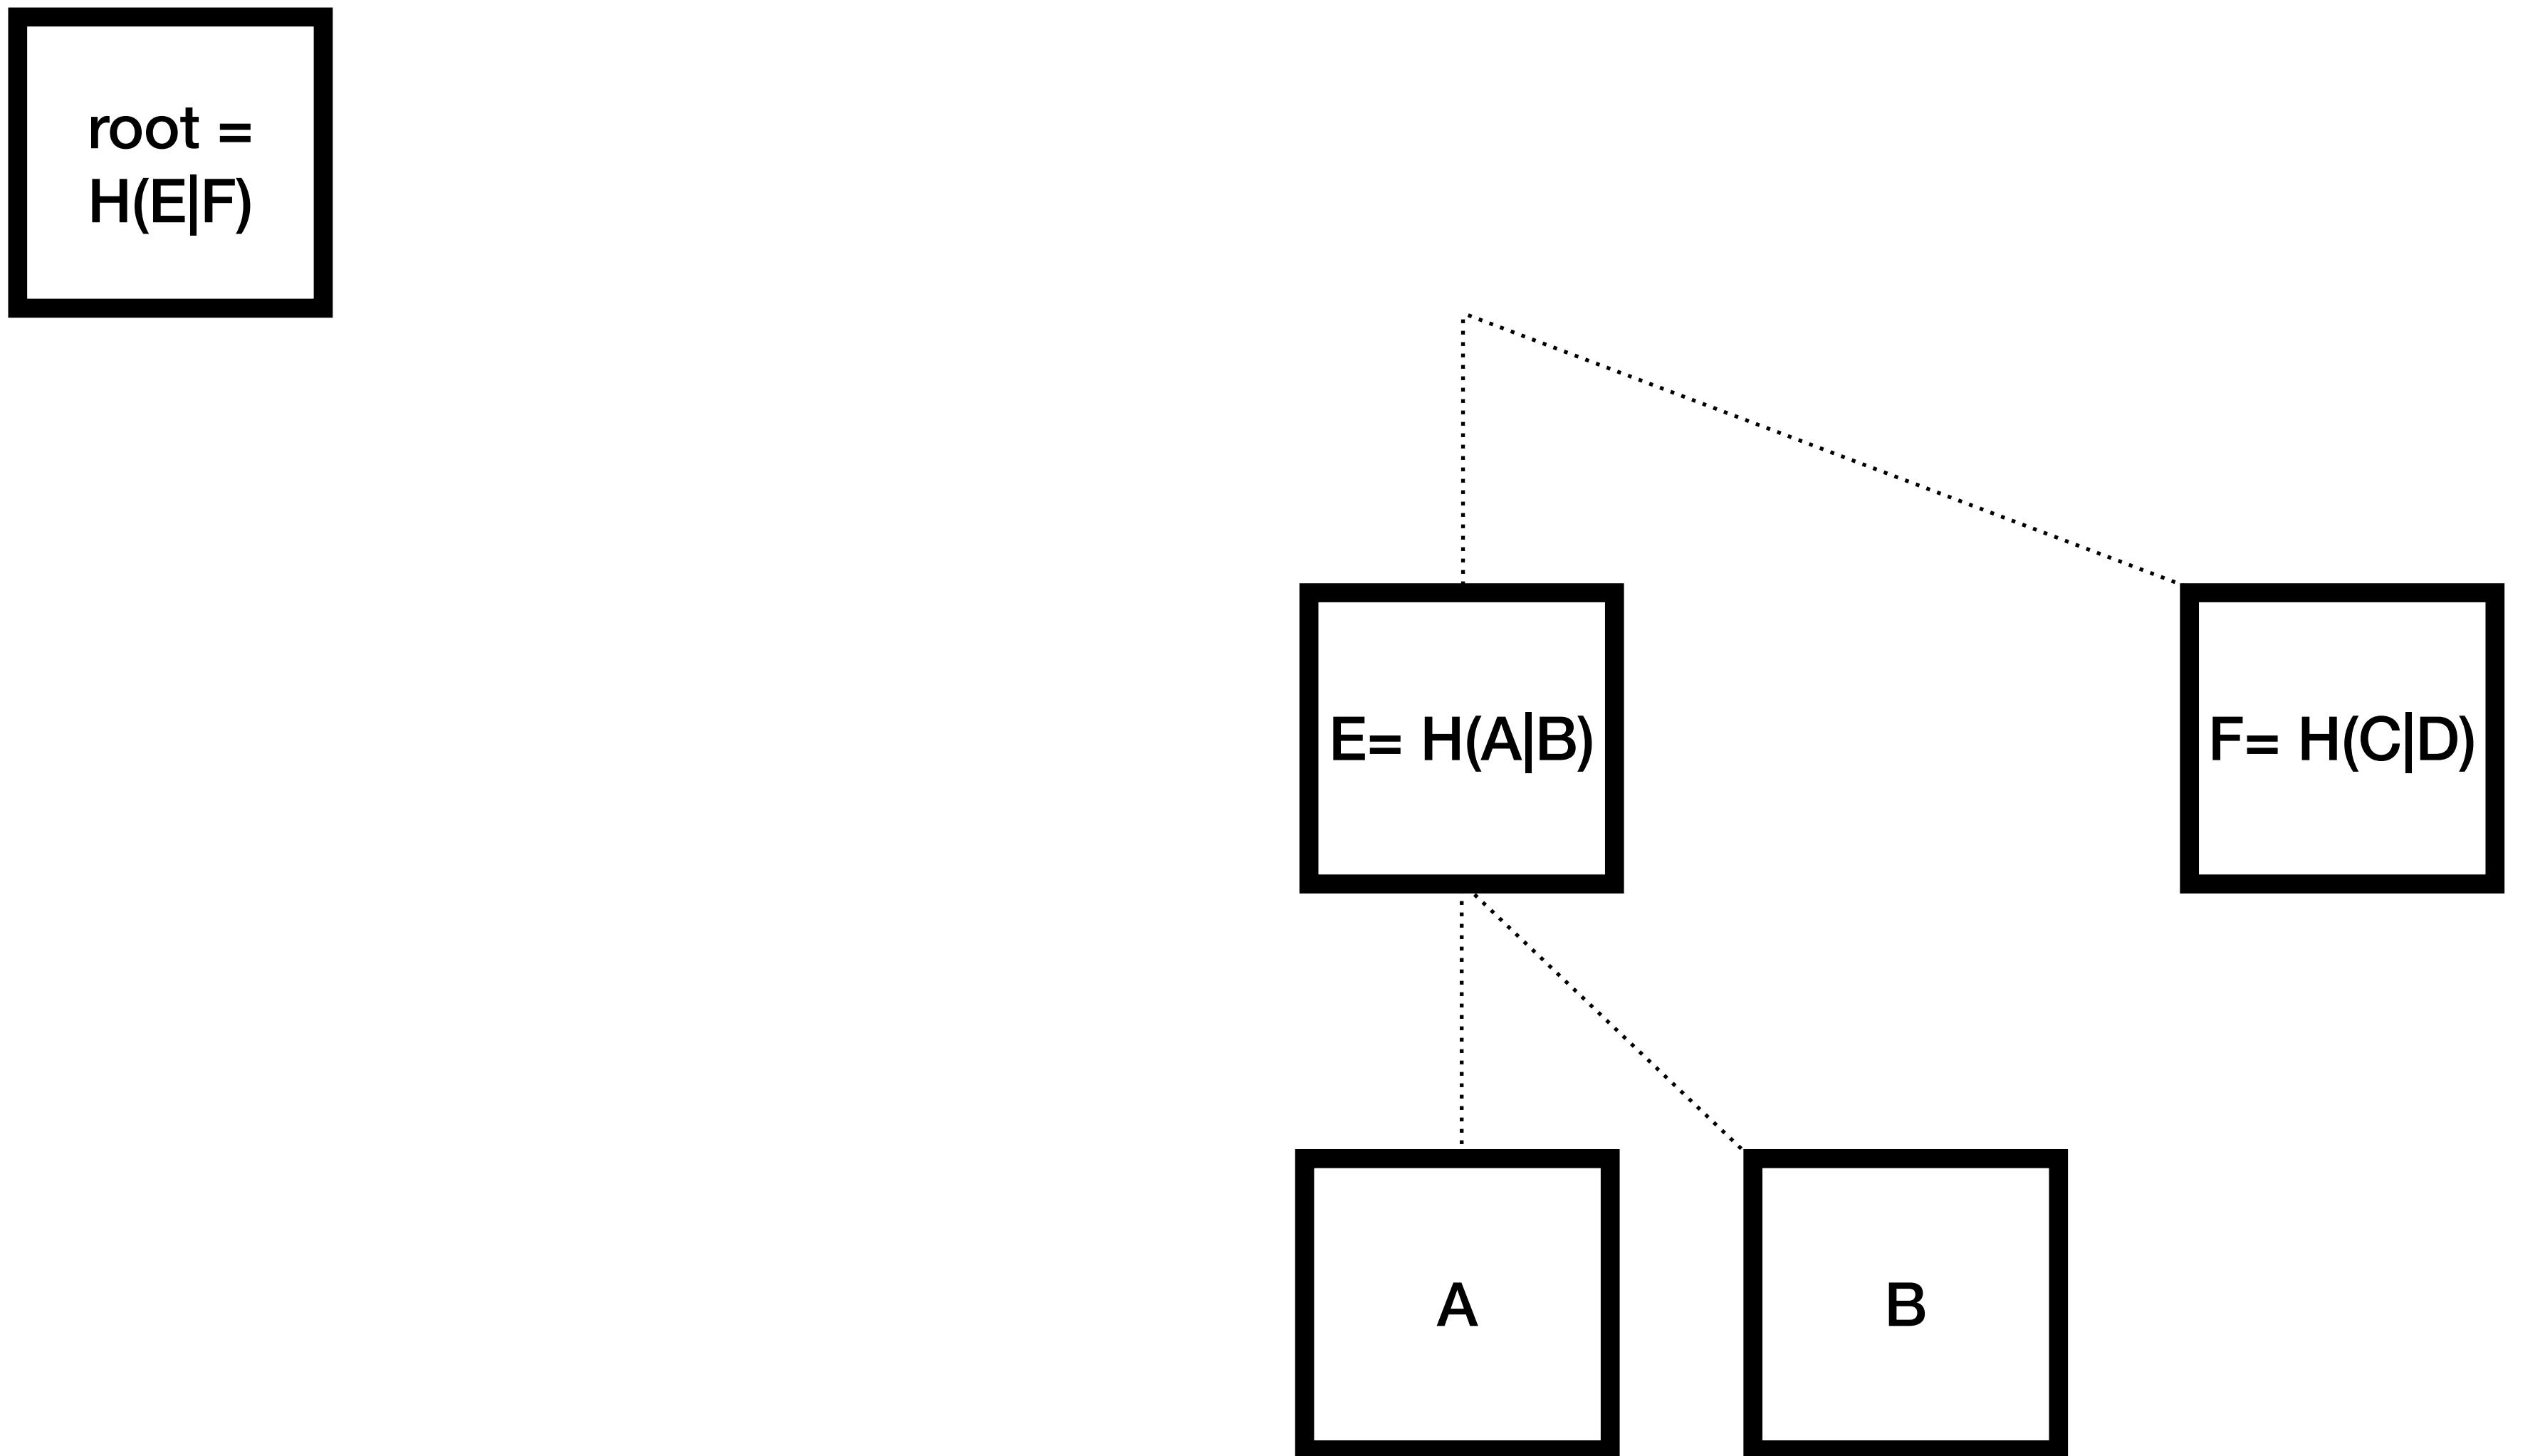
root =
 $H(E|F)$



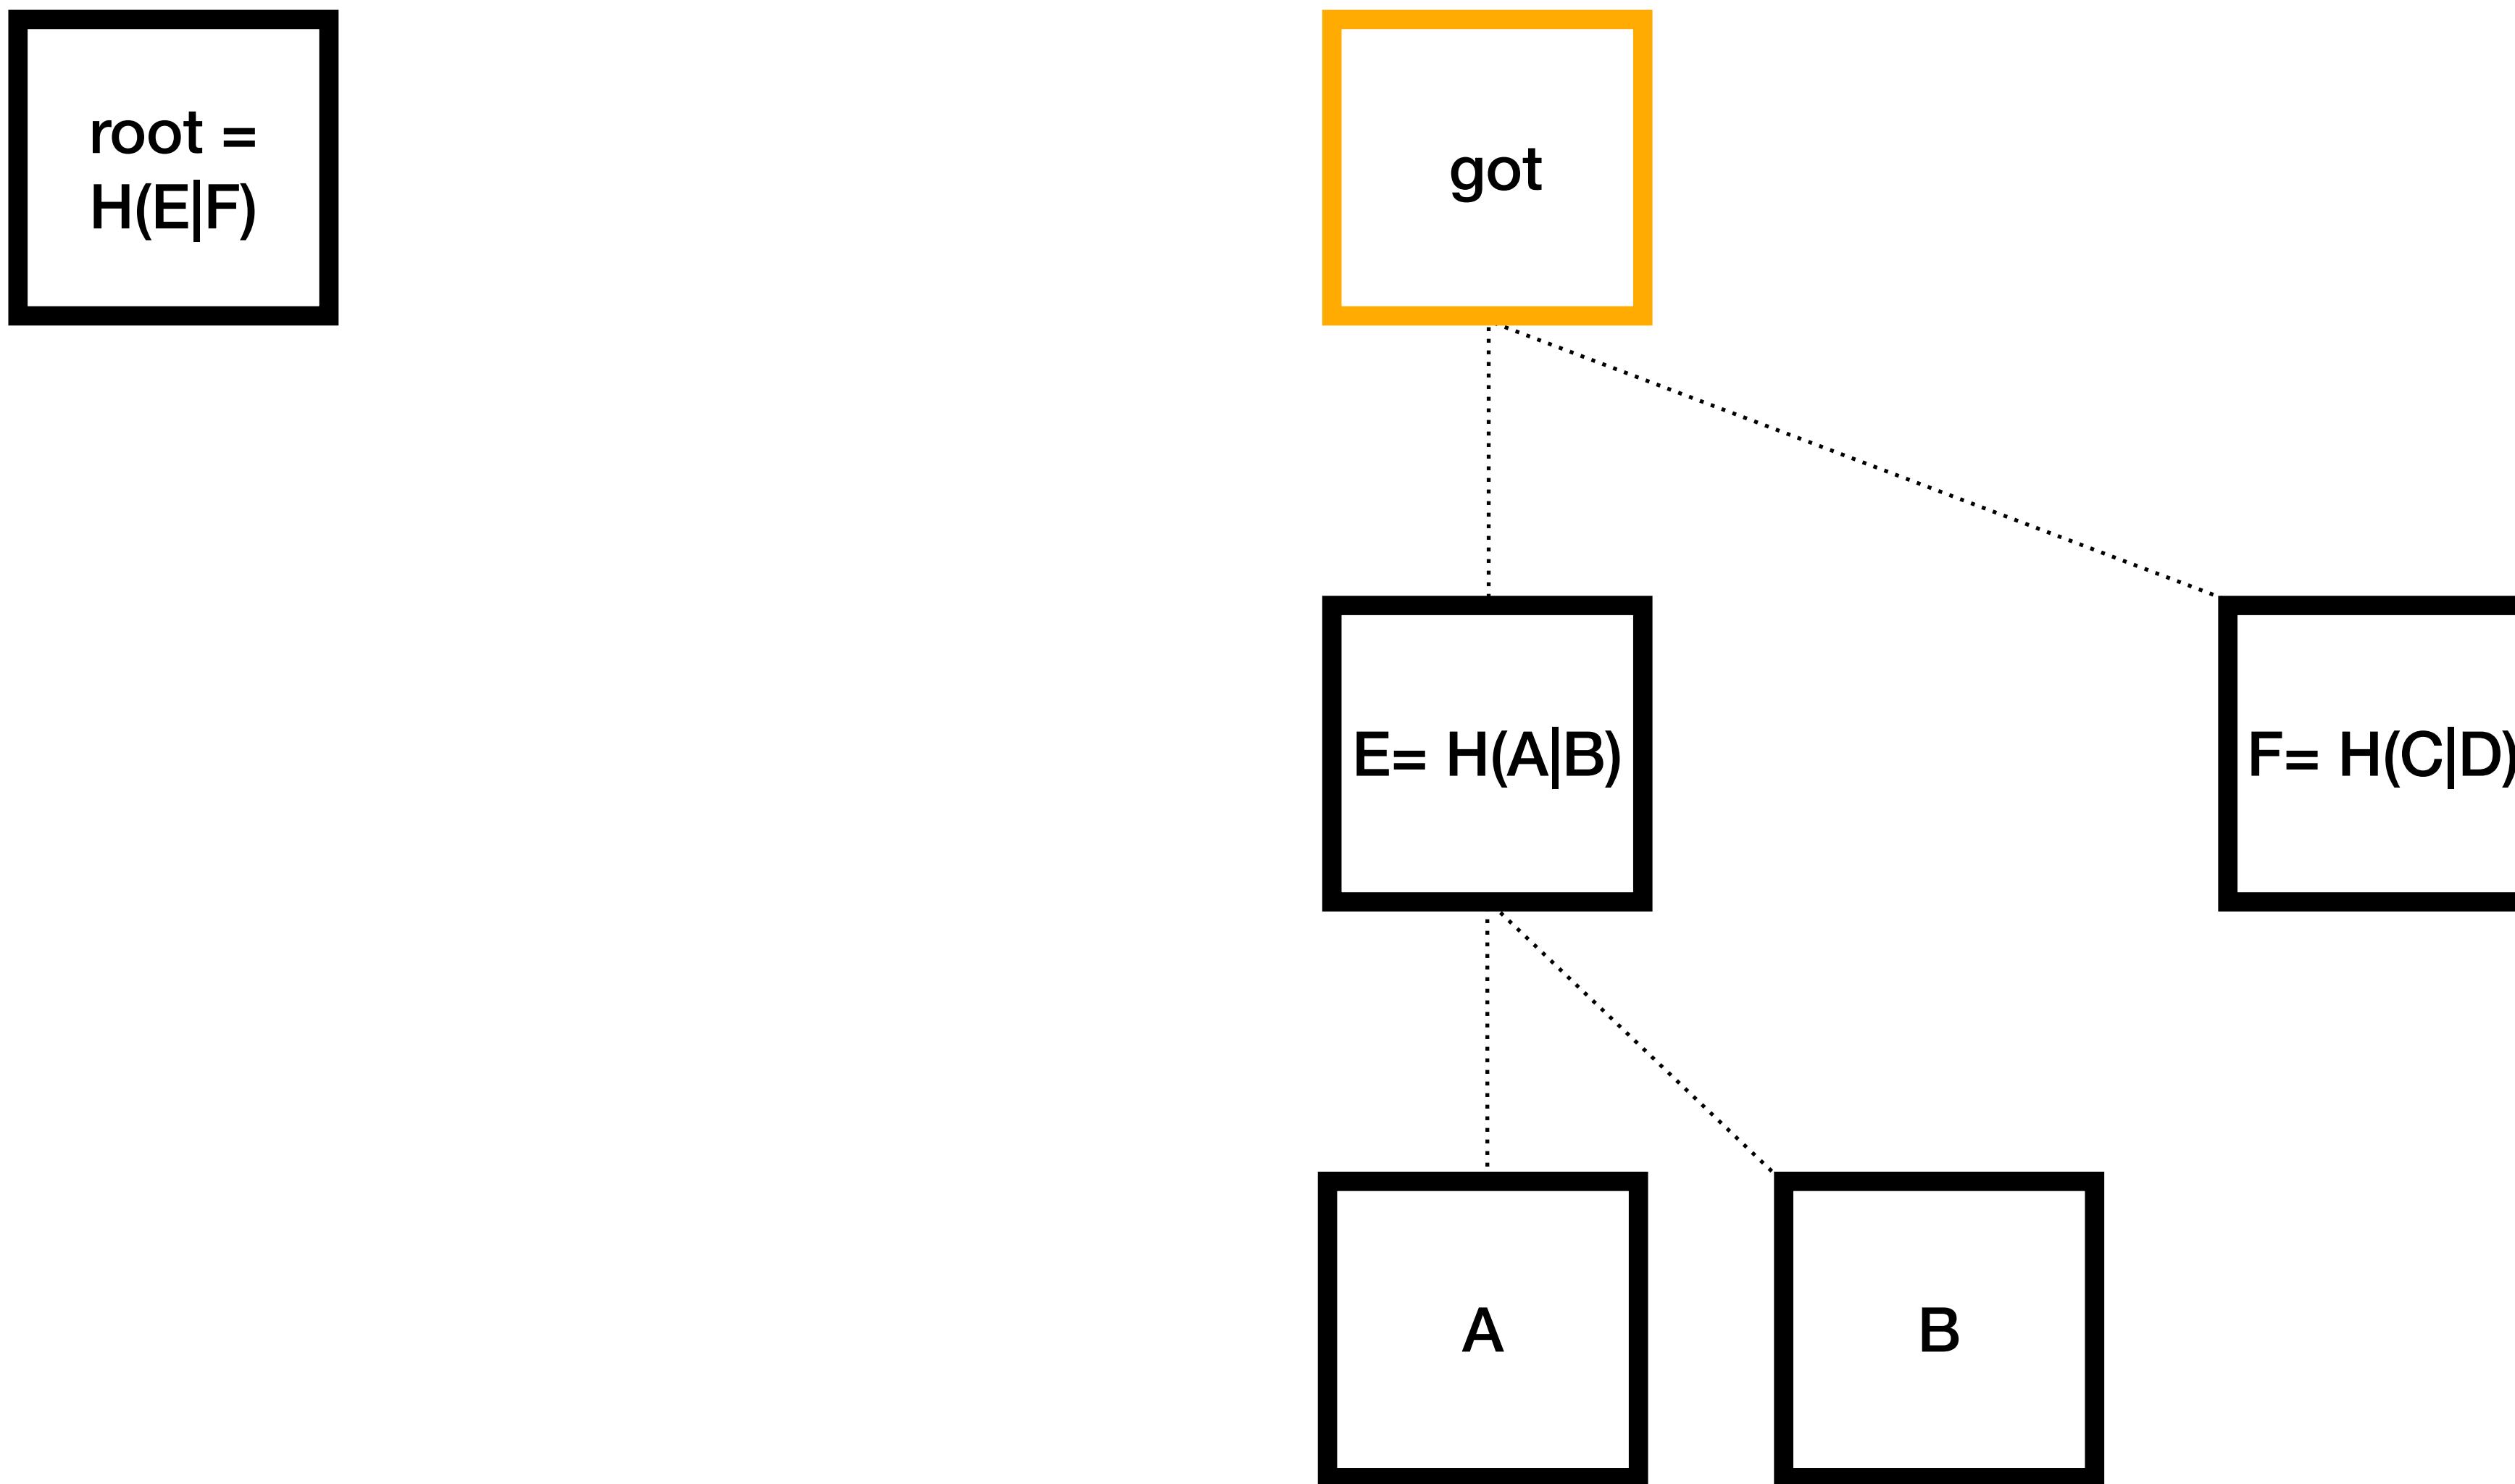
Fetch A



Calculate E



Calculate root



Compare roots

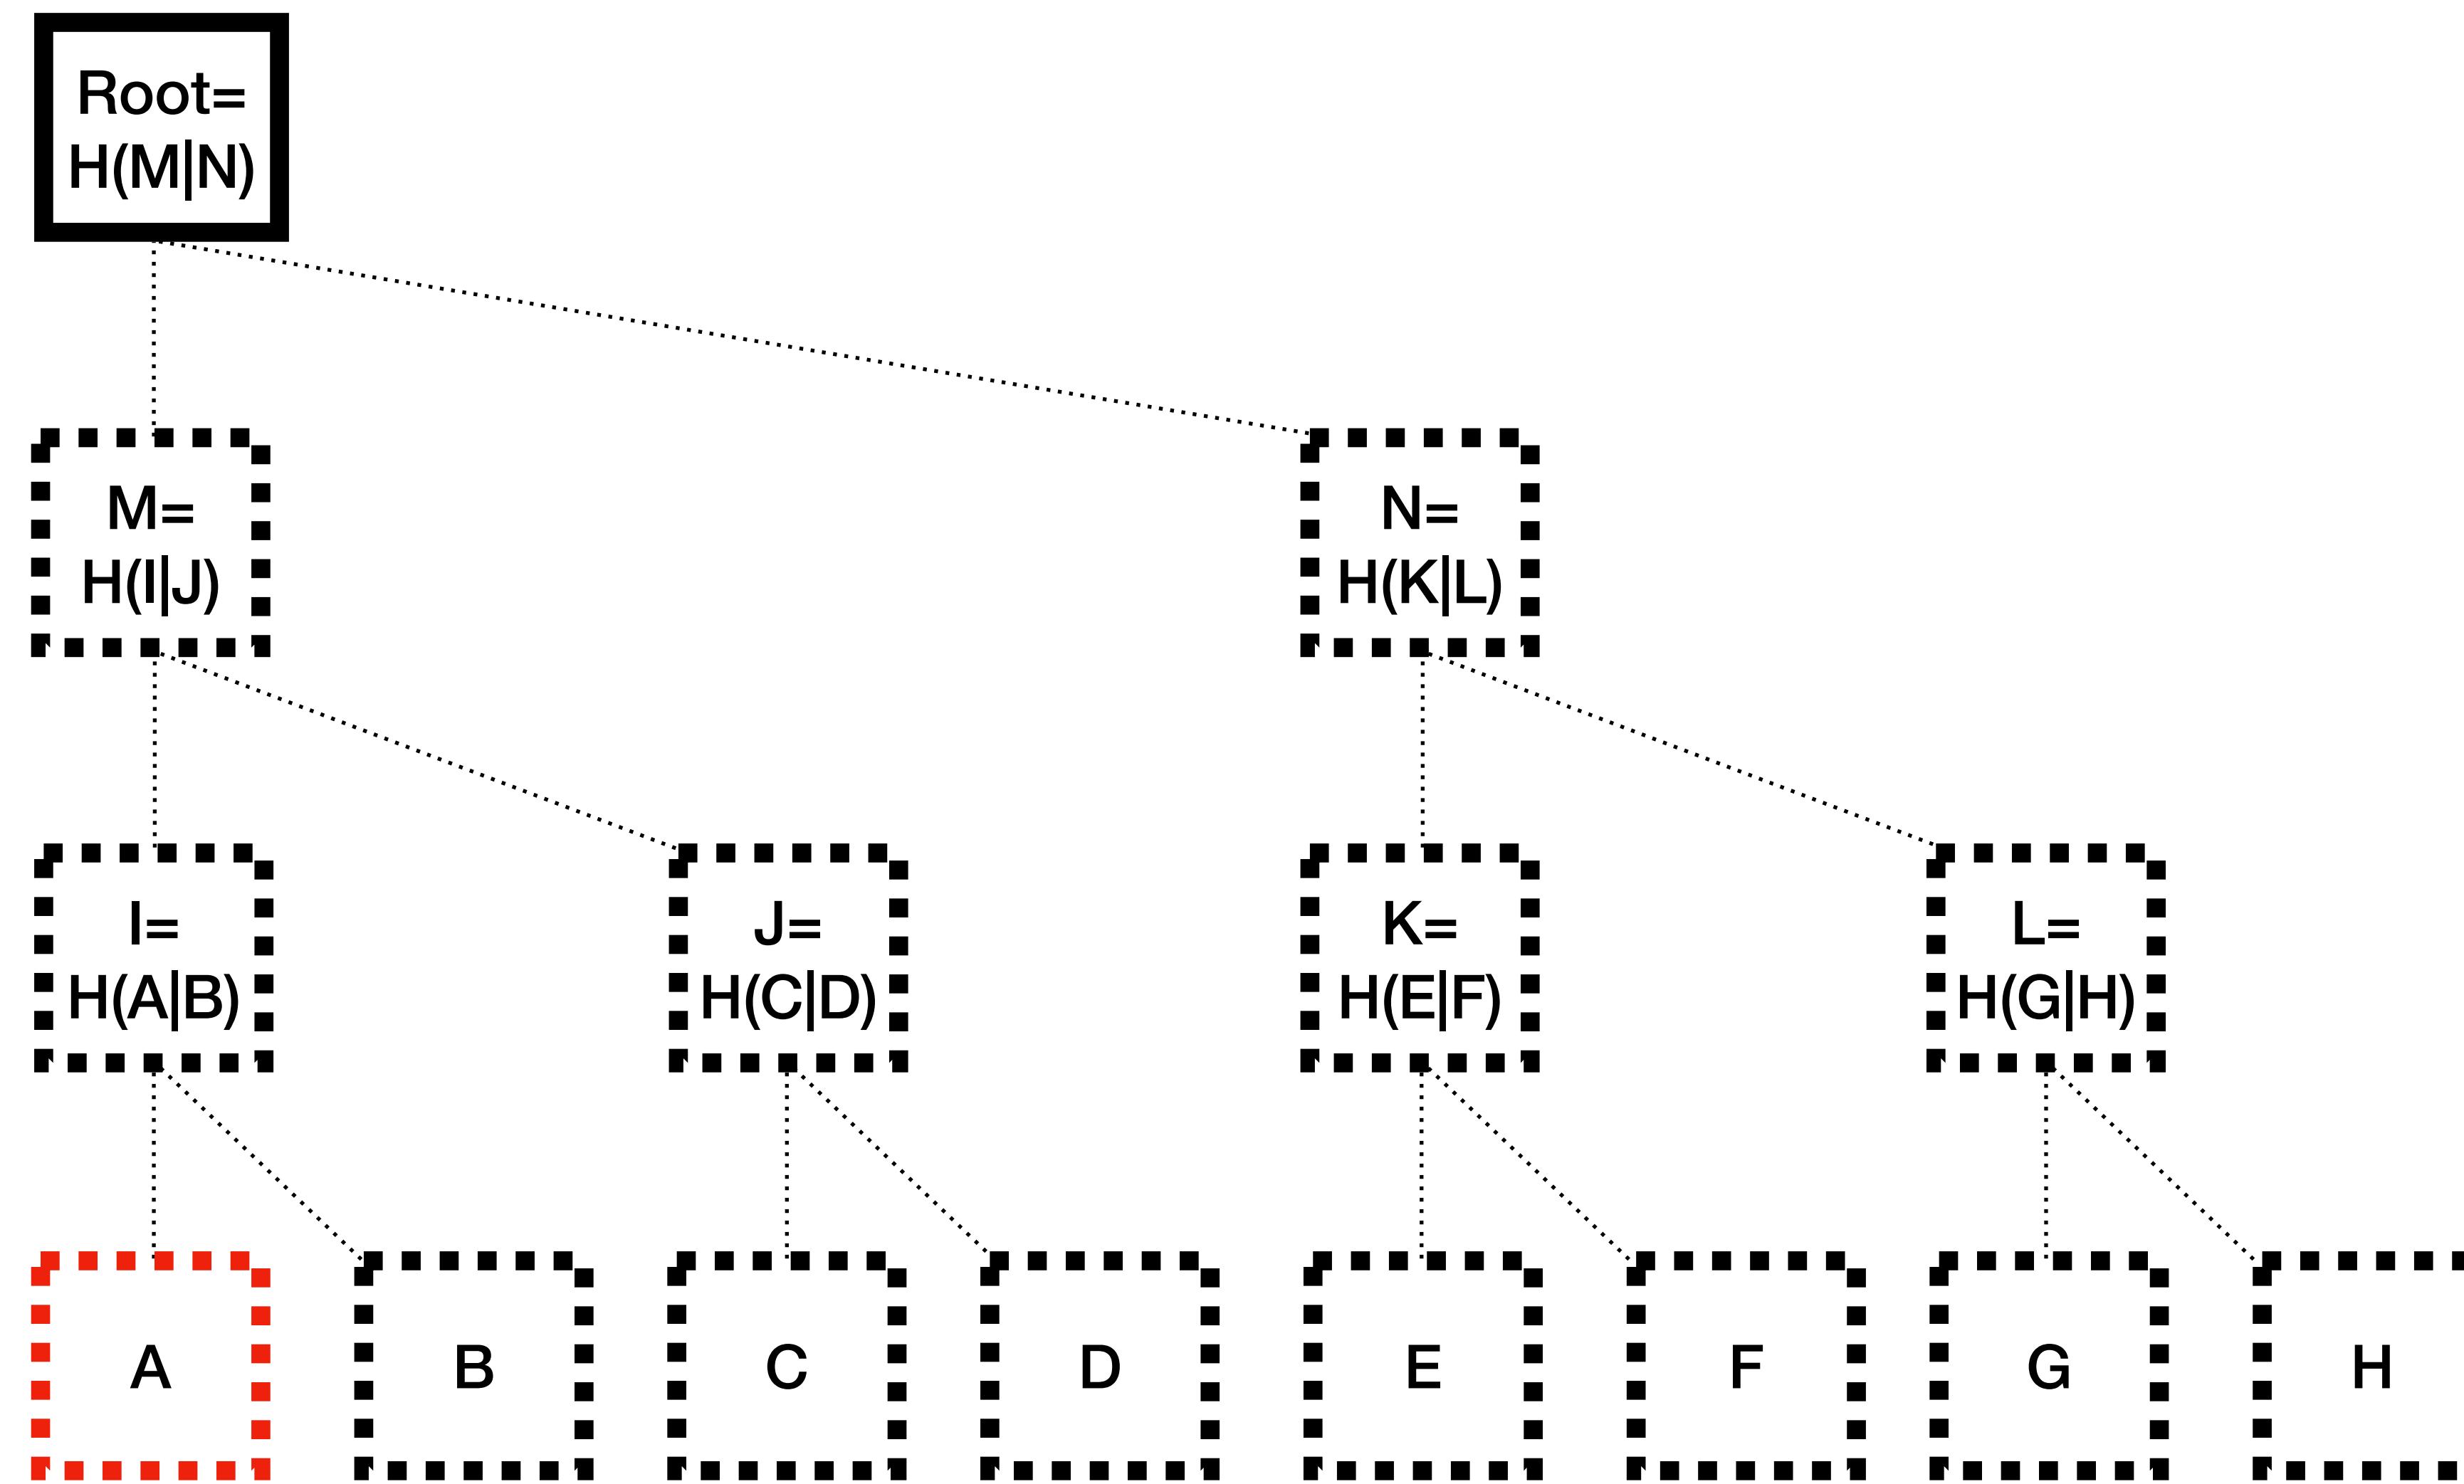
Proof is valid if root is the same

root= $H(E|F)$

==

Got Root

Deleting A

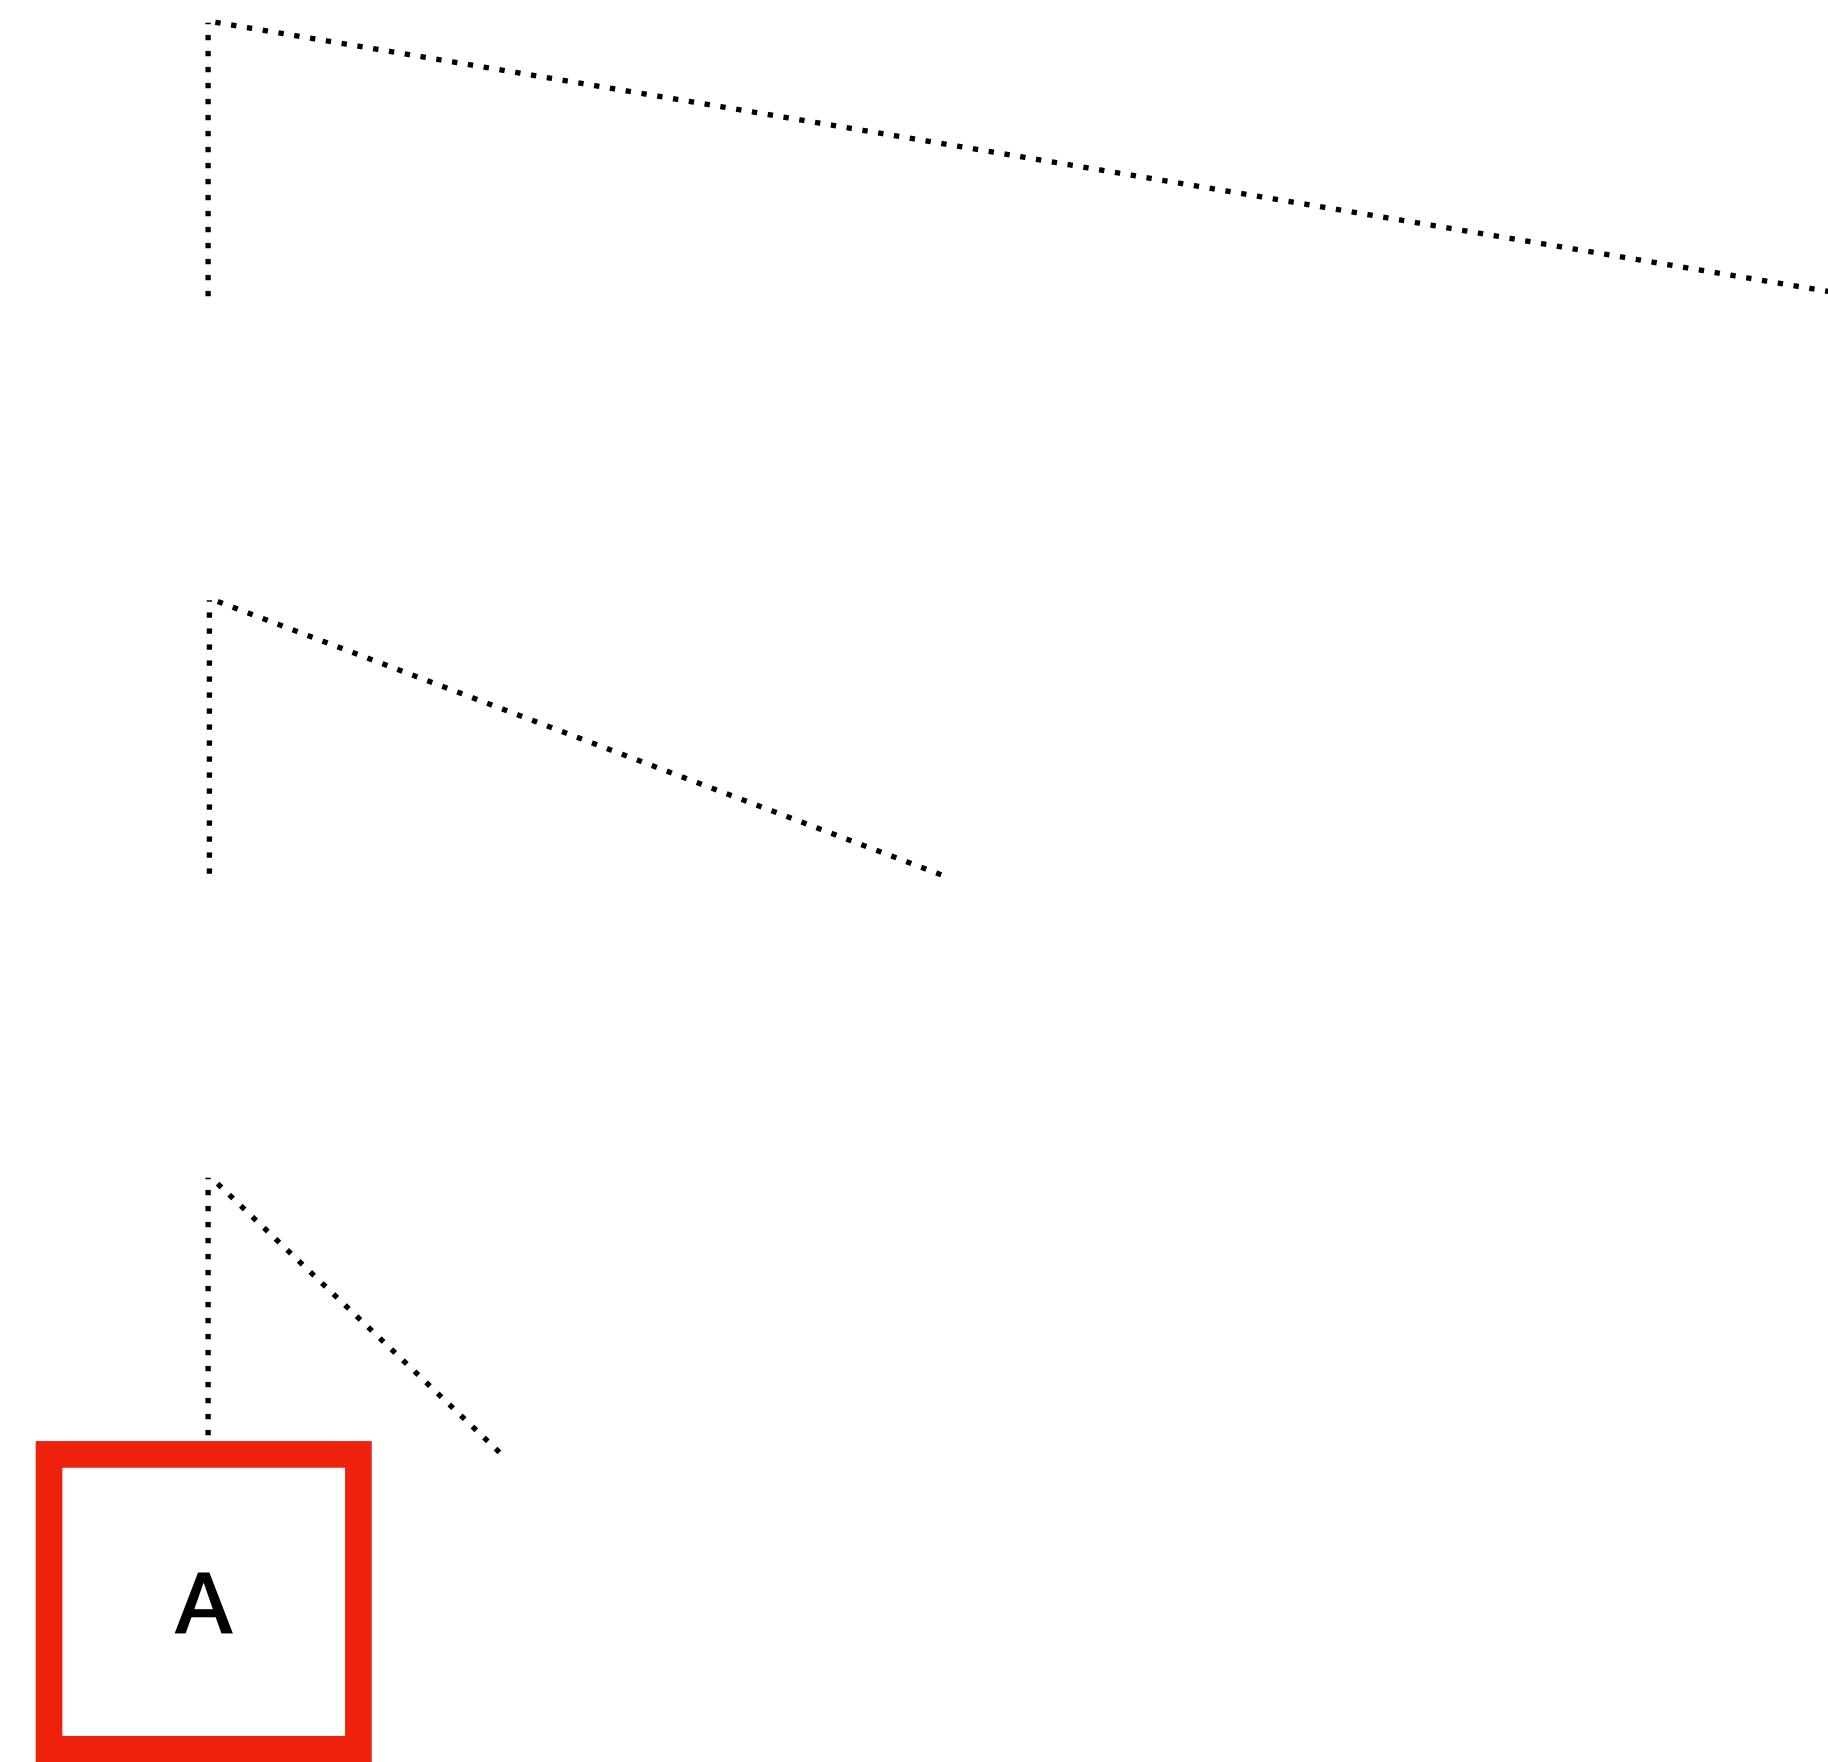


The only requirement is the roots

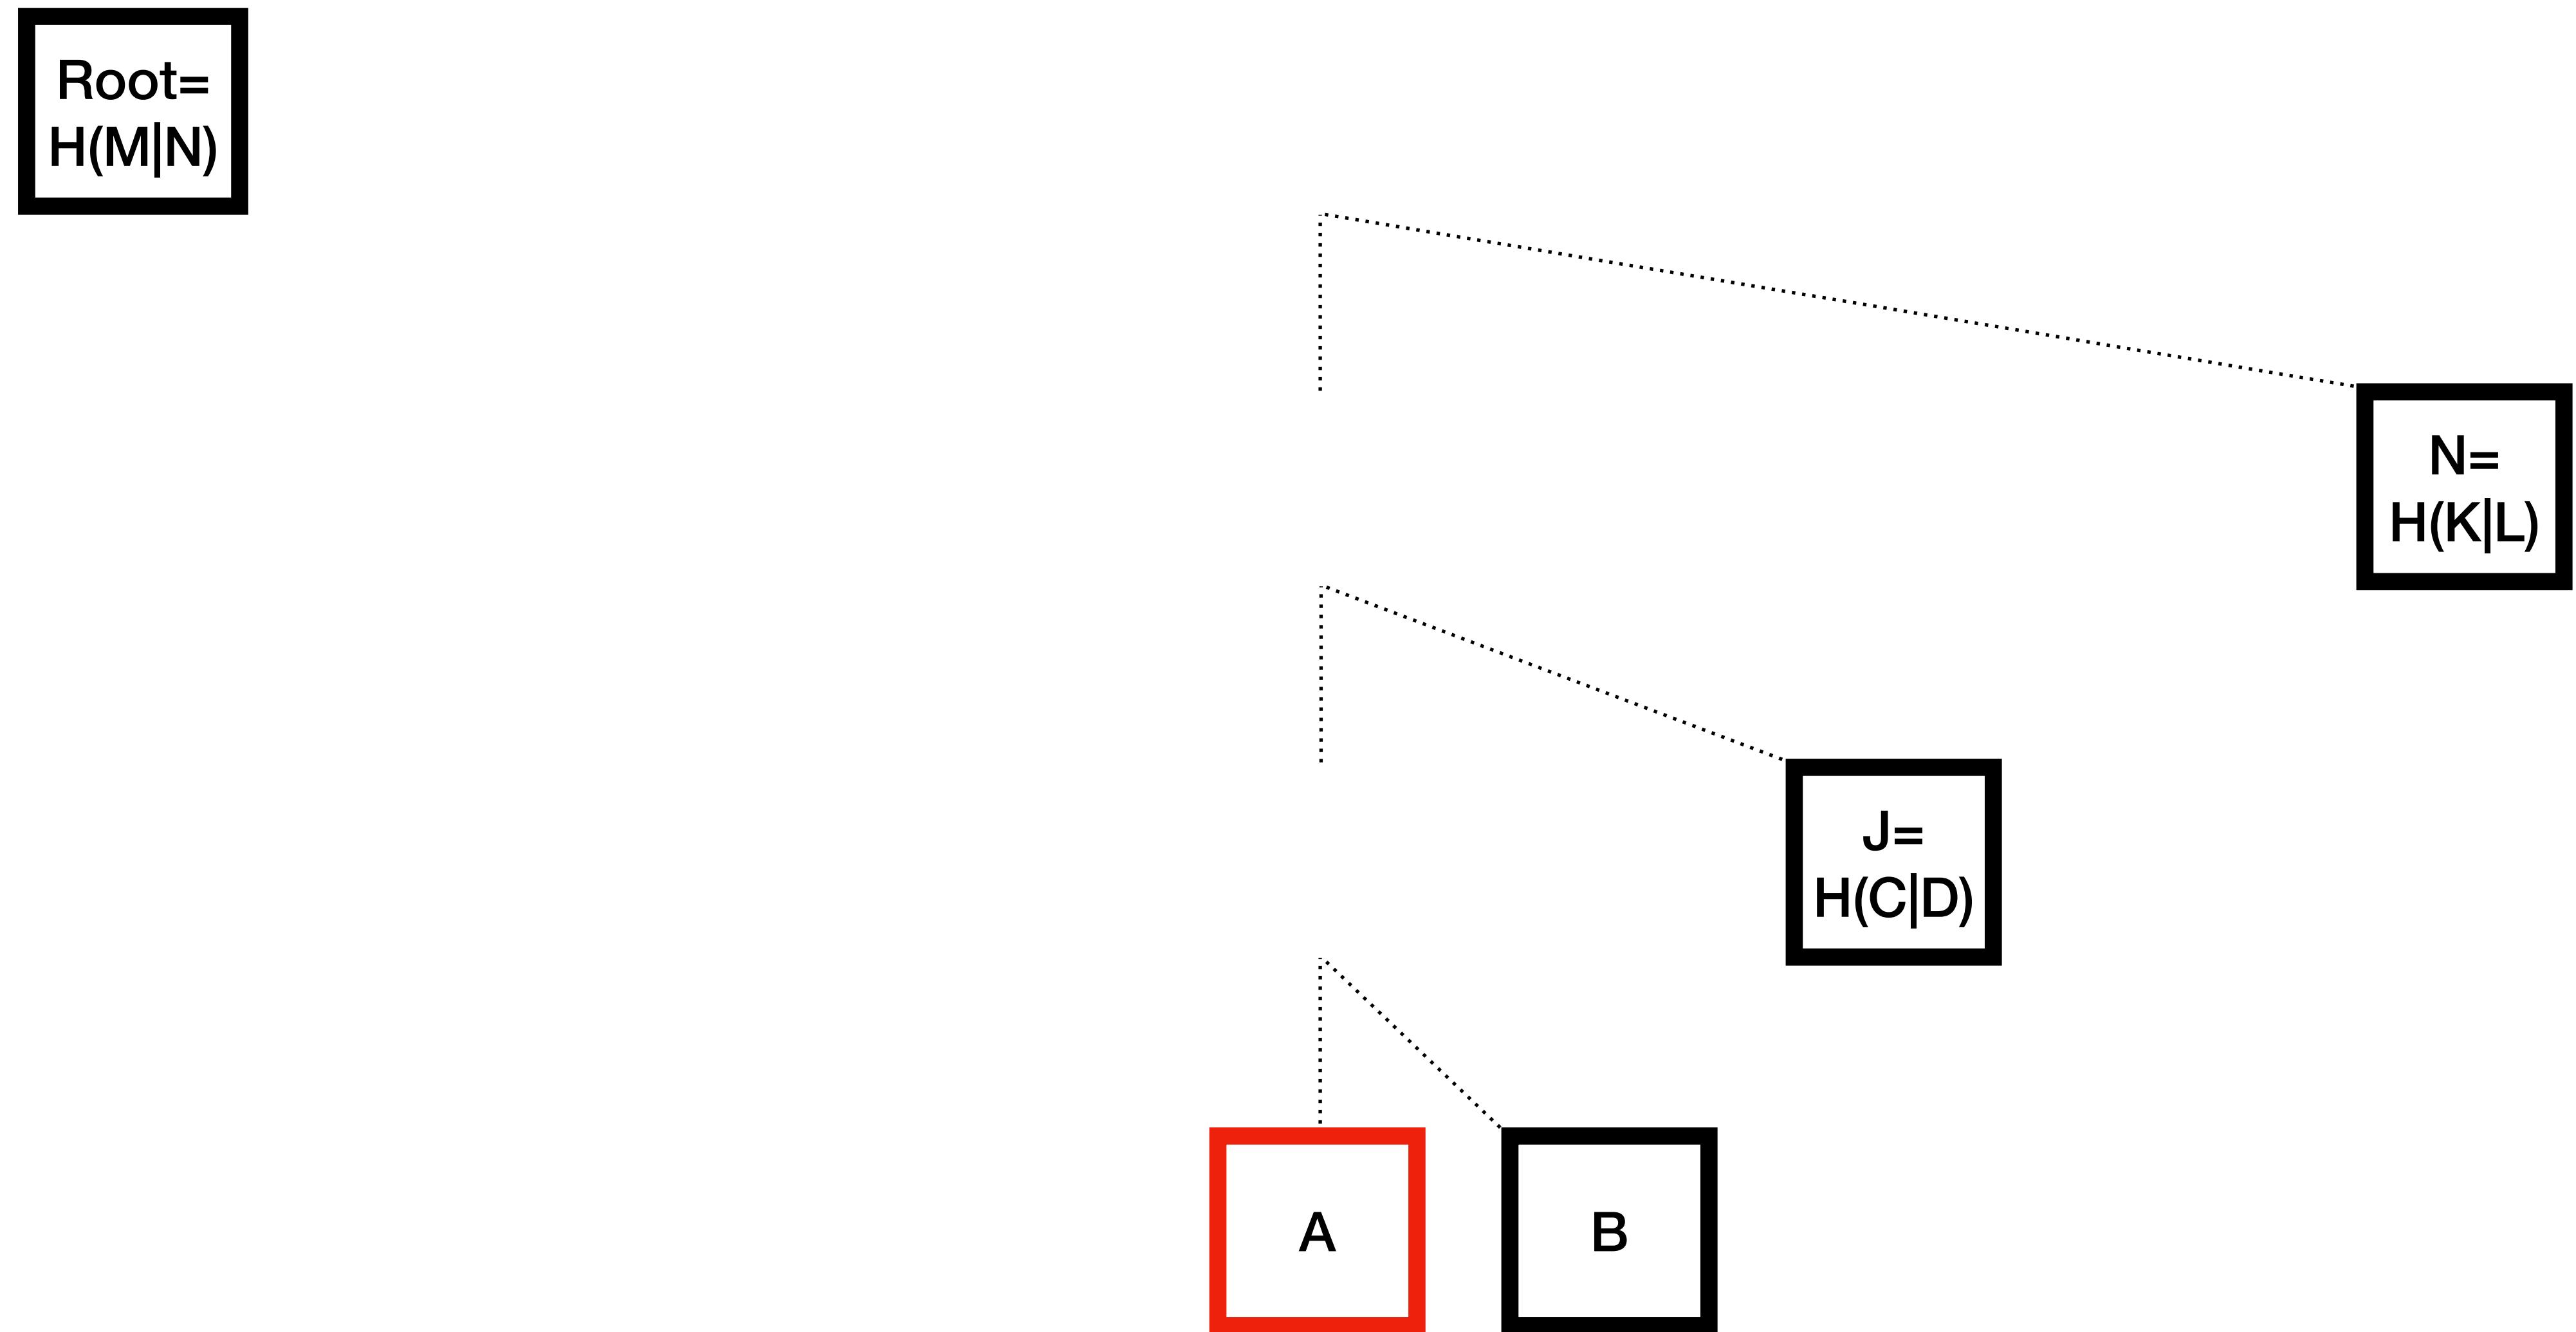
Root=
 $H(M|N)$

Calculate A from the Bitcoin Block

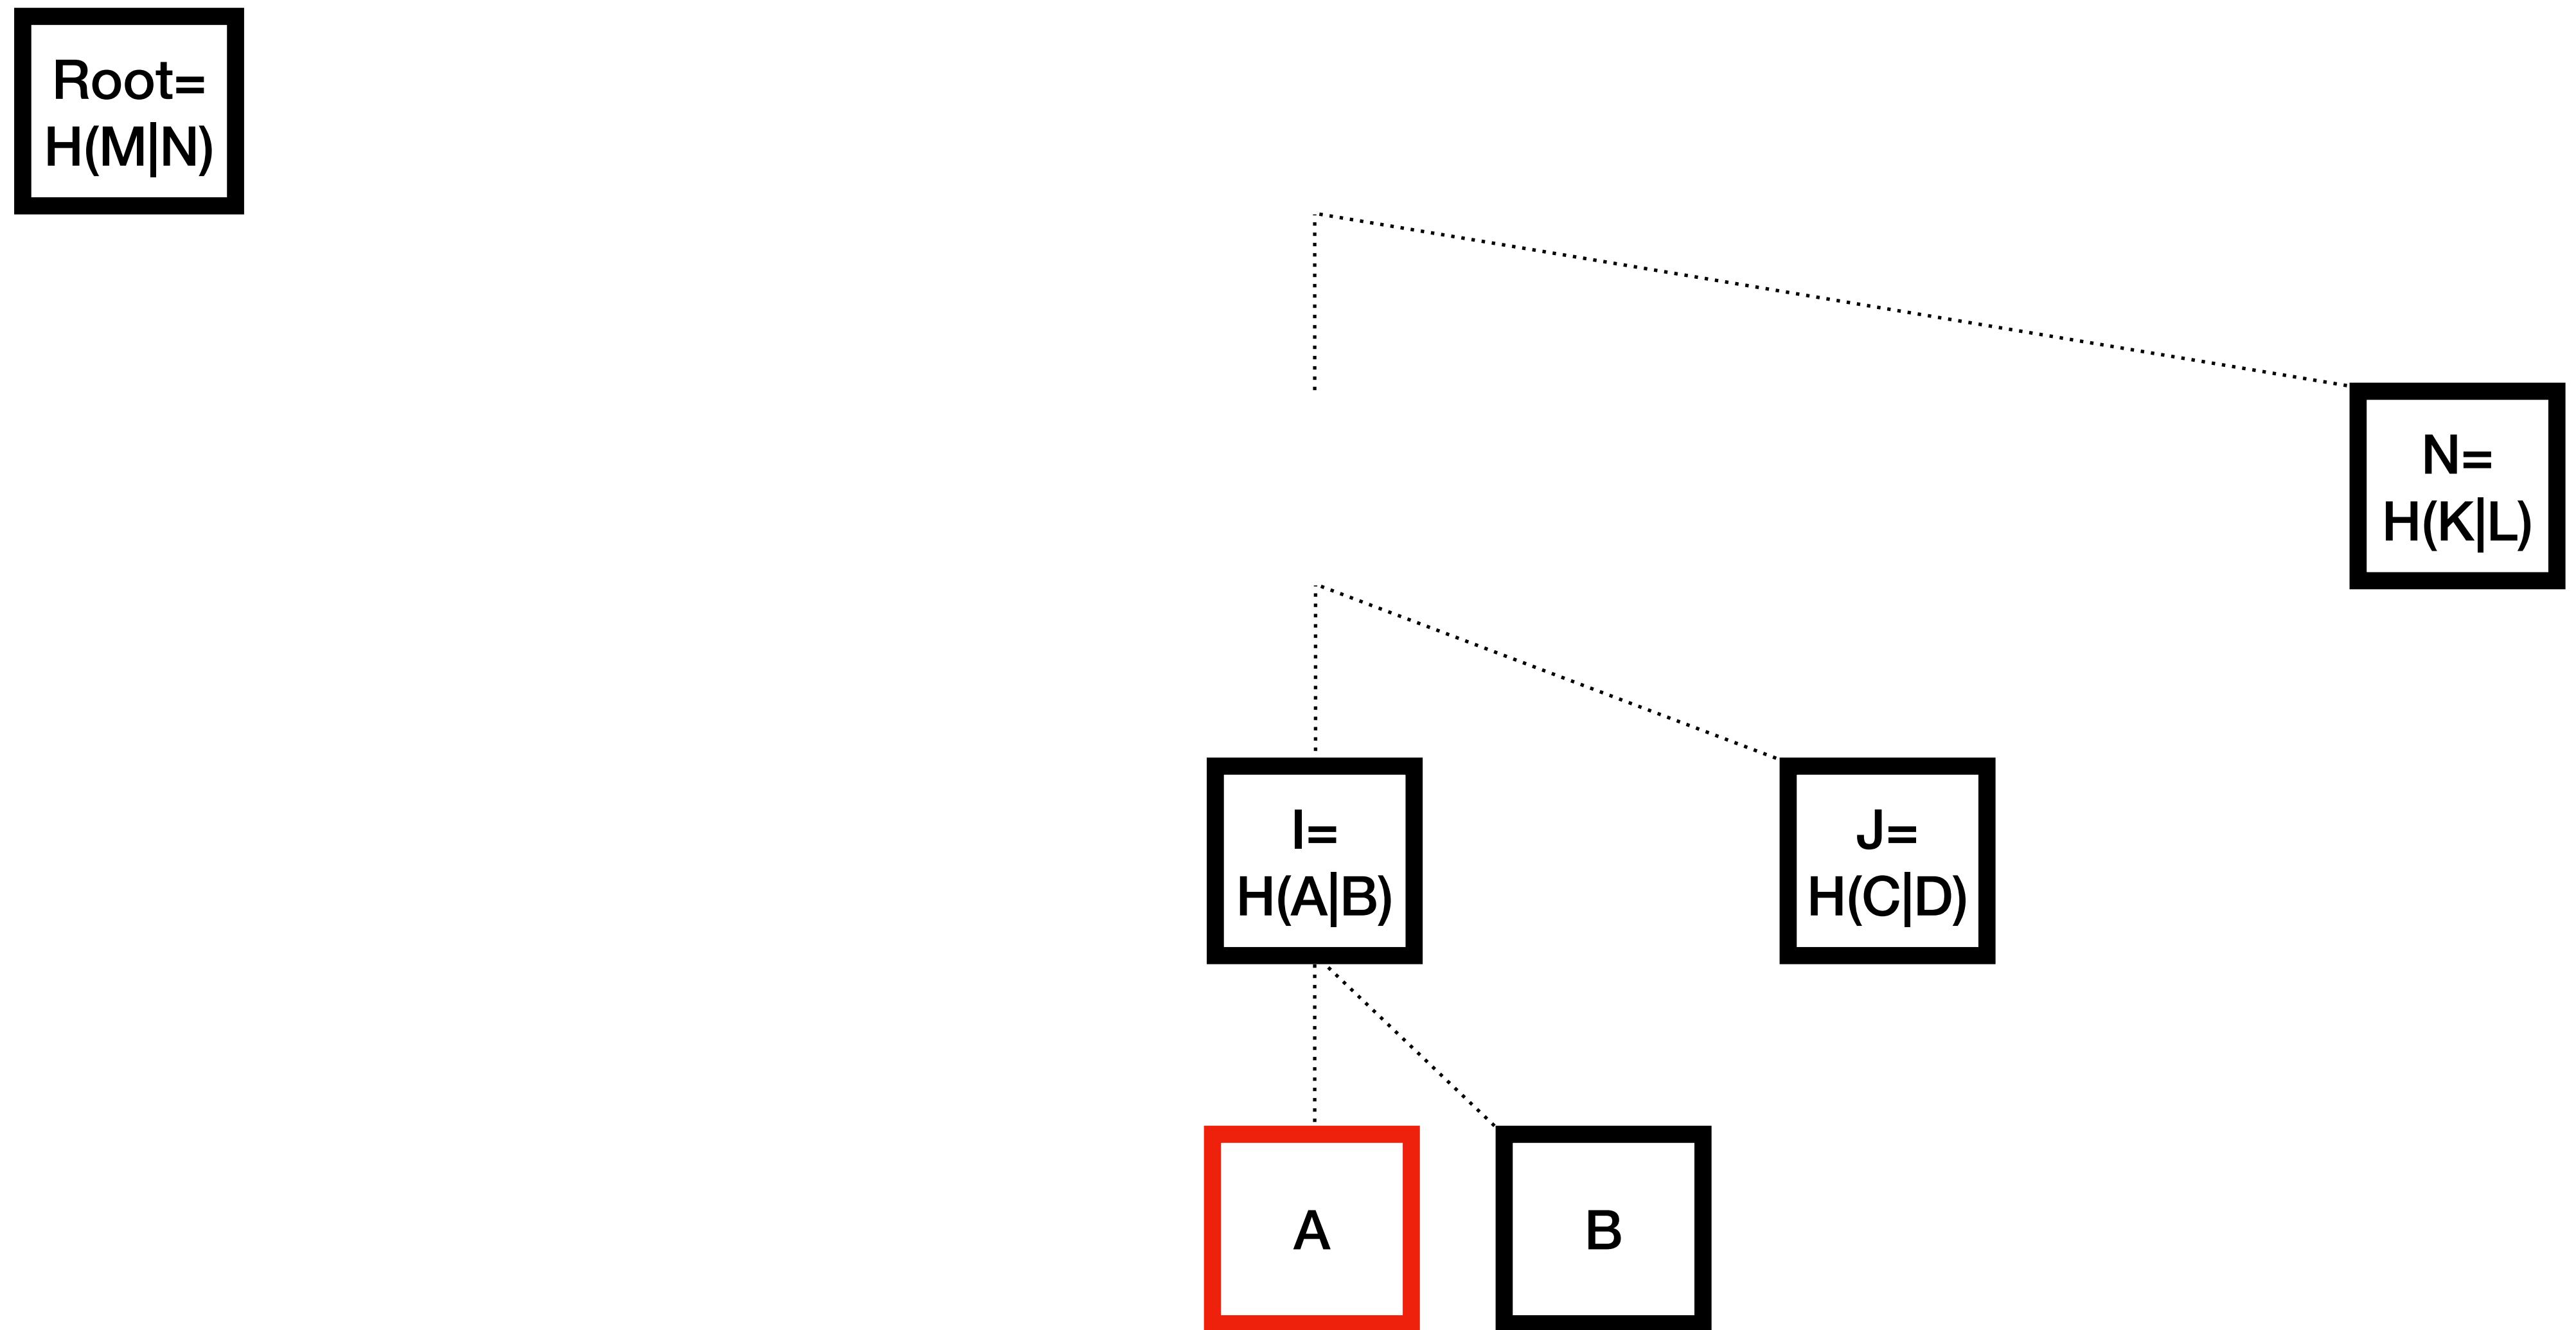
Root=
 $H(M|N)$



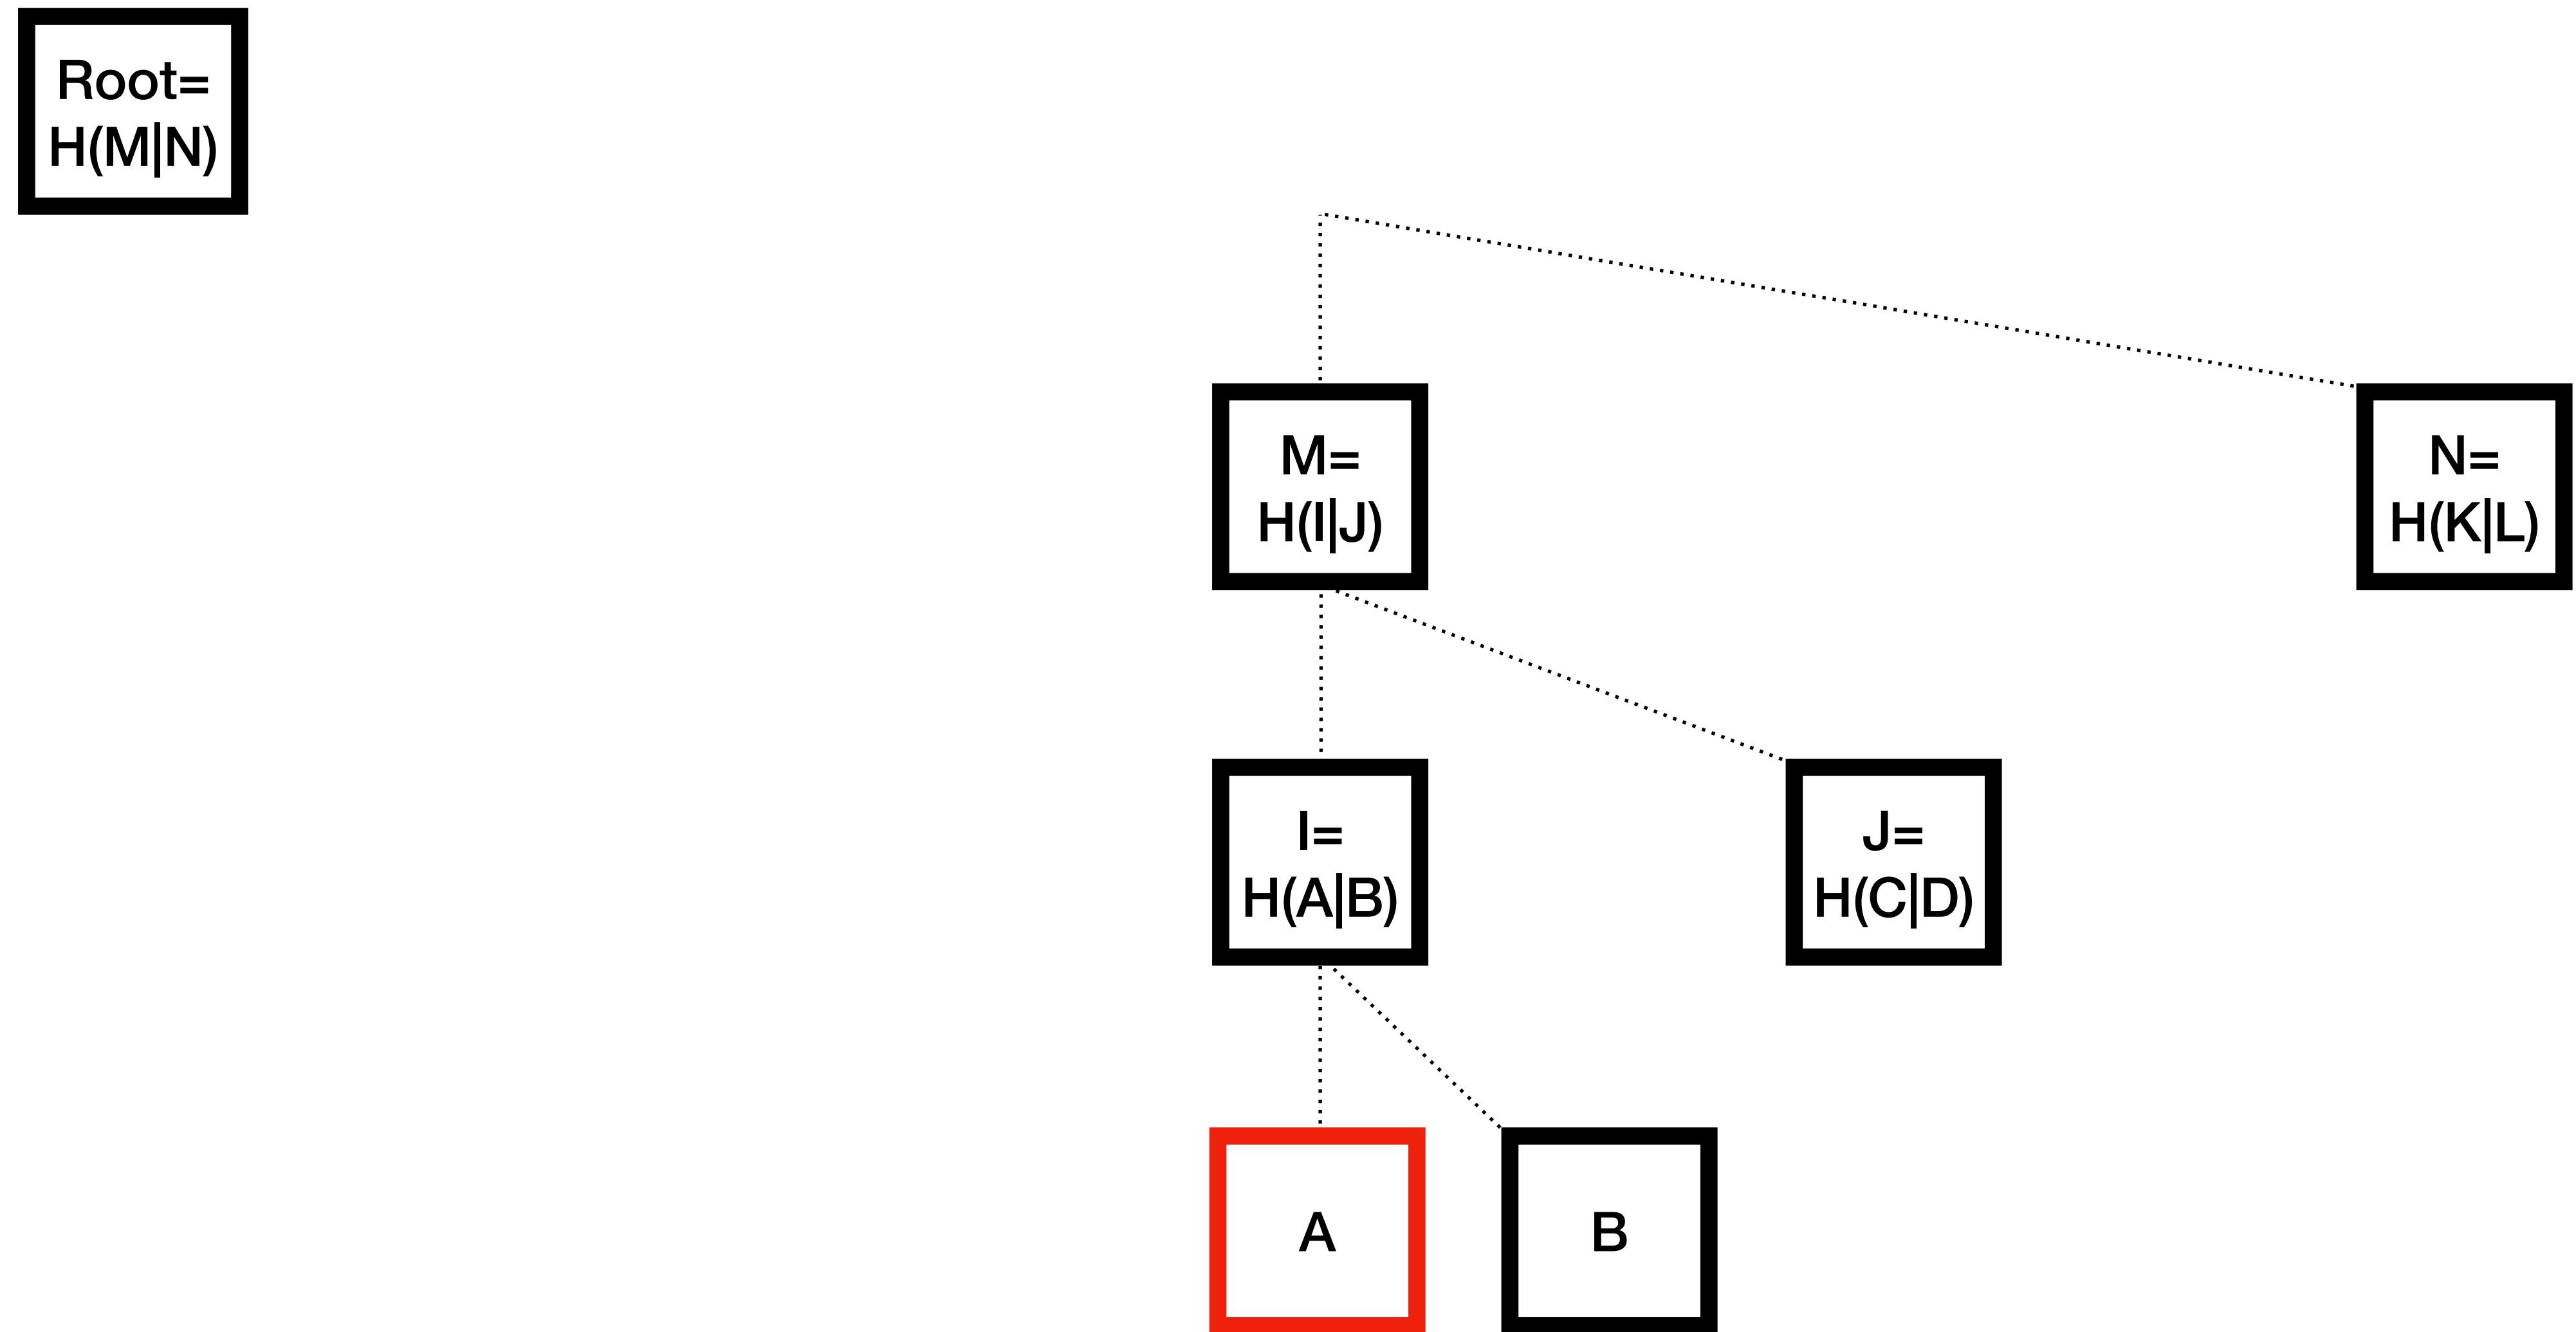
Receive the proof: B, J, N



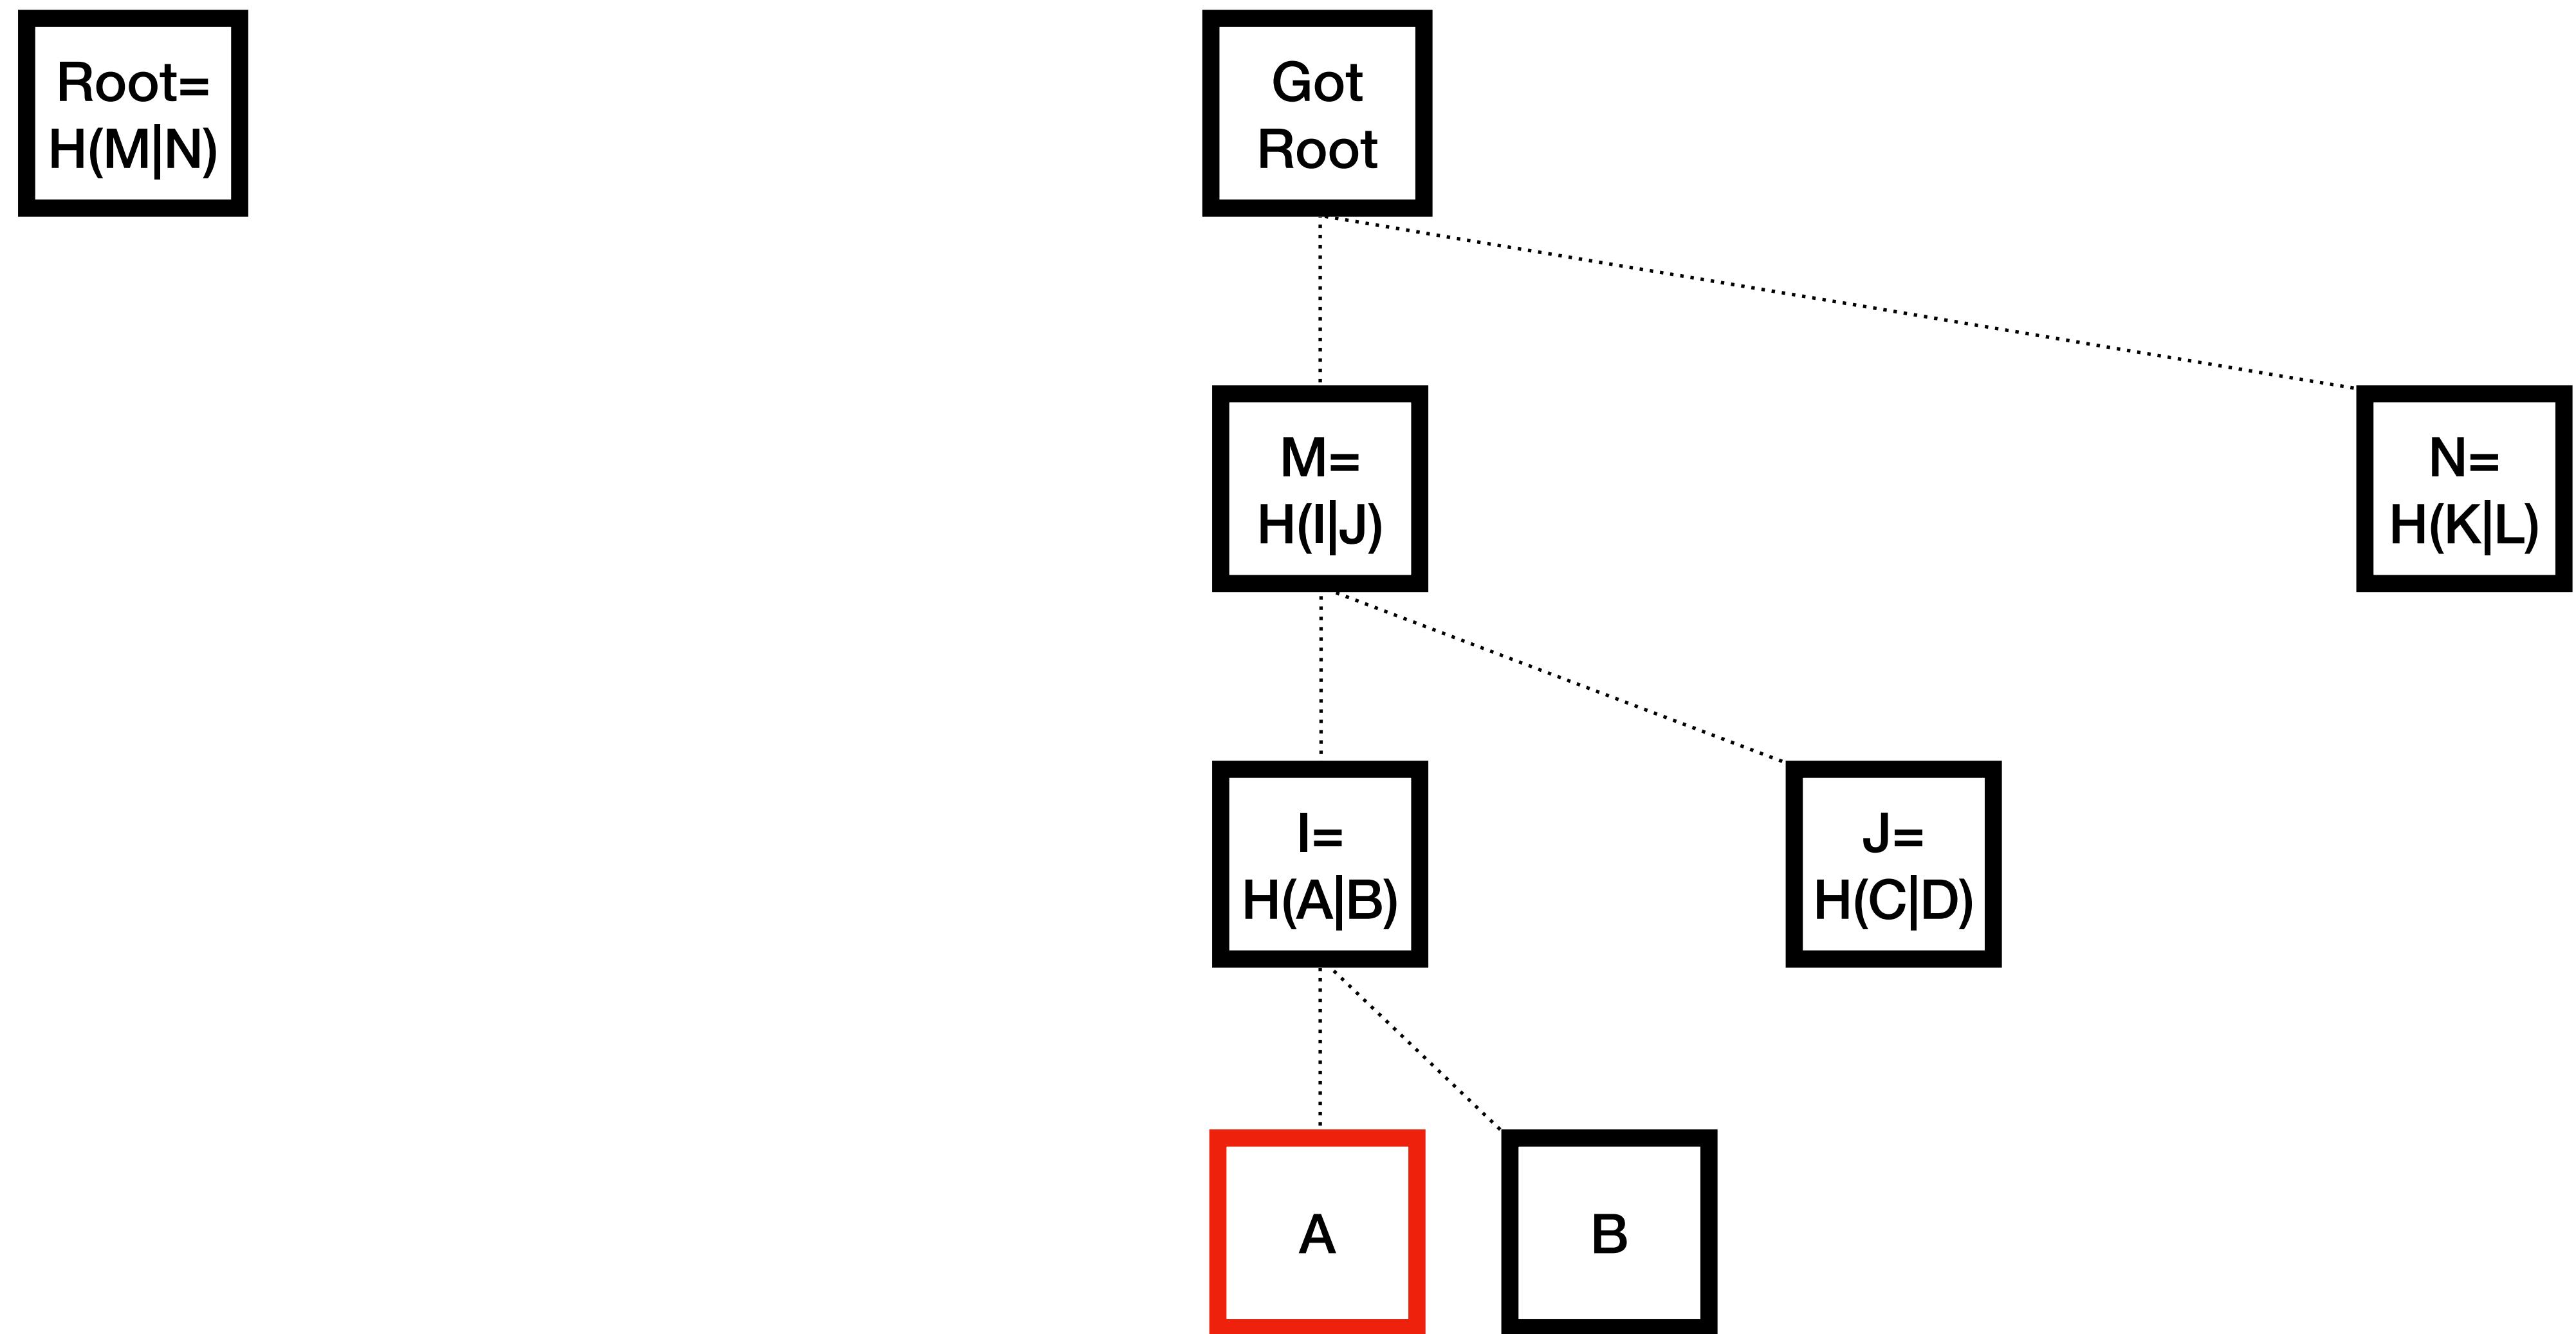
Calculate the hash for I



Calculate the hash for M



Calculate Root



Compare roots

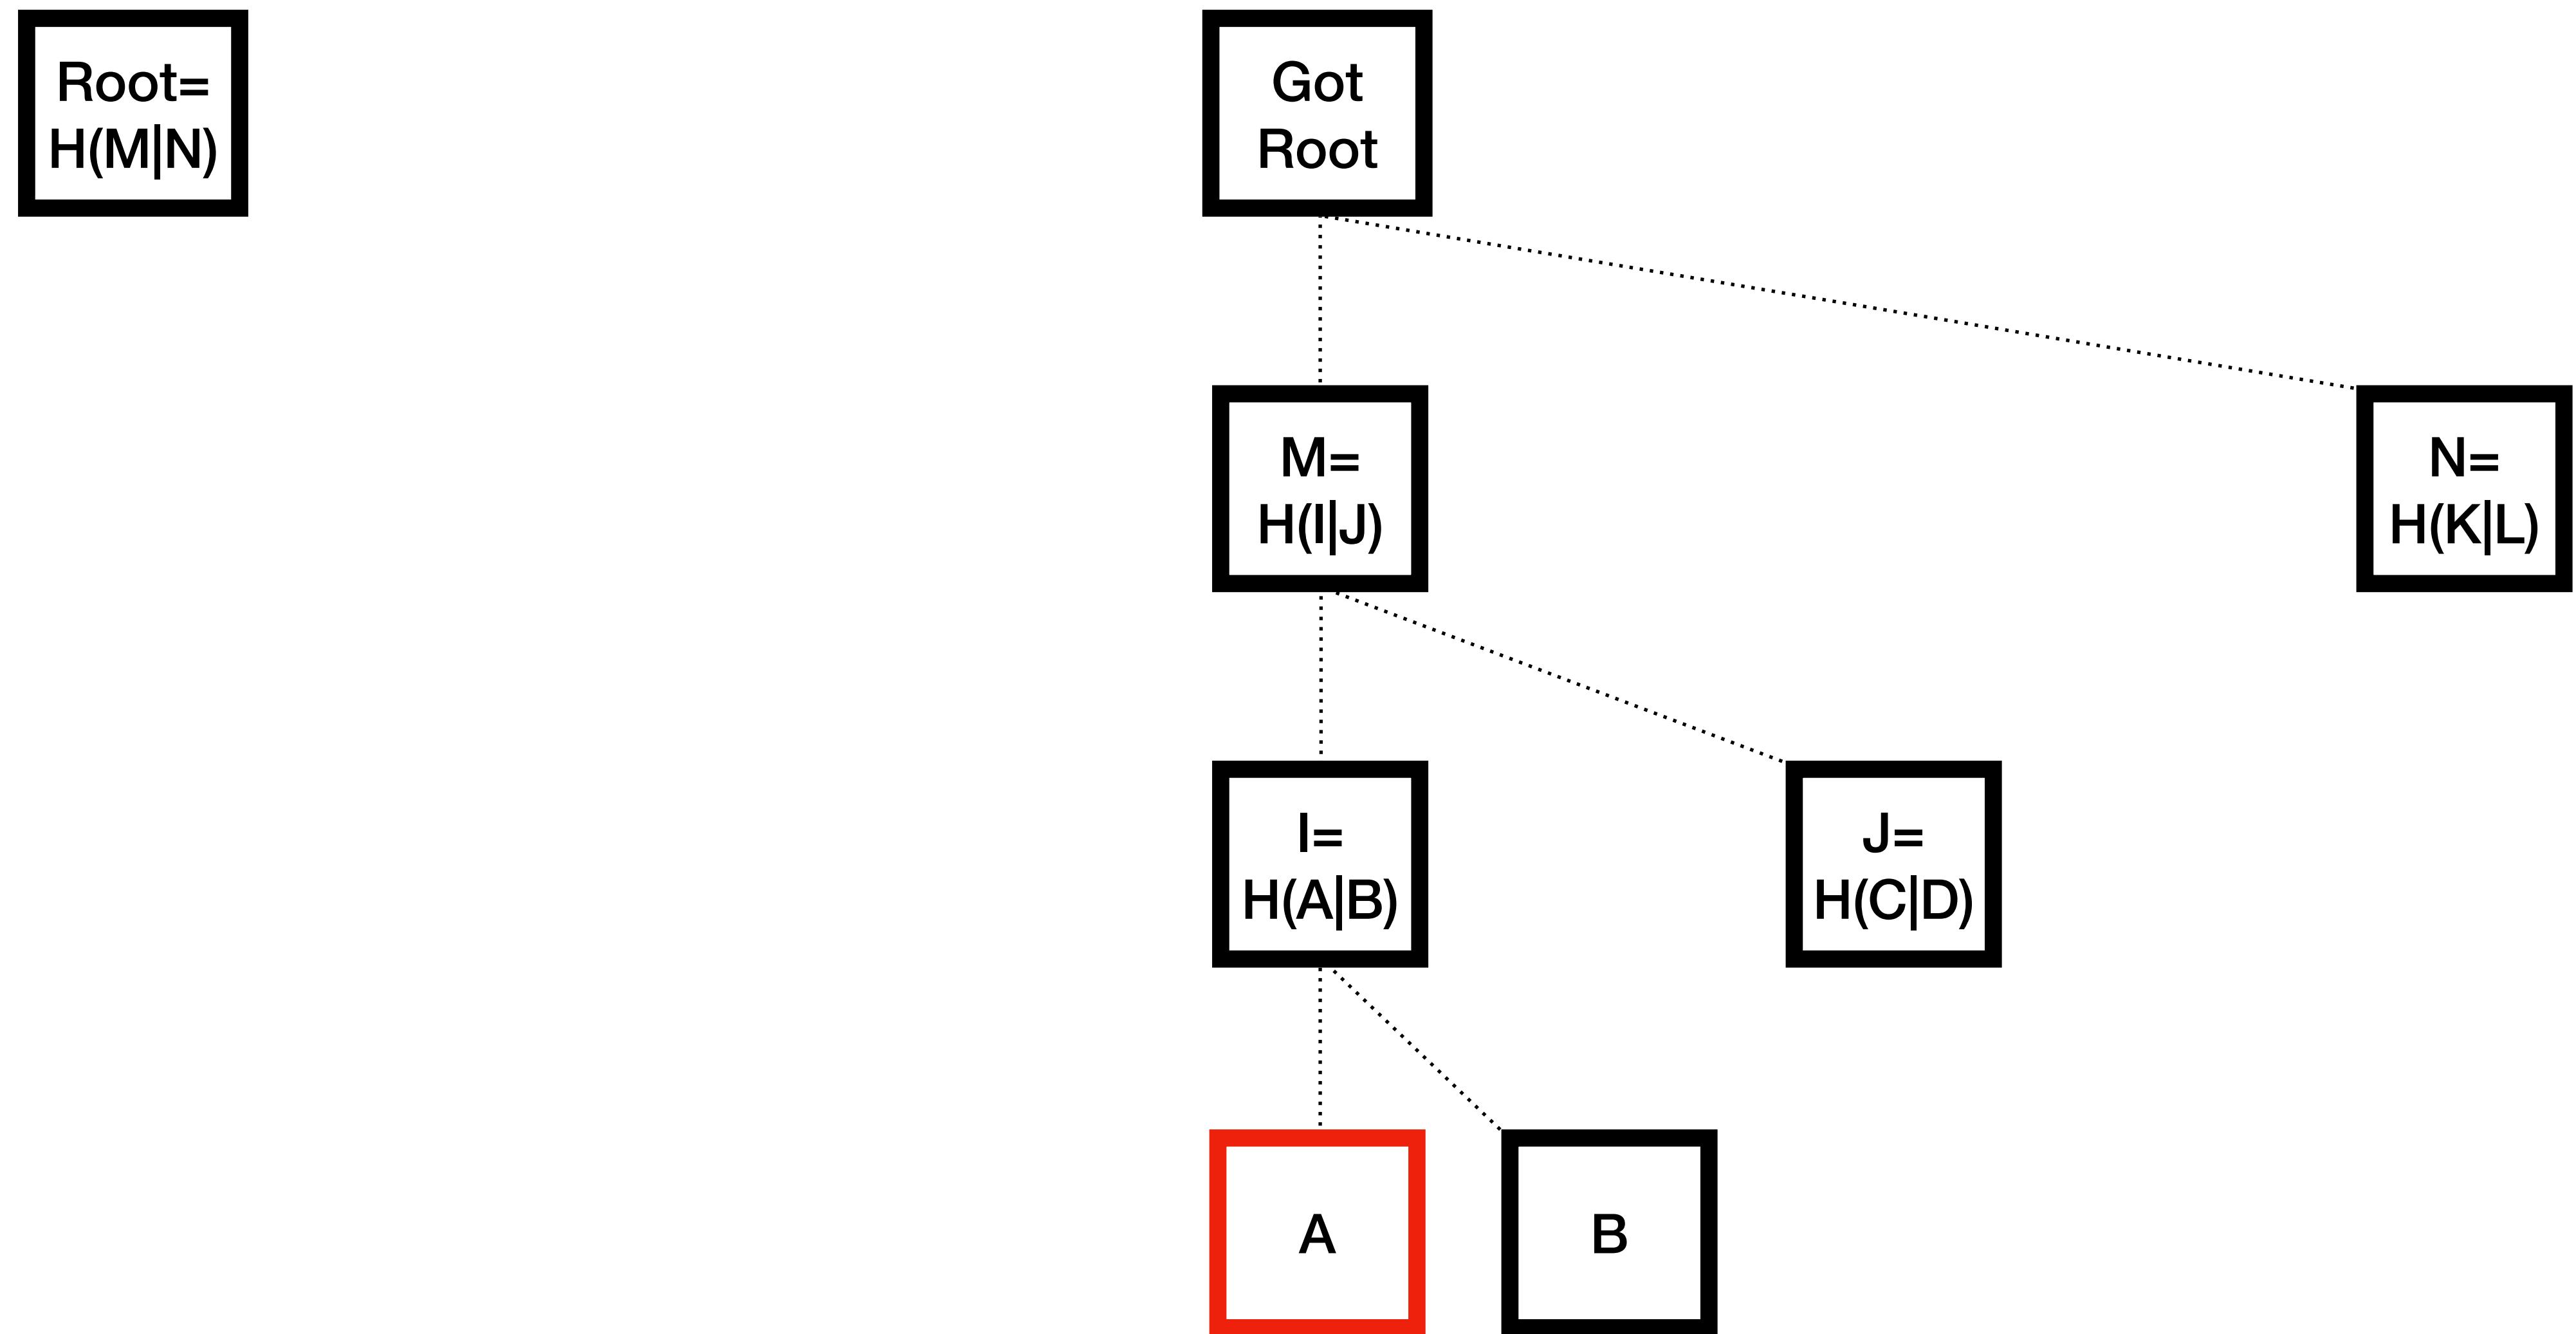
Continue if the roots are equal. Ban peer if not equal

Root= $H(M|N)$

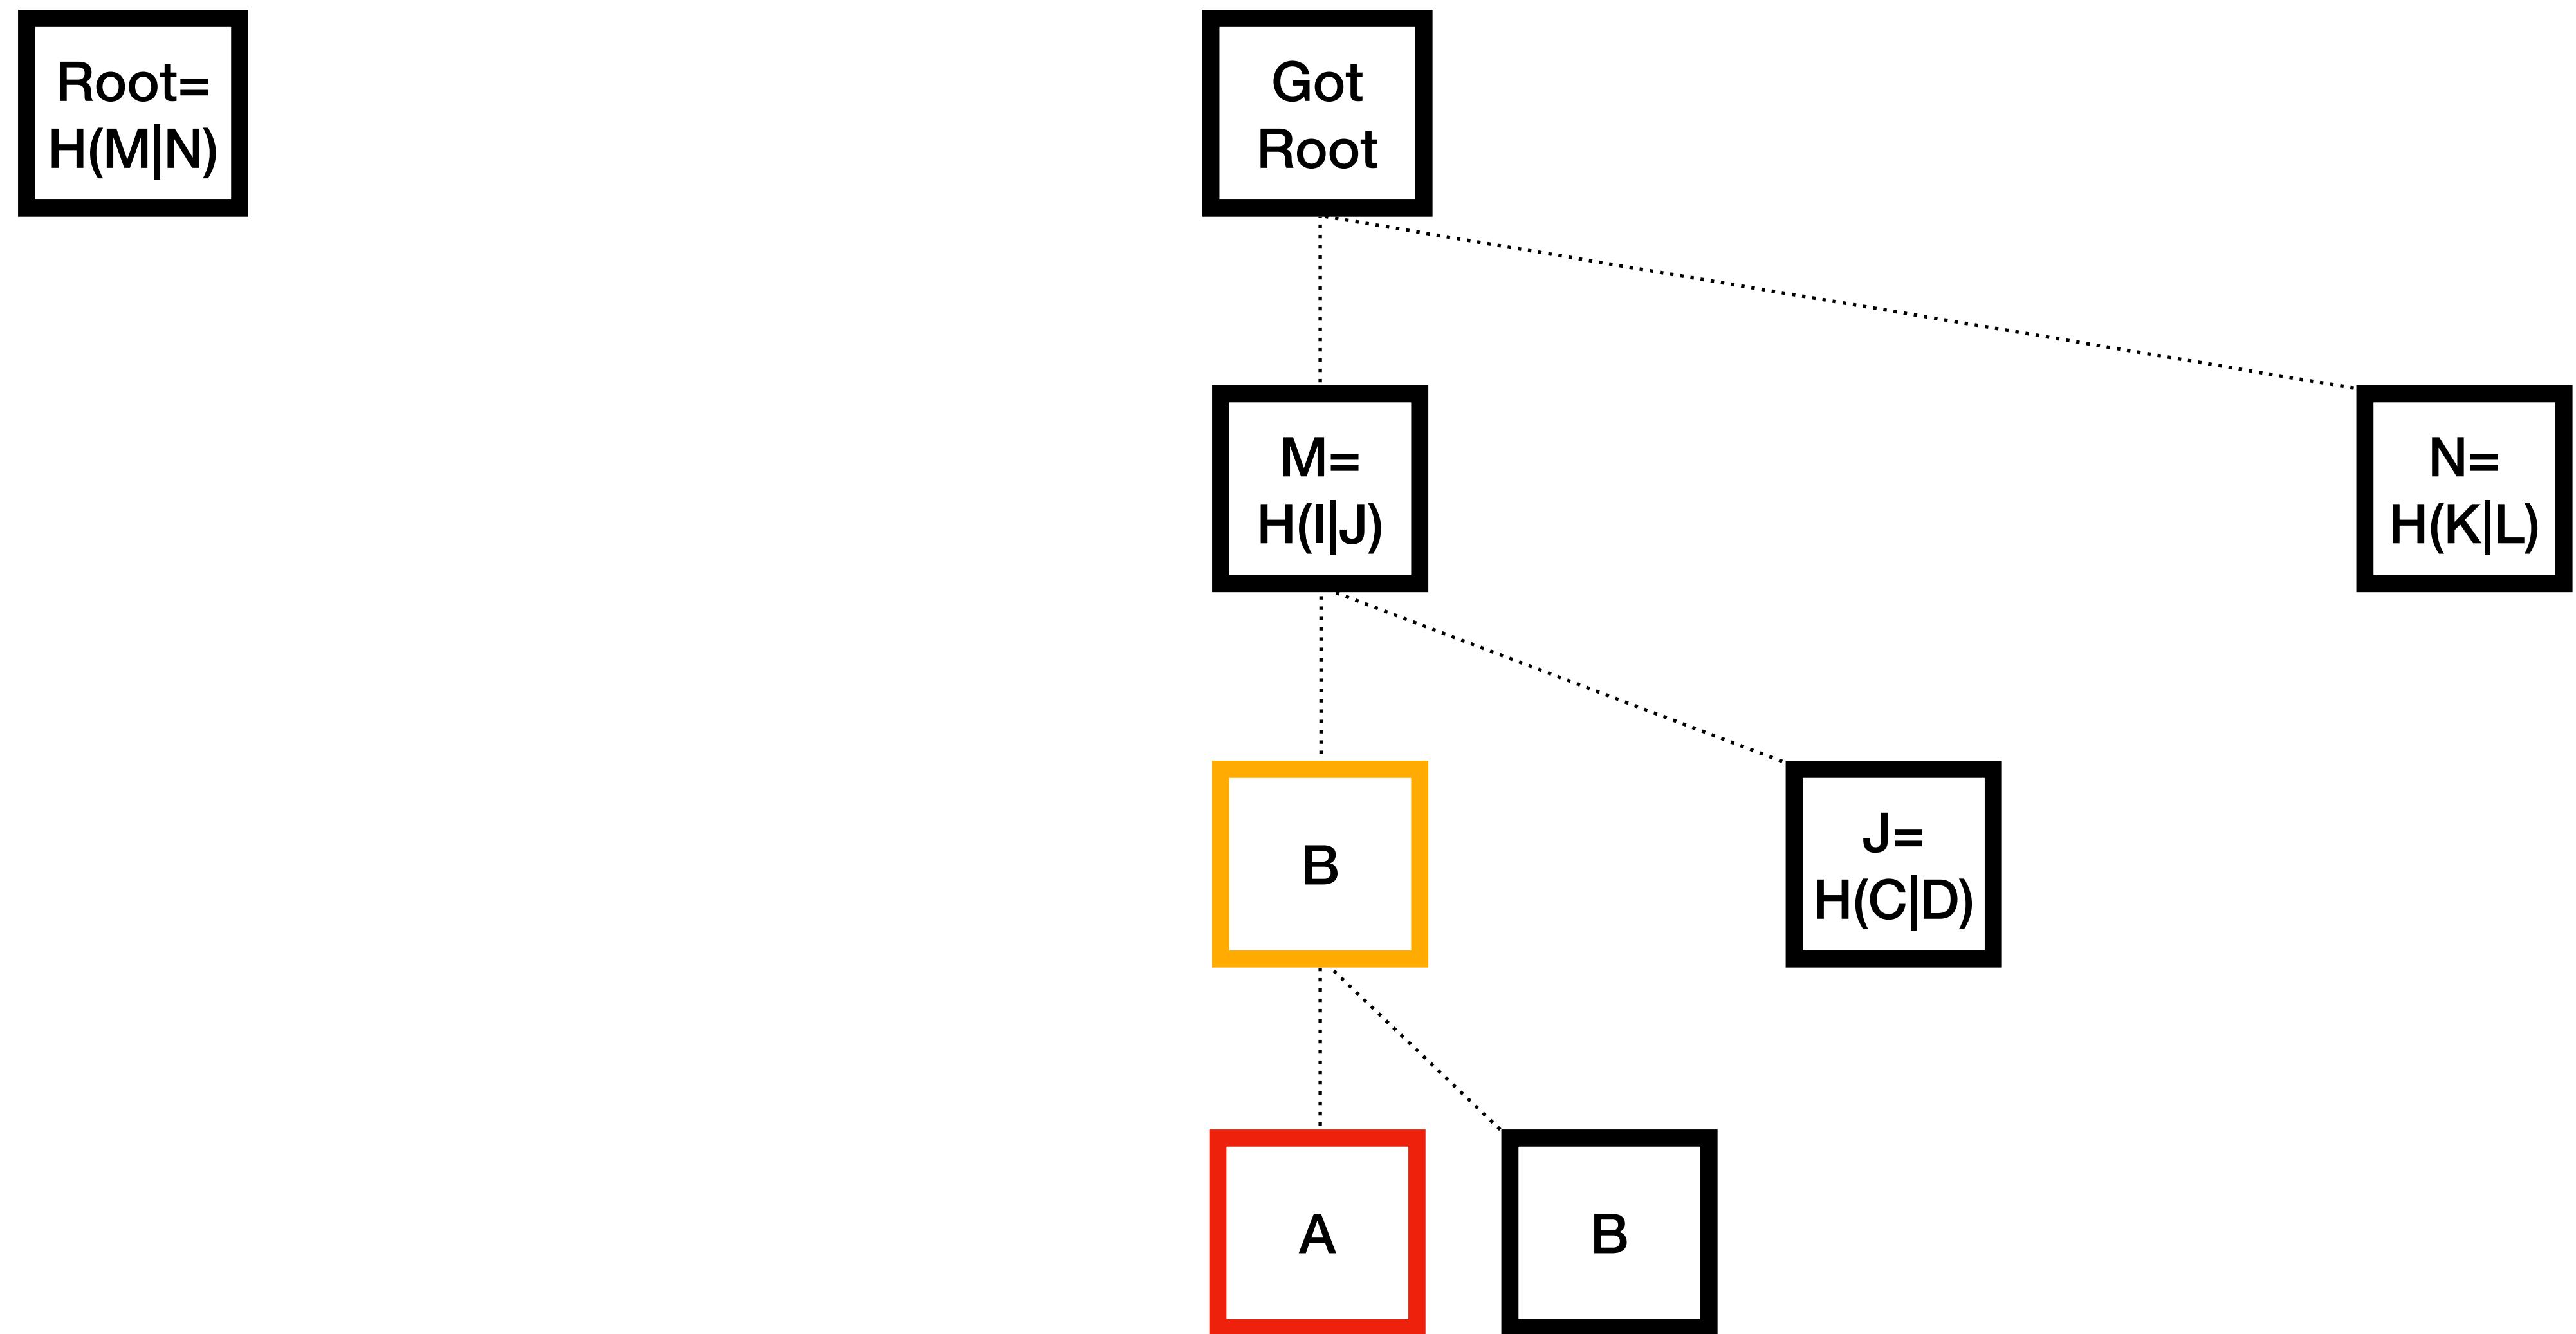
==

Got Root

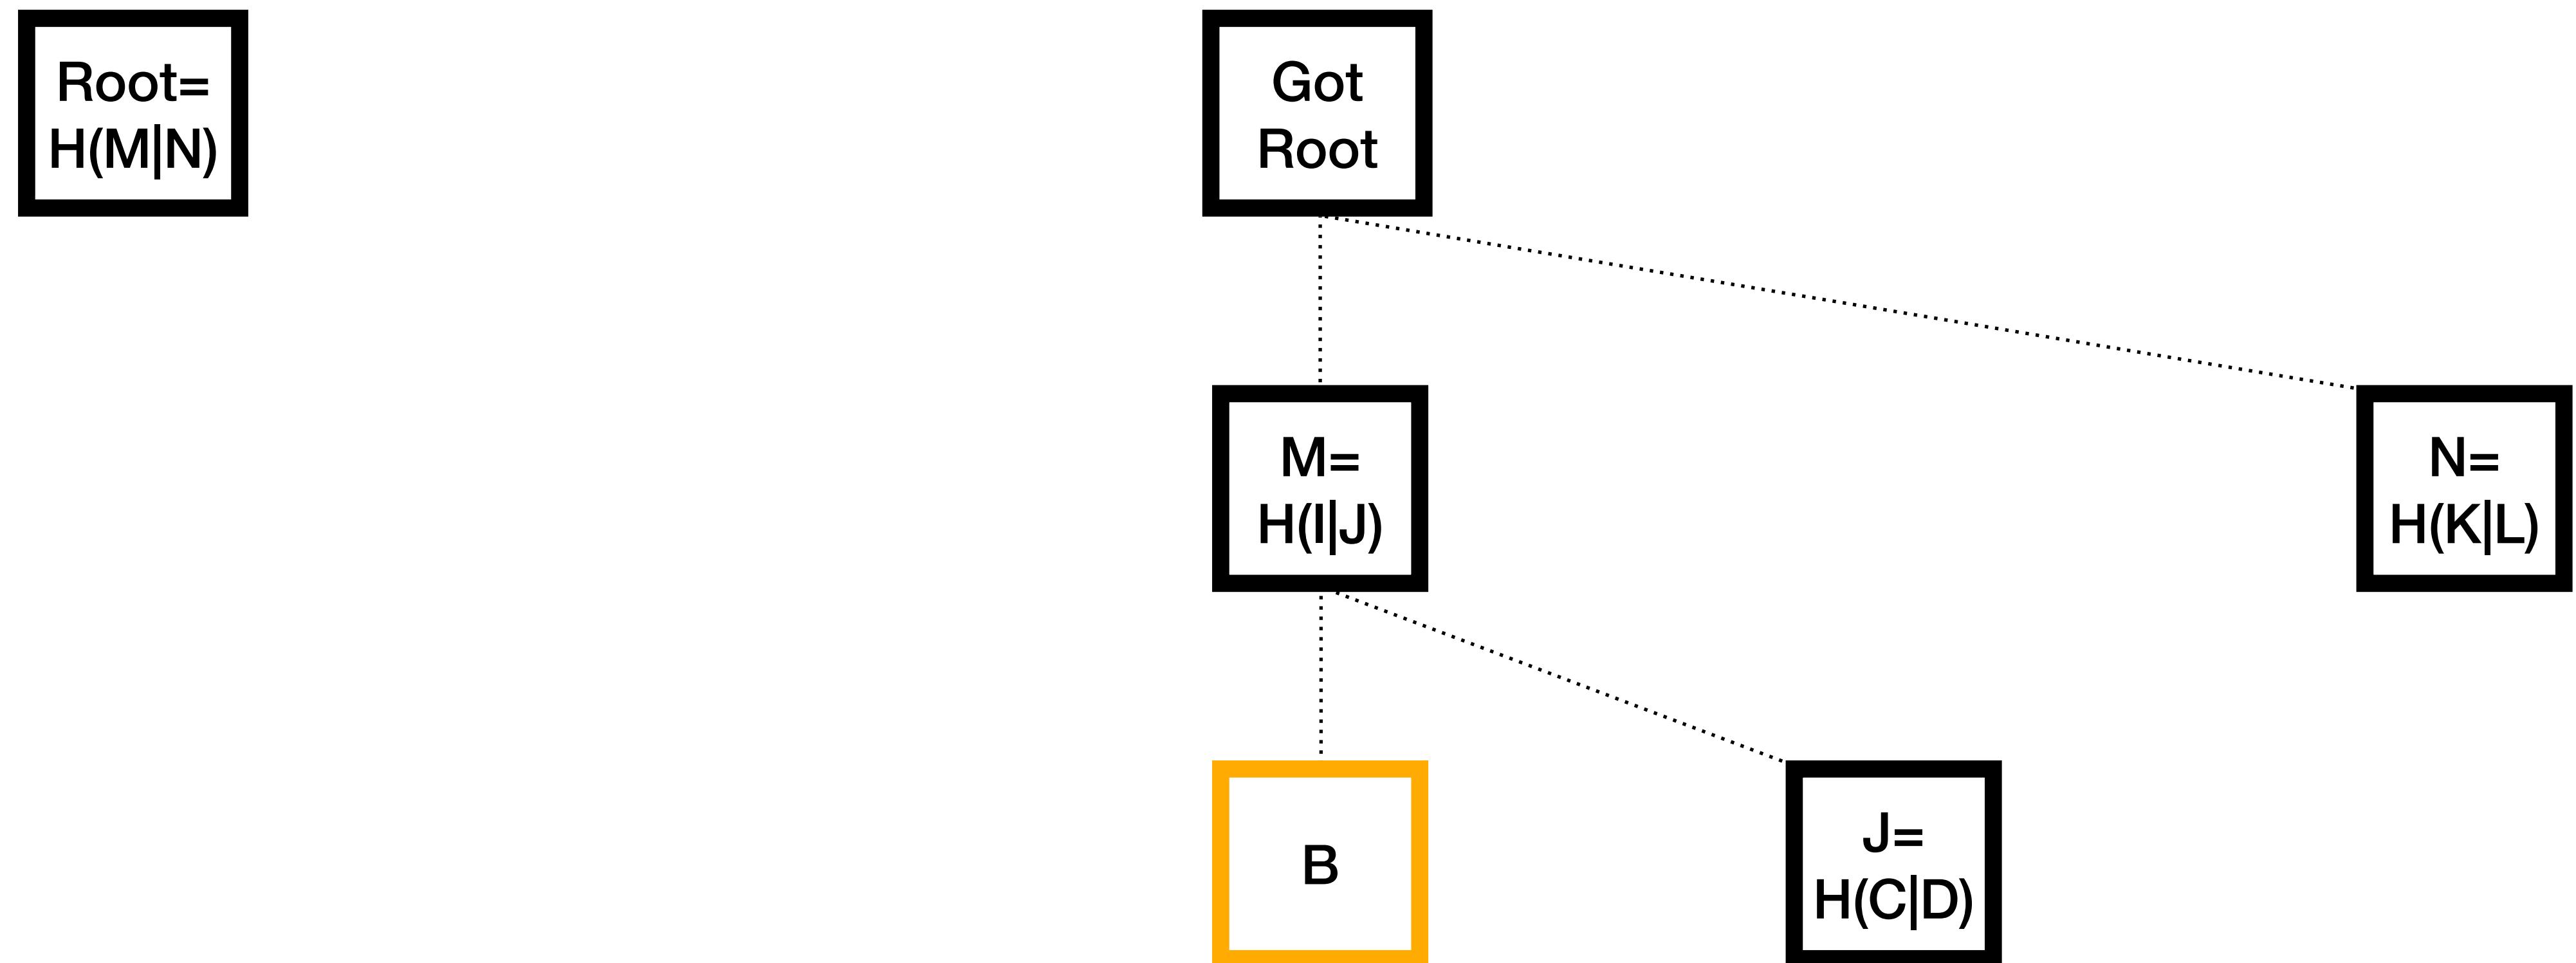
Calculate root after deletion



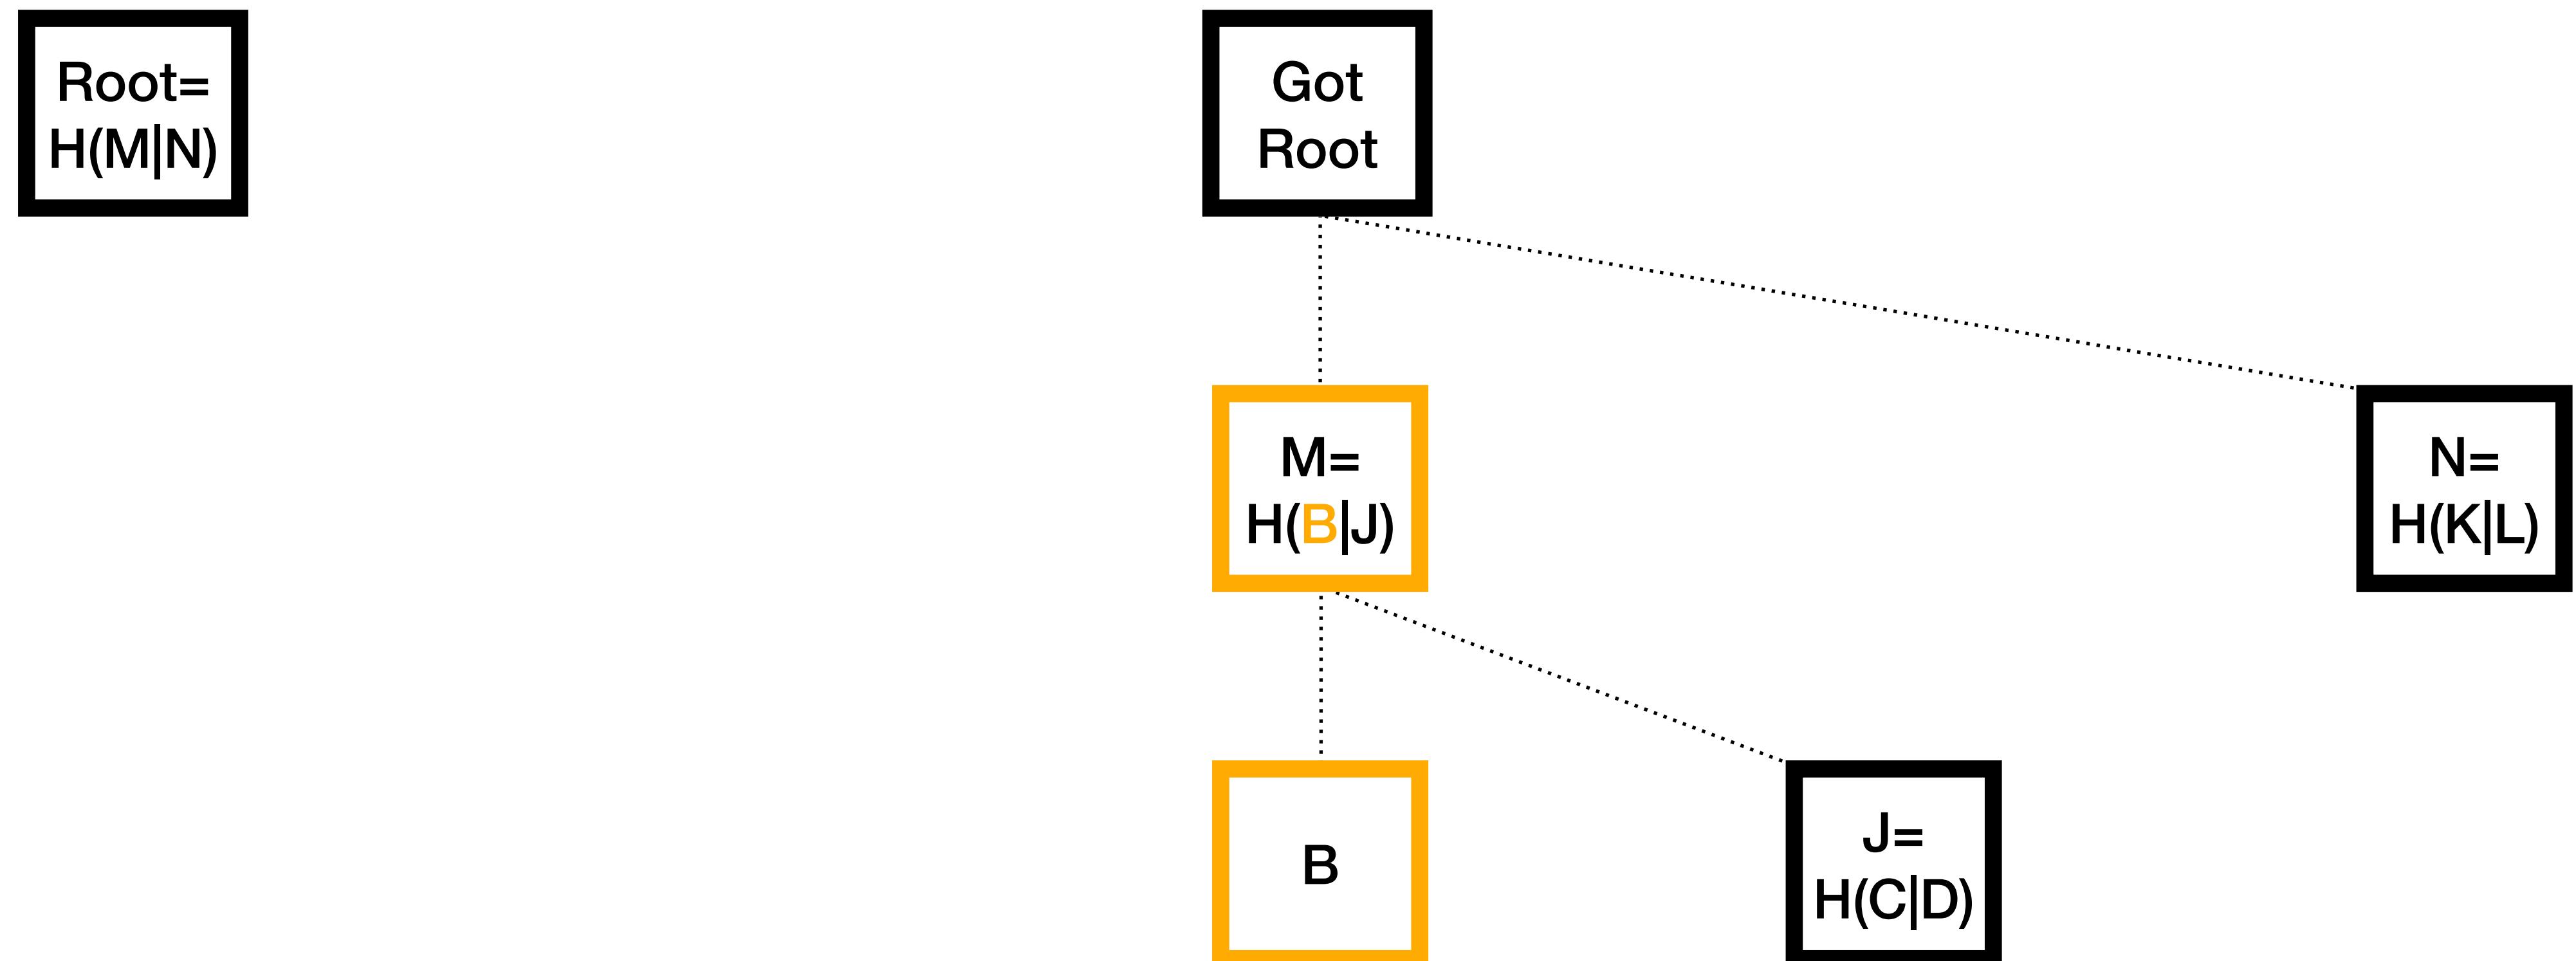
Move up B to I



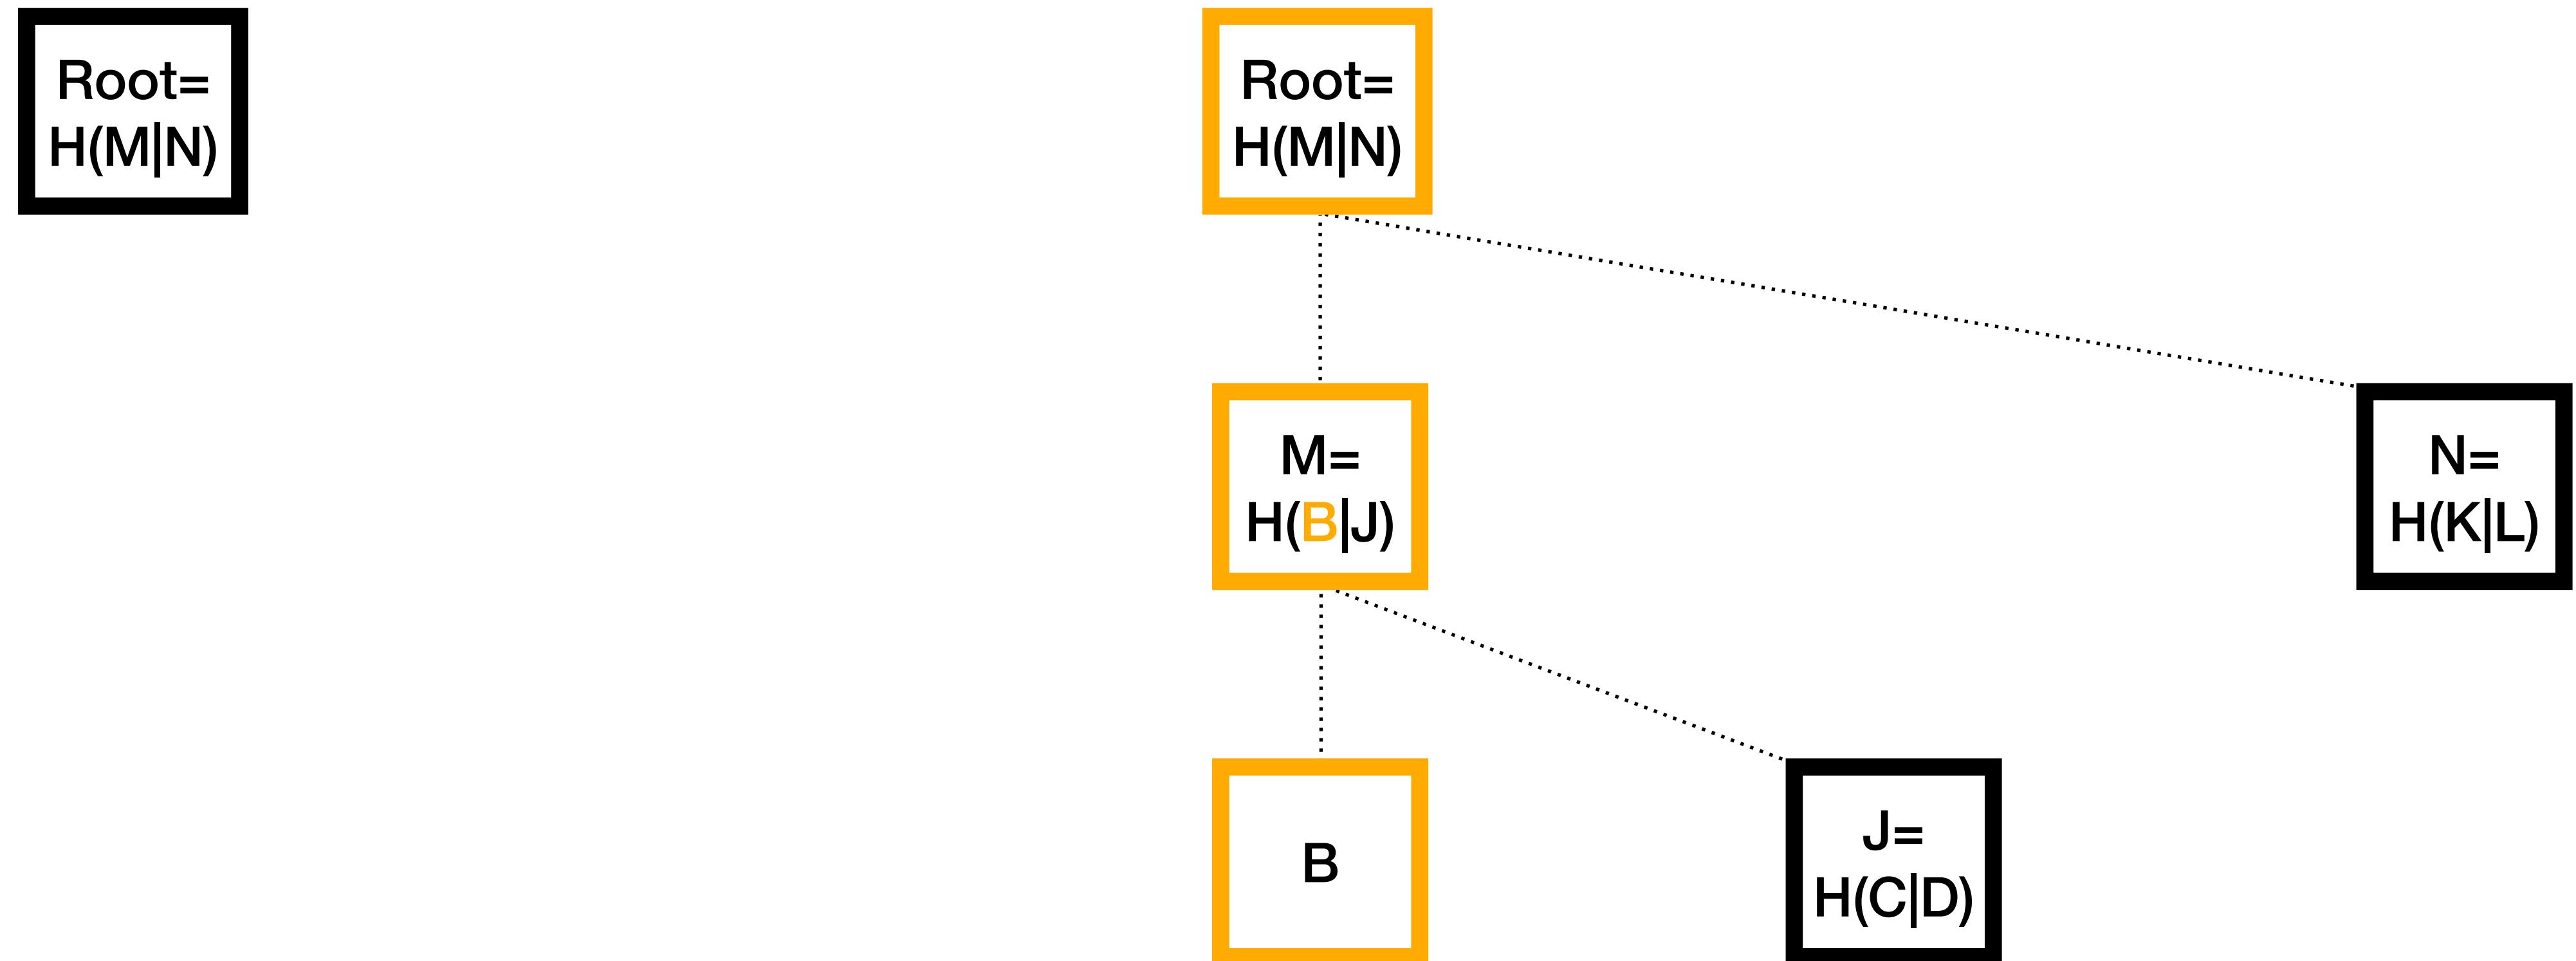
Remove old nodes for A&B



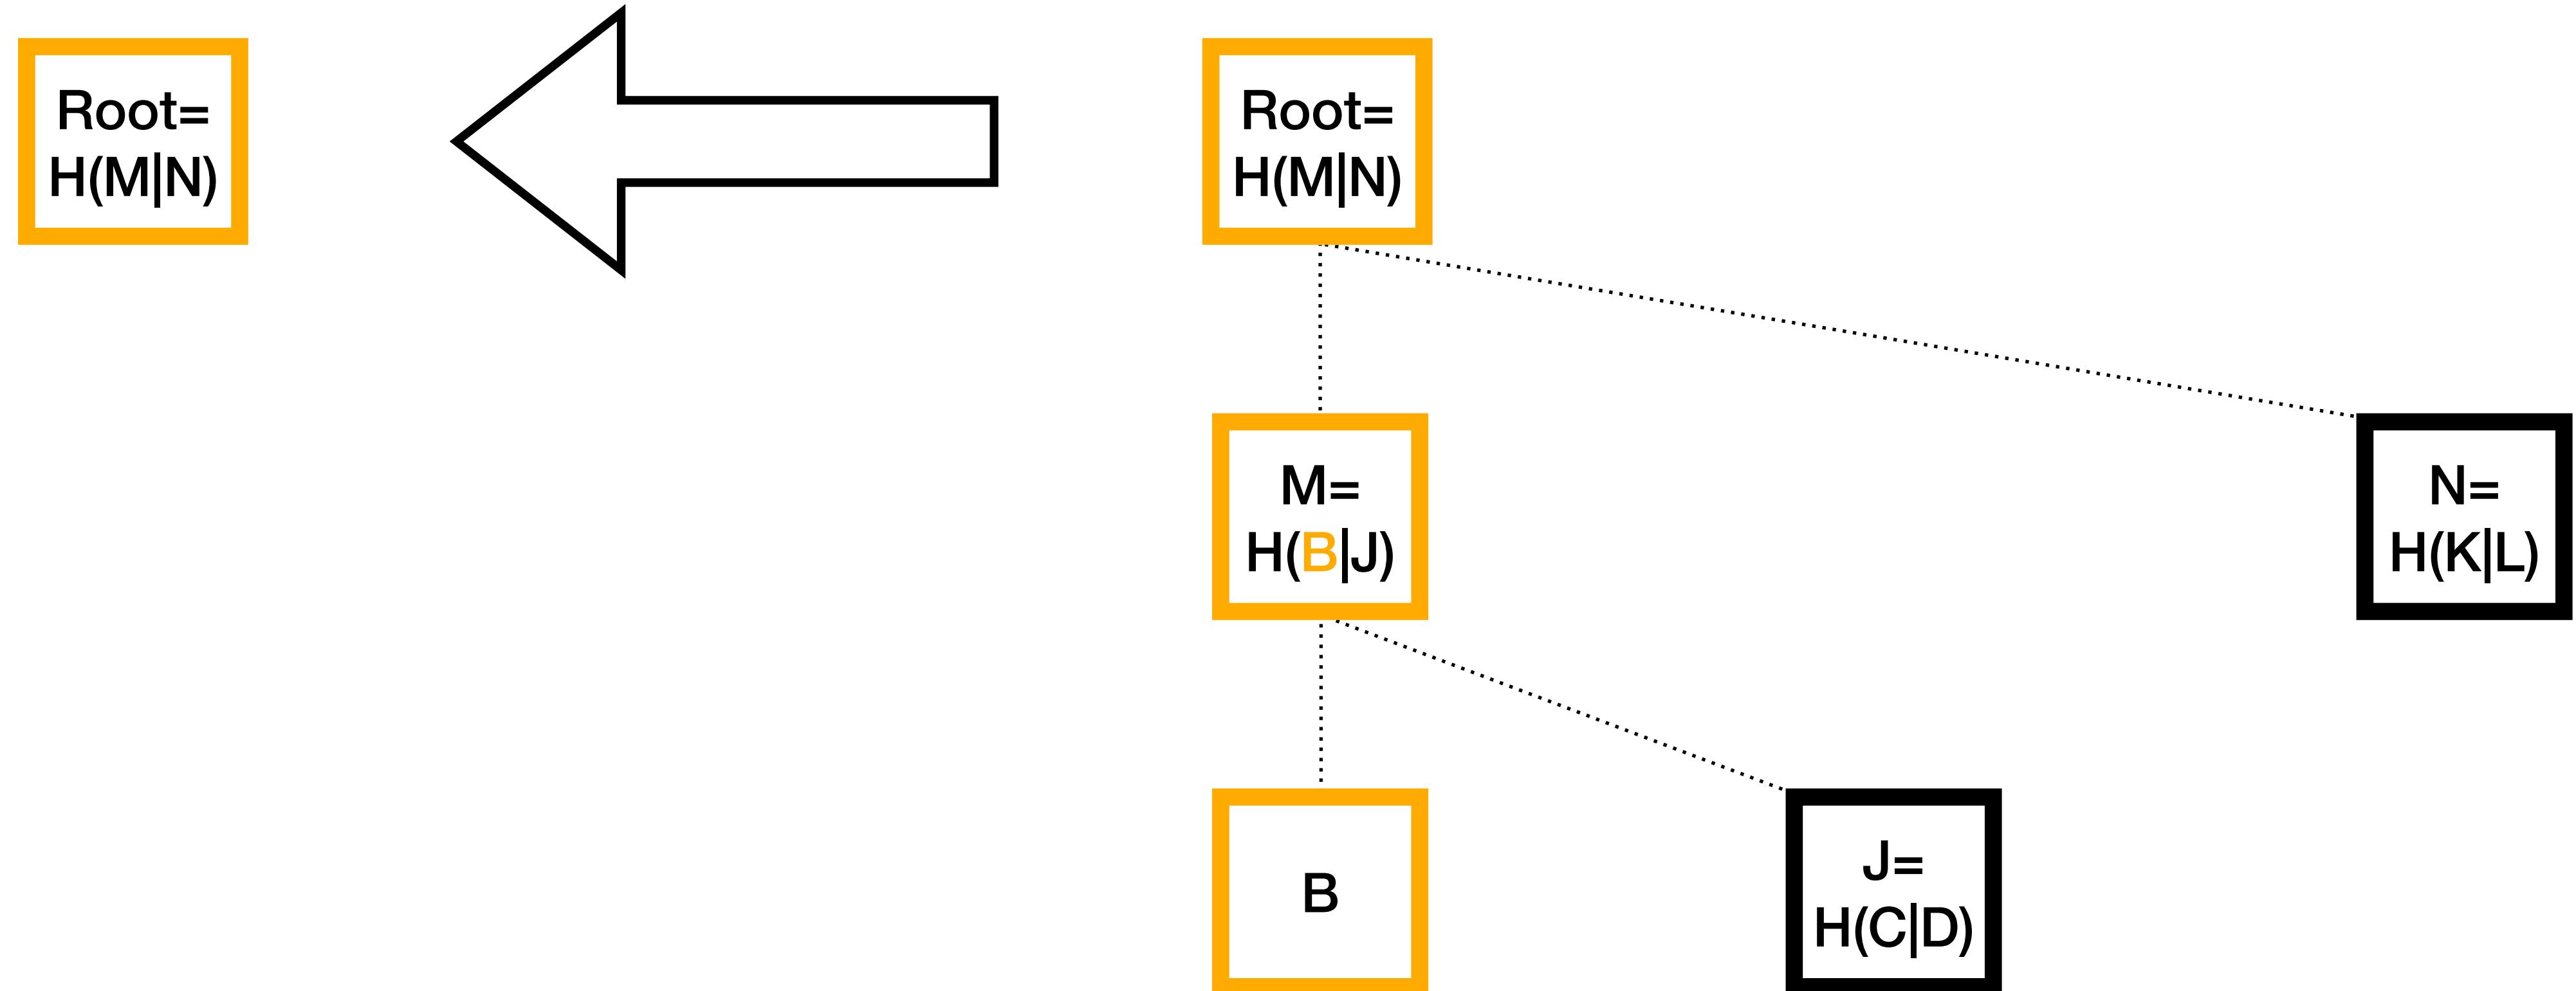
Calculate $H(B|J)$



Calculate new root



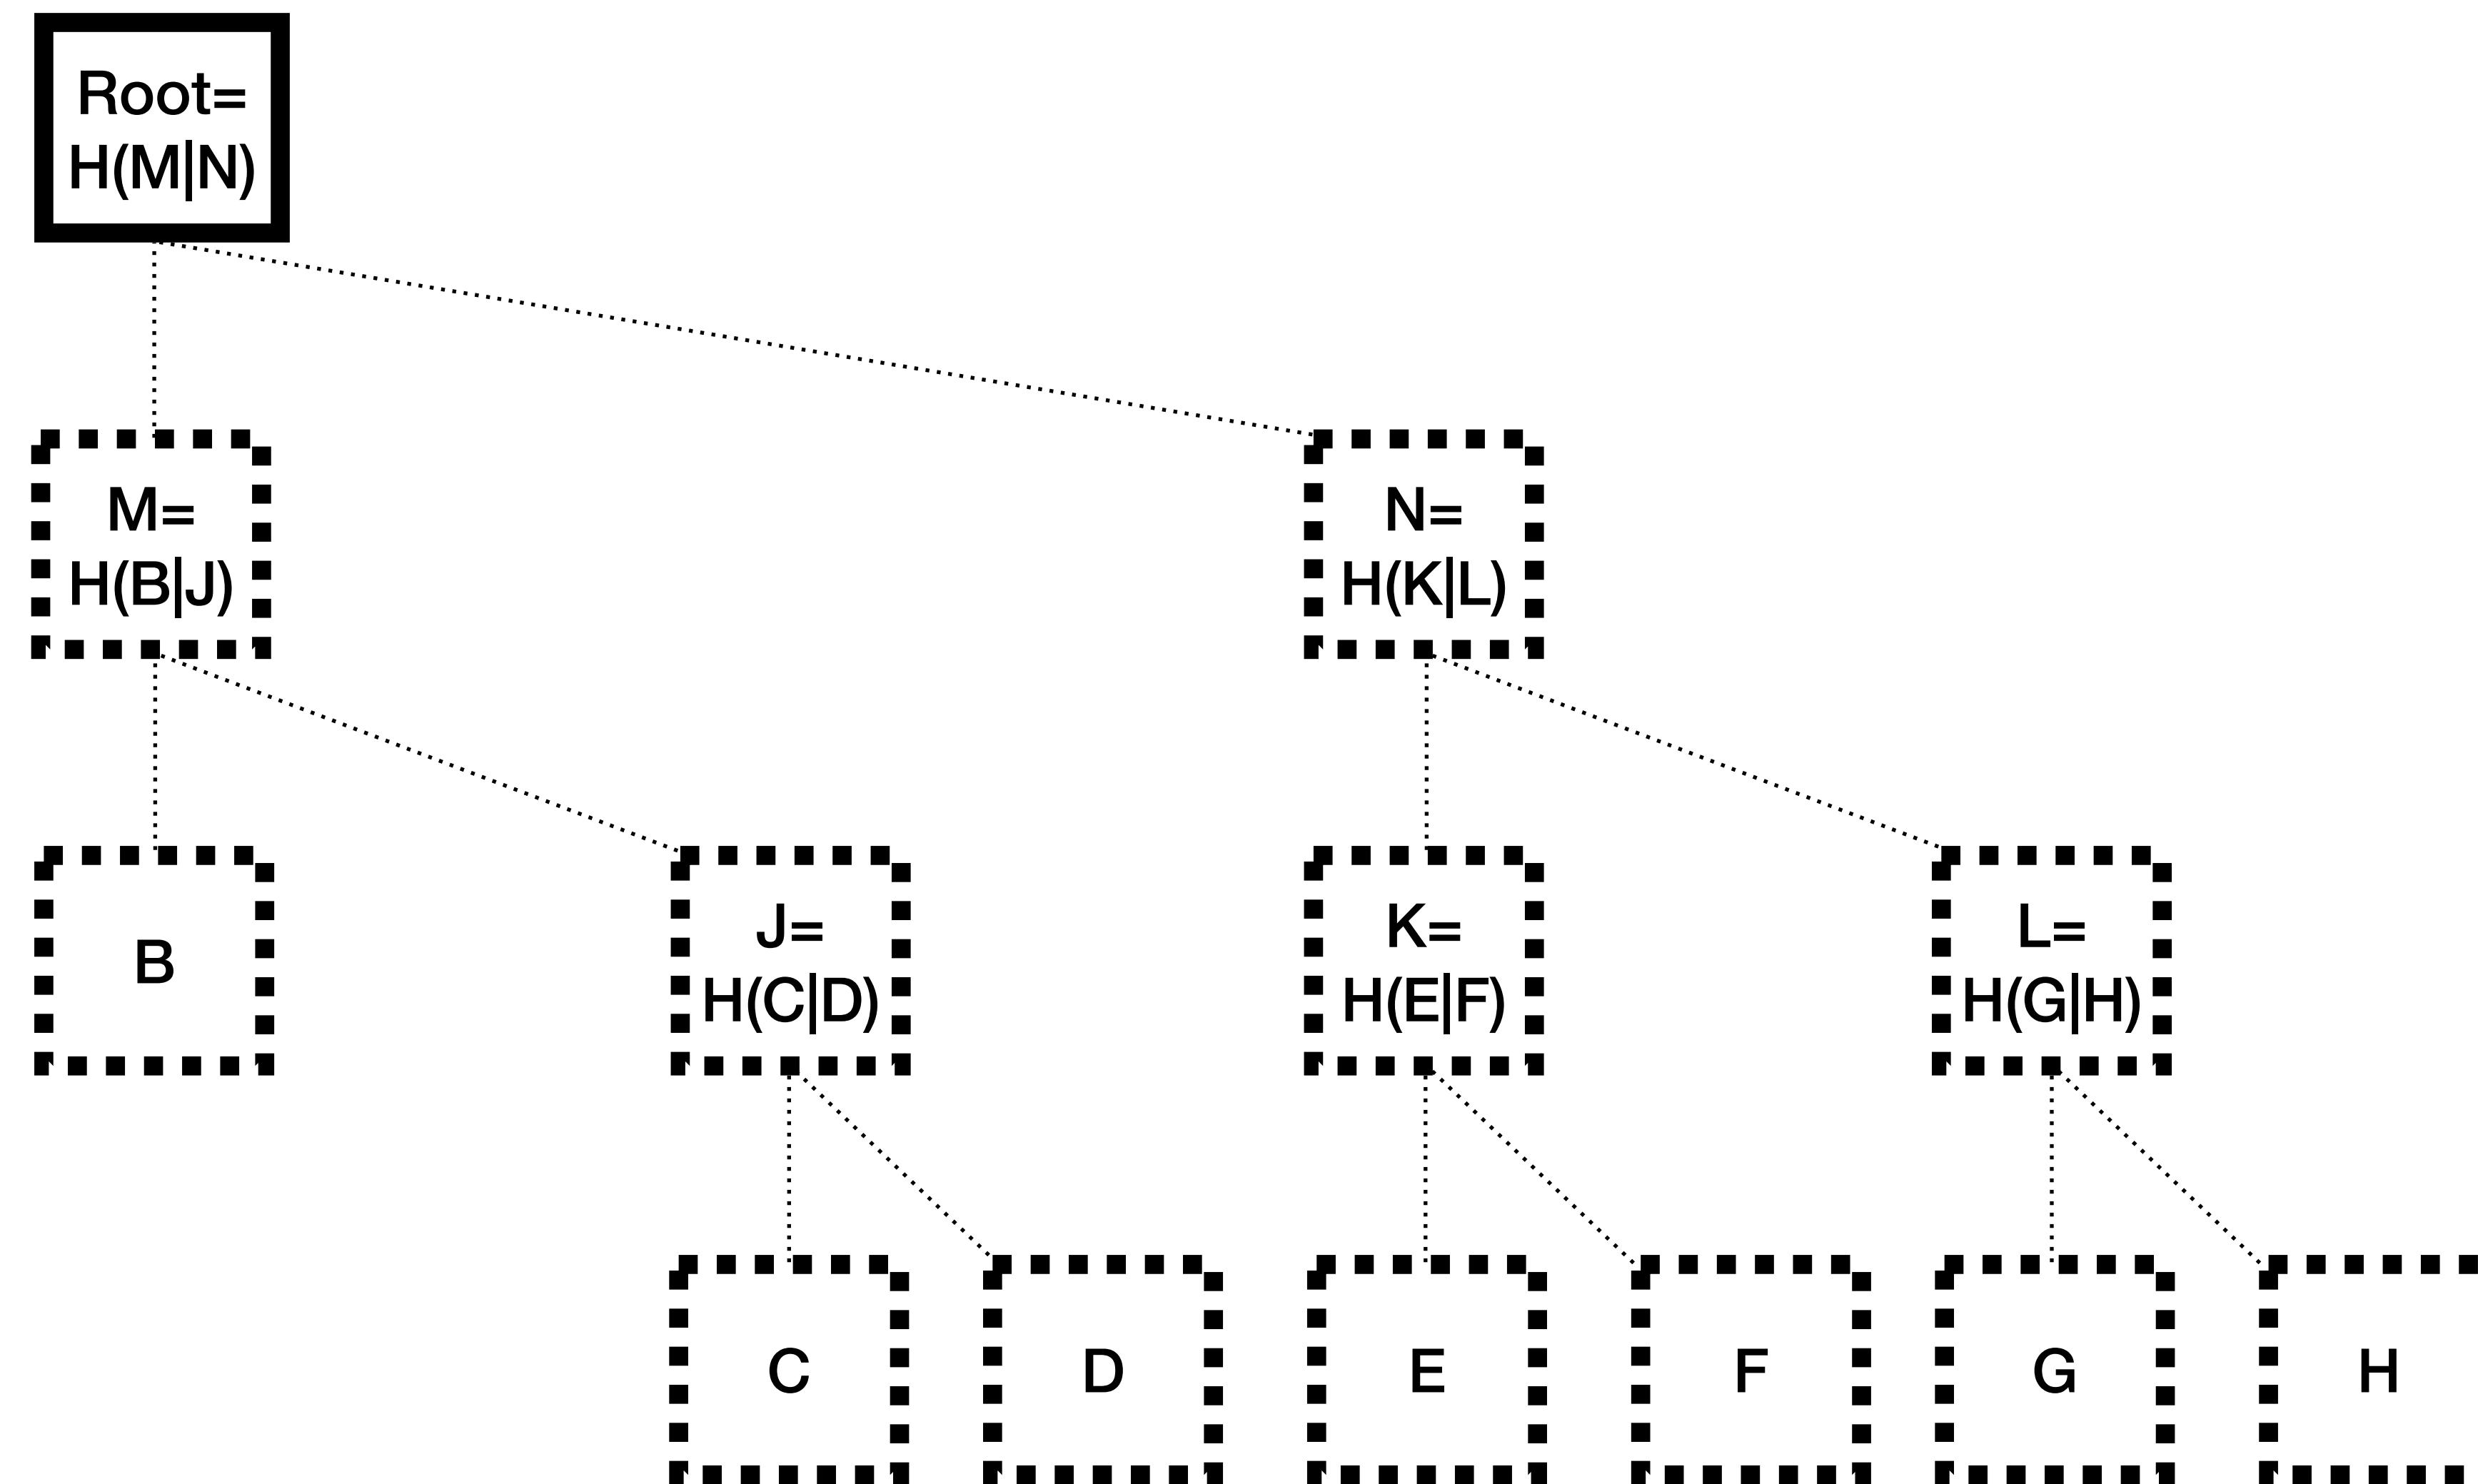
Copy over the new root to be saved



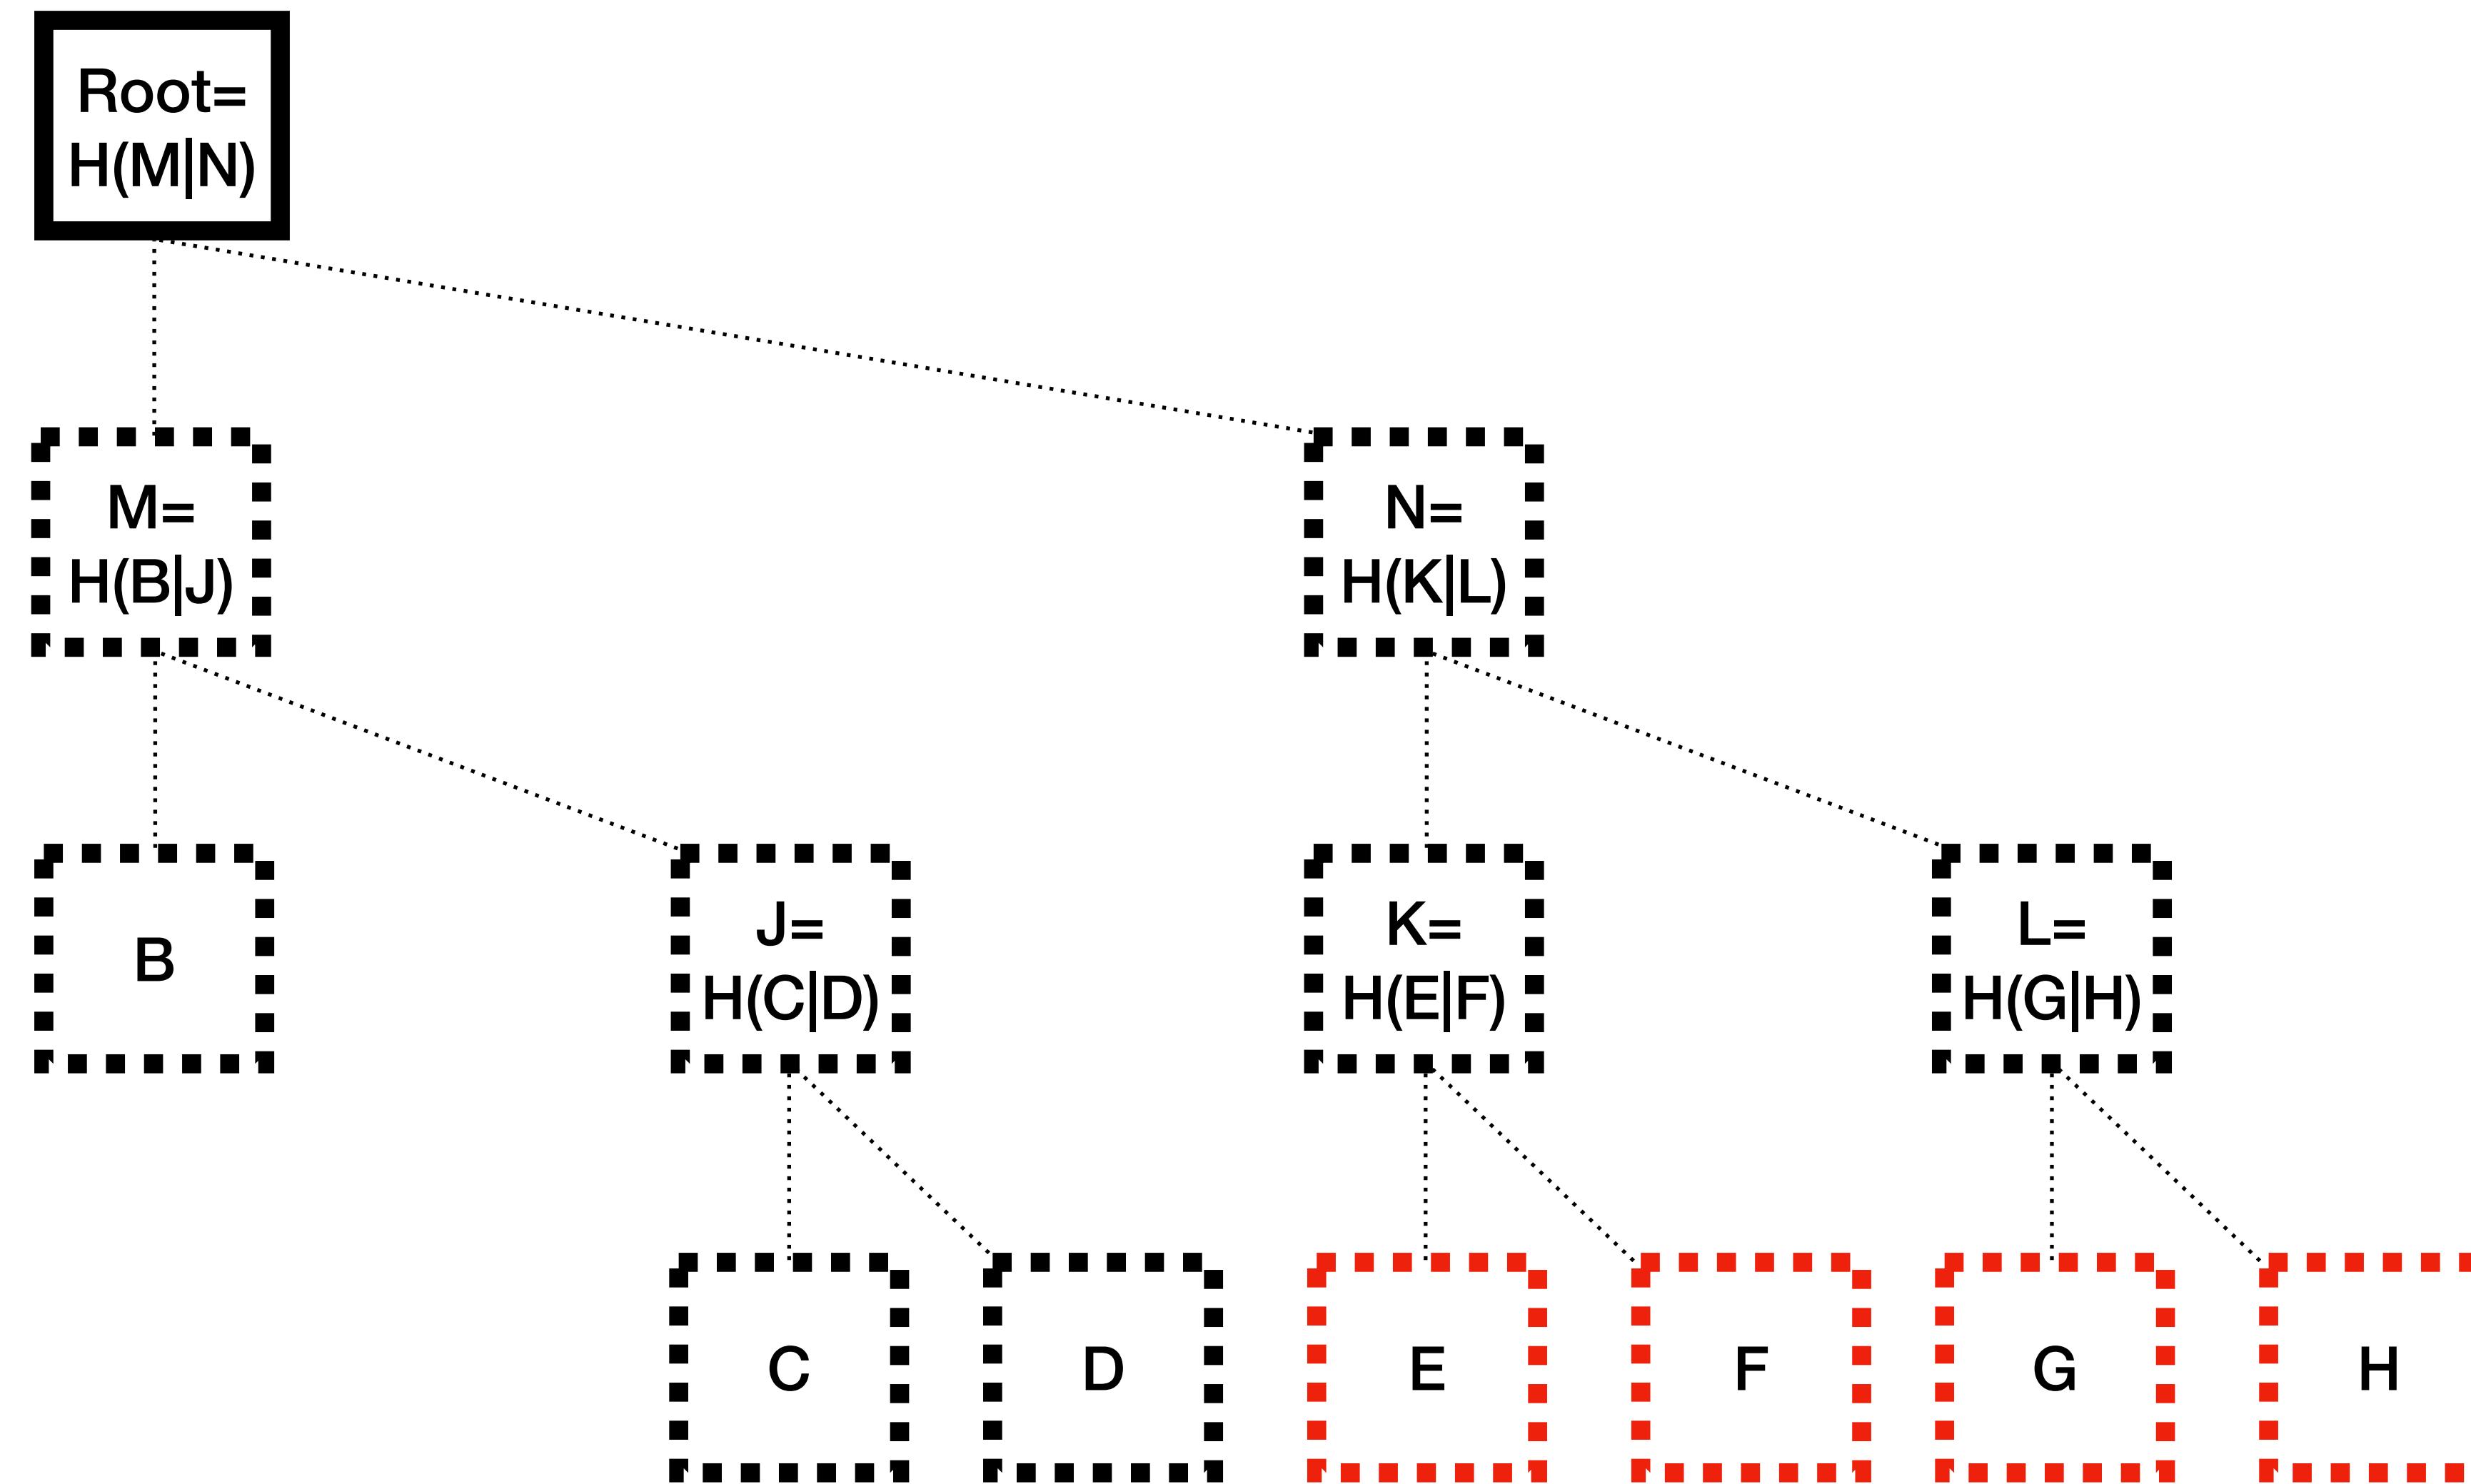
Done

Root=
 $H(M|N)$

After deleting A



Delete E, F, G, H. (Batched deletions)

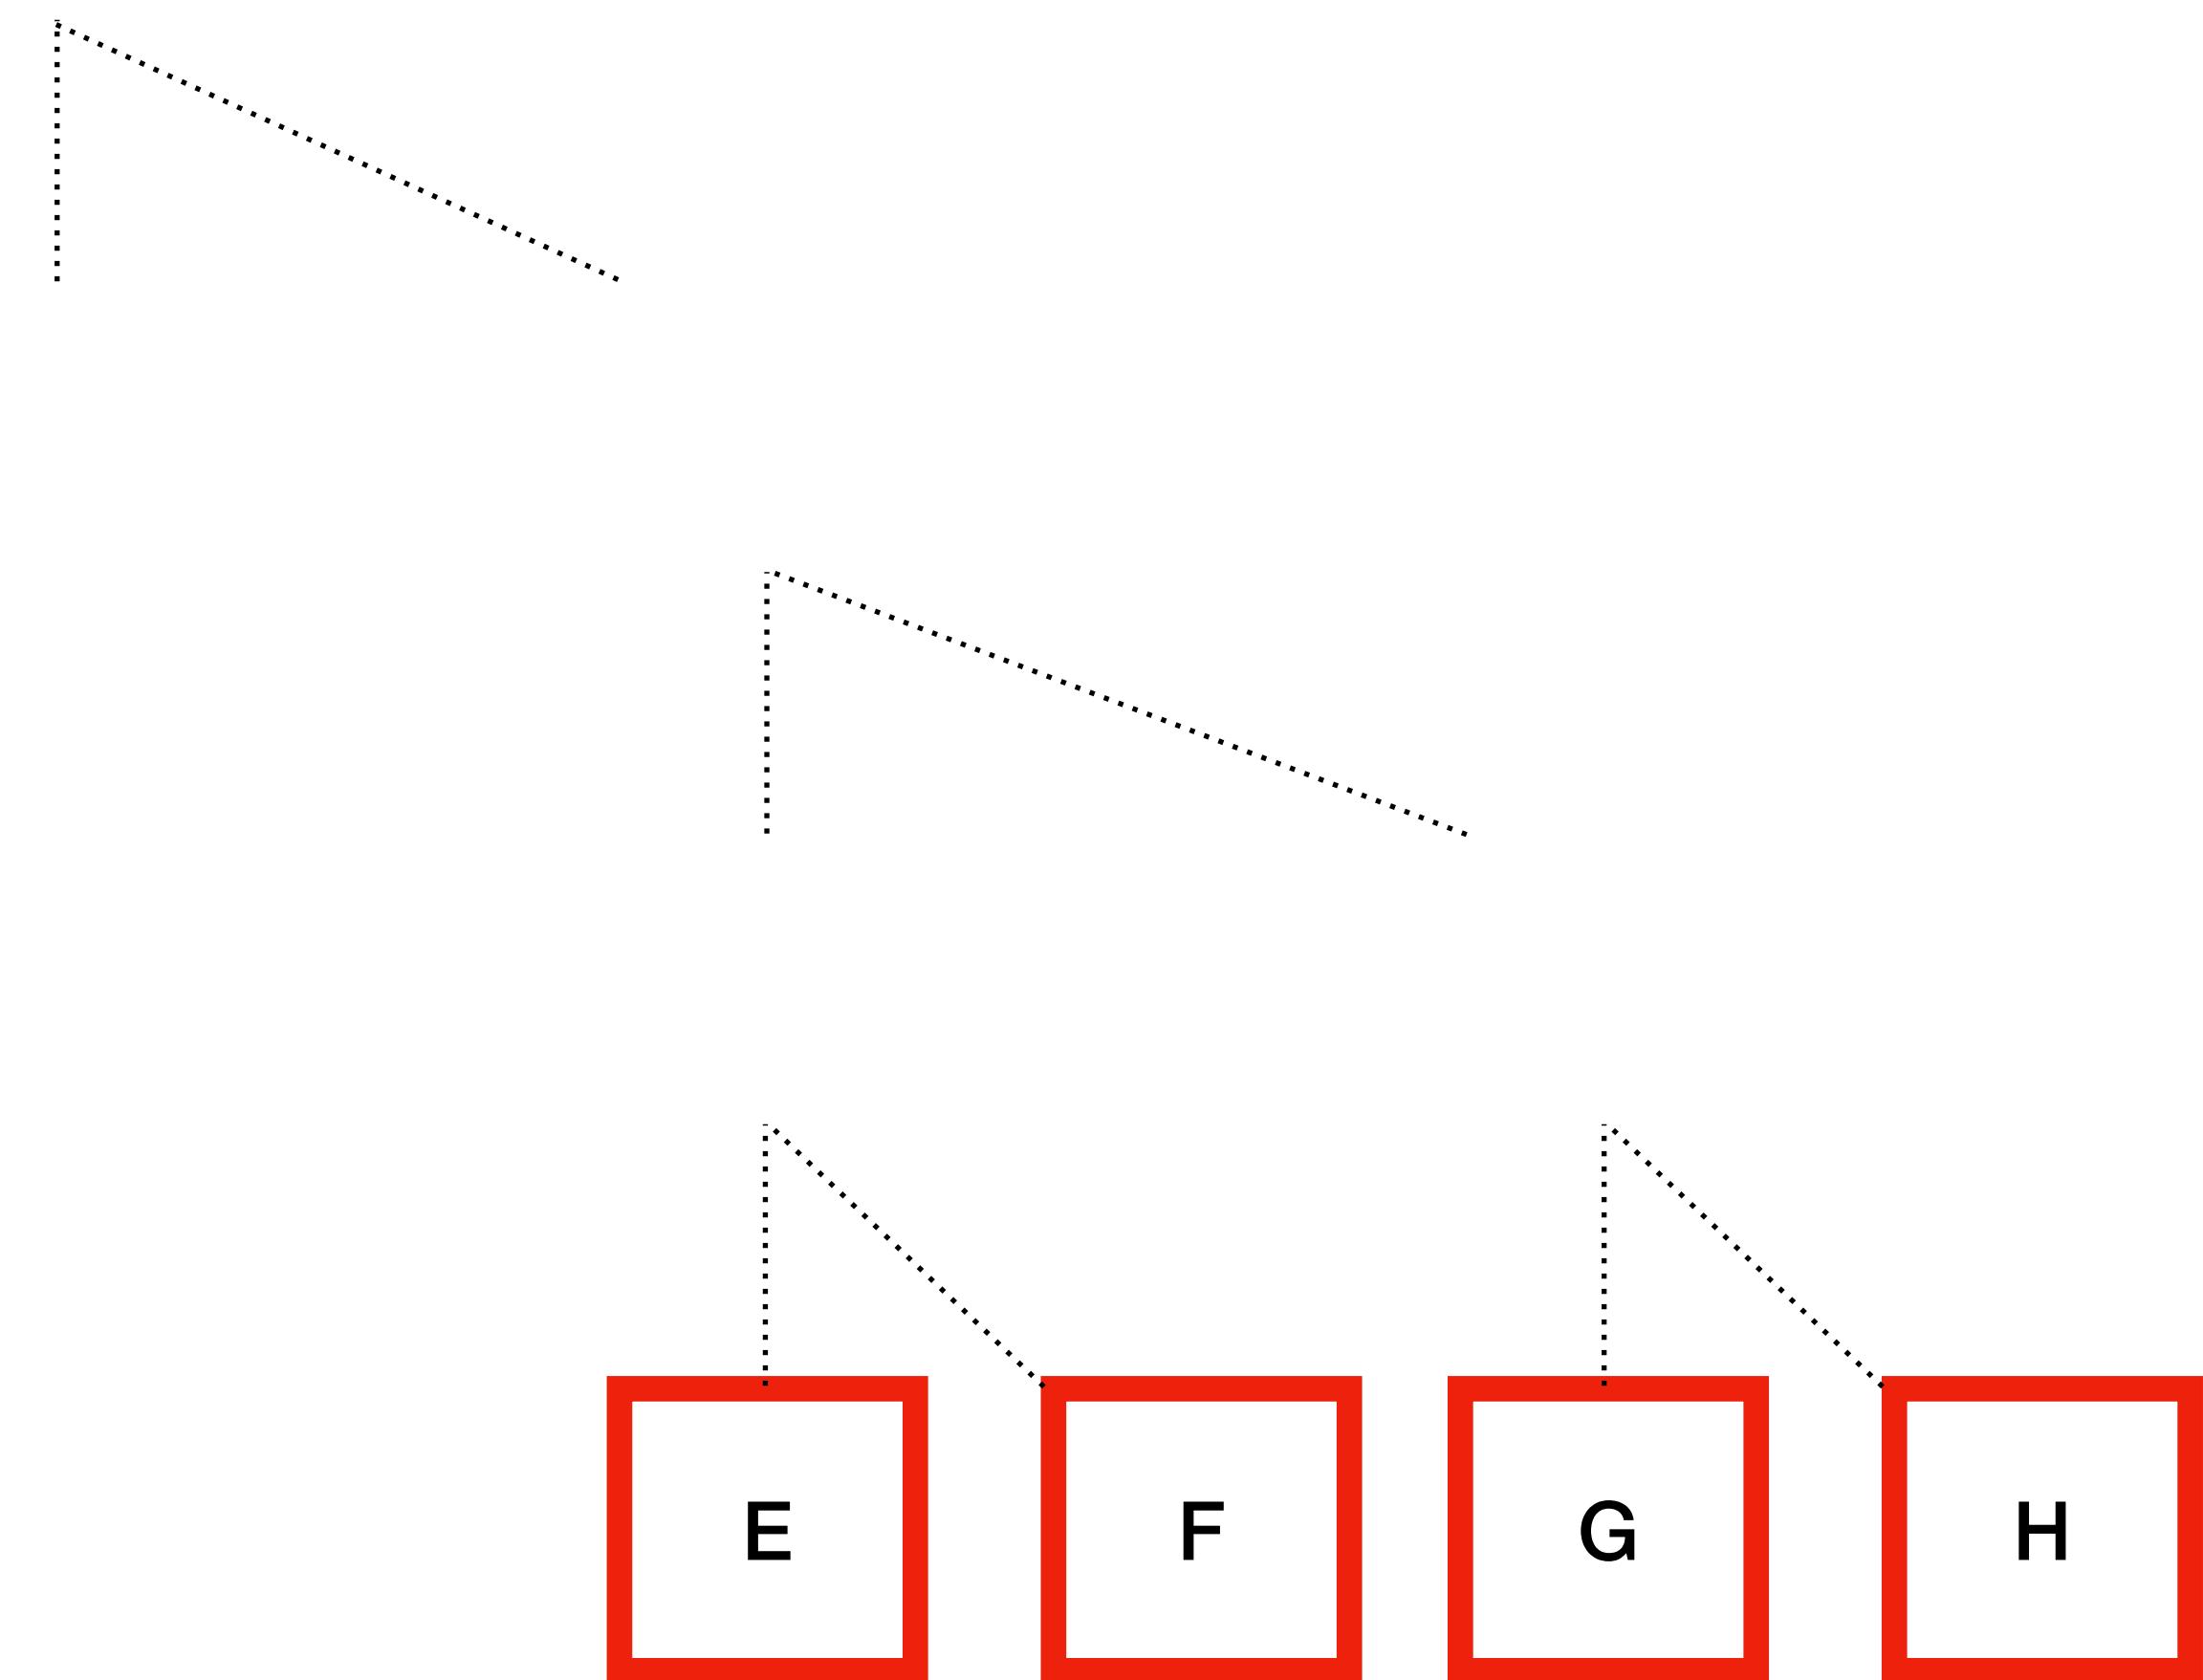


Start off with only the root

Root=
 $H(M|N)$

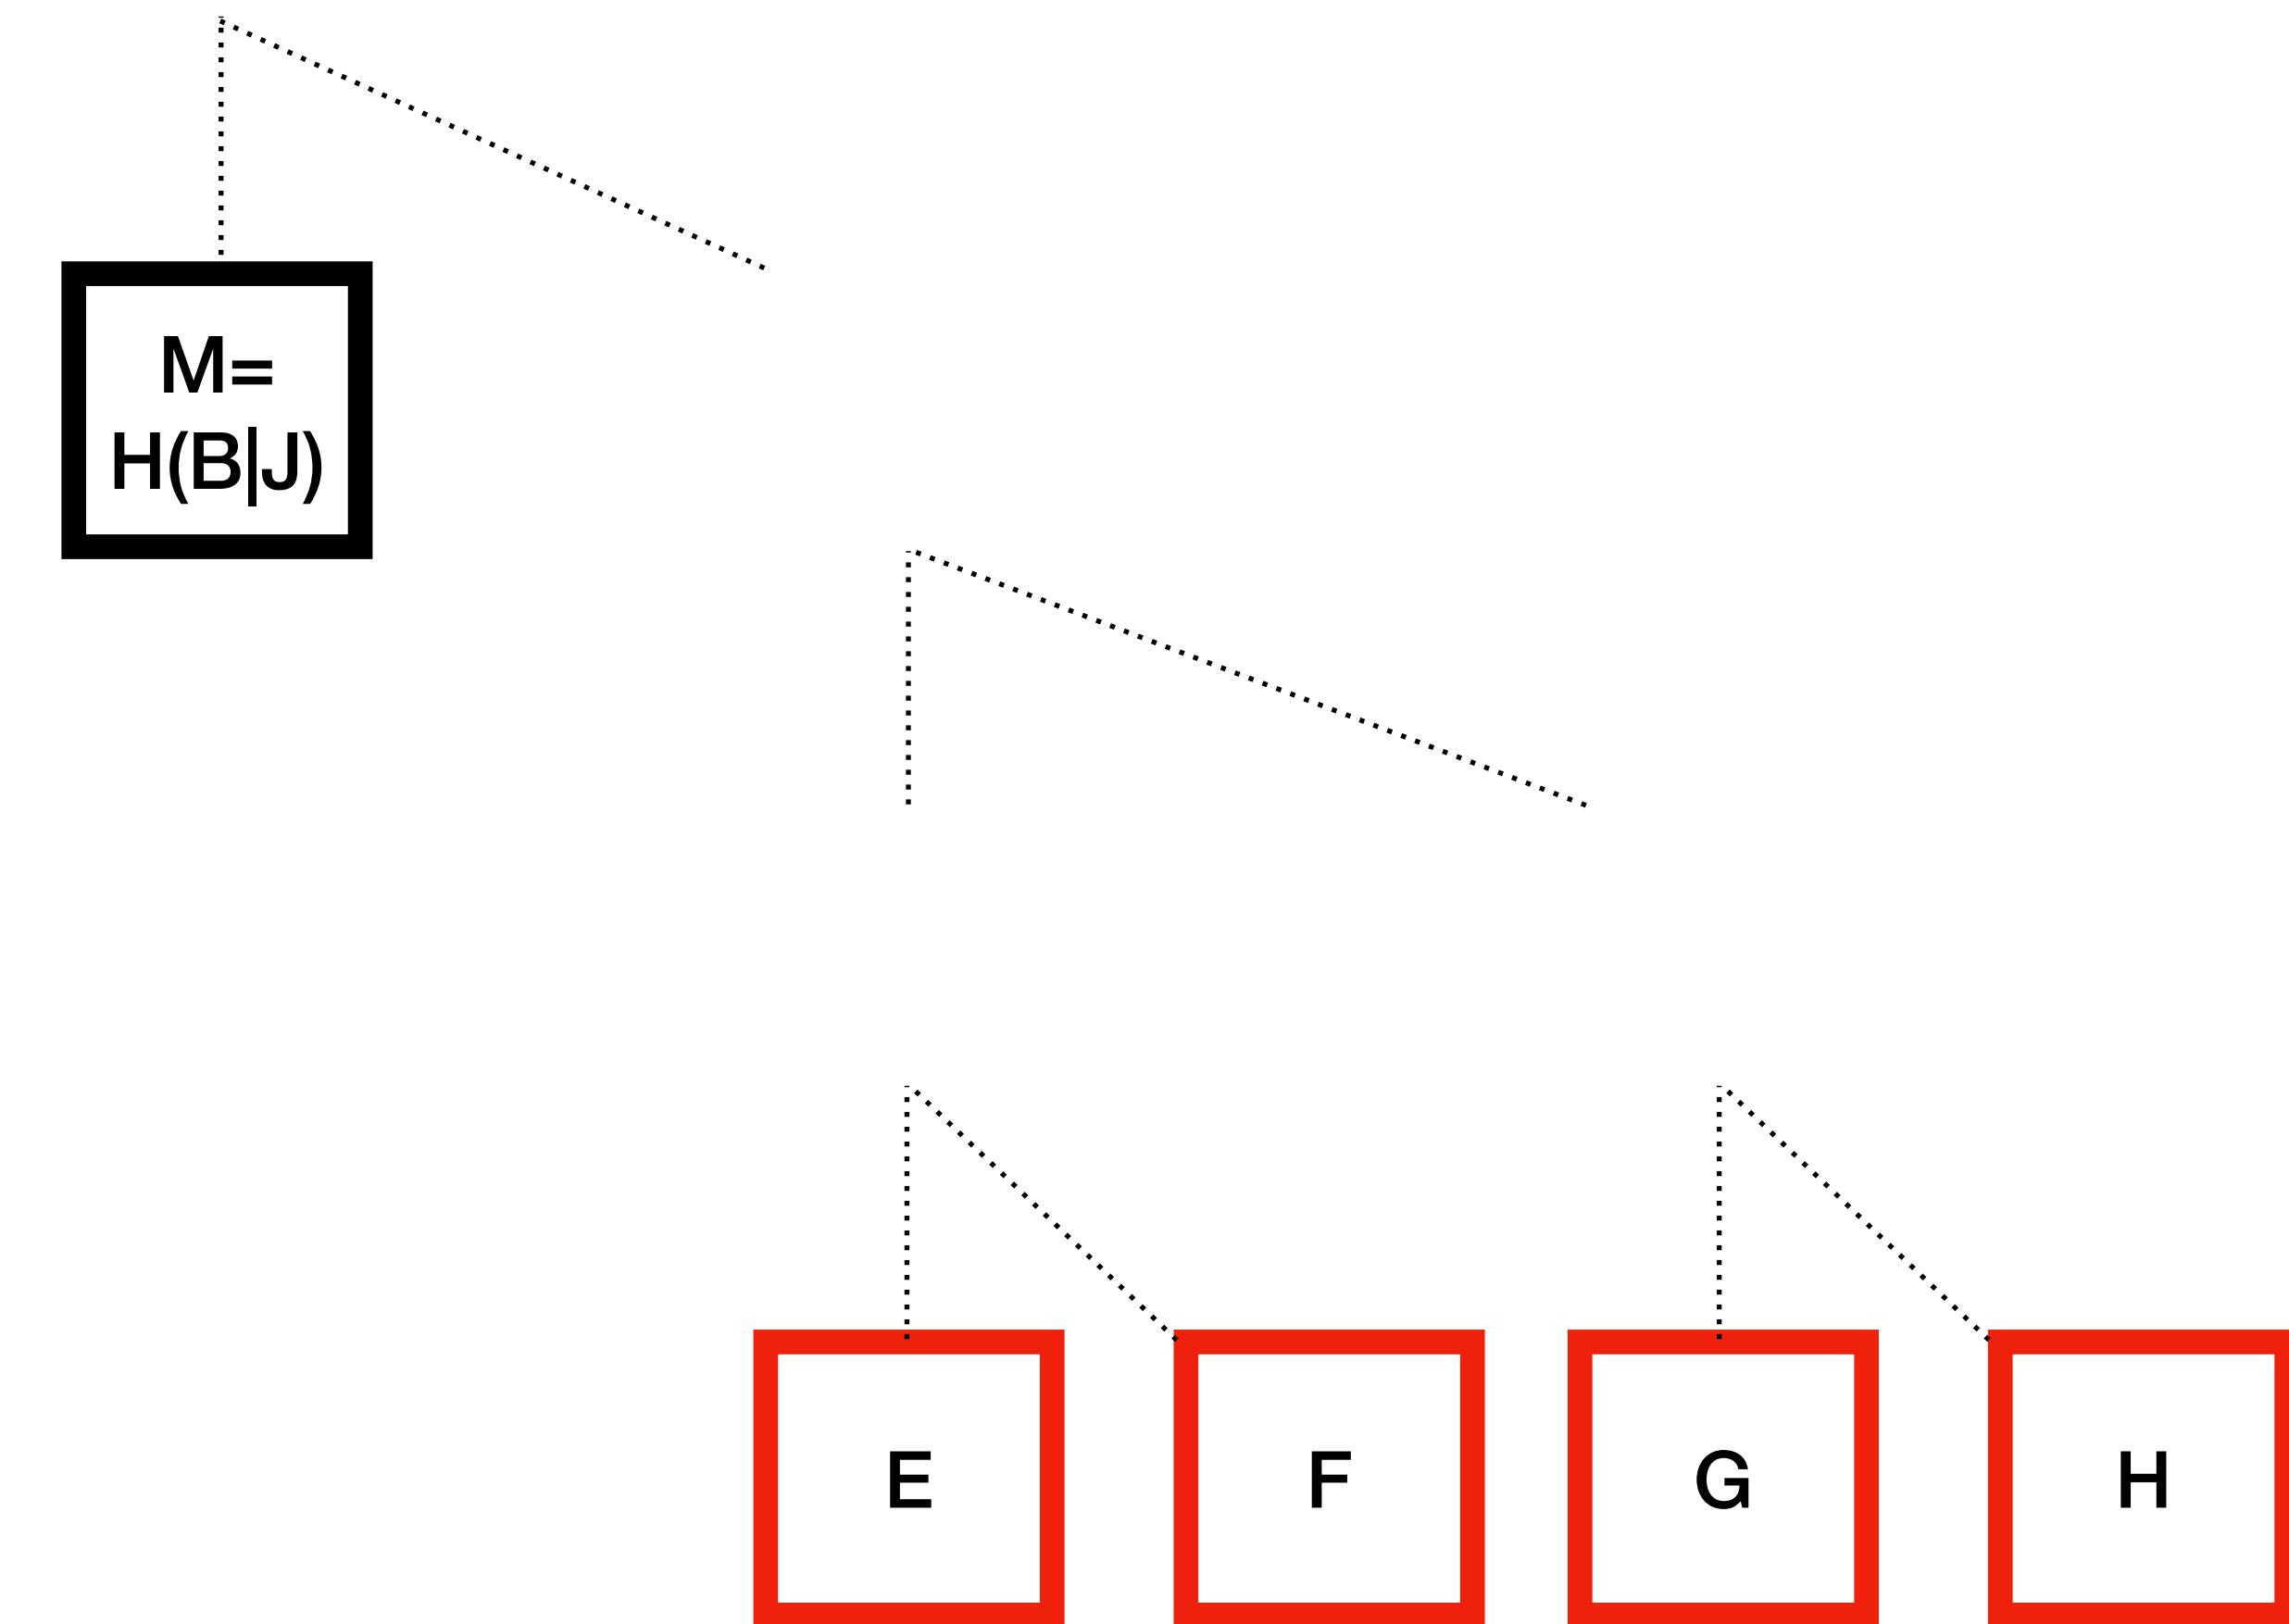
Calculate E, F, G, H from the Bitcoin Block

Root=
 $H(M|N)$

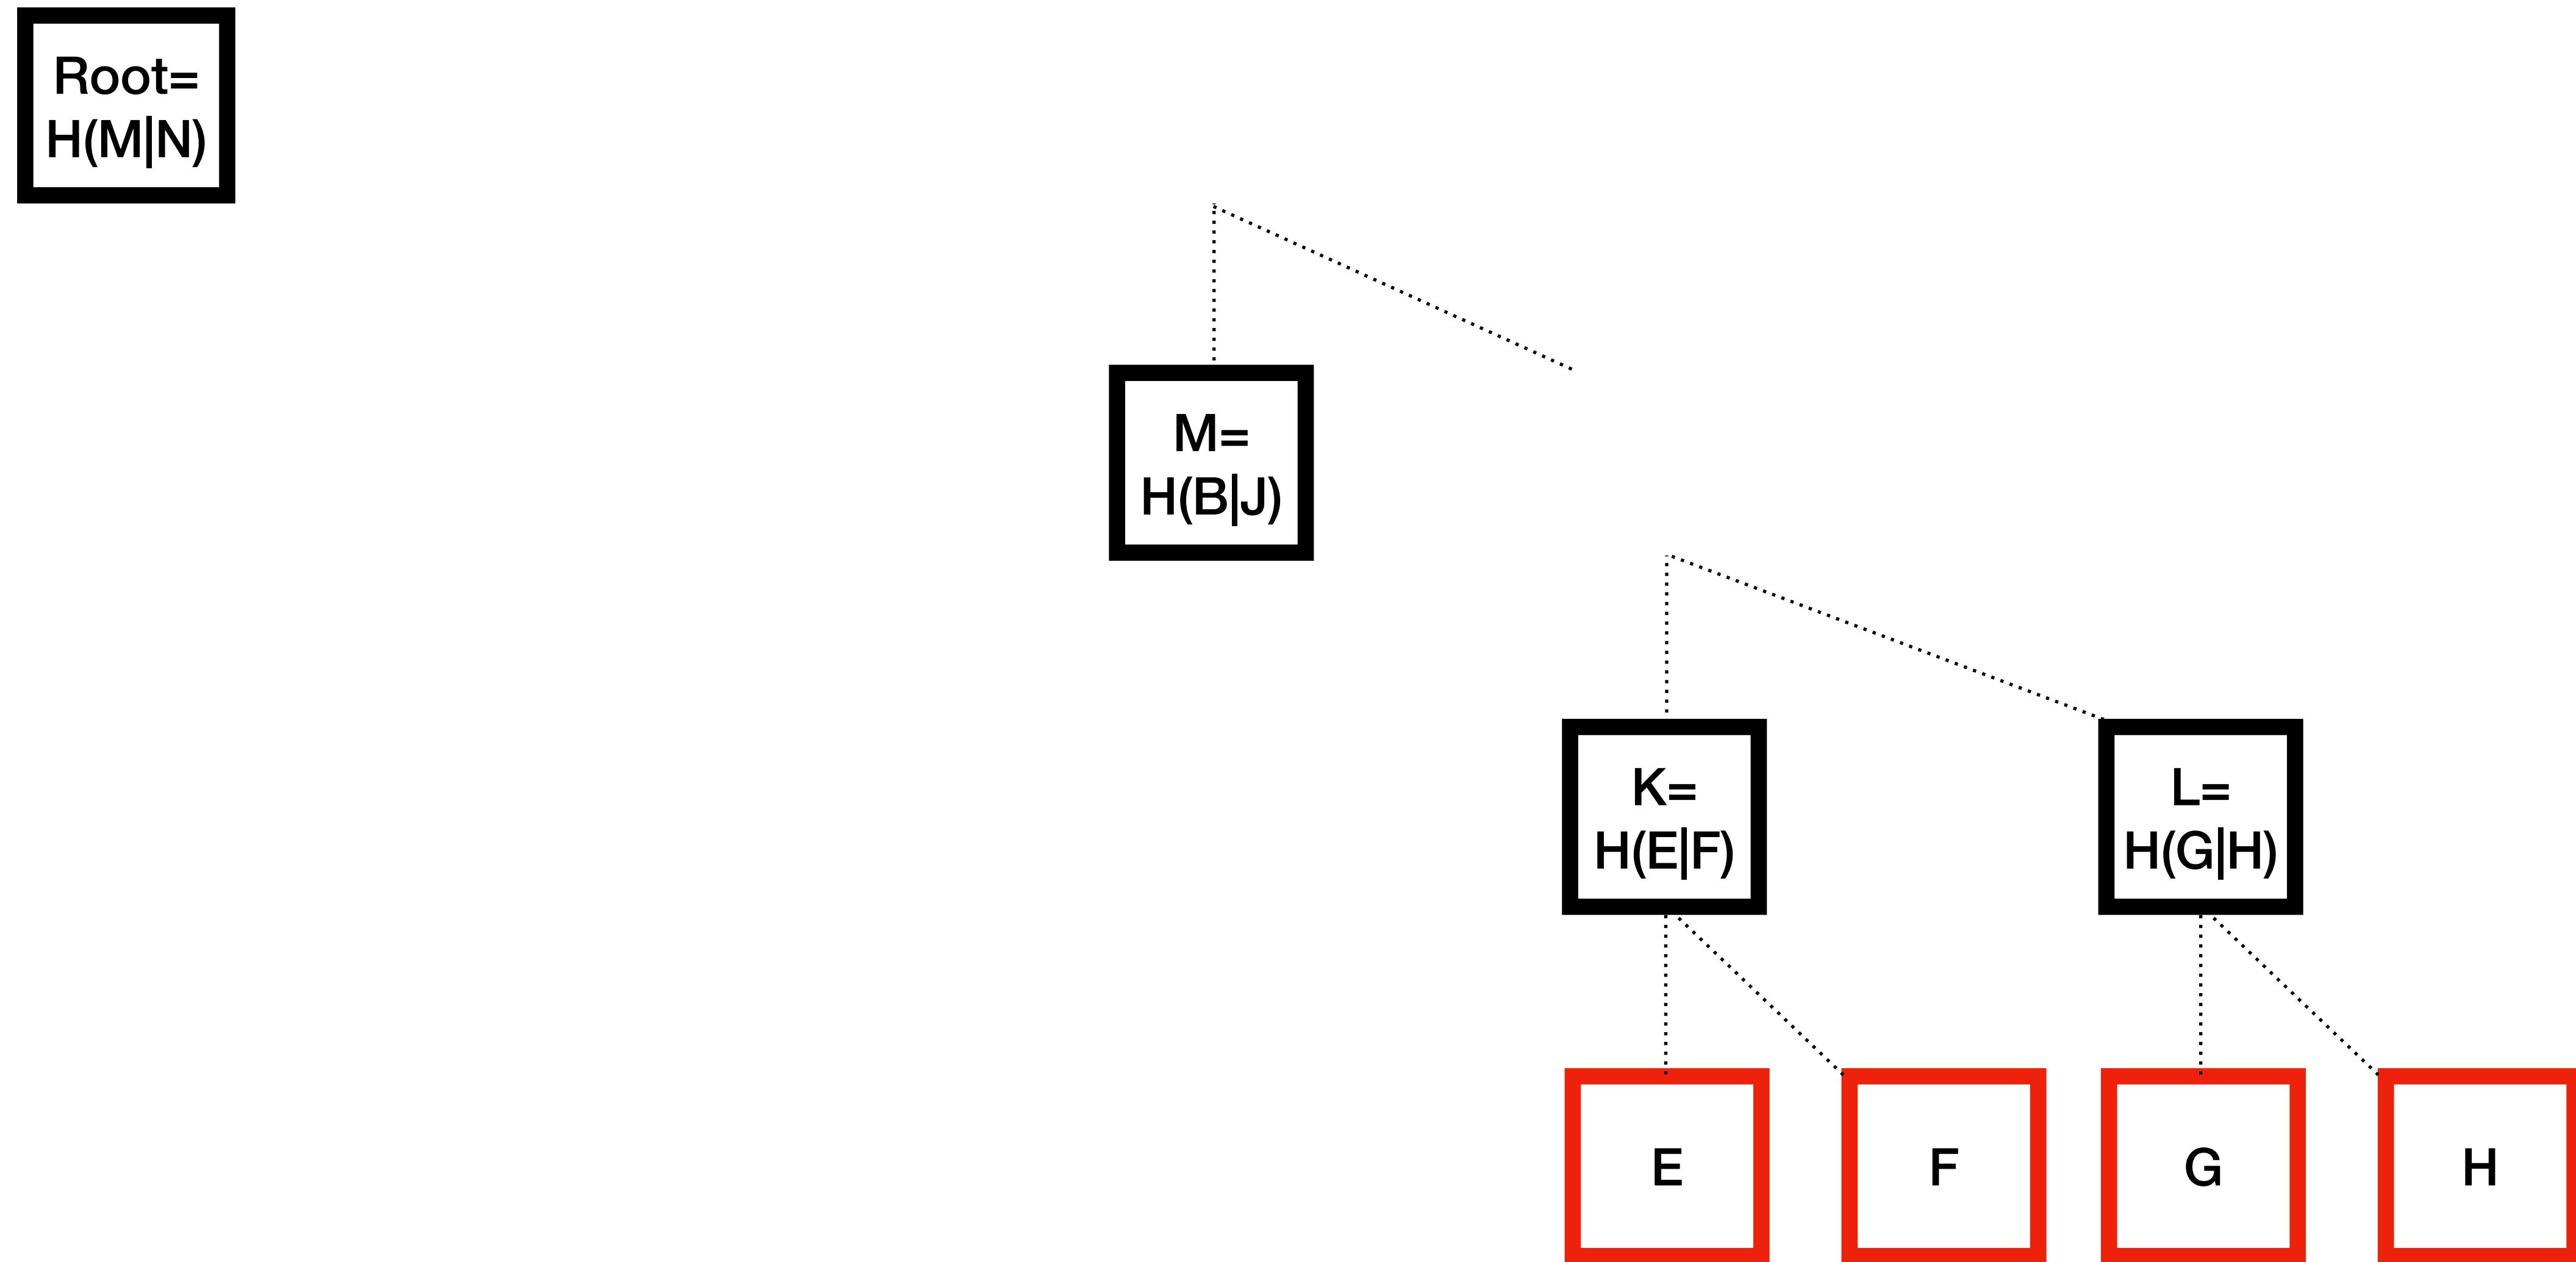


Receive the proof: M

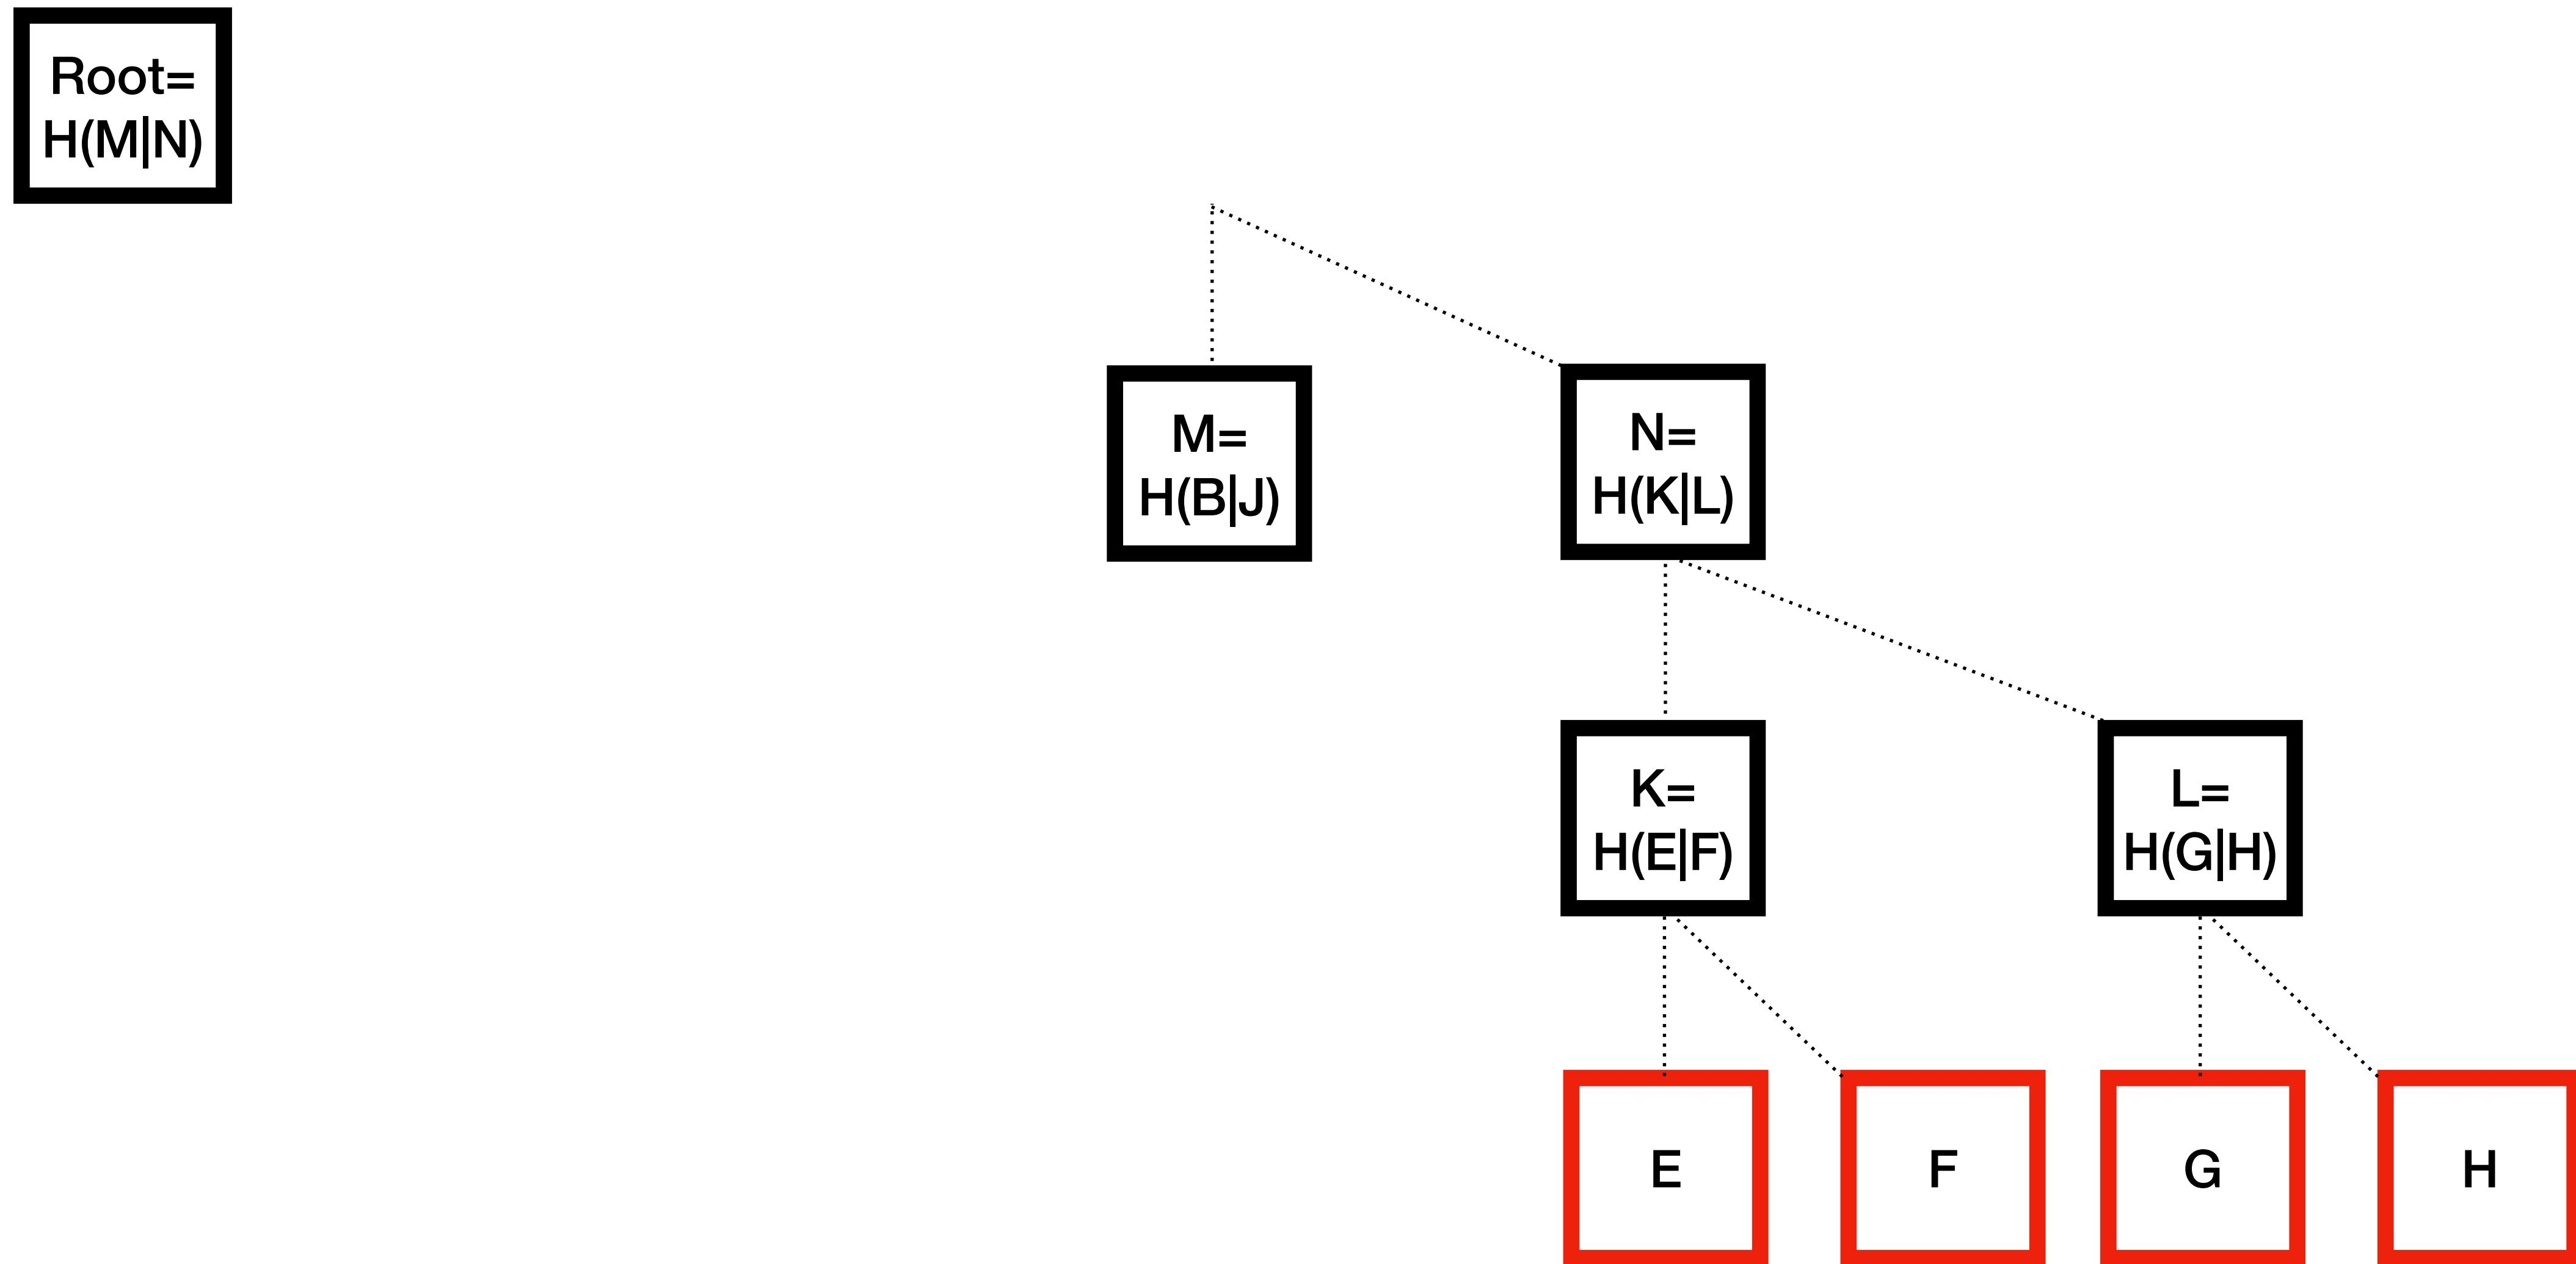
Root=
 $H(M|N)$



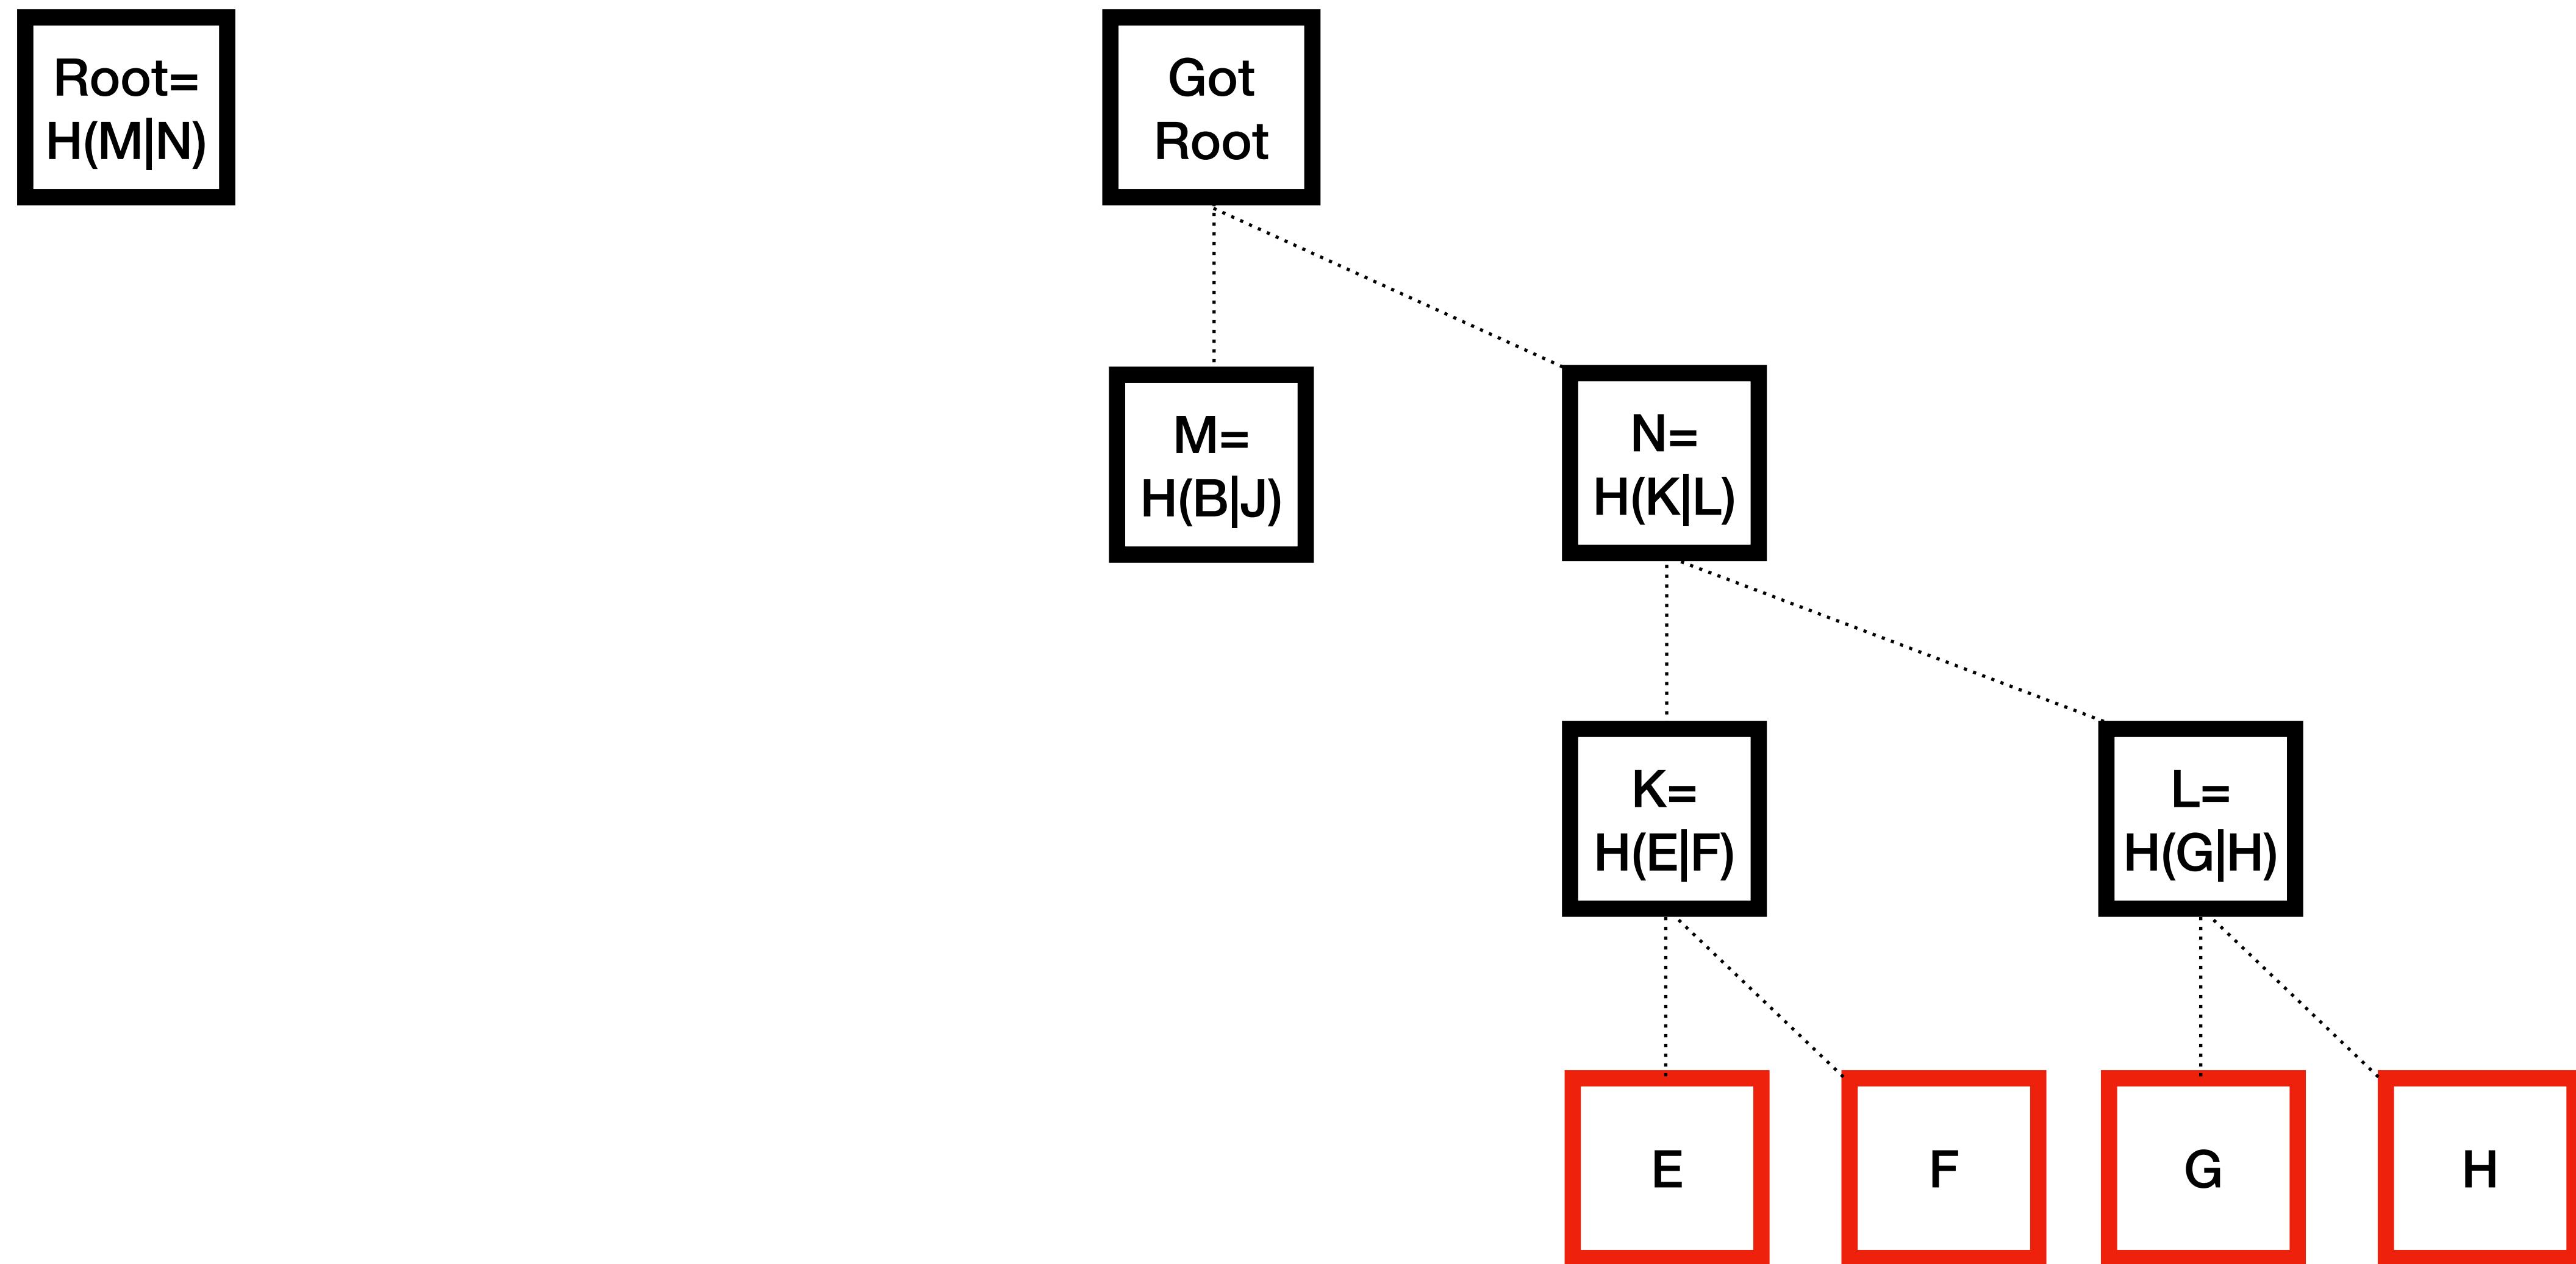
Calculate K and L



Calculate N



Calculate the Root



Compare roots

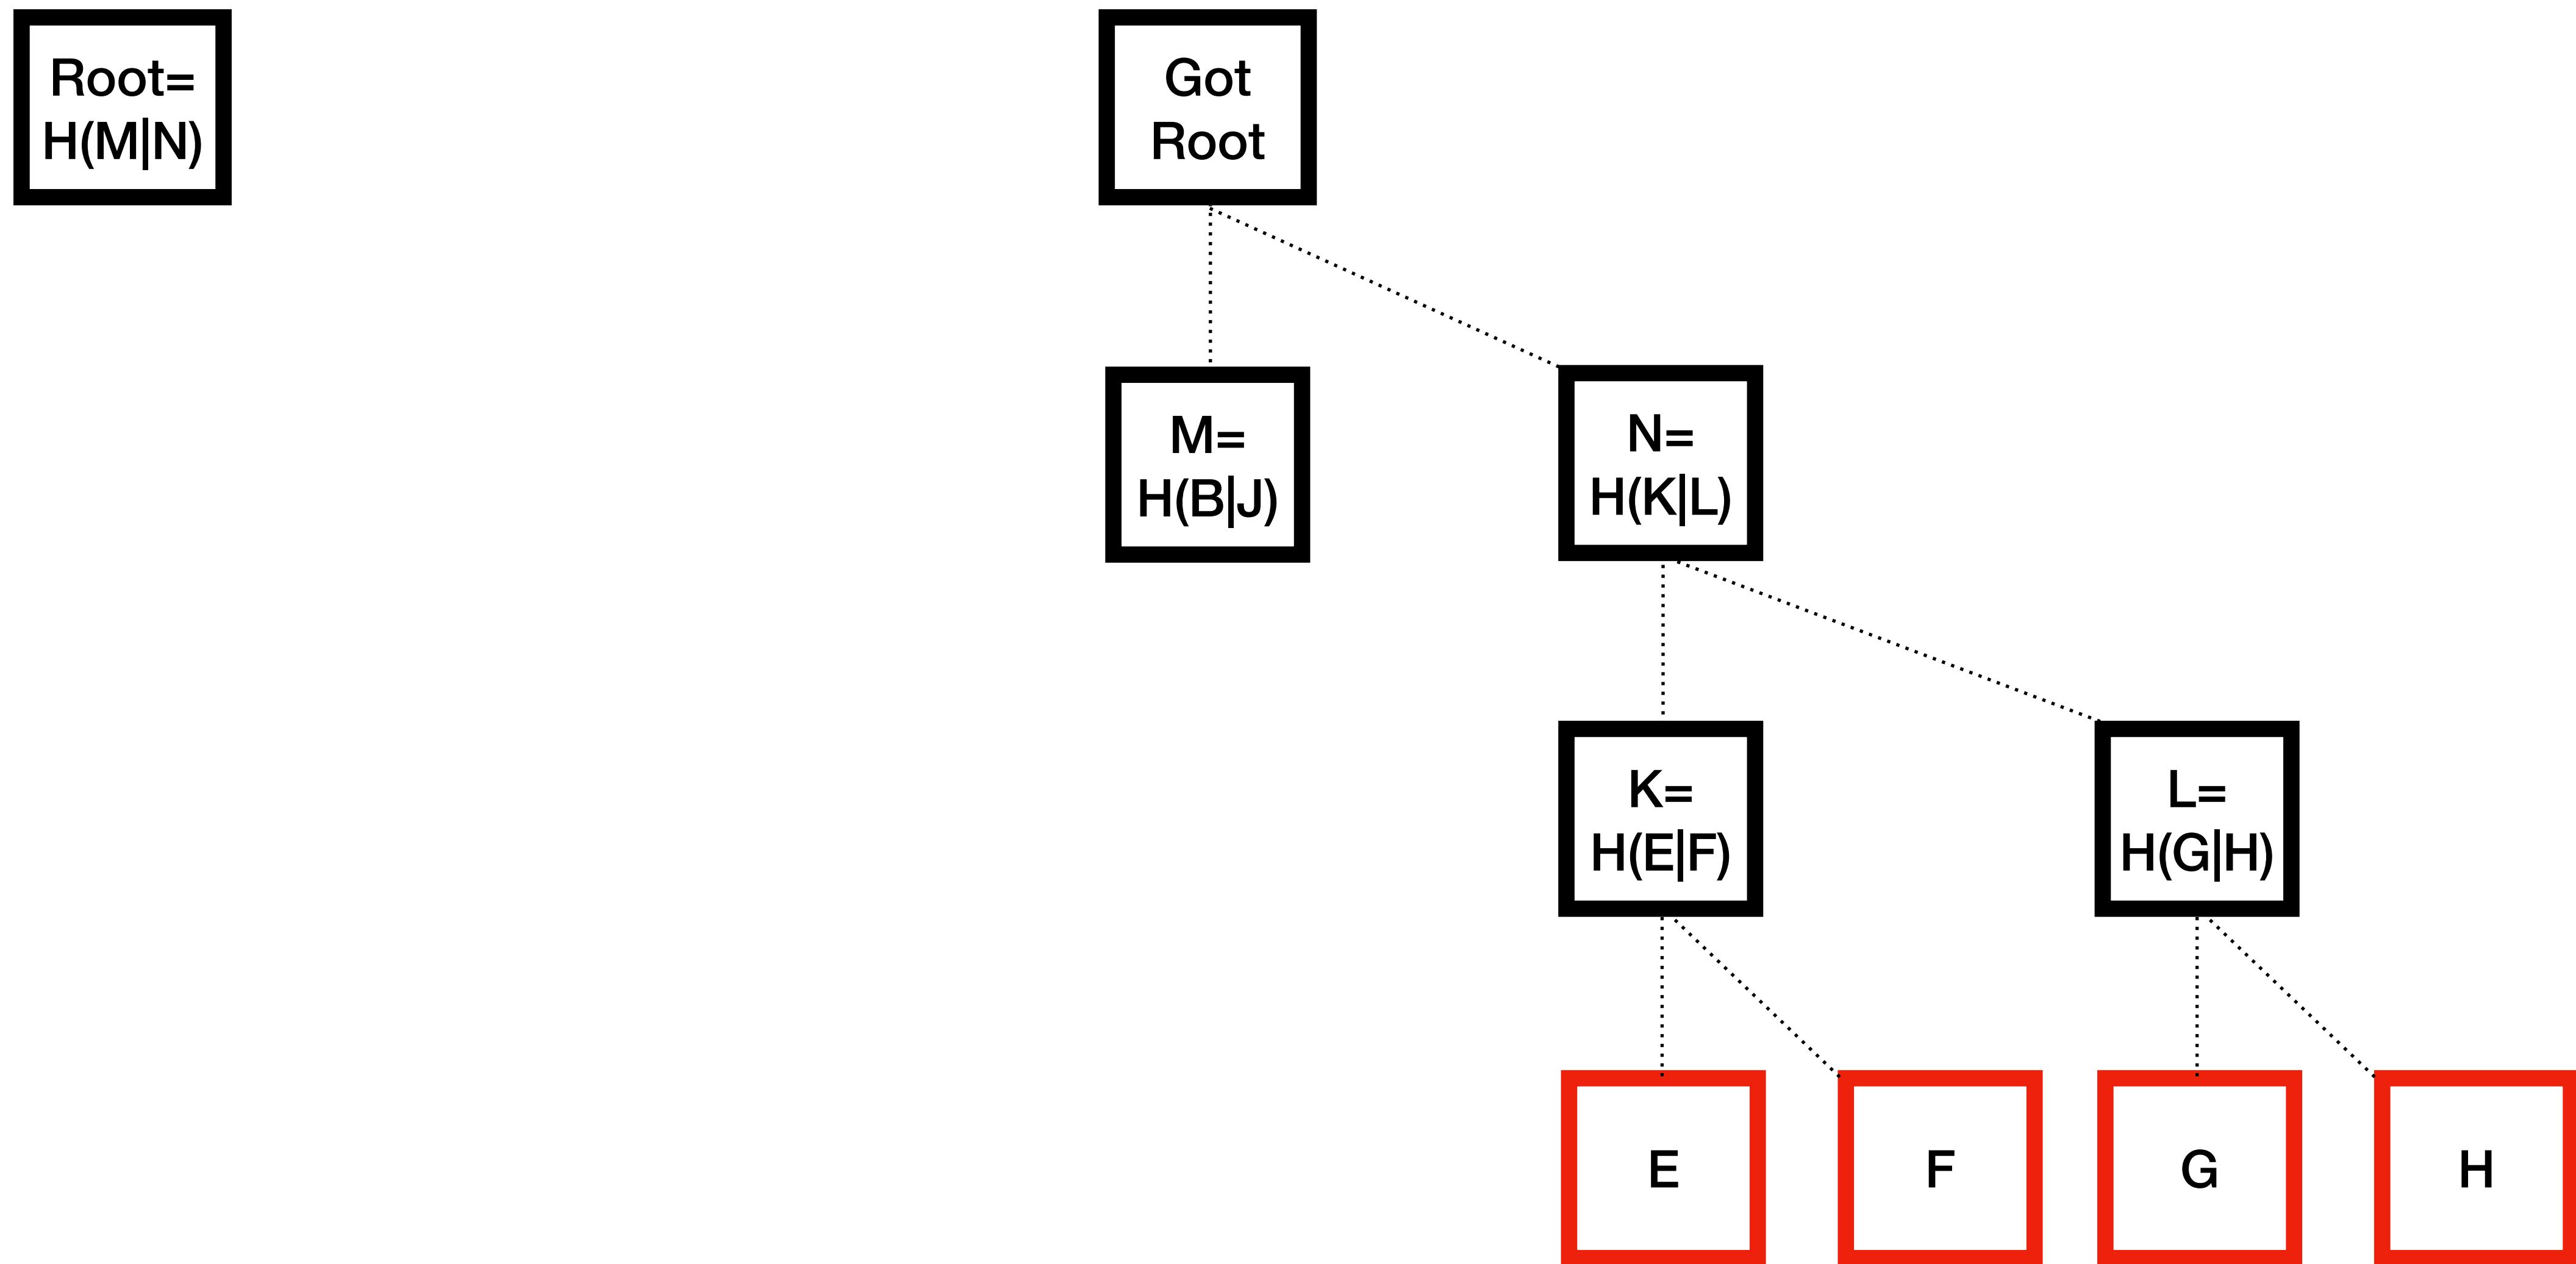
Continue if the roots are equal. Ban peer if not equal

Root= $H(M|N)$

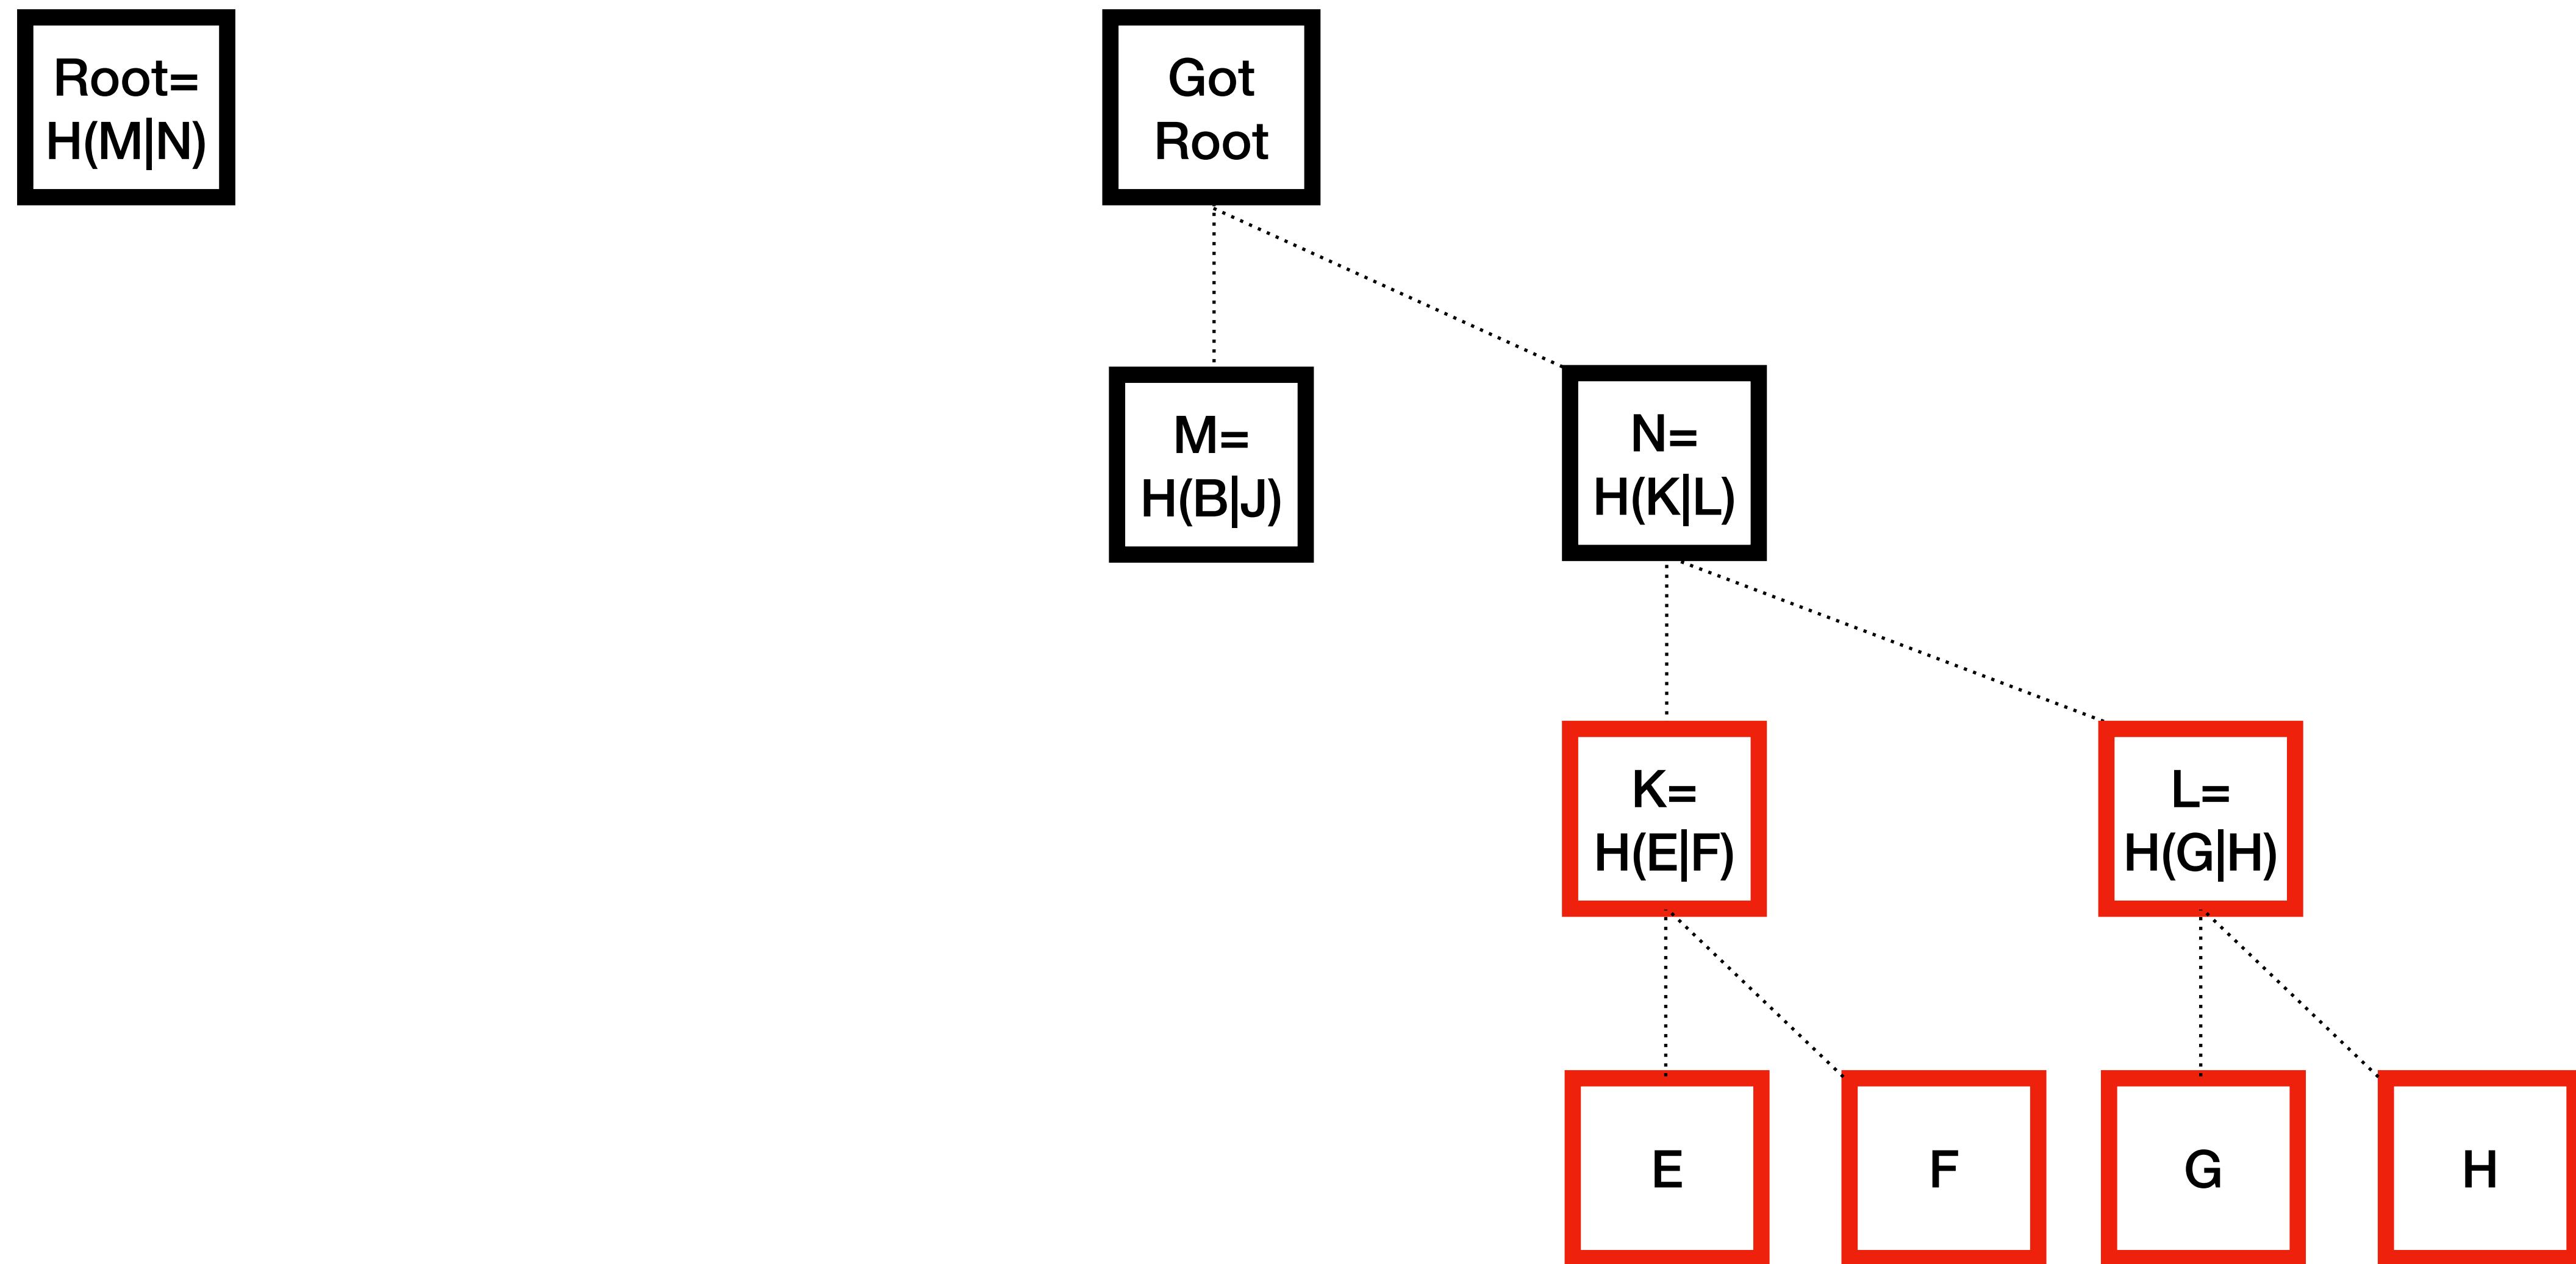
==

Got Root

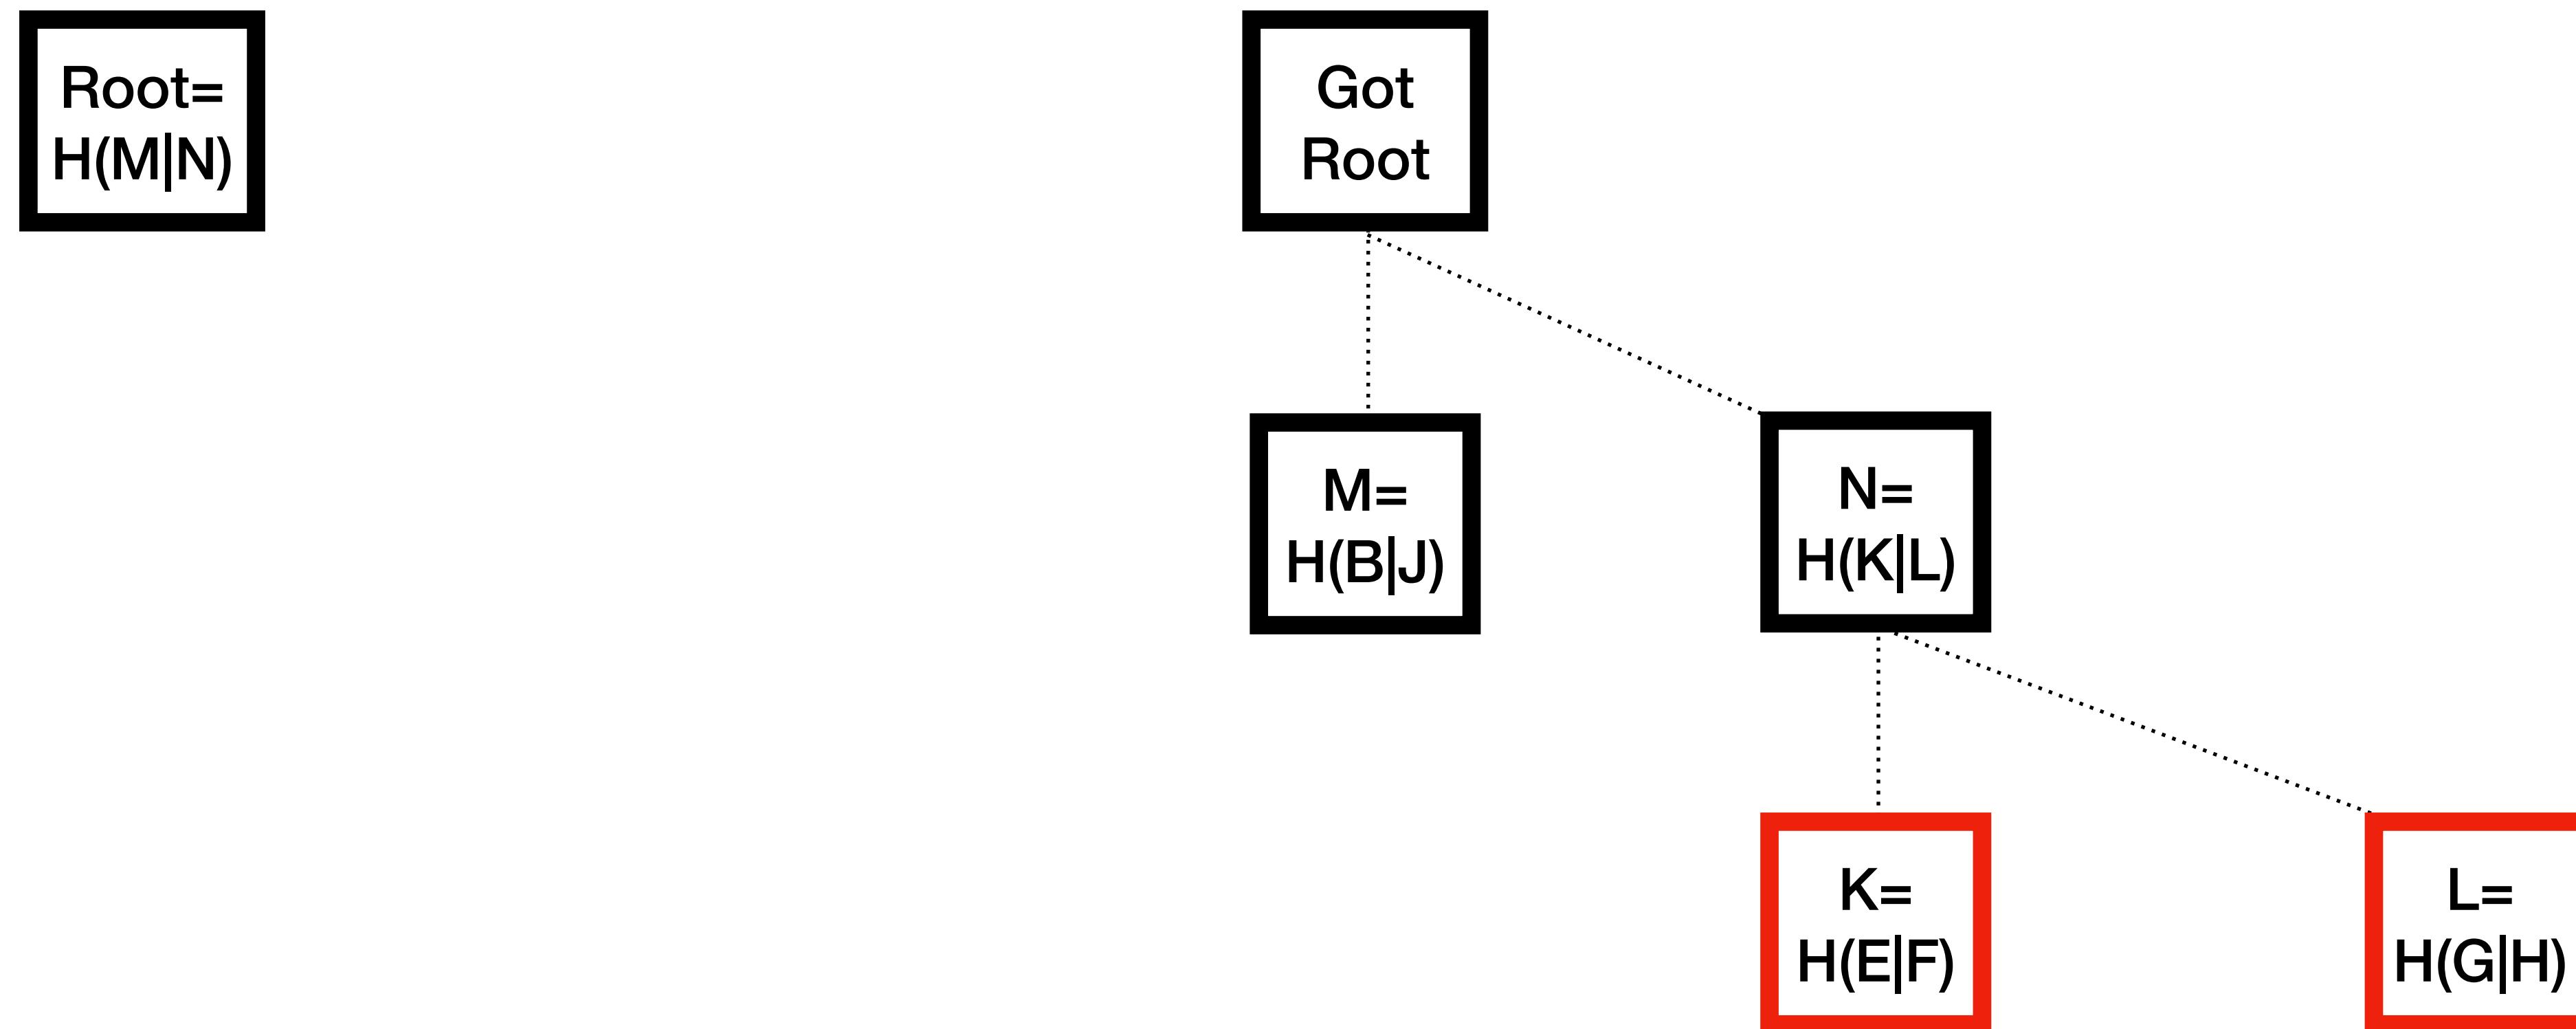
Calculate new root



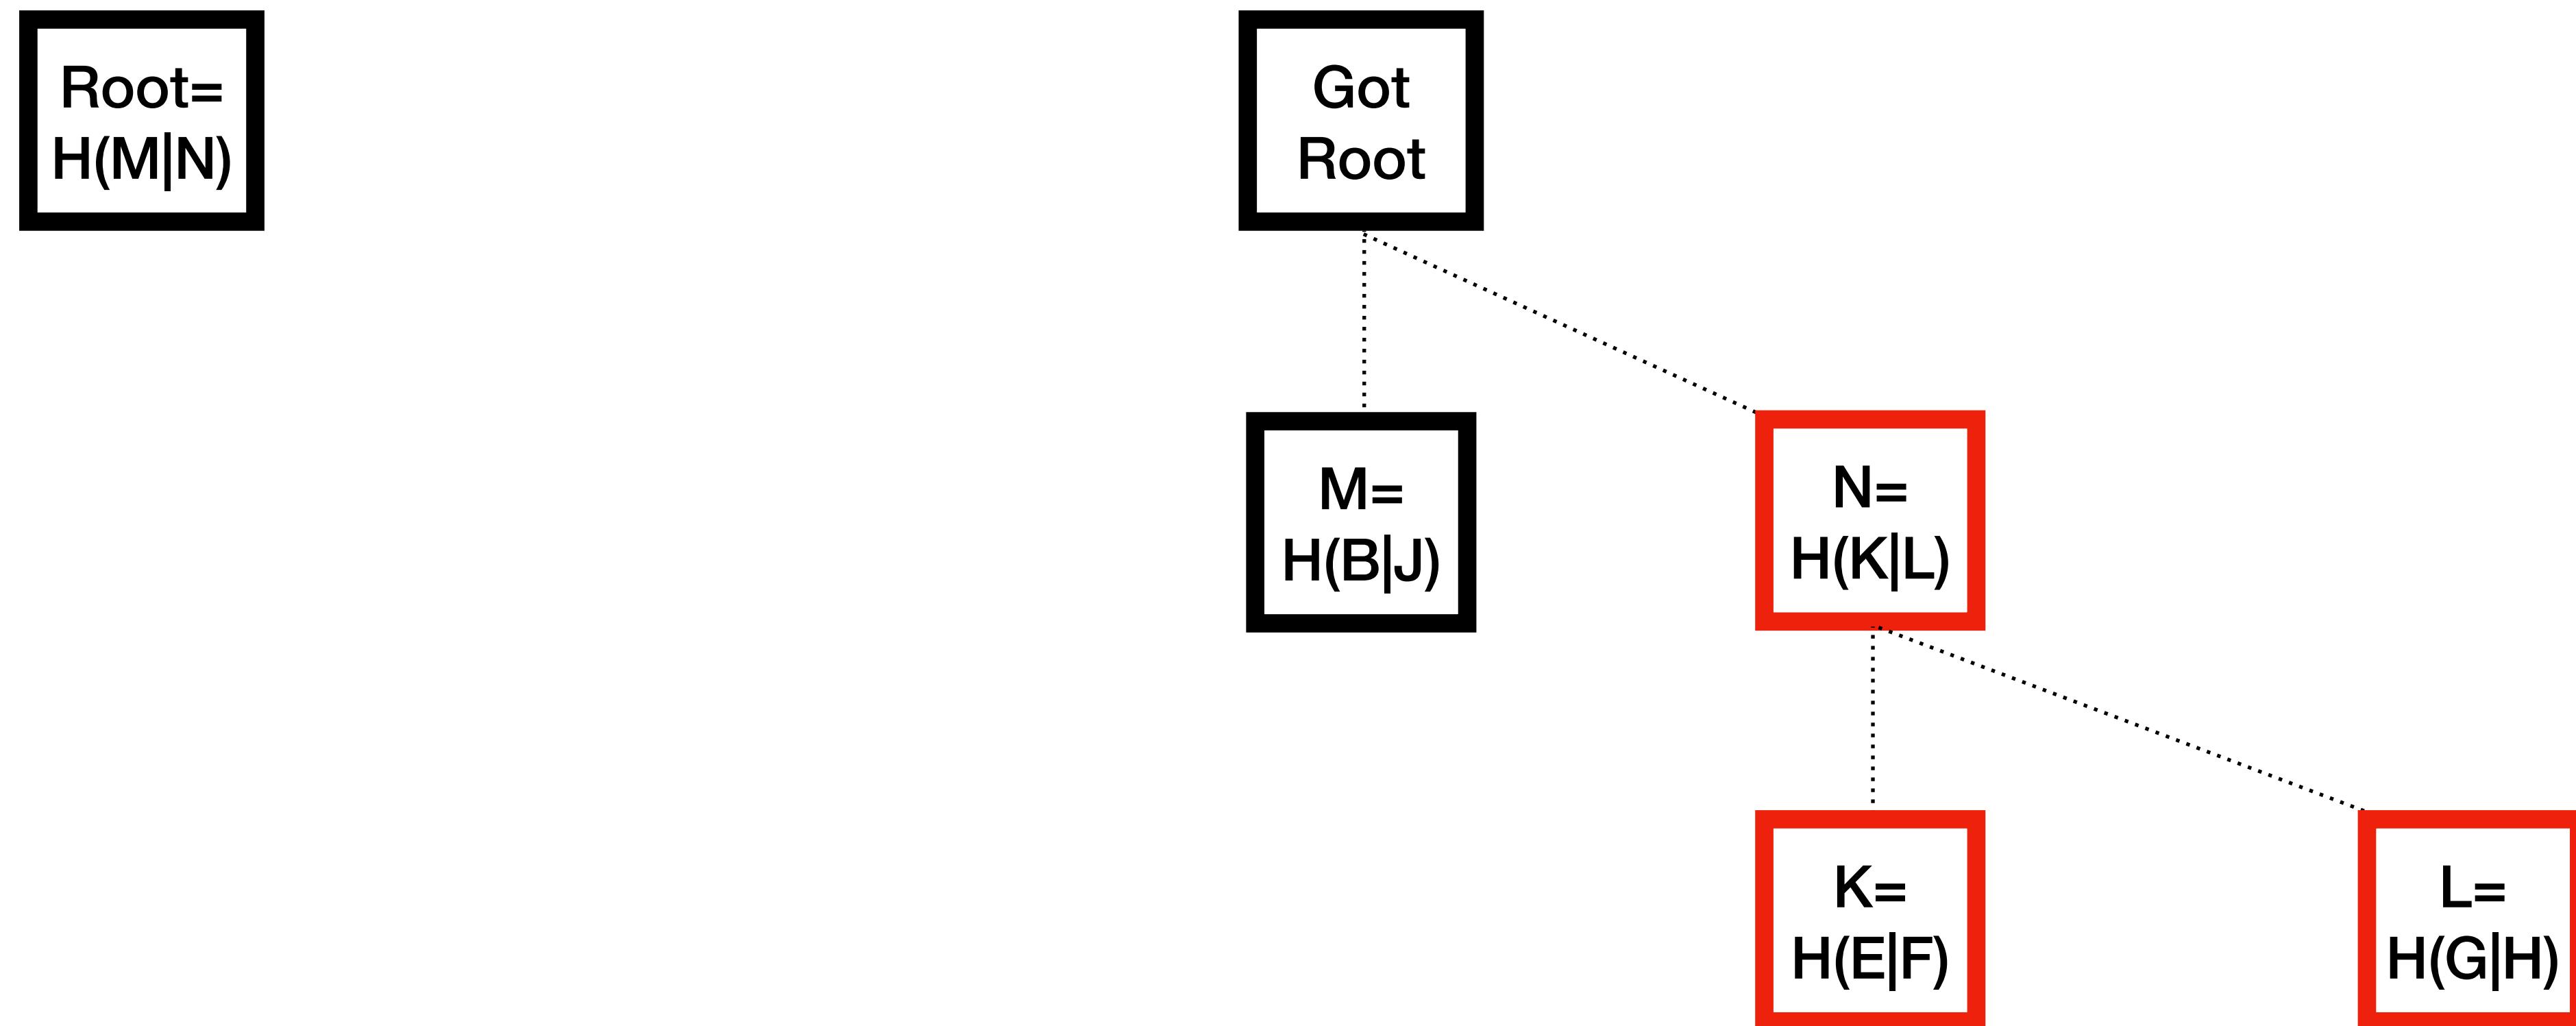
Mark K & L as node to delete



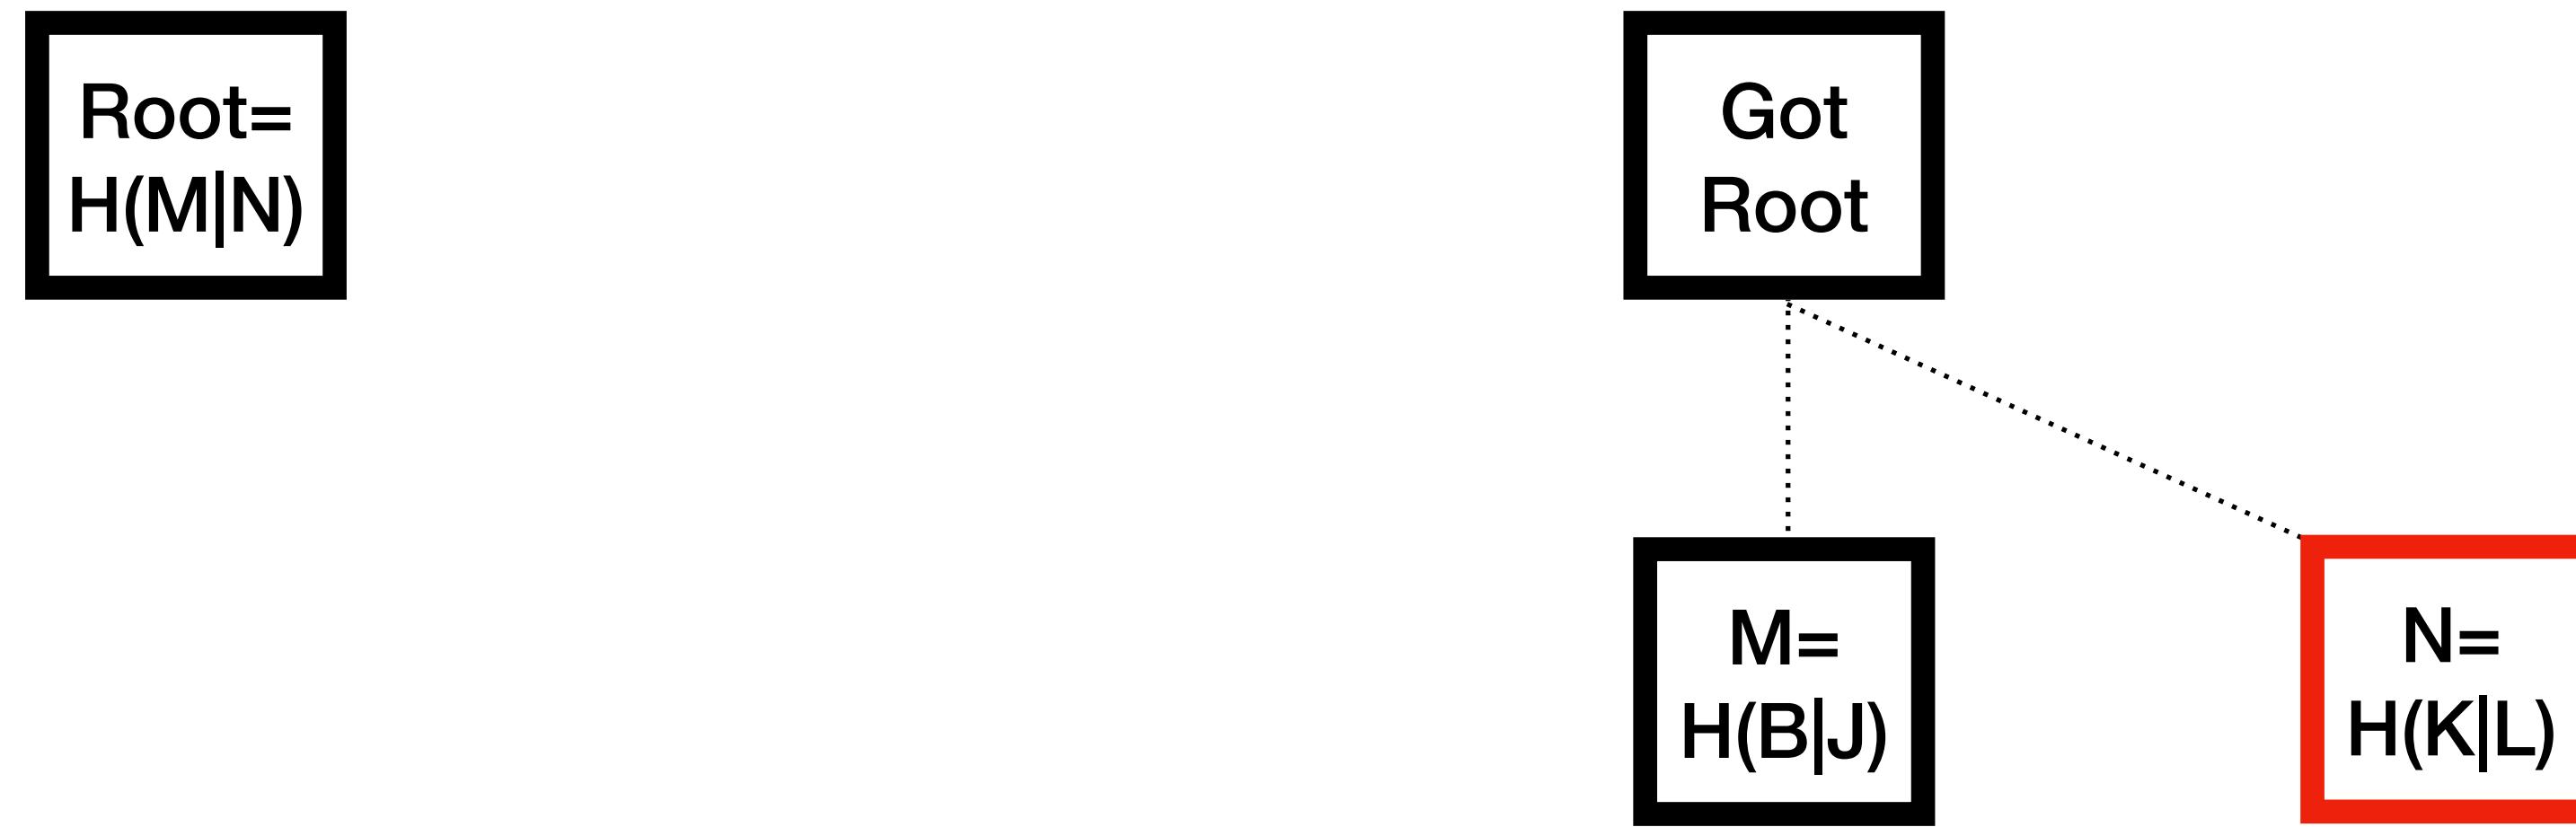
Remove old nodes for E, F, G, H



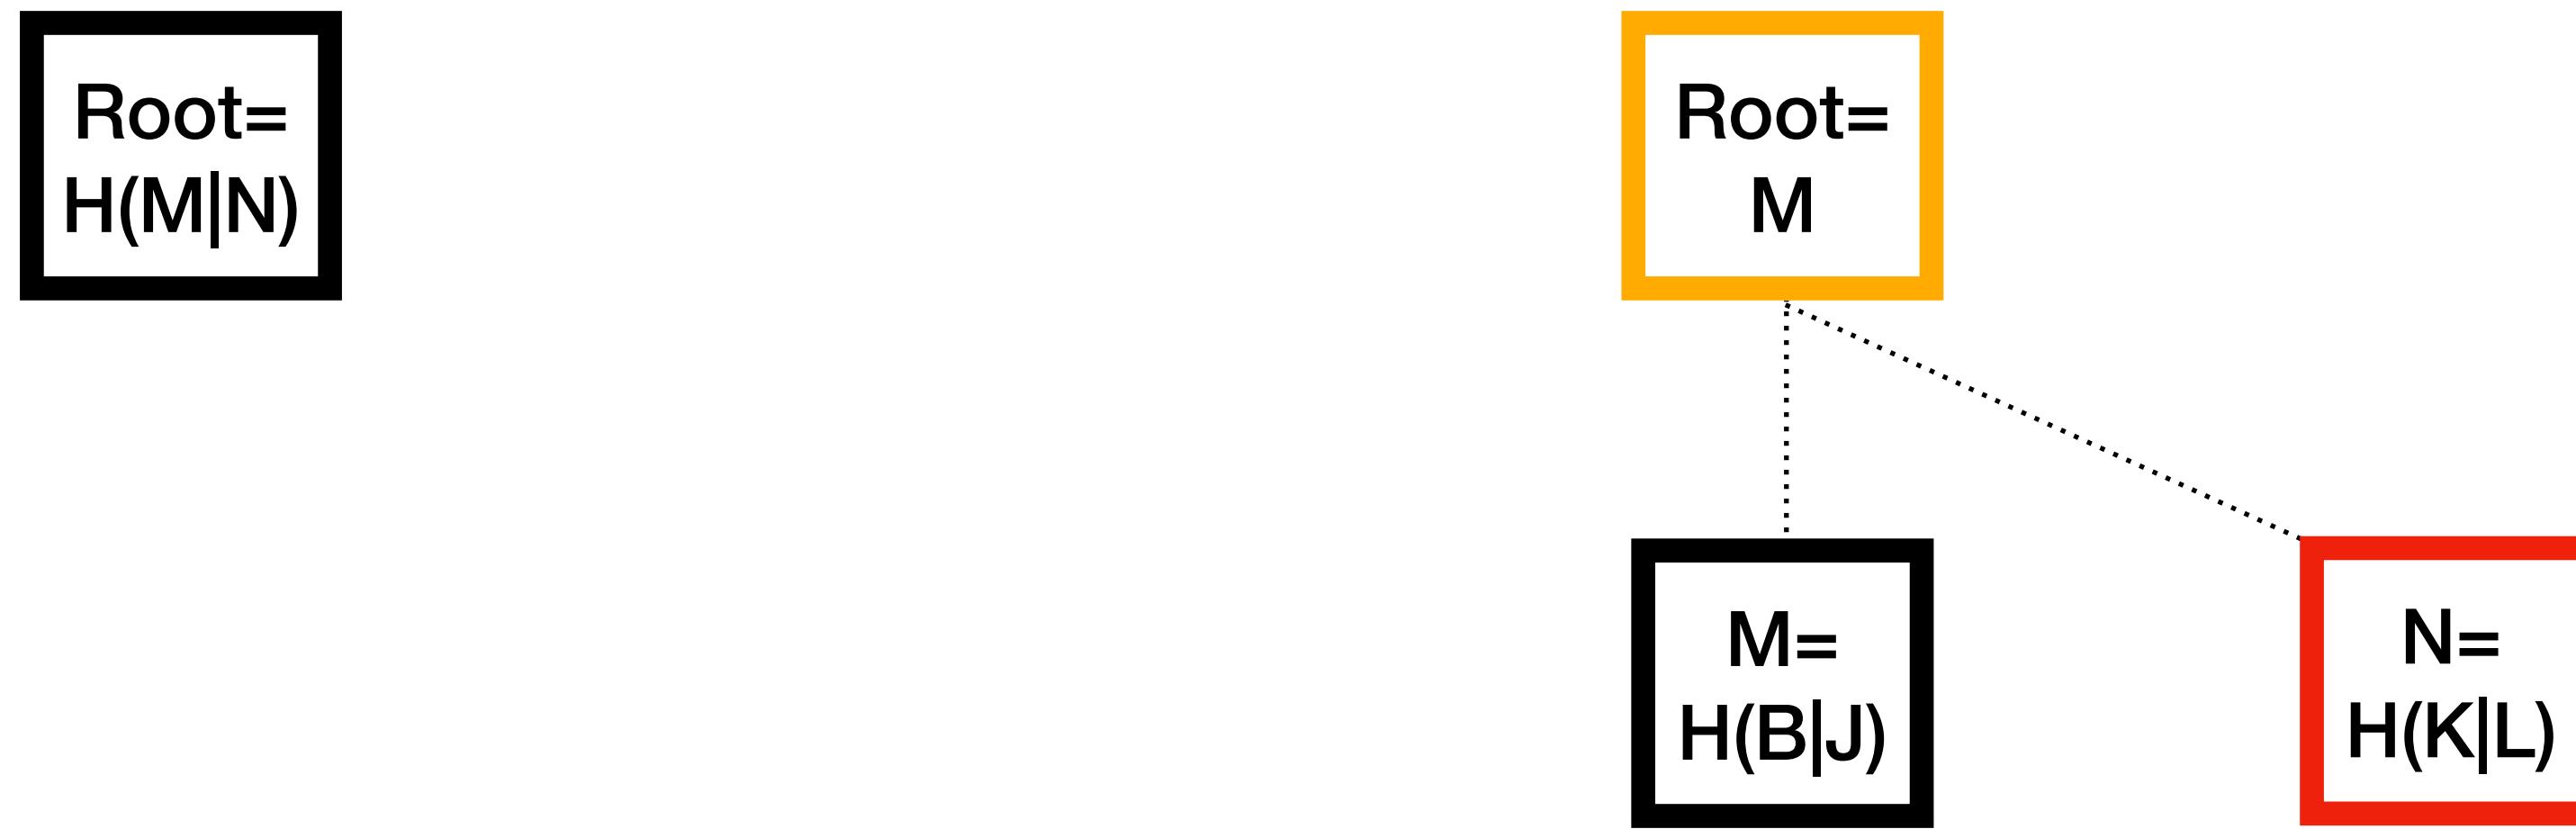
Mark N as node to delete



Remove old nodes for K & L



Move up M

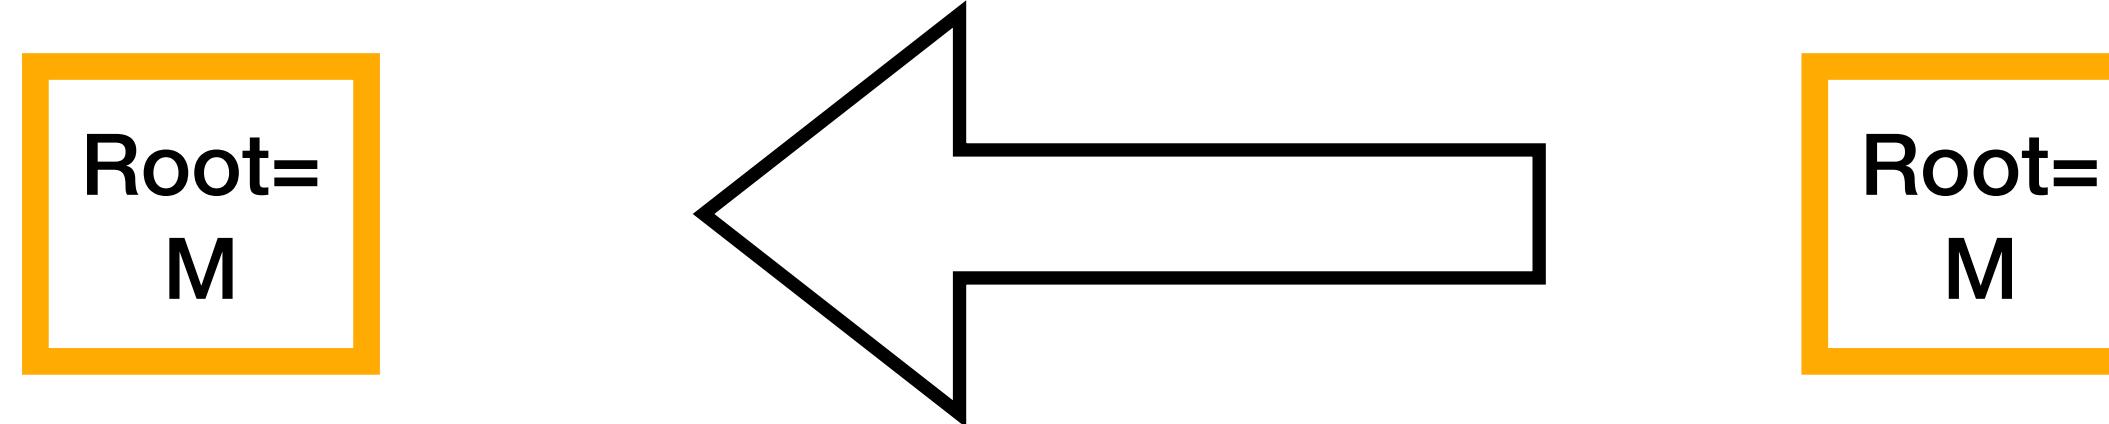


Remove old nodes for M & N

Root=
 $H(M|N)$

Root=
M

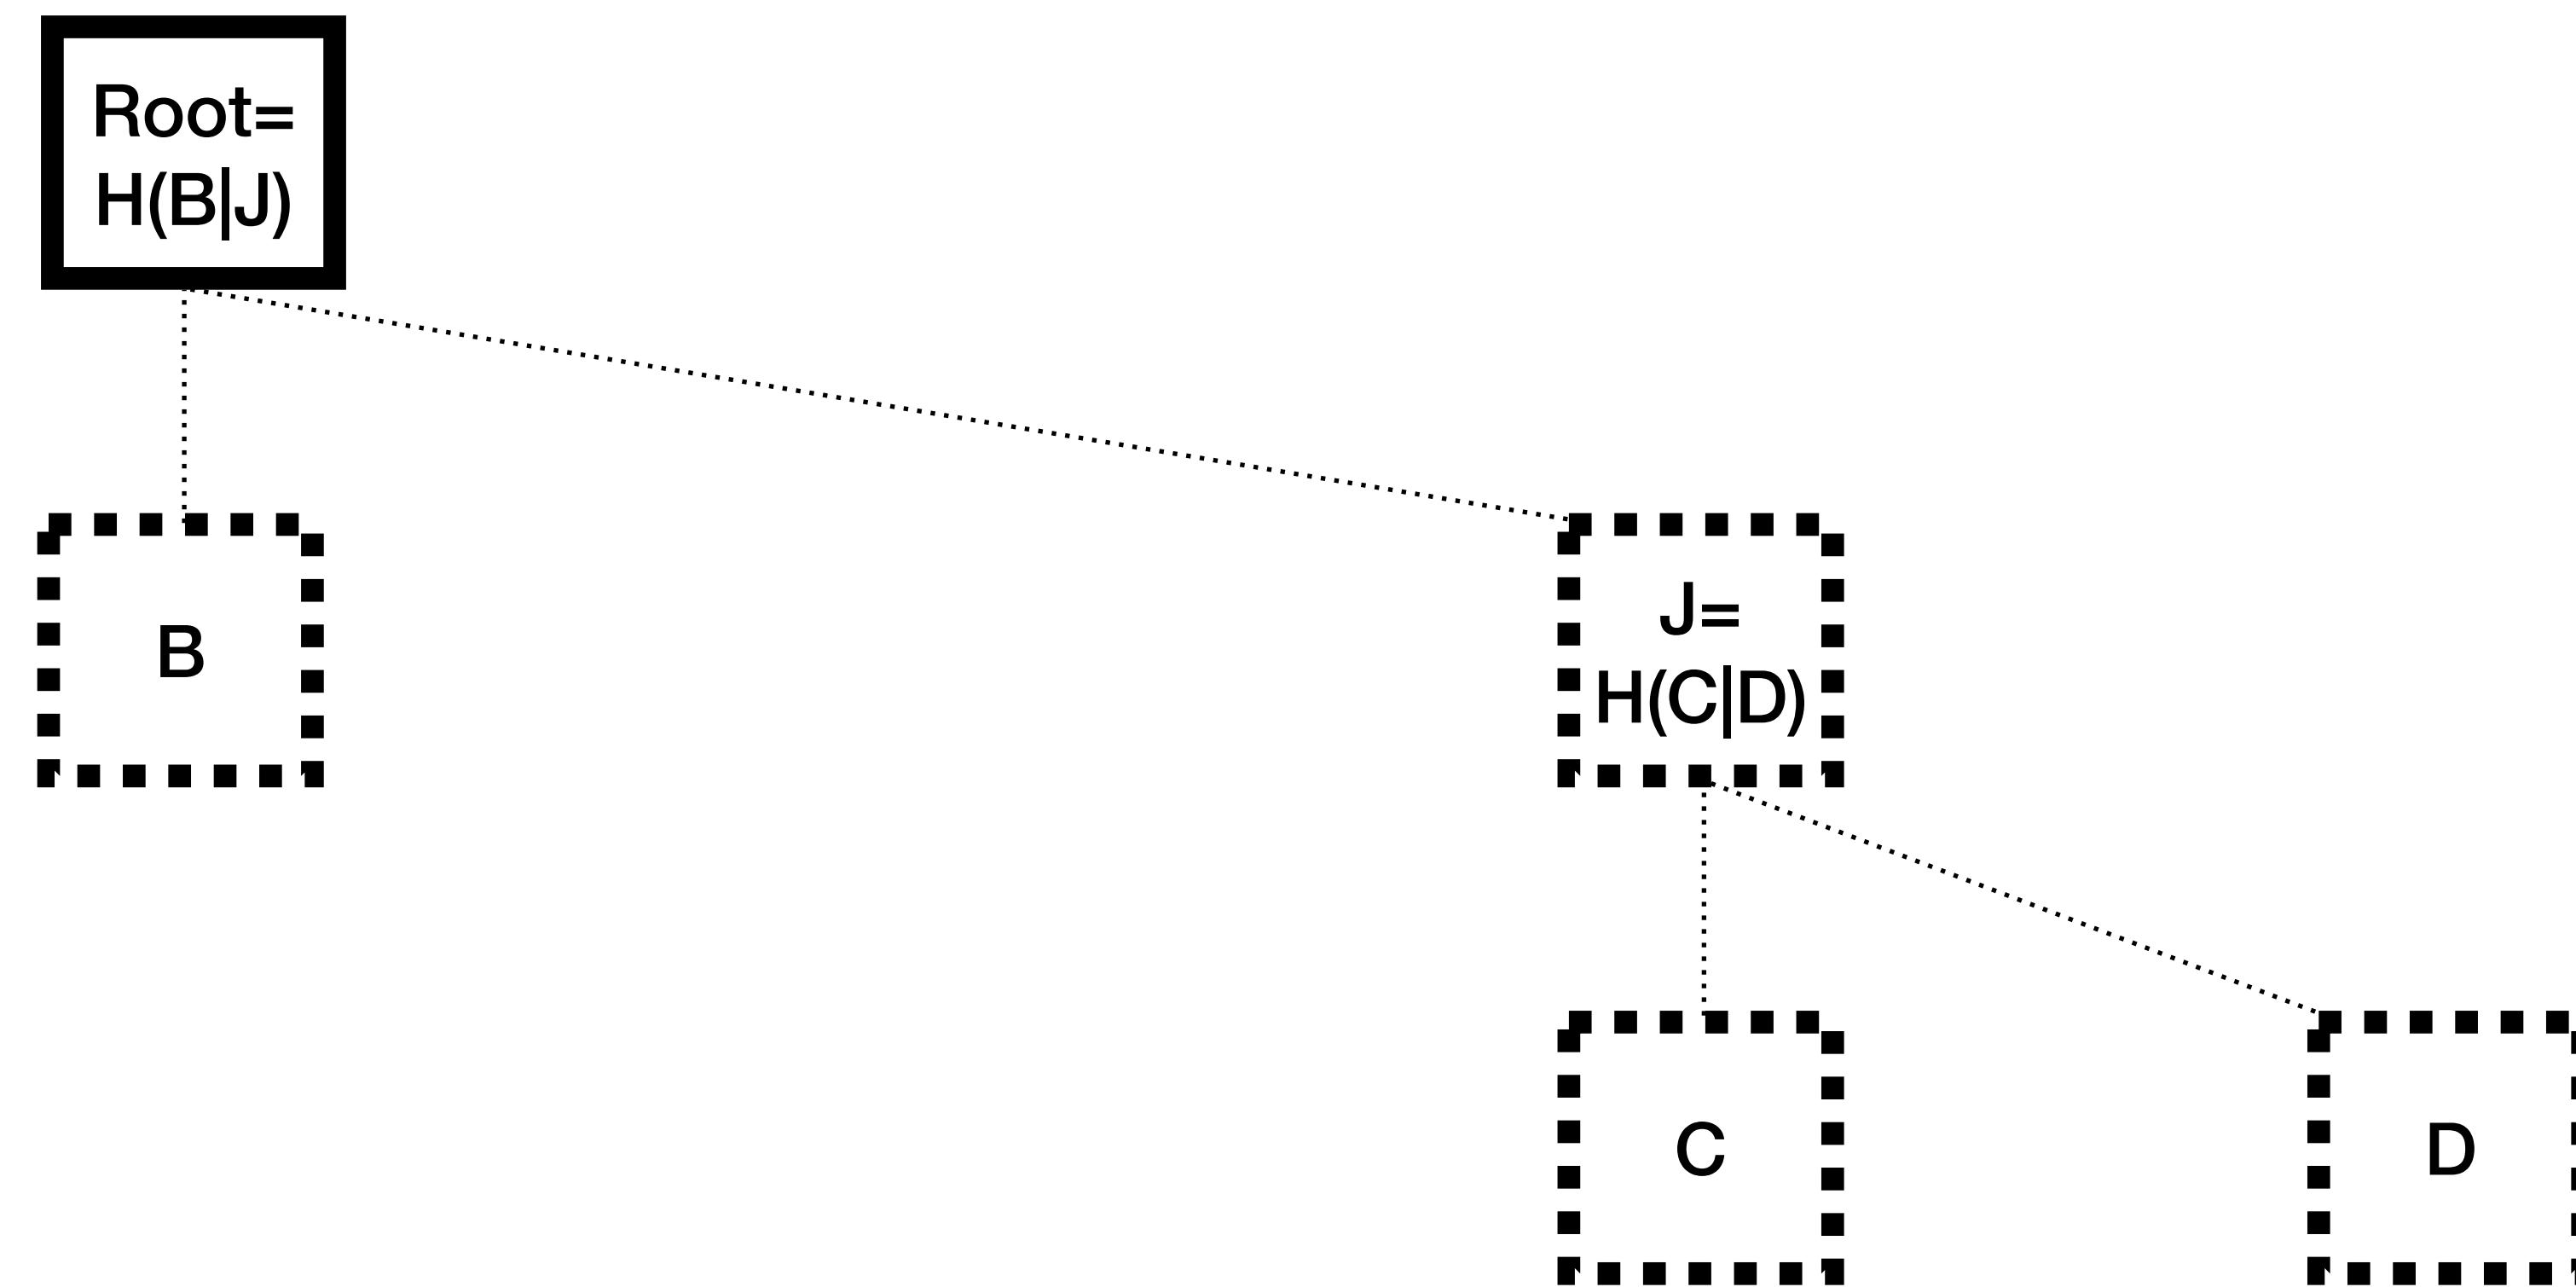
Copy over the new root to be saved



Done

Root=
M

After deleting E, F, G, H



Consensus

Add, Verify, Delete

- That's all the relevant algorithm for consensus
- Implemented in 174 lines of Python:
github.com/utreexo/pytreexo



github.com/kcalvin/alvin/slides-for-tabconf-2023